

Commvault Certified Associate

Interview Preparation Guide

300 Interview Questions & Answers • 7 Topic Categories • CCA Exam Focused

300

Q&A Pairs

7

Categories

CCA

Exam Prep

What's Covered

- Commvault Fundamentals & Architecture
- Installation & Configuration
- Troubleshooting
- CCA Exam Preparation
- Backup & Recovery Operations
- CommCell Administration
- Security & Compliance

Table of Contents

Commvault Certified Associate – 300 Interview Questions

1	Commvault Fundamentals & Architecture <i>CommServe, MediaAgent, client architecture, CommCell, storage policies, agents overview</i>	56 Questions
2	Backup & Recovery Operations <i>Full, incremental, synthetic full, GRT, subclient, deduplication, restore workflows</i>	45 Questions
3	Installation & Configuration <i>CommServe install, MediaAgent setup, agent deployment, storage library configuration</i>	40 Questions
4	CommCell Administration <i>Schedules, alerts, reports, user groups, RBAC, workflow engine, capacity licensing</i>	40 Questions
5	Troubleshooting <i>Job failures, media errors, event viewer, log analysis, performance debugging</i>	35 Questions
6	Security & Compliance <i>Encryption, RBAC, GDPR, ransomware protection, audit trail, MSP multi-tenancy</i>	30 Questions
7	CCA Exam Preparation <i>Key concepts, scenario questions, architecture review, exam tips & practice questions</i>	54 Questions

300 Total Questions

7 Categories

CCA Certified Ready

Commvault Fundamentals & Architecture

Q1

What is Commvault and what problem does it solve?

Answer:

Commvault is an enterprise data protection and information management platform that provides unified backup, recovery, replication, archiving, and cloud management capabilities. It solves the problem of data sprawl across hybrid environments by delivering a single platform to protect, manage, and recover data across on-premises, cloud, and SaaS environments. It eliminates the need for multiple point solutions by centralising all data management operations.

Example:

A hospital with 500 VMs, 200 physical servers, Office 365, and AWS workloads uses Commvault as a single pane of glass – one platform managing all backup policies, retention, recovery, and compliance reporting across every environment.

Q2

What are the core components of the Commvault architecture?

Answer:

Commvault's architecture consists of: (1) CommServe – the central management server and SQL database, the brain of the entire environment; (2) MediaAgent – the data mover that manages storage and moves backup data; (3) Client – any system being protected (VM, physical server, application); (4) CommCell – the collective term for all components managed by a single CommServe; (5) iDataAgent – the software agent installed on clients to facilitate backup and recovery; (6) Storage – disk libraries, tape libraries, cloud storage targets connected to MediaAgents.

Example:

In a 3-tier deployment: CommServe at headquarters manages the CommCell. MediaAgents at each remote site move backup data to local storage. Clients (Windows servers, Linux servers, VMs) have iDataAgents installed. CommCell Console provides unified management.

Q3

What is the CommServe and what are its responsibilities?

Answer:

The CommServe is the central management server of the Commvault environment. It hosts the CommServe database (Microsoft SQL Server) that stores all configuration, job history, and metadata. Responsibilities: (1) Scheduling and orchestrating all backup, restore, and auxiliary copy jobs; (2) Managing policies (storage, schedule, retention); (3) Authentication and authorisation; (4) Reporting and audit logging; (5) License management; (6) Communications hub for all MediaAgents and clients. The CommServe does not move backup data – it only coordinates operations.

Example:

CommServe schedules a nightly backup of 300 VMs at 10 PM. It sends job instructions to MediaAgents and clients, monitors job progress, updates the SQL database with job results, and sends failure alerts to administrators – without touching a single byte of backup data.

Q4

What is a MediaAgent in Commvault?

Answer:

A MediaAgent (MA) is the data movement engine in Commvault. It reads data from clients, deduplicates and compresses it, and writes it to storage (disk, tape, or cloud). It manages DDB (Deduplication Database) for dedup-enabled storage. Key roles: (1) Data mover between clients and storage; (2) DDB management; (3) Encryption of backup data; (4) Auxiliary copy operations; (5) Data verification. Multiple MediaAgents provide load balancing and redundancy. MediaAgents can be physical or virtual.

Example:

During backup of 100 VMs: MediaAgent receives backup streams from ESXi hosts via Commvault VSA proxy, deduplicates the data (50% reduction), compresses it (30% additional), encrypts with AES-256, and writes to a disk library. DDB on the MediaAgent tracks all unique blocks.

Q5

What is a CommCell and how is it different from CommServe?

Answer:

A CommCell is the logical grouping of all Commvault components (CommServe, MediaAgents, Clients) managed by a single CommServe instance. CommServe is the physical server and software; CommCell is the administrative domain it manages. A CommCell has one CommServe (can have standby). Large enterprises may have multiple CommCells managed through CommCell Console's multi-CommCell view. The CommCell ID uniquely identifies the environment.

Example:

Enterprise with global operations: CommCell-NA (CommServe in New York) manages 500 clients in North America. CommCell-EMEA (CommServe in London) manages 300 EMEA clients. Both CommCells visible from a single Command Center using multi-CommCell management.

Q6

What is an iDataAgent (iDA) in Commvault?

Answer:

An iDataAgent is the client-side software module installed on protected systems that provides application-aware backup and recovery. Different iDataAgents handle different workloads: (1) File System iDA – protects files and folders; (2) Virtual Server Agent (VSA) – protects VMware, Hyper-V, Azure VMs; (3) SQL Server iDA – application-aware SQL backup; (4) Exchange iDA – Exchange mailbox protection; (5) Oracle iDA – Oracle RMAN integration; (6) SAP iDA – SAP landscape protection. Each iDA understands its application's API for consistent backups.

Example:

SQL Server with SQL iDA: Commvault communicates with SQL VSS writer for backup, creates application-consistent snapshots, backs up transaction logs every 15 minutes for point-in-time recovery, and enables granular restore of individual databases or tables – capabilities unavailable with a generic file backup.

Q7

What is a Storage Policy in Commvault?

Answer:

A Storage Policy defines where backup data is stored and how it is managed. It contains: (1) Primary copy – where initial backup data lands; (2) Secondary copies – for replication or tape copies; (3) Retention rules – how long data is kept; (4) Deduplication settings; (5) Encryption settings; (6) Copy precedence – which copy is used for restores. Multiple subclients can share the same storage policy. Storage policies link clients to MediaAgents and storage libraries.

Example:

Storage Policy 'Finance-SP': Primary copy on local disk library (30-day retention, dedup enabled). Secondary copy on tape (7-year retention for SOX compliance). Tertiary copy on Azure Blob (1-year retention for DR). All Finance subclients assigned to this policy – unified retention and copy management.

Q8

What is a Subclient in Commvault?

Answer:

A Subclient is the fundamental backup unit in Commvault – it defines what data is backed up, when, and how. A subclient belongs to an iDataAgent and specifies: (1) Content – files, databases, VMs included in backup; (2) Storage policy – where data goes; (3) Schedule policy – when backups run; (4) Backup type – full, incremental, differential; (5) Pre/post scripts; (6) Filters – exclusions. Default subclients are created automatically; custom subclients can be created for granular control.

Example:

SQL Server with 4 databases: Default subclient backs up all databases. Custom subclient 'HRDB' backs up only HRDatabase with 4-hour log backup frequency (vs 24-hour for other DBs). Different retention policies can be applied to each subclient independently.

Q9

What is the Deduplication Database (DDB) in Commvault?**Answer:**

The Deduplication Database (DDB) is a specialised database maintained by the MediaAgent that stores unique block signatures (hashes) for all deduplicated backup data. When new backup data arrives, each block is hashed and compared against the DDB – if the hash exists, only a reference pointer is stored (not the data again). This eliminates redundant data across multiple backups and clients. The DDB resides on the MediaAgent's local storage and must be on fast I/O (SSD recommended).

Example:

100 VMs with Windows OS share the same OS blocks. Without dedup: 100 copies of Windows system files (e.g., 20GB × 100 = 2TB). With DDB: Windows blocks stored once, 99 copies are pointers □ 20GB stored instead of 2TB. Dedup ratio: 100:1 for identical OS content.

Q10

What are backup types supported by Commvault?**Answer:**

Commvault supports: (1) Full backup – complete copy of all selected data; baseline for all other types; (2) Incremental backup – only data changed since the last backup (full or incremental); smallest backup, fastest; (3) Differential backup – data changed since last full backup; larger than incremental but simpler restore; (4) Synthetic Full – created at the MediaAgent by merging previous full + incrementals without re-reading client data; (5) Selective Full – full backup run at specific intervals (e.g., weekly); (6) Transaction Log backup – for databases (SQL, Oracle) for point-in-time recovery.

Example:

Backup schedule for SQL Server: Full backup Sunday 10 PM (baseline). Differential backup Mon–Sat 10 PM (captures week's changes relative to Sunday full). Transaction log every 15 minutes (point-in-time recovery to within 15 minutes). Monthly synthetic full created without touching the production server.

Q11

What is Commvault Command Center?**Answer:**

Command Center is the modern web-based management interface for Commvault (introduced in v11). It replaces/supplements the legacy Java-based CommCell Console with a responsive, browser-accessible UI. Features: (1) Dashboard with health, jobs, alerts; (2) Configuration of clients, storage, policies; (3) Monitoring and reporting; (4) Self-service recovery portal for end users; (5) Multi-CommCell management; (6) Role-based access; (7) Mobile-friendly responsive design. CommCell Console (Java) is still available for advanced configurations.

Example:

Enterprise admin accesses Command Center from any browser without Java. End users access a self-service portal to restore accidentally deleted files from their laptops – without contacting IT. Backup health dashboards give management real-time SLA compliance visibility.

Q12

What is the Virtual Server Agent (VSA) in Commvault?**Answer:**

The Virtual Server Agent (VSA) is the Commvault iDataAgent for virtualisation platforms. It provides agentless backup of VMs by interfacing with hypervisor APIs (VMware VADP, Microsoft VSS for Hyper-V). Key features: (1) Image-level backup of entire VMs; (2) Application-consistent snapshots (VMware quiescing, VSS); (3) Granular file-level restore from VM backups; (4) Live recovery (IntelliSnap + instant boot); (5) Cross-hypervisor recovery (restore VMware VM as Hyper-V); (6) Cloud replication to AWS/Azure. The VSA proxy can be deployed on Windows or Linux.

Example:

VMware environment with 200 VMs: VSA proxy communicates with vCenter to create crash-consistent/quiesced snapshots of each VM. Backup data sent to MediaAgent. From a single VM backup, admin can: restore entire VM, restore specific files, recover a deleted Exchange email, or mount the backup for instant access.

Q13

What is IntelliSnap in Commvault?

Answer:

IntelliSnap is Commvault's hardware snapshot integration framework. It coordinates storage array-level snapshots (NetApp SnapVault, EMC RecoverPoint, Pure Storage, HPE Nimble, etc.) with Commvault's backup workflow. IntelliSnap: (1) Triggers array snapshot for near-instant, low-impact backup; (2) Catalogues the snapshot in CommServe for searchable recovery; (3) Optionally copies snapshot to backup storage (backup copy job); (4) Enables instant VM recovery from array snapshots; (5) Reduces backup window to near-zero. Array snapshots remain accessible for fast local recovery.

Example:

200TB Oracle database on NetApp SAN. Traditional backup: 18-hour backup window. IntelliSnap: triggers NetApp snapshot in 2 seconds (near-zero impact on Oracle), catalogues it in CommServe, then background backup copy job copies snapshot data to tape for long-term retention. Backup window: 2 seconds vs 18 hours.

Q14

What is Commvault HyperScale X?

Answer:

HyperScale X is Commvault's hyperconverged backup infrastructure — a software-defined, scale-out backup storage and management platform. It runs on commodity hardware (or purpose-built nodes) and provides: (1) Built-in deduplication and compression; (2) Scale-out storage (add nodes to expand); (3) Integrated MediaAgent and storage in one appliance; (4) Built-in Commvault software; (5) Simplified deployment and management; (6) Global deduplication across all nodes. Target: enterprises wanting all-in-one backup appliance without separate storage arrays.

Example:

Retail chain replacing aging tape infrastructure: Deploy 3 HyperScale X nodes (each with 96TB raw, ~400TB usable after dedup). Commvault software pre-installed. All 500 stores backup to HyperScale X nodes. Easy scale-out: add node 4 when capacity needed. No separate storage management required.

Q15

What is Commvault's approach to cloud backup and management?

Answer:

Commvault provides comprehensive cloud support: (1) Backup to cloud — use AWS S3, Azure Blob, GCP Storage as backup targets (cloud library); (2) Backup from cloud — protect AWS EC2, Azure VMs, GCP instances; (3) SaaS protection — Microsoft 365, Google Workspace, Salesforce; (4) Cloud replication — replicate backups to cloud for DR; (5) Cloud migration — move workloads to/from cloud; (6) Cloud cost optimisation — intelligent tiering to cheaper cloud storage (Glacier, archive tiers); (7) Air-gap — immutable cloud storage for ransomware protection.

Example:

Enterprise cloud strategy with Commvault: On-prem VMs backed up to local MediaAgent (30-day retention). Auxiliary copy to AWS S3 Standard (90-day retention). Lifecycle policy moves to S3 Glacier after 90 days (7-year retention). Microsoft 365 data backed up directly from cloud to cloud. Single CommServe manages all.

Q16

What is the Commvault CommCell Console?

Answer:

CommCell Console is the traditional Java-based rich-client management interface for Commvault. It provides comprehensive configuration and management capabilities including: (1) Client and agent configuration; (2) Storage library management; (3) Policy creation and management; (4) Job monitoring; (5) Alert configuration; (6) Reports; (7) Advanced troubleshooting. While Command Center (web) is the modern interface, CommCell Console provides deeper configuration options. It requires Java Runtime Environment and can connect to any CommServe.

Example:

Advanced configuration requiring CommCell Console: configuring NDMP tape libraries, setting up disk library properties with specific block sizes, configuring auxiliary copy with advanced throttling, setting up CommServe DR, and advanced deduplication database management — features not yet available in the web-based Command Center.

Q17

What is Commvault's retention policy and how does it work?

Answer:

Retention in Commvault is managed at the storage policy copy level. Two retention criteria work together – BOTH must be satisfied before data is pruned: (1) Retention Days – data retained for minimum X days; (2) Retention Cycles – minimum number of full backup cycles retained. Data is pruned only when BOTH criteria are met. This ensures you never lose data prematurely. Extended retention can be applied to specific jobs (infinite retention for compliance). Retention is managed by the Data Aging job on the CommServe.

Example:

Retention: 30 days AND 2 cycles. Scenario: full backup runs monthly. After 30 days, first criterion met. But only 1 cycle has elapsed – second criterion not met. Data retained until a second full cycle completes (at 60 days). Pruning occurs only when 60 days AND 2 cycles both satisfied.

Q18

What is a Schedule Policy in Commvault?

Answer:

A Schedule Policy is a reusable template that defines when and how frequently backup jobs run. Instead of configuring individual schedules per subclient, a schedule policy is created once and applied to many subclients. Schedule types: (1) Automatic – continuous incremental triggered by changed blocks; (2) Full – scheduled full backup frequency; (3) Incremental – frequency of incremental backups; (4) Differential – frequency of differential backups; (5) Log – frequency of transaction log backups. Schedules can be time-based (specific hours) or interval-based.

Example:

Schedule Policy 'Production-Standard': Full backup Sunday 10 PM, Incremental Mon-Sat 10 PM, Transaction logs every 30 minutes. Applied to 50 SQL Server subclients – all follow the same schedule. Changing the policy changes all 50 subclients simultaneously. Individual overrides allowed per subclient if needed.

Q19

What is auxiliary copy in Commvault?

Answer:

Auxiliary copy is the process of replicating backup data from one storage policy copy to another (e.g., disk to tape, disk to cloud, local to remote). It runs as a separate job managed by the MediaAgent. Use cases: (1) Disk-to-tape for long-term retention; (2) Local-to-remote for DR; (3) Disk-to-cloud for cloud archival; (4) Selective copy – copy only specific jobs (e.g., monthly fulls only to tape); (5) Synchronous copy – inline copy during backup for continuous data protection. Auxiliary copy ensures multiple copies of data exist at different locations.

Example:

Storage policy with 3 copies: Primary (local disk, 30 days) □ Auxiliary copy to Secondary (remote disk, 90 days, runs nightly) □ Auxiliary copy to Tertiary (AWS Glacier, 7 years, runs weekly for monthly fulls only). Data automatically flows through the copy hierarchy.

Q20

What is data deduplication in Commvault and what are the types?

Answer:

Deduplication eliminates redundant data to reduce storage consumption. Commvault supports: (1) Source-side deduplication – dedup happens on the client/proxy before data is sent to MediaAgent; reduces network bandwidth consumption; (2) MediaAgent-side deduplication – dedup happens at the MediaAgent when writing to storage; (3) Inline deduplication – dedup during the backup job itself; (4) Post-process deduplication – dedup after backup data is written. The DDB (Deduplication Database) stores block hashes. Dedup ratios of 10:1 to 50:1 are common in enterprise environments.

Example:

Source-side dedup for WAN backup: 100 branch offices backing up over WAN to central MediaAgent. Without source dedup: 500GB/day per branch × 100 = 50TB WAN traffic. With source dedup: each branch sends only unique blocks – WAN traffic reduced to ~5TB/day (10:1 ratio). Critical for WAN-constrained environments.

Q21

What is a disk library in Commvault?**Answer:**

A disk library is a storage resource configured in Commvault that uses disk-based storage (local disk, NAS, SAN) as backup target. Types: (1) Magnetic disk library – traditional HDD-based; (2) DDB-enabled disk library – supports Commvault deduplication; (3) Cloud library – cloud object storage presented as a disk library; (4) HyperScale – software-defined scale-out disk library. Disk libraries are managed by MediaAgents and can have multiple mount paths across different storage systems. They support simultaneous read/write for concurrent jobs.

Example:

Disk library 'Primary-DL': mount path 1 on NetApp NFS share (50TB), mount path 2 on Windows share (30TB). Total capacity: 80TB. MediaAgent 'MA-01' manages this library. Deduplication enabled – effective capacity 800TB with 10:1 ratio. 200 clients write to this library simultaneously via multiple mount paths.

Q22

What is Commvault's approach to ransomware protection?**Answer:**

Commvault ransomware protection features: (1) Air-gap – immutable backup copies stored offline or in cloud with Object Lock (S3, Azure Blob); (2) Anomaly detection – ML-based detection of unusual backup patterns (spike in changed data indicating encryption); (3) Honeypot files – Commvault places decoy files; if modified, ransomware alert fires; (4) Immutable storage – WORM-enabled storage targets; (5) Isolated recovery environment – spin up clean VMs from backup without connecting to infected network; (6) Point-in-time recovery – restore to any backup point before infection; (7) Threat scan – scan backup data for malware before recovery.

Example:

Ransomware event at 3 AM. Commvault anomaly detection fires alert at 3:05 AM (200% spike in changed data rate). Admin isolates infected VMs. Threat scan identifies clean restore point from 2:45 AM (before encryption). Isolated recovery environment: VMs restored in clean VLAN, validated, then reconnected to production. 15-minute data loss vs hours or days with traditional backup.

Q23

What is a Storage Policy Copy in Commvault?**Answer:**

A storage policy copy is a specific instance of backed-up data within a storage policy. A storage policy contains one or more copies: (1) Primary copy – the initial backup destination; always present; (2) Secondary copy – created by auxiliary copy; can be selective or synchronous; (3) Tertiary copy – another level for tape or cloud. Each copy has its own: retention settings, storage target (disk/tape/cloud), MediaAgent assignment, deduplication settings, and encryption settings. Copies enable the 3-2-1 backup rule within a single storage policy.

Example:

Storage policy 'CriticalData-SP' copies: Primary copy (local SSD disk library, 14-day retention), Secondary copy (remote NAS via auxiliary copy, 90-day retention), Tertiary copy (AWS S3, selective – monthly fulls only, 7-year retention). Three copies, two media types, one off-site – 3-2-1 rule satisfied.

Q24

What is Commvault's CommServe Disaster Recovery (DR)?**Answer:**

CommServe DR protects the CommServe database – the most critical component. Methods: (1) CommServe DR backup – automatic daily export of CommServe DB to a protected location; (2) Standby CommServe – passive CommServe with current DB replica, activated in emergency; (3) Live Sync – real-time CommServe DB replication to a standby; (4) CommServe DR restore – restoring CommServe from DR backup. If CommServe fails without DR: all backup metadata lost, jobs cannot run, restore operations blocked. CommServe DR is mandatory in production deployments.

Example:

CommServe DR configuration: Daily CommServe DB export to: (1) Local disk path on CommServe, (2) Remote UNC path on MediaAgent, (3) FTP/cloud location. If CommServe hardware fails: restore Windows OS on new server, install Commvault CommServe, restore DB from DR backup. CommCell operational in 2-4 hours with all job history and config intact.

Q25

What are the Commvault ports used for communication?**Answer:**

Key Commvault communication ports: (1) TCP 8400 – default tunnel/daemon port between all components; (2) TCP 8401-8403 – additional service ports for CommServe services; (3) TCP 8600-8699 – MediaAgent data transfer ports; (4) TCP 443 – Command Center (HTTPS); (5) TCP 80 – Command Center (HTTP, redirects to 443); (6) TCP 1433 – SQL Server (CommServe database); (7) TCP 8403 – Communication Manager. All communication between CommServe, MediaAgents, and clients uses these ports. Port 8400 is the primary tunnel – must be open in firewalls.

Example:

Firewall rules for Commvault: Clients ↔ CommServe: TCP 8400-8403. Clients ↔ MediaAgent: TCP 8400, 8600-8699. CommServe ↔ SQL Server: TCP 1433. Admin workstations ↔ Command Center: TCP 443. Missing port 8600-8699 between client and MediaAgent: backup jobs fail at data transfer phase despite management connectivity.

Q26

What is the Commvault Workflow Engine?**Answer:**

The Workflow Engine is a process automation framework built into Commvault that allows creating automated, conditional workflows for complex operations. Workflows can: (1) Automate multi-step DR processes; (2) Create approval workflows (e.g., restore requires manager approval); (3) Send notifications at specific steps; (4) Integrate with external systems via REST API; (5) Conditionally execute steps based on outcomes; (6) Automate routine tasks (daily health checks, cleanup). Workflows are created visually in CommCell Console or via XML.

Example:

Restore approval workflow: End user requests file restore via Command Center ↔ Workflow triggers ↔ Manager receives email approval request ↔ If approved: restore automatically executed ↔ Completion notification sent to user. No admin involvement after workflow designed. Audit trail of every restore request, approval, and execution.

Q27

What is Commvault Live Sync?**Answer:**

Live Sync is Commvault's replication feature that maintains an up-to-date copy of VMs or data at a DR site by continuously replicating changes after each backup job. It provides: (1) Near-current VM copies at DR site (RPO = backup interval); (2) Automatic synchronisation after each backup; (3) Recovery point selection at DR site; (4) Test failover capability; (5) Support for VMware, Hyper-V, and Azure targets. Live Sync differs from traditional replication – it replicates within the backup framework, leveraging existing backup data.

Example:

Exchange server backed up hourly by Commvault. Live Sync configured to AWS: after each hourly backup, Commvault automatically syncs the latest backup to AWS EC2 instance (pre-provisioned). At DR, recovered Exchange is always within 1 hour of production. Test failover: boot DR Exchange, validate email delivery, power off – production unaffected.

Q28

What is the Commvault Content Store?**Answer:**

Content Store (formerly Content Store Mailbox) is Commvault's archiving solution for email, files, and SharePoint. It captures content from mailboxes, file servers, and SharePoint into an indexed archive for long-term retention, compliance search, and eDiscovery. Features: (1) Single-instance storage (stores one copy of duplicated attachments); (2) Full-text search and eDiscovery; (3) Legal hold; (4) Stubbing – replaces archived content with a small stub link; (5) Compliance-friendly WORM storage; (6) User self-service archive search.

Example:

Legal requirement: retain all financial emails for 7 years. Content Store archives all Finance department emails: removes attachments from Exchange (stubs), stores single-instance copies in WORM storage. Exchange storage reduced by 60%. Legal hold applied to CFO's mailbox for ongoing investigation. eDiscovery search for 'acquisition' emails returns results in seconds from 7-year archive.

Q29

What is a Commvault client group?

Answer:

A client group is a logical collection of clients (servers, VMs, desktops) in Commvault used for policy management and operations. Benefits: (1) Apply storage policies, schedule policies, and settings to all group members at once; (2) Run reports on a group; (3) Execute bulk operations (install agents, run backups); (4) Dynamic groups — clients auto-added based on rules (OS type, location, tags); (5) Smart client groups — based on Active Directory OUs. Client groups simplify management of large environments with hundreds of clients.

Example:

Client groups: 'Production-SQL-Servers' (45 SQL clients), 'Dev-VMs' (200 dev VMs), 'Remote-Offices' (150 laptop clients). Storage policy change: update 'Production-SQL-Servers' group □ all 45 SQL servers updated simultaneously. Dev VMs get different (less expensive) storage policy. Remote offices get WAN-optimised policy with source dedup.

Q30

What is Commvault's support for Microsoft 365?

Answer:

Commvault protects Microsoft 365 (Exchange Online, SharePoint, OneDrive, Teams) using native Microsoft Graph API and Exchange Web Services. Features: (1) Full mailbox backup (email, calendar, contacts, tasks); (2) SharePoint site and document library backup; (3) OneDrive for Business backup; (4) Teams conversations and files; (5) Granular restore — single email, attachment, file version; (6) eDiscovery and compliance search; (7) Legal hold; (8) Data can be backed up to on-premises storage or cloud. No agents required — cloud-to-cloud or cloud-to-on-prem backup.

Example:

Microsoft 365 backup for 5,000-user enterprise: All Exchange Online mailboxes backed up daily (4 versions retained). Deleted mailboxes retained 90 days after M365 purges. SharePoint sites backed up weekly. User accidentally deletes an entire SharePoint document library: Commvault restores it in 15 minutes from yesterday's backup — Microsoft's 90-day recycle bin not required.

Q31

What are Commvault alerts and how are they configured?

Answer:

Commvault alerts notify administrators of events requiring attention. Alert categories: (1) Job alerts — backup completed, failed, succeeded with errors; (2) Media alerts — library offline, media full, media errors; (3) System alerts — CommServe unreachable, disk space low, service failure; (4) SLA alerts — clients missed backup SLA; (5) Security alerts — failed login attempts, unauthorised access. Alerts can trigger: email, SNMP trap, RSS feed, Event Log, pager. Alerts configured in CommCell Console □ Alerts or Command Center □ Manage □ Alerts.

Example:

Alert configuration: Job failure □ email to backup-team@company.com + PagerDuty webhook. SLA miss □ email to manager@company.com. Disk library >80% full □ email + SNMP to SolarWinds. CommServe unreachable □ SMS to on-call. Each alert type has configurable thresholds and recipient lists for granular notification routing.

Q32

What is Commvault's global deduplication?

Answer:

Global deduplication in Commvault means all clients sharing the same DDB (Deduplication Database) benefit from deduplication across all of them — not just within a single client's backups. If 100 clients have Windows Server OS files, those blocks are stored once in the global DDB, and all 100 clients point to the same blocks. Global dedup maximises storage savings in environments with many similar clients (VDI, Windows servers, standard OS images). One DDB can be shared by multiple MediaAgents (requires DDB synchronisation).

Example:

VDI environment: 500 Windows 10 virtual desktops. Each desktop: 40GB base OS (same image). Without global dedup: $500 \times 40\text{GB} = 20\text{TB}$ OS data. With global dedup: 40GB stored once, 499 copies as pointers = 40GB + metadata. Dedup ratio: 500:1 for OS content. User data (unique per desktop) deduplicated separately — overall ratio: 50:1 across all 500 desktops.

Q33

What is Commvault's role-based access control (RBAC)?**Answer:**

Commvault RBAC provides granular control over who can perform what operations on which data: (1) Roles — define permissions (backup operator, restore operator, admin, report viewer); (2) User groups — collections of users with the same role; (3) Capability groups — fine-grained permissions within roles; (4) Scope — limit access to specific client groups, subclients, or CommCells; (5) Active Directory integration — map AD groups to Commvault roles; (6) Local users — Commvault-native user accounts. Principle of least privilege: users see and operate only what they need.

Example:

RBAC setup: AD group 'Backup-Admins' □ Commvault Admin role (full access). AD group 'App-Owners' □ Restore-Only role scoped to their application clients (can restore own apps, cannot see others). AD group 'Audit-Team' □ Report-Viewer role (view reports only, no operations). User offboarding: disable AD account □ Commvault access revoked immediately.

Q34

What is the Commvault Job Controller?**Answer:**

The Job Controller is the CommServe service responsible for scheduling, queuing, and monitoring all backup, restore, and auxiliary copy jobs. Functions: (1) Evaluates scheduled jobs and queues them at run time; (2) Manages resource allocation (which MediaAgent handles which job); (3) Monitors running jobs and updates job status; (4) Handles job retries on failure; (5) Enforces resource throttling (limits simultaneous jobs); (6) Provides job history and audit trail. Job Controller status visible in CommCell Console □ Job Controller and Command Center □ Jobs.

Example:

Job Controller managing 500-client environment: 50 backup jobs running simultaneously (configured max). 100 more scheduled but queued (resource limit). MediaAgent-01 handling 30 jobs, MediaAgent-02 handling 20 jobs. Job prioritisation: production clients get higher priority than dev clients. Failed jobs automatically retry after 30-minute delay (configurable).

Q35

What is Commvault's support for physical servers?**Answer:**

Commvault protects physical servers through: (1) File System iDA — backup of files and folders on Windows, Linux, AIX, Solaris, HP-UX; (2) Bare Metal Recovery (BMR/1-Touch) — full system recovery to same or dissimilar hardware; (3) Application iDAs — SQL, Oracle, SAP, Exchange agents for app-aware protection; (4) System State backup — Windows system state (registry, boot files, AD); (5) Block-level backup — faster than file-level for large volumes; (6) Live File Recovery — mount backup for instant access. Physical servers require Commvault agent (iDA) installation.

Example:

Critical physical database server (Oracle on AIX): Oracle iDA provides RMAN-integrated backup, incremental-forever strategy, archive log backup every hour. Physical server failure: 1-Touch recovery boots from a minimal OS, connects to MediaAgent, restores the full server to identical or different hardware in 3 hours vs 1 day for manual rebuild.

Q36

What is Commvault's tape library support?**Answer:**

Commvault manages tape libraries (autoloaders, enterprise tape libraries) through: (1) Tape library configuration — add tape drives and slots to CommCell; (2) Tape media management — barcode labelling, cleaning, media replacement; (3) Tape encryption — AES encryption of data written to tape; (4) Multiplexing — multiple backup streams to one tape drive; (5) Media rotation policies — manage tape sets for offsite rotation; (6) NDMP — for NAS backup directly to tape; (7) Tape import/export — move media between sites for DR. Commvault uses ACSLS, Scalar libraries, StorageTek, IBM, and HP libraries.

Example:

Tape library setup: 2 × LTO-8 drives, 50 tape slots. Commvault manages media labelling (CVLT001-CVLT050). Auxiliary copy from disk to tape runs nightly (2-hour window). Tapes rotated offsite weekly (Tapes 1-10 offsite, 11-20 on-site, alternating). Commvault tracks tape location — can request specific tape from vault when needed for restore.

Q37

What is the Commvault Data Classification feature?**Answer:**

Data Classification (DC) in Commvault uses the built-in index and content analysis engine to classify data based on: (1) File type and extension; (2) Content patterns (regex, keywords — e.g., credit card numbers, SSNs, PHI); (3) Age and last access; (4) Ownership; (5) Location. Use cases: (1) Identify sensitive data (GDPR, PII, PHI, PCI); (2) Storage optimisation — identify stale data for deletion or archiving; (3) eDiscovery — locate files matching legal requirements; (4) Compliance reporting — document what data exists where.

Example:

GDPR compliance project: Data Classification scans 500 file servers for files containing EU citizen PII (name + email + date of birth patterns). Report identifies 15,000 files with PII spread across 200 shares. IT classifies and restricts access, legal reviews data, and GDPR register is updated. Ongoing: new PII files flagged automatically weekly.

Q38

What is a Commvault plan and how does it simplify management?**Answer:**

Plans (introduced in Commvault v11) are simplified policy objects that combine storage, schedule, and retention settings into a single, easy-to-apply template. They replace the older multi-object model (subclient □ storage policy □ schedule policy) with a streamlined approach: (1) One plan defines backup frequency, retention, and storage destination; (2) Plans applied to workloads directly; (3) Simpler than traditional storage + schedule policy separation; (4) Server plans, laptop plans, database plans — tailored templates; (5) Override allowed per workload if needed.

Example:

Traditional Commvault: create storage policy + schedule policy + assign to subclient = 5 configuration steps per workload. With plans: create 'SQL-Standard-Plan' (daily full, 30-day retention, disk library) □ assign to SQL Server □ done in 2 steps. 50 SQL servers assigned to same plan in minutes. Changing plan changes all 50 simultaneously.

Q39

What is Commvault's Hedvig distributed storage?**Answer:**

Hedvig (acquired by Commvault) is a software-defined distributed storage platform that provides scale-out block, file, and object storage. Integration with Commvault: (1) Hedvig can serve as a backup target for Commvault MediaAgents; (2) Provides scale-out capacity without proprietary hardware; (3) Supports deduplication and compression; (4) Erasure coding for data protection; (5) Multi-site replication; (6) Cloud tiering to S3/Azure. Hedvig replaces traditional NAS/SAN for backup storage in Commvault deployments.

Example:

Organisation replacing aging NAS with Hedvig: 4 commodity x86 servers with NVMe SSDs form Hedvig cluster (400TB raw, 200TB usable with erasure coding). Commvault MediaAgents write backup data to Hedvig NFS mount paths. Hedvig replicates data between nodes for resilience. Scale-out: add a 5th node to increase capacity without downtime.

Q40

What is the Commvault Web Console?**Answer:**

The Commvault Web Console (also called End User Console or Self-Service Console) is a simplified browser-based interface designed for end users — not administrators. It allows: (1) End users to browse and restore their own files; (2) Browse email archives (Content Store); (3) View backup status of their devices; (4) Request restores (with workflow approval if configured); (5) Download backed-up files. The Web Console requires no client software — purely browser-based. Administrators configure what each user can access through RBAC.

Example:

Employee accidentally deletes project folder from file server. Instead of logging a ticket: (1) Opens Web Console in browser; (2) Browses to file server backup; (3) Navigates to deleted folder; (4) Selects version from yesterday; (5) Restores directly — completed in 5 minutes. IT team uninvolved. Web Console self-service reduces IT helpdesk tickets by 40%.

Q41

What is the Commvault Metrics Reporting Server?

Answer:

The Metrics Reporting Server (MRS) is a centralised reporting component that aggregates data from multiple CommCells for consolidated enterprise reporting. Features: (1) Multi-CommCell consolidated reports; (2) SLA compliance dashboards; (3) Capacity planning reports; (4) Job success/failure trends; (5) License usage tracking; (6) Chargeback reports per department; (7) Scheduled report delivery by email; (8) Custom report builder. MRS collects data from CommCells and provides executive-level reporting without accessing each CommCell individually.

Example:

Enterprise with 5 CommCells (Americas, EMEA, APAC, Dev, DR): MRS aggregates all data. Monthly executive report: 99.3% backup success rate globally, 245TB protected, 3 SLA breaches (identified root causes). Chargeback: Finance department consumed 45TB (billed accordingly). MRS generates this report automatically on the 1st of each month.

Q42

What is Commvault's support for SAP HANA?

Answer:

Commvault protects SAP HANA through: (1) SAP HANA iDA — native integration using SAP HANA backup API (BBAPIClient); (2) Full database backup — complete HANA system backup; (3) Incremental backup — only changed data since last full; (4) Log backup — continuous transaction log backup for point-in-time recovery; (5) Tenant database support — HANA multi-tenant deployments; (6) Cross-instance restore — recover to different HANA instance; (7) Integration with SAP HANA Studio and HANA Cockpit for recovery initiation.

Example:

SAP HANA 2.0 SPS05, 8TB database: Weekly full backup Sunday via Commvault HANA iDA (4-hour window). Incremental Mon-Sat (45-minute window). Log backups every 15 minutes (RPO target). March month-end data corruption: Commvault restores HANA to point-in-time 10:15 AM from 11:00 AM corruption event. 45-minute data loss — within SLA.

Q43

What are Commvault's hypervisor integration capabilities?

Answer:

Commvault supports multiple hypervisors through the VSA framework: (1) VMware vSphere — via VADP (vStorage APIs for Data Protection); (2) Microsoft Hyper-V — via VSS and WMI; (3) Microsoft Azure — Azure VM backup via Azure API; (4) Amazon AWS — EC2 backup using EBS snapshots; (5) Google Cloud Platform — GCP VM backup; (6) Nutanix AHV — AHV native API; (7) Red Hat OpenShift Virtualization; (8) Oracle VM Server. All hypervisors managed through the same VSA framework with consistent policies and recovery options.

Example:

Mixed-hypervisor environment: 300 VMware VMs, 50 Hyper-V VMs, 100 Azure VMs, 30 Nutanix AHV VMs. Single Commvault CommServe manages all. Same storage policy applied across hypervisors. Cross-hypervisor recovery: Commvault restores VMware VM as an Azure VM for cloud DR scenario — disk format conversion handled automatically.

Q44

What is Commvault's encryption capability?

Answer:

Commvault provides end-to-end encryption: (1) In-flight encryption — data encrypted during network transfer between client and MediaAgent (SSL/TLS or Blowfish); (2) At-rest encryption — data encrypted before writing to storage (AES-256, 3DES); (3) Tape encryption — hardware or software encryption for tape media; (4) Key management — Commvault Key Management Server (KMS) or external KMS (KMIP compatible); (5) Encryption granularity — per storage policy copy, per subclient; (6) Pass-phrase or certificate-based key management. Encryption does not impact deduplication ratios if applied correctly.

Example:

Healthcare environment (HIPAA): PHI backup data encrypted with AES-256 (at-rest) and TLS 1.2 (in-flight). Keys managed by Commvault KMS on dedicated server. Tape libraries use hardware LTO encryption. Cloud copies (Azure Blob) encrypted with customer-managed keys. HIPAA audit: all backup data cryptographically protected throughout lifecycle — finding: no gaps.

Q45

What is the Commvault one-pass backup?**Answer:**

One-pass backup (also called SnapProtect One-pass) combines backup and archiving in a single scan. In a traditional approach, backup and archiving scan the same data twice (two passes). One-pass: (1) Scans file system once; (2) In the same pass: backs up data AND archives eligible content (based on age, type, size rules); (3) Eliminates duplicate scanning overhead; (4) Reduces backup window; (5) Creates search index in one pass. Significant for large file servers where scanning takes hours.

Example:

Large file server with 50 million files. Traditional approach: backup scan = 6 hours + archive scan = 6 hours = 12 hours total window.

One-pass: single scan backs up recent files AND archives old files in 6 hours total – 50% reduction in backup window. Critical for servers where backup and archive windows would otherwise overlap.

Q46

What is Commvault's support for Oracle database?**Answer:**

Commvault protects Oracle through: (1) Oracle iDA – integrates with Oracle RMAN (Recovery Manager); (2) Online backup – backup running Oracle database without shutdown; (3) Incremental backup – block change tracking for efficient incrementals; (4) Archive log backup – continuous log backup for point-in-time recovery; (5) RAC support – Oracle Real Application Clusters backup through any node; (6) Oracle DataGuard integration – backup from standby to reduce primary impact; (7) Cross-platform recovery – restore Oracle on different OS/hardware; (8) Tablespace and data file-level restore.

Example:

Oracle 19c RAC (4-node, 5TB database): Commvault Oracle iDA backs up from Node 1 (standby-preferred). RMAN incremental block-change tracking: daily incremental (only changed blocks). Archive logs backed up every 30 minutes. Monthly full backup. Corruption on tablespace USERS: Commvault restores just the USERS tablespace from 6-hour-old backup + archive logs applied for point-in-time to exact transaction – production DBs on other tablespaces unaffected.

Q47

What is Commvault's data verification feature?**Answer:**

Data verification (DV) in Commvault validates that backed-up data is readable and intact. Types: (1) Quick DV – verifies media headers and data checksums without reading all data (fast); (2) Full DV – reads all backup data blocks from storage and verifies integrity (thorough); (3) On-demand – triggered manually for specific jobs or media; (4) Scheduled – automatic regular verification. Results logged in CommServe. Failed DV identifies corrupt or damaged backup data before a recovery is attempted in a real DR situation.

Example:

Weekly data verification schedule on all primary copies: Quick DV on all jobs older than 7 days. Monthly full DV on all monthly full backups. October verification: full DV on September monthly backup reports 3 blocks unreadable (NAS hardware issue identified).

MediaAgent re-reads data from a secondary copy to repair – backup integrity restored before it was needed for a real recovery.

Q48

What is the Commvault Index Server?**Answer:**

The Index Server maintains a searchable index of all backed-up files, emails, and content across the CommCell. Functions: (1) Enables content search across backup data (find a specific file across all clients); (2) Supports eDiscovery and legal hold searches; (3) Enables file restore from search results; (4) Supports Content Store (email archive) search; (5) Powers the Web Console browse-and-restore capability; (6) Global search across multiple CommCells (with multi-CommCell index). Index data is stored on a dedicated server for performance.

Example:

Legal hold scenario: Legal team needs all emails containing 'Project Phoenix' from Finance team for last 5 years. Index Server search: returns 12,000 emails across 15 mailboxes in 45 seconds. Legal hold applied to these results – they cannot be deleted by normal retention policies. Results exported to PST for attorney review. Without Index Server: manual tape restore of 5 years of data – days of work.

Q49

What is Commvault Activate?

Answer:

Commvault Activate is an advanced data intelligence module that goes beyond backup to provide: (1) Sensitive data discovery – automatically identify PII, PHI, PCI data across all protected data; (2) Data risk analysis – identify who has access to sensitive data; (3) Compliance reporting – GDPR, CCPA, HIPAA compliance dashboards; (4) Threat analysis – identify potential insider threats from access patterns; (5) Entity extraction – extract and analyse specific content types (contracts, invoices). Activate uses AI/ML to analyse content indexed by Commvault.

Example:

GDPR compliance with Activate: scans 2 million files across 100 file servers. Identifies 45,000 files with EU PII, 8,000 with financial data, 1,200 with health information. Risk score assigned to each data store. Identifies a public file share accessible by 3,000 users containing 500 sensitive HR documents – immediate action taken. GDPR DPA report generated automatically.

Q50

What is Commvault's Backup Copy feature?

Answer:

Backup Copy is a specific auxiliary copy mode where Commvault creates a traditional backup copy from an IntelliSnap (hardware snapshot) source. After IntelliSnap creates an array snapshot: (1) Backup Copy job mounts the snapshot; (2) Streams data to MediaAgent; (3) Writes to backup storage (disk library, tape, cloud); (4) Enables long-term retention beyond array snapshot limits; (5) Allows transport to offsite/cloud for DR; (6) Data deduplicated and compressed during the copy. Backup Copy decouples the production-facing snapshot (fast) from the storage-efficient long-term copy.

Example:

NetApp array with IntelliSnap: snapshot created in 2 seconds (production zero-impact). Backup Copy runs at 2 AM: mounts NetApp snapshot □ streams data through MediaAgent □ writes to disk library with dedup. Production impact: zero during primary snapshot, minimal during background copy. Long-term copy now available on disk/tape/cloud independent of array snapshot schedule.

Q51

What is the difference between Commvault backup and archiving?

Answer:

Key differences: Backup – (1) Retains all data regardless of age/access frequency; (2) Multiple versions for point-in-time recovery; (3) Short-to-medium retention (days to years); (4) Original data remains on source; (5) Primary goal: recovery from loss. Archiving – (1) Moves infrequently accessed data from source to archive storage; (2) Source storage freed (stub may replace original file); (3) Long-term retention (years to decades); (4) Single copy of archived content; (5) Primary goal: compliance, storage optimisation. Commvault handles both in a unified platform.

Example:

File server management: Files <90 days old □ backed up (recovery focus, retain 90 days, original stays on server). Files >90 days old, unaccessed □ archived (moved to cheap object storage, stub left on file server, 7-year retention for compliance). File server storage freed by 60%. Users access archived files seamlessly via stub (transparent to user). IT saves 40TB of expensive SAN storage.

Q52

What is the Commvault REST API used for?

Answer:

Commvault's REST API enables programmatic access to virtually all CommCell functions: (1) Job management – trigger, monitor, cancel jobs; (2) Client management – add/remove clients, configure agents; (3) Restore operations – initiate file or VM restores; (4) Report generation – pull job history, SLA reports; (5) Storage management – configure libraries, policies; (6) Alert management; (7) RBAC – create users and roles; (8) Monitoring – poll job and system status. API supports JSON responses and uses token-based authentication. Swagger documentation available on CommServe.

Example:

DevOps integration: CI/CD pipeline calls Commvault REST API to: (1) Trigger backup of dev database before deployment; (2) Monitor job completion; (3) If deployment fails: trigger restore from pre-deployment backup automatically. All via REST API – no manual admin action. Backup integrated into DevOps workflow transparently.

Q53

What is Commvault's Continuous Data Replicator (CDR)?**Answer:**

Continuous Data Replicator (CDR) is Commvault's replication feature for continuous, near-real-time replication of data changes from source to target. CDR uses: (1) Journal-based change tracking — captures every write at the block level; (2) Near-zero RPO — changes replicated within seconds; (3) Multiple recovery points at target; (4) Failover and failback orchestration; (5) Support for physical servers, VMs, and file systems. CDR provides RPO in seconds vs backup-based protection (minutes to hours). Used for most critical tier-0 workloads.

Example:

Core banking application on physical RHEL server (RPO requirement: <30 seconds). CDR installed: captures every file system write, replicates to DR server in real time. RPO: 8 seconds average. In disaster: admin initiates failover in CDR console □ DR server mounts replicated volumes □ application starts. Total downtime: <5 minutes. Traditional backup RPO for same system: 4 hours (hourly backup schedule).

Q54

What is a Commvault library and what types exist?**Answer:**

A library in Commvault is a storage resource managed by a MediaAgent. Types: (1) Disk library — disk-based storage (local, NAS, SAN); most common; supports dedup; (2) Cloud library — cloud object storage (S3, Azure Blob, GCP); acts like a disk library; (3) Tape library — physical tape autoloader or enterprise tape library; (4) Standalone tape drive — single tape drive (no automation); (5) Magnetic tape library — legacy classification; (6) Air-gap library — isolated storage with no direct network connection for ransomware protection. Libraries linked to storage policy copies.

Example:

Tiered library strategy: Disk library 'Primary-SSD' (fast, expensive, 30TB, 30-day retention) □ Disk library 'Secondary-NAS' (slower, cheaper, 200TB, 90-day retention) □ Cloud library 'AWS-S3' (Glacier-tiered, unlimited, 7-year retention). Data flows through tiers automatically via auxiliary copy. Each tier optimised for cost and performance.

Q55

What is Commvault Threat Scan?**Answer:**

Threat Scan analyses backup data for malware/ransomware before recovery, ensuring clean data is restored. Process: (1) Mounts a specific backup version from storage; (2) Scans mounted data using integrated AV engine or ClamAV; (3) Reports infected files; (4) Allows recovery of only clean backup versions; (5) Automated scanning of all new backups (optional); (6) Integration with external SIEM/SOAR for incident response. Particularly valuable for ransomware recovery — validates the recovery point is clean before committing to production.

Example:

Ransomware recovery scenario: 10 backup versions available (every hour). Admin initiates Threat Scan on each version from newest to oldest. Versions 1-4 (last 4 hours): infected (ransomware active). Version 5 (5 hours ago): CLEAN. Admin recovers from Version 5 — validated clean recovery point. Without Threat Scan: risk of restoring ransomware-infected backup, re-infecting production immediately.

Q56

What is Commvault's Disaster Recovery Orchestration?**Answer:**

DR Orchestration in Commvault provides automated, documented, and testable DR runbooks: (1) Define VM boot order and dependencies for recovery; (2) Pre/post scripts for each VM or group; (3) IP address remapping (production □ DR IP); (4) Network configuration automation; (5) Failover test capability (non-disruptive); (6) Automated RTO tracking; (7) DR documentation generation; (8) Recovery plan versioning. Integrates with VMware vCenter, Hyper-V, and AWS/Azure APIs. Complements Live Sync for automated failover.

Example:

DR runbook for 3-tier CRM application: Phase 1 — boot Oracle DB VMs (2 VMs, 5-minute wait). Phase 2 — boot JBoss App servers (3 VMs, validate DB connectivity, 3-minute wait). Phase 3 — boot Web servers + F5 load balancer VIP update. RTO target: 25 minutes. Quarterly DR test: runbook executed in isolated environment — all phases automated, RTO 22 minutes achieved — documented and filed.

Backup & Recovery Operations

Q57

What is the difference between a full backup and a synthetic full backup in Commvault?

Answer:

Full backup: reads all data from the source client and sends to MediaAgent for storage — consumes client CPU/IO resources and network bandwidth. Synthetic full: created entirely at the MediaAgent by merging the last full backup with all subsequent incremental backups — does NOT read from the client. Benefits of synthetic full: (1) Zero client impact; (2) No backup window impact on production; (3) Reduced network traffic; (4) Restores from synthetic full are as fast as a regular full. Synthetic full is created as a scheduled operation on the MediaAgent.

Example:

SQL Server backup strategy: Sunday traditional full backup (5-hour window, production impact). Mon–Sat: incremental backups (20 minutes, minimal impact). Saturday night: synthetic full created from Sunday full + Mon–Sat incrementals — client not touched, no extra backup window. Next Sunday's traditional full again provides a fresh baseline. Synthetic full used for all daily restores.

Q58

How does Commvault perform VMware backup?

Answer:

Commvault VMware backup process: (1) VSA proxy communicates with vCenter to initiate a snapshot; (2) VM snapshot created (quiesced for application consistency if VMware tools installed); (3) Commvault reads changed blocks via VMware's Changed Block Tracking (CBT); (4) Data streamed from ESXi host to VSA proxy; (5) VSA proxy sends data to MediaAgent; (6) MediaAgent deduplicates and stores on backup storage; (7) VM snapshot deleted after backup completes. CBT enables efficient incrementals — only changed blocks since last backup are read.

Example:

200 VM environment (vSphere 7.0): VSA proxy (Windows Server) communicates with vCenter. VMware snapshot created, CBT reads only 5GB changed blocks (instead of 200GB full VM). Data flows: ESXi → VSA proxy → MediaAgent → Disk Library. Backup completes in 20 minutes. Full VM backup (no CBT first time): 2 hours. CBT incremental 10× faster after first full.

Q59

What is Changed Block Tracking (CBT) in VMware backup?

Answer:

CBT is a VMware feature that tracks which disk blocks have changed since a reference point (last backup). Commvault uses CBT via the VADP (vStorage APIs for Data Protection) to perform efficient incremental backups — reading only changed blocks instead of the entire VM disk. CBT benefits: (1) Dramatically reduced backup window (read only changed blocks); (2) Reduced network traffic; (3) Reduced storage consumption (only changed data). CBT issues: CBT must be reset after vSphere upgrades, snapshots, or storage migrations. Commvault has automated CBT reset capability.

Example:

1TB VM with 10GB changed daily. Without CBT (full backup): read 1TB, transfer 1TB, 2-hour backup. With CBT incremental: read 10GB, transfer 10GB, 12-minute backup. CBT reduces daily backup time by 90%. Monthly full still reads 1TB — CBT incremental used for all other days. Commvault monitors CBT validity and resets automatically if VADP returns an error.

Q60

How do you perform a full VM restore in Commvault?

Answer:

Full VM restore steps: (1) CommCell Console or Command Center □ Restore; (2) Select the backed-up VM and browse to the required version/date; (3) Choose restore type: Restore to original location (same vCenter/ESXi/datastore) or alternate location; (4) Configure: target datastore, resource pool, network mapping; (5) Power on VM after restore option; (6) Submit restore job; (7) Monitor in Job Controller. Options include: overwrite existing VM, restore to new VM name, restore with different hardware specs, cross-hypervisor restore.

Example:

VM 'HR-AppServer' corrupted by failed OS update. Restore steps: Command Center □ VMs □ HR-AppServer □ Restore □ select yesterday's backup □ Restore to original location □ overwrite VM □ Power on after restore. Job completes in 25 minutes. VM powers on with yesterday's clean state. Total downtime: 25 minutes vs 2-hour rebuild from scratch.

Q61

What is granular recovery in Commvault and what can be recovered?

Answer:

Granular recovery allows restoring specific items from a backup without restoring the entire VM or system: (1) Files and folders – from VM or file system backups; (2) Exchange – individual emails, mailboxes, calendar items, contacts; (3) SharePoint – sites, lists, documents, versions; (4) SQL – individual databases, tables (using table-level restore); (5) Active Directory – individual objects (users, groups, OUs); (6) Oracle – specific tablespaces, data files; (7) OneDrive/SharePoint files; (8) Virtual disk files from within VM backups. Granular recovery eliminates the need for full system restores for small recoveries.

Example:

User accidentally deletes 50 emails from Outlook. Without granular recovery: restore entire Exchange server (hours). With Commvault Exchange iDA granular: browse backed-up mailbox □ locate deleted folder □ restore 50 emails directly to mailbox or PST file – 5 minutes. Production Exchange never taken offline. This is the most common recovery operation in enterprise environments.

Q62

What is Commvault's out-of-place restore and when is it used?

Answer:

Out-of-place restore recovers data to a location different from the original source. Use cases: (1) Source system is still running (avoid overwriting production); (2) Testing recovery without impacting production (restore to alternate server); (3) Migrating data to new location; (4) Forensic investigation (restore to isolated environment); (5) DR scenario (restore to DR site). For VMs: restore to different ESXi host, datastore, vCenter, or cloud. For files: restore to different folder path or different server. For databases: restore to different SQL instance with different name.

Example:

Production SQL Server still running, but last month's database state needed for audit. Out-of-place restore: restore 'Finance_DB' from 60-day-old backup □ to SQL Server 'AUDIT-SQL-01' □ new database name 'Finance_DB_Nov_Audit'. Production SQL unaffected. Auditors access the restored database on AUDIT-SQL-01 independently. Production continues normal operations.

Q63

How does Commvault handle backup of large file servers?

Answer:

Large file server optimisation: (1) Subclient content filtering – exclude temp files, recycle bin, specified extensions; (2) Use-based backup – only back up recently modified files in incrementals; (3) Block-level backup – backup at disk block level (faster than file-level for large volumes); (4) Parallel backup streams – multiple data readers simultaneously; (5) Accelerated File System Scanner (ACFS) – pre-scans file system in parallel to identify changed files; (6) One-pass backup – backup and archive in single scan; (7) Live File Recovery – instant access to mounted backup for browse/restore.

Example:

File server with 50 million files, 100TB data: (1) ACFS pre-scan runs 30 minutes before backup window. (2) Incremental identifies 500,000 changed files (1% daily change). (3) 8 parallel backup streams process changed files. (4) Files >2 years, unaccessed □ archived in same pass (one-pass). Backup window: 2 hours vs 8 hours without optimisation. Storage savings: 30% from archiving old files.

Q64

What is Commvault's Instant Recovery (IR) and how does it work?**Answer:**

Instant Recovery (IR) allows a backed-up VM to be booted directly from backup storage (without full restore to production storage first). Process: (1) Commvault presents backup data to ESXi as a virtual disk (NFS/iSCSI); (2) VM boots from this 'virtual' location within minutes; (3) Writes go to a journal layer (changes don't affect backup data); (4) In the background, data migrated from backup storage to production storage (Storage vMotion-like); (5) VM continues running during migration; (6) Migration complete: VM running entirely on production storage. RTO: minutes instead of hours.

Example:

500GB VM needs urgent recovery. Traditional restore: copy 500GB from backup to production datastore = 3 hours before VM boots.

Instant Recovery: VM boots from backup storage in 4 minutes. Users resume work immediately. Background migration: Commvault copies data to production storage over next 2 hours. Users experience no interruption – RTO reduced from 3 hours to 4 minutes.

Q65

How does Commvault recover a SQL Server database to a specific point in time?**Answer:**

SQL point-in-time recovery with Commvault: (1) Requires SQL iDA with transaction log backup enabled; (2) Select the SQL database subclient for restore; (3) Choose restore type: Point in Time; (4) Specify exact date and time for recovery (e.g., 10:47 AM); (5) Commvault: restores latest full backup before the time point, then replays transaction logs up to the exact time; (6) Options: recover with/without full recovery mode, leave in restoring state for additional log apply. SQL must be in FULL recovery model for log backups.

Example:

Logical corruption in SQL database at 3:27 PM (bad stored procedure executed). Transaction log backups run every 15 minutes.

Point-in-time recovery: (1) Restore full backup from Sunday. (2) Apply incremental (Monday-Thursday). (3) Apply transaction logs up to 3:26 PM exactly. (4) DB recovers to 3:26 PM – 1 minute before corruption. Data loss: 1 minute. Without log backups: would restore to previous night's backup – 15 hours of data loss.

Q66

What is the Commvault Browse and Restore feature?**Answer:**

Browse and Restore allows administrators and end users to browse the backup data like a file system and select specific items for recovery. Features: (1) Point-in-time browse – view backed-up data as it existed at any backup date; (2) File-level browsing within VM backups; (3) Browse without mounting the backup (Commvault reads from index); (4) Multiple browse views: by client, by date, by content type; (5) Drag-and-drop restore from browse results; (6) Available in CommCell Console, Command Center, and Web Console (end user). Enables precise recovery without scripting.

Example:

File server browse: Admin selects file server client □ browse □ navigates to /finance/reports/ □ views versions from last 30 days □ identifies that the March_Report.xlsx was correct on March 15 (corrupted on March 20) □ right-clicks March 15 version □ restore to original location. Entire operation: 2 minutes. Backup job ID or specific version knowledge not required.

Q67

How does Commvault handle backup job failures and retries?**Answer:**

Commvault handles failures through: (1) Automatic retry – failed jobs retry automatically (configurable number of attempts and interval); (2) Partial success – jobs with some successes and some failures log as 'Completed w/ one or more errors'; (3) Kill and restart – if job gets stuck, admin can kill and restart; (4) Orphan job cleanup – CommServe detects and cleans stuck jobs at startup; (5) Pre/post scripts – custom error handling; (6) Alerts – email/SNMP notification on failure; (7) Job Priority – failed jobs can be immediately re-queued at high priority. Root cause analysis via job events/logs.

Example:

Backup job for SQL server fails at 11 PM (SQL VSS writer timeout). Commvault: (1) Logs failure event with error code. (2) Waits 30 minutes. (3) Automatic retry at 11:30 PM. (4) Retry succeeds. (5) Email alert sent: 'Job completed on retry'. Morning review: admin investigates the VSS writer timeout – SQL maintenance job was running at 11 PM. Fix: stagger SQL maintenance to start at midnight. No data loss – retry succeeded.

Q68

What is a backup copy job in Commvault?

Answer:

Backup Copy is an auxiliary copy job specifically designed to copy IntelliSnap (hardware snapshot) data to backup storage. After IntelliSnap creates a snapshot: (1) Backup Copy job mounts the snapshot on a proxy server; (2) Reads snapshot data; (3) Streams data to MediaAgent; (4) Writes to disk library, tape, or cloud with dedup/compression. Backup Copy decouples the fast snapshot operation from the slower long-term copy operation. Backup Copy enables: tape copies from snapshots, cloud copies from snapshots, long-term retention beyond array snapshot limits.

Example:

VMware IntelliSnap schedule: snapshot every 2 hours (24-hour retention on array). Backup Copy schedule: runs at 2 AM daily, copies the most recent snapshot to disk library (14-day retention) and to tape (7-year retention). Result: 6 array snapshots (fast restore) + 14 disk copies (medium-term) + monthly tape copies (long-term). All from a single production snapshot framework.

Q69

How do you back up and restore Active Directory with Commvault?

Answer:

Active Directory protection: (1) Windows System State iDA — backs up AD DS (Active Directory Domain Services) including the Active Directory database (NTDS.dit), SYSVOL, and system state; (2) Granular AD recovery — using Commvault's AD granular recovery, individual objects (users, groups, OUs, GPOs) restored without full AD restore; (3) Authoritative restore — recover specific objects as authoritative (overrides replication); (4) Non-authoritative restore — recover a DC and let replication sync from other DCs; (5) Forest recovery — full AD forest recovery for disaster scenarios.

Example:

50,000-user AD environment. Junior admin accidentally deletes 'Finance' OU (5,000 user accounts and 200 groups). Without granular recovery: authoritative restore of entire AD from backup (risk: lose 8 hours of changes to all other OUs). With Commvault granular AD recovery: restore only the 'Finance' OU and its objects from backup — authoritative for Finance OU only. Other OUs: unaffected. Recovery time: 45 minutes. Change loss: zero outside Finance OU.

Q70

What is the Commvault 3-2-1 backup strategy implementation?

Answer:

3-2-1 rule implementation in Commvault: (1) 3 copies of data — primary backup + 2 secondary copies in storage policy; (2) 2 different media types — e.g., disk (primary) + tape or cloud (secondary); (3) 1 offsite copy — auxiliary copy to remote disk, cloud, or offsite tape. Commvault storage policy with 3 copies automatically enforces this: Primary copy (on-site disk) □ Secondary copy (offsite disk/cloud via auxiliary copy) □ Tertiary copy (tape/archive). Retention rules applied independently per copy.

Example:

3-2-1 implementation: Primary copy □ on-site NAS disk library (30 days retention). Secondary copy □ Azure Blob cloud library via auxiliary copy (90 days, different data type = cloud). Tertiary copy □ LTO-8 tape, offsite storage facility (7 years, different media). All managed by one Commvault storage policy. Ransomware hits on-site disk: Azure copy + tape copy both untouched and available.

Q71

How does Commvault perform Exchange Server backup?

Answer:

Exchange Server backup with Commvault: (1) Exchange iDA — application-aware agent using VSS for consistent backups; (2) Full database backup — entire Exchange database (EDB); (3) Incremental — transaction logs backed up and truncated; (4) Database availability group (DAG) — backup from passive copy to reduce primary server impact; (5) Granular recovery — individual mailboxes, emails, calendar items, public folders; (6) Mailbox-level restore — restore to same or different mailbox; (7) In-place restore — restore deleted item back to original location; (8) PST export — export mailbox backup to PST file.

Example:

Exchange 2019 with 3-node DAG: Commvault backs up from Node 3 (passive copy). Passive backup: zero impact on active mail flow. Database backup nightly (log truncation). Transaction logs backed up every 30 minutes. CEO accidentally deletes sent items from 3 months ago: granular restore □ 15 minutes □ 3-month-old emails back in CEO's Outlook. Entire Exchange offline: database restore from last night's backup □ 2 hours. Both scenarios handled by same Exchange iDA.

Q72

What is Commvault's Live Browse feature?

Answer:

Live Browse allows mounting a backup and browsing/recovering files directly from the mounted backup without a full restore. It presents the backup as a live file system: (1) Commvault mounts the backup volume via NFS or block storage presentation; (2) Admin browses using Windows Explorer or file manager; (3) Copy specific files directly from the mounted backup; (4) No full restore required; (5) Fast access — no need to extract backup archives; (6) Supported for VM backups, file system backups, and application backups. Particularly useful for granular recovery of specific files from large backups.

Example:

100GB VM backup from yesterday. Admin needs to recover 3 specific log files from /var/log/. Traditional restore: extract 100GB backup to staging server (45 minutes). Live Browse: mount backup □ navigate to /var/log/ □ copy 3 files □ unmount. Total time: 4 minutes. Live Browse saved 41 minutes. Production VM doesn't need to be stopped. Original backup data unmodified.

Q73

How does Commvault handle backup of Nutanix AHV VMs?

Answer:

Commvault protects Nutanix AHV through: (1) Native Nutanix API integration — no hypervisor-level agents required; (2) VM snapshot via Nutanix snapshot API; (3) Changed Region Tracking (CRT) — Nutanix equivalent of VMware CBT for efficient incrementals; (4) VSA proxy deployed on or near Nutanix cluster; (5) Application consistency — VMware Tools equivalent scripts via Nutanix guest scripts; (6) Granular file restore from AHV VM backups; (7) Cross-hypervisor restore — recover AHV VM as VMware or Azure VM. All managed through the same VSA framework as VMware.

Example:

Nutanix AHV cluster with 80 VMs: VSA proxy (Linux VM) on the Nutanix cluster. Daily backup: Nutanix snapshot API creates crash-consistent VM snapshots. CRT reads only 8GB of changed blocks per VM (vs 200GB full VM disk). Full backup Sunday, incremental Mon–Sat. Recovery of AHV VM to VMware: Commvault converts QCOW2 disk to VMDK during restore — cross-platform recovery in 30 minutes.

Q74

What is a CommVault backup job report and what does it contain?

Answer:

Commvault backup job reports provide operational visibility: (1) Job summary — total jobs run, success/failure/partial counts; (2) Per-client status — which clients backed up successfully; (3) Data backed up — GB backed up per client; (4) Dedup ratio — storage savings from deduplication; (5) Duration — backup window per job; (6) SLA compliance — clients that missed scheduled backup; (7) Media usage — disk and tape consumption; (8) Errors and warnings — specific failure reasons; (9) Trend analysis — comparison over time. Reports available in Command Center, Metrics Reporting, and via API.

Example:

Weekly backup report emailed to IT Manager every Monday: 485 successful jobs (97%), 12 completed with errors (2.4%), 3 failed (0.6%). Failed: 2 SQL servers (VSS writer issue), 1 VM (snapshot error). Data protected this week: 2.8TB. Dedup ratio: 12:1 (saved 30.8TB of storage). SLA miss: 3 clients (scheduled but not run — power issue during maintenance window). Actionable intelligence in one report.

Q75

How do you restore a VM to a different location (cross-cluster or cross-vCenter) in Commvault?

Answer:

Cross-location VM restore steps: (1) CommCell Console □ Restore □ select VM □ restore to alternate location; (2) Select target vCenter (can be different from source); (3) Select target ESXi host and cluster; (4) Select target datastore; (5) Map source networks to target networks (network mapping); (6) Configure power-on state after restore; (7) Optionally rename VM; (8) Submit job. Commvault handles cross-vCenter authentication independently. Cross-hypervisor: select different hypervisor type (VMware □ Hyper-V or Azure).

Example:

DR scenario: production vCenter unreachable (site failure). Restore VMs to DR vCenter: CommCell Console □ each critical VM □ Restore to Alternate Location □ DR vCenter □ DR cluster □ DR datastore □ map production network to DR network. Network mapping: VLAN100 (prod) □ VLAN200 (DR). Power on: yes. 20 critical VMs restored and running in DR in 45 minutes.

Q76

What is Commvault's data aging job?

Answer:

Data Aging is a scheduled CommServe job that identifies and prunes backup data that has exceeded its retention policy. Process: (1) Commvault evaluates each backup job against retention rules (days AND cycles); (2) Jobs past retention: media (disk blocks, tape files) marked for reuse; (3) DDB entries for aged blocks pruned; (4) Disk space reclaimed; (5) Tape media marked as reusable. Data Aging runs daily by default (usually late night). It never deletes data within retention period. Extended retention jobs (compliance) exempt from normal aging.

Example:

Data Aging runs nightly at 3 AM. Retention policy: 30 days, 2 cycles. Tonight's aging: identifies backup jobs from 31+ days ago that have had 2+ full cycles. Marks 500GB of disk blocks as reusable. DDB updated: 50,000 hash entries removed. Disk library shows 500GB freed in morning. Tape media with all aged jobs: entire tape marked scratch (reusable). Audit log records every deletion for compliance.

Q77

How does Commvault back up and restore Kubernetes workloads?

Answer:

Commvault Kubernetes protection: (1) Application-aware protection — backs up applications as logical units (namespace level); (2) Etcd backup — backs up Kubernetes control plane; (3) Persistent Volume (PV) backup — protects persistent volume data; (4) ConfigMaps and Secrets — captured as part of namespace backup; (5) Helm chart aware — restores applications with proper configuration; (6) Cross-cluster restore — restore to different Kubernetes cluster; (7) Cloud-native: supports EKS, AKS, GKE, and on-premises Kubernetes. Uses Commvault Kubernetes operator deployed in the cluster.

Example:

MongoDB StatefulSet in Kubernetes (EKS): Commvault Kubernetes operator deployed. Namespace 'mongodb-prod' backed up daily: captures persistent volumes (1TB data) + ConfigMaps + StatefulSet definition. Cluster failure: restore MongoDB namespace to new EKS cluster □ Commvault provisions PVs, restores data, recreates StatefulSet. MongoDB operational in 20 minutes on new cluster.

Q78

What is Commvault's wildcard backup and how does it help?

Answer:

Wildcard backup uses pattern matching to define subclient content, automatically including new data that matches the pattern without manual subclient updates. Examples: (1) Include all directories matching /data/project_*. (2) Include all databases on SQL instance (auto-discovers new databases); (3) Include all VMs with tag 'backup-required'; (4) Include all mailboxes in an Exchange DAG. Wildcard subclients auto-protect new resources as they are provisioned — critical for dynamic cloud environments and growing datasets.

Example:

Cloud environment: new VMs provisioned daily with tag 'env:production'. Commvault wildcard subclient: include all VMs with tag 'env:production' in the vSphere inventory. New VM provisioned Monday □ Tuesday morning: already in backup scope □ Tuesday night: first backup completed automatically. No admin action required. Prevents 'forgotten VM' backup gaps that manual subclient management causes.

Q79

How does Commvault handle incremental forever backup strategy?

Answer:

Incremental forever (continuous incremental): Only one full backup is taken initially; all subsequent backups are incremental. Benefits: (1) Minimal backup window (only changed data); (2) Reduced storage consumption; (3) Consistent RPO; (4) No disruptive weekly full backup windows. Challenges: long restore time if full + many incrementals must be merged. Commvault addresses with synthetic full — periodically merges full + incrementals into a new full at MediaAgent without touching source. Result: fast restore + incremental forever strategy.

Example:

File server with 10TB data: Initial full backup Sunday (8-hour window). Every subsequent day: 30-minute incremental. Monthly: synthetic full created at MediaAgent (no client impact). Restore performance: from synthetic full + few incrementals (fast). Client impact: 8 hours once, then 30 minutes daily forever. Without synthetic full: restore would require replaying months of incrementals — slow.

Q80

What is the CommCell Migration feature in Commvault?**Answer:**

CommCell Migration allows moving clients and their backup history from one CommCell (CommServe) to another. Use cases: (1) CommServe consolidation – merge multiple CommCells into one; (2) CommServe upgrade – migrate from old CommServe hardware; (3) Regional reorganisation – move clients between regional CommCells; (4) M&A – integrate acquired company's CommCell. Migration includes: client configuration, job history, backup data references (data physically moved or re-catalogued). Requires matching Commvault versions between CommCells.

Example:

M&A scenario: Acquired company has 50 clients on their CommCell (v11 SP20). Parent company runs v11 SP24. Migration plan: (1) Upgrade acquired CommCell to SP24. (2) Use CommCell Migration wizard to transfer 50 clients + job history to parent CommCell. (3) Backup data on acquired MediaAgents re-associated to parent CommServe. (4) Acquired CommServe decommissioned. One CommCell now managing 550 clients total.

Q81

How does Commvault manage backup jobs during peak business hours?**Answer:**

Resource management during business hours: (1) Activity control – pause or throttle backup jobs during business hours; (2) Schedule windows – configure backups to run only during off-peak hours; (3) Throttling – limit network bandwidth and data transfer rates; (4) Priority management – production backup jobs de-prioritised during peak; (5) Blackout windows – prevent backups from running in specific periods (e.g., batch processing); (6) Smart bandwidth throttling – automatically adjust to available bandwidth; (7) Stream limits – limit concurrent backup streams to reduce host impact.

Example:

Business-hours policy: 8 AM–6 PM blackout window for backup jobs. Exception: critical SQL transaction logs every 15 minutes (very low impact). 6 PM: backups resume at full throughput. WAN throttle: business hours (9 AM–6 PM) limit WAN replication to 20Mbps; overnight (6 PM–9 AM) uncapped. Result: no backup impact on business operations; all SLAs met overnight.

Q82

What is the Commvault 'kill job' and 'suspend job' function?**Answer:**

Job management operations in Commvault: (1) Suspend Job – pauses a running job; data already backed up is saved; job can be resumed later (useful during peak hours or resource contention); (2) Kill Job – terminates a job immediately; partially completed data depends on job type (file backups: data backed up so far is valid; DB backups: may be incomplete); (3) Hold – similar to suspend but prevents job from re-running by schedule; (4) Resume – restart a suspended job from where it paused; (5) All job management from CommCell Console Job Controller or Command Center Jobs.

Example:

Large file server backup starts at 10 PM, still running at 8 AM (business hours). Action: Suspend Job at 8 AM □ job pauses, retains all data backed up so far. Resume at 6 PM □ job continues from the last file processed, not from beginning. Without suspend: kill job and lose all progress. With suspend: 16-hour backup completed in two sessions without restarting from scratch.

Q83

How does Commvault handle network failures during a backup job?**Answer:**

Network failure handling: (1) Checkpointing – Commvault checkpoints job progress at regular intervals; (2) If network drops: job enters a waiting state (not failed immediately); (3) Configurable timeout – wait X minutes for network to recover before failing; (4) Auto-restart – on network recovery, job resumes from last checkpoint (not beginning); (5) Partial data validity – data backed up before failure is valid and recoverable; (6) Retry – if job fails, automatic retry at configured interval; (7) Bandwidth throttling with QoS – prevents Commvault backup from consuming all network, reducing failure likelihood.

Example:

Large backup running over WAN link when network drops for 15 minutes at 2 AM. Commvault: waits 20 minutes (configured timeout). Network recovers at 2:15 AM. Commvault resumes from last checkpoint (file 450,000 of 600,000 – not starting over). Job completes successfully at 3:30 AM. Without checkpointing: restart from file 1 □ 5-hour restart. With checkpointing: 1-hour resume.

Q84

What are backup streams in Commvault and how do they affect performance?

Answer:

Backup streams are parallel data transfer channels between a client and MediaAgent. Multiple streams process data in parallel, improving throughput. Settings: (1) Maximum streams per MediaAgent – limits concurrent load; (2) Streams per subclient – controls parallelism for a specific backup; (3) Database streams – for SQL/Oracle, multiple database files backed up in parallel streams; (4) More streams = faster backup but more resource consumption; (5) Too many streams = resource contention, slower backup. Optimal streams depend on: client CPU/IO, MediaAgent capacity, network bandwidth, storage throughput.

Example:

SQL Server with 8 databases and 4 backup streams configured: 4 databases backed up simultaneously (4 streams). Total backup time: 2 hours. With 1 stream: 8 hours (sequential). With 8 streams: 1.5 hours. But with 16 streams: 2.5 hours (storage I/O bottleneck – too many parallel reads from SAN). Optimal for this environment: 4-6 streams. Performance testing determines the right value.

Q85

How does Commvault's SLA (Service Level Agreement) monitoring work?

Answer:

SLA monitoring in Commvault tracks whether protected workloads receive backup within the configured SLA window. Configuration: (1) SLA period – how often backup should run (daily, every 12 hours); (2) SLA window – acceptable backup completion time; (3) Missed SLA – client that didn't receive a backup within the period; (4) SLA reports – dashboard showing protected vs unprotected workloads; (5) SLA alerts – notification when clients miss SLA; (6) Command Center SLA dashboard – visual overview of compliance %. SLA monitoring enables proactive backup management.

Example:

SLA configured: all 500 clients must have at least one successful backup in the last 24 hours. Monday morning report: 487 clients met SLA (97.4%), 13 missed SLA (8 failed, 5 never started). Failed: 8 clients had VSS errors. Never started: 5 clients offline during backup window (weekend patching). Actions: fix VSS issues (5 clients), schedule re-try for patched clients. SLA report emailed to IT manager automatically.

Q86

What is a CommVault protected object?

Answer:

A protected object is any entity that has been configured for data protection in Commvault: VMs, physical servers, databases, mailboxes, files, SharePoint sites, cloud resources. Protected object management includes: (1) Discovery – automatic discovery of new VMs, databases, mailboxes; (2) Protection status – actively protected or unprotected; (3) Last backup time; (4) Next scheduled backup; (5) Associated storage policy and plan; (6) Protection gap analysis – identify unprotected resources. Command Center provides a Protected Objects view for visibility into all assets and their protection status.

Example:

Command Center Protected Objects view: 1,250 VMs total. Protected: 1,180 (94.4%). Unprotected: 70 (5.6%). Investigation: 40 are dev/test VMs (intentionally excluded), 15 are new VMs not yet assigned to a plan, 15 were removed from subclient by mistake. Action: assign 15 new VMs to plan, re-add 15 removed VMs to subclient. Coverage improves to 98%.

Q87

What is Commvault's support for NAS backup?

Answer:

Commvault NAS backup approaches: (1) NDMP (Network Data Management Protocol) – standard NAS backup protocol for NetApp, EMC Isilon, HPE; (2) Network File System – mount NAS share on proxy, back up as file system; (3) IntelliSnap NAS – array-level snapshots of NetApp/EMC volumes via storage API; (4) NDMP 3-way copy – NAS → MediaAgent → tape without backup server processing data; (5) NAS-to-NAS copy – direct NAS replication for DR; (6) Cloud tiering – archive old NAS data to S3/Azure Blob. NDMP preferred for large NAS systems – efficient and low network overhead.

Example:

NetApp ONTAP cluster (500TB NFS file shares): Commvault NDMP backup via NDMP agent on NetApp. Data flows directly from NetApp to tape library (NDMP 3-way) – CommServe coordinates but data bypasses Windows MediaAgent servers. 100TB/night backup window. NetApp IntelliSnap snapshots hourly for fast local recovery. Daily NDMP backup to tape for offsite retention. Combined: hourly local recovery + daily tape copies.

Q88

How does Commvault handle backup during VM snapshots in VMware?

Answer:

Snapshot management during VMware backup: (1) Commvault creates a VMware snapshot before reading changed blocks (CBT); (2) Snapshot holds writes in a delta file during backup; (3) After data is read: Commvault automatically removes the snapshot; (4) If snapshot removal fails: Commvault detects and alerts (orphaned snapshots cause VM performance issues); (5) Quiesced snapshots — VMware Tools quiesces the VM before snapshot for application consistency; (6) Non-quiesced — crash-consistent snapshot for VMs without VMware Tools; (7) Commvault manages snapshot lifecycle automatically.

Example:

Backup of 50 VMs simultaneously: 50 VMware snapshots created (quiesced where Tools available). Backup reads CBT data while snapshot holds in-progress writes. All 50 backups complete in 2 hours. Snapshots automatically deleted by Commvault after each VM's backup completes. One VM: snapshot deletion fails (vCenter API timeout). Alert fires. Admin manually deletes orphaned snapshot. Commvault's snapshot tracking ensures no orphaned snapshots go undetected.

Q89

What is Commvault's archive pruning and how does it work?

Answer:

Archive pruning removes archived data that has exceeded its configured retention period. Process: (1) Commvault evaluates archived content age against archive storage policy retention; (2) Stub files on source servers: if archived data pruned, stub file optionally removed or replaced with 'not available' notice; (3) Archived data in Content Store: deleted from archive storage; (4) Legal hold: archived content under legal hold is NEVER pruned regardless of retention; (5) Pruning schedules — Data Aging job handles archive pruning alongside backup pruning. Critical: ensure legal holds are applied before pruning deletes compliance-relevant content.

Example:

Email archive retention: 7 years. After 7 years: email archive pruning runs. Except: 200 mailboxes under legal hold — those archived emails protected indefinitely. For the other 4,800 mailboxes: emails older than 7 years pruned from archive storage. Stub files on Exchange: updated to 'item not available — retention expired'. Storage freed: 2TB of archive space reclaimed. Audit log records every pruning action for compliance.

Q90

What is the difference between a full backup and a synthetic full backup in Commvault?

Answer:

A full backup copies all selected data from the source, while a synthetic full backup is constructed by merging the latest full backup with subsequent incremental backups — without reading from the source client. Synthetic fulls reduce backup window and network load while maintaining a complete restore point.

Example:

In a weekly schedule, you run a true full on Sunday and incrementals Monday–Saturday. On the following Sunday, Commvault creates a synthetic full from Sunday's original full + 6 incrementals, reducing backup time from 8 hours to under 1 hour.

Q91

What is Granular Recovery Technology (GRT) in Commvault?

Answer:

GRT allows recovery of individual items — such as emails, mailboxes, files, or VM disks — from application-consistent backups without restoring the entire backup set. It eliminates the need for item-level backup agents for each application by parsing backup data at recovery time.

Example:

Instead of restoring an entire Exchange database to recover one deleted email, GRT lets you browse the backup, locate the specific email, and restore it directly to the mailbox in minutes.

Q92

Explain the role of the MediaAgent in the backup data path.

Answer:

The MediaAgent is responsible for moving data between the client (or data mover) and storage. It manages data deduplication, compression, encryption, and writing to media. It acts as the bridge between CommServe-controlled jobs and the physical storage targets.

Example:

When a VM backup runs, the VSA proxy reads VM data and sends it to the MediaAgent. The MediaAgent deduplicates and compresses the stream, then writes the result to a disk library or cloud target.

Q93

What is a Subclient in Commvault and why is it important?

Answer:

A Subclient is the fundamental unit of data protection in Commvault. It defines the content to be backed up, the schedule, the storage policy, and backup options for a specific set of data. Each client can have multiple subclients, allowing granular control over different data sets.

Example:

A SQL Server client may have separate subclients for the production database (backed up every 4 hours) and the log database (backed up every 30 minutes), each pointing to different storage policies.

Q94

What is Inline Deduplication in Commvault?

Answer:

Inline deduplication analyzes and removes duplicate data blocks during the backup stream, before writing to storage. This reduces storage consumption and network bandwidth in real time, as opposed to post-process deduplication which deduplicates after writing.

Example:

During a backup of 100 VMs sharing a common OS template, inline deduplication identifies duplicate OS blocks across all VMs and writes each unique block only once to the SIBB, reducing storage use by up to 90%.

Q95

How does Commvault handle backup of virtual machines — what is the VSA agent?

Answer:

The Virtual Server Agent (VSA) is Commvault's agent for protecting hypervisor environments such as VMware vSphere, Microsoft Hyper-V, and Nutanix AHV. It uses hypervisor-level APIs (like VMware VADP) to create application-consistent snapshots and back up VM data without installing agents inside each VM.

Example:

For VMware, VSA uses VADP Changed Block Tracking (CBT) to identify changed blocks since the last backup, enabling fast incremental backups. A 500GB VM may transfer only 5GB of changed data per incremental job.

Q96

What is the difference between a hardware snapshot and a Commvault snapshot?

Answer:

A hardware snapshot is created by the storage array (e.g., NetApp Snapshot, EMC SnapVX) and lives on the same array as the primary data. A Commvault snapshot is a logical snapshot managed by CommServe that can catalog, replicate, and back up hardware snapshots to secondary storage — adding management, scheduling, and retention on top of array capabilities.

Example:

Commvault IntelliSnap creates a NetApp Snapshot, catalogues it in CommServe, then optionally replicates the snapshot to a secondary NetApp system and backs it up to tape — all in one automated workflow.

Q97

What is Auxiliary Copy in Commvault?**Answer:**

Auxiliary Copy is the process of copying backup data from primary storage to secondary or tertiary storage. It is used for creating additional copies for disaster recovery, long-term retention, or offsite storage. It is scheduled independently of the backup job itself.

Example:

Primary backups are written to a local disk library. An Auxiliary Copy job runs nightly to copy all new backup data to a remote cloud library, ensuring a geographically separated copy exists for DR.

Q98

Explain the concept of a Deduplication Database (SIDB) in Commvault.**Answer:**

The Scale-out Intelligent Deduplication Database (SIDB) stores the signature (hash) of every unique data block backed up. Before writing a new block, Commvault computes its hash and checks the SIDB — if the hash exists, the block is not written again, only a reference is stored. The SIDB must be protected and should reside on fast local disk for optimal performance.

Example:

A 1TB backup environment with 70% deduplication ratio stores only 300GB of unique data in the disk library. The SIDB contains millions of hash entries pointing to those 300GB of unique blocks.

Q99

What happens during a CommCell Disaster Recovery if the CommServe is lost?**Answer:**

Commvault's Disaster Recovery (DR) backup automatically backs up the CommServe database, configuration, and policies to a DR folder. If CommServe is lost, it can be rebuilt by installing CommServe on a new server and restoring the DR backup — recovering all client configurations, schedules, job histories, and storage policies. The DR backup should be copied offsite.

Example:

After a server failure, the admin installs CommServe on a new VM, runs the DR restore wizard, and points it to the latest DR backup on a secondary site. Within 2 hours, the full CommCell environment with all 500 clients is operational.

Q100

What is a Storage Policy and how does it relate to backup data flow?**Answer:**

A Storage Policy defines the path backup data takes from the client to storage, including the MediaAgent to use, the primary library, copy rules, retention settings, and deduplication configuration. Every subclient must be associated with a Storage Policy — it is the central data management policy in Commvault.

Example:

A Storage Policy named 'VM-Prod-Policy' routes VM backup data through MediaAgent-01 to a local disk library, retains data for 30 days, and has an auxiliary copy rule to replicate to Azure Blob after 24 hours.

Q101

What is the purpose of the Job Controller in CommServe?**Answer:**

The Job Controller is a core CommServe service that schedules, initiates, monitors, and manages all backup, restore, and auxiliary copy jobs across the CommCell. It allocates resources, handles job queuing, manages concurrency limits, and reports job status.

Example:

When 50 backup jobs are scheduled at midnight, the Job Controller queues them based on resource availability (MediaAgent streams, client connections) and starts them in order — preventing system overload while ensuring all jobs complete within the backup window.

Installation & Configuration

Q102

How do you configure a Commvault disaster recovery (DR) site?

Answer:

DR site configuration: (1) Install CommServe at DR site (standby) – Live Sync replicates DB in real time; (2) Install MediaAgent(s) at DR site for local backup processing; (3) Configure auxiliary copy from primary site disk library to DR site disk library; (4) Configure Live Sync for critical VMs: protects and replicates VM backups to DR site; (5) DR network mapping: map production network to DR network in restore settings; (6) DR runbook: define VM boot order, scripts, IP remapping; (7) Test DR: quarterly non-disruptive test restore to DR site; (8) Documentation: maintain DR runbook with RTOs and RPOs per application.

Example:

DR site setup (primary: Mumbai, DR: Pune): (1) CommServe standby at Pune synced via Live Sync. (2) MediaAgent at Pune with 50TB disk library. (3) Auxiliary copy: nightly, Mumbai disk to Pune disk library (30-day retention). (4) Live Sync: 50 critical VMs replicated to Pune after each backup. (5) DR runbook: Tier-1 apps (Oracle, SAP) boot first, Tier-2 second. (6) Test: quarterly, VMs booted at Pune in isolated VLAN, validated, deleted. RTO: 2 hours for Tier-1.

Q103

What are pre-installation checks for Commvault CommServe?

Answer:

Pre-installation checklist: (1) OS: Windows Server 2016/2019/2022, all patches applied; (2) SQL Server: installed and running (or install with Commvault); (3) .NET Framework 4.8+ installed; (4) Hostname: static IP, DNS entry resolves from all clients; (5) Ports: 8400-8403, 443, 1433 open in local firewall; (6) Disk: OS drive separate from SQL data drive (SSD for SQL); (7) Antivirus: exclude Commvault installation path and SQL data directories; (8) Time: NTP configured (time skew causes certificate issues); (9) User account: local admin for installation; (10) Disable UAC during installation.

Example:

Pre-install validation script output: OS: Windows Server 2022, SQL 2019 installed, .NET 4.8, Hostname COMMSERVE01 resolves, Port 8400 open, SQL drive: D:\ SSD 500GB, AV exclusions: C:\Program Files\Commvault\ excluded, NTP synced (drift <1s). All checks passed – proceed with CommServe installation. Installation completed in 45 minutes.

Q104

How do you configure a Linux MediaAgent in Commvault?

Answer:

Linux MediaAgent configuration: (1) Supported OS: RHEL, CentOS, Ubuntu, SLES; (2) Prerequisites: glibc 2.17+, libstdc++6, required packages installed; (3) Download Unix installer from CommServe software store; (4) Run: ./cvpkgadd -install -silent -cs [CommServe_IP] -cspassword [password] -role MA; (5) Mount paths: configure NFS/NAS mounts on Linux for disk library paths; (6) DDB: SSD mount for DDB path (/ddb partition); (7) Service: Commvault services start automatically, managed via commvault start/stop; (8) Verify: CommCell Console MediaAgents Linux MA shows Online.

Example:

RHEL 8 MediaAgent: prerequisites installed (gcc-c++, libstdc++). Commvault installer downloaded. Run: sudo ./cvpkgadd -install -silent -cs 10.10.1.50 -cspassword Admin@123 -role MA. Installation: 12 minutes. NFS mounts added to /etc/fstab: NAS01:/backup /backup/vol1 nfs defaults. Disk library created on /backup/vol1 (100TB NFS). DDB: /dev/nvme0n1 (1TB NVMe, mounted /ddb). Linux MA online – 300MB/s throughput achieved.

Q105

What is the Commvault service account and what permissions does it need?

Answer:

Commvault requires service accounts for: (1) Commvault services on CommServe/MediaAgent/clients: local system or dedicated domain account; (2) Domain account permissions needed: (a) Log on as a service right; (b) Local admin on CommServe and MediaAgent servers; (c) Read access to backed-up data (file servers); (d) SQL sysadmin for CommServe SQL instance; (3) Client push installation: domain admin temporarily needed; (4) vCenter account: Commvault needs vCenter role with snapshot, datastore, and VApp permissions; (5) Cloud service principal: specific API permissions for AWS/Azure/GCP; (6) Exchange: Commvault service account needs Exchange Organization Admin role for mailbox backup.

Example:

Service account setup: AD account svc-commvault@company.local created. Permissions: (1) Local admin on CommServe, both MediaAgents. (2) SQL sysadmin on CommServe SQL instance. (3) Log on as service right via Group Policy. (4) vCenter role 'CommvaultBackup' with custom permissions. (5) Azure SP: Reader + Storage Blob Data Contributor on backup storage account. (6) Exchange: Organization Management role. Password: 24-char random, stored in CyberArk, 90-day rotation.

Q106

How do you configure CommCell firewall settings for NAT environments?

Answer:

CommCell Firewall configuration for NAT/segmented networks: (1) CommCell Console □ Administration □ CommCell Firewall; (2) Configure pre-defined routes — tell CommServe how to reach clients behind NAT (via tunnel through gateway); (3) Tunnel port: all communication through port 8400 (single firewall rule); (4) One-way firewall: client-initiated connection to CommServe (CommServe cannot reach client directly); (5) Persistent tunnel: Commvault maintains persistent connection to clients behind NAT; (6) Proxy gateway: use a Commvault proxy as communication relay for isolated segments; (7) DMZ deployment: for cloud-hosted clients communicating back to on-premises CommServe.

Example:

Branch office NAT scenario: Branch VMs behind NAT router (private IP 192.168.1.x). CommServe at HQ: public IP 203.0.113.5. Configuration: (1) Commvault firewall enabled. (2) Branch clients: CommServe address set to 203.0.113.5:8400. (3) Clients initiate outbound connection to CommServe (NAT allows outbound). (4) CommServe maintains persistent tunnel through NAT. (5) NAT router: allows outbound TCP 8400 from branch network. No inbound rules on branch firewall needed. Backups work through single outbound port.

Q107

What is Commvault's CommCell migration and when is it used?

Answer:

CommCell migration moves clients, their backup history, and associated data from one CommServe to another. Use cases: (1) Hardware refresh: old CommServe □ new CommServe; (2) CommCell consolidation: merge 3 regional CommCells into 1; (3) M&A: integrate acquired company's CommCell; (4) Version upgrade via parallel migration. Process: (1) Ensure source and target CommCells on same Commvault version; (2) Export client from source CommCell; (3) Import client into target CommCell; (4) Reassociate backup data on MediaAgents; (5) Validate backup jobs on new CommCell; (6) Decommission from source. Backup data optionally migrated (disk) or left on original MediaAgent.

Example:

Consolidation project: 3 CommCells (Region1: 100 clients, Region2: 80 clients, Region3: 60 clients) □ 1 Central CommCell. Process: all 3 CommCells upgraded to SP32 (match Central CommCell). 240 clients exported and imported to Central CommCell over 3 weekends (80/weekend). Backup data: MediaAgents from all 3 regions connected to Central CommServe (no physical data movement needed). Result: 1 CommCell managing 240 clients. 3 old CommServe servers decommissioned. Annual license saving: 3 CommServe licenses.

Q108

How do you configure Commvault for SAP backup?

Answer:

SAP backup configuration: (1) Install SAP iDA on SAP application server; (2) Configure BACKINT interface (SAP's standard backup API): SAP calls Commvault BACKINT during backup jobs; (3) Create Commvault subclient for SAP: specify SAP SID, database, instance; (4) Configure storage policy with appropriate retention; (5) SAP R/3 (ABAP): BACKINT integration with SAP database (MaxDB, Oracle, SQL Server); (6) SAP HANA: separate HANA iDA using HANA backup API; (7) Schedule: align with SAP maintenance windows; (8) BR*Tools integration for Oracle-based SAP; (9) Test: trigger backup from SAP, verify in CommCell Console.

Example:

SAP ECC on Oracle configuration: SAP server (SAPAPP01) with Commvault SAP iDA installed. BACKINT configured: `/usr/sap/SID/SYS/profile/initSID.utl` points to Commvault BACKINT executable. SAP DBA COCKPIT: backup job triggers ☐ Commvault BACKINT receives call ☐ creates subclient backup in CommCell ☐ data to MediaAgent. Schedule: SAP DB backup daily 11 PM (Commvault schedule mirrors SAP). Restore: SAP BRRESTORE calls Commvault BACKINT for file restore. Audit: all SAP backup jobs visible in CommCell Console job history.

Q109

How is Commvault licensed and what are the license types?

Answer:

Commvault licensing models: (1) Capacity-based: licensed by TB of data protected (most common); (2) Per-socket/server: license per physical server socket; (3) VM license: per VM (agent-based or agentless); (4) SaaS license: Microsoft 365, Salesforce per user; (5) Perpetual: one-time purchase + annual maintenance; (6) Subscription: annual or 3-year subscription; (7) Commvault Go: simplified all-in-one subscription (common for SMB); (8) License types: CommServe, MediaAgent, iDataAgent specific. License consumed when client configured for backup. View usage: CommCell Console ☐ License ☐ License Summary.

Example:

Enterprise license audit: 1TB capacity license (covers 1TB of deduplicated protected data). Current usage: 850GB deduplicated protected data. Headroom: 150GB. Growth rate: 50GB/month. Renewal planning: current license expires in 90 days. Action: purchase 2TB capacity license (room for 2 years growth). Savings: negotiate 3-year subscription (15% discount vs annual). License applied: CommCell Console ☐ License ☐ Activate.

Q110

What is the process to configure HyperScale X in Commvault?

Answer:

HyperScale X deployment: (1) Rack and power HyperScale X nodes (minimum 3 nodes for production); (2) Connect to management network and backup network; (3) Run HyperScale X setup wizard (web-based); (4) Configure node cluster: provide IPs, hostnames, network settings; (5) CommServe: either built-in (HyperScale X has option) or external; (6) Storage configuration: Commvault configures distributed storage across nodes automatically; (7) Disk library: automatically created on the scale-out storage; (8) CommCell Console: HyperScale X appears as MediaAgent(s) + disk library; (9) Assign to storage policies. Scale-out: add new nodes via wizard.

Example:

HyperScale X 3-node cluster: Node1 (48 cores, 192GB RAM, 24x12TB SAS, 2x2TB NVMe for DDB+cache). Node2, Node3 same specs. Total raw: $3 \times 288\text{TB} = 864\text{TB}$ raw. Usable after erasure coding (3-way): $\sim 288\text{TB}$. With 15:1 dedup ratio: 4.3PB effective. Setup: 3 hours end-to-end (rack to first backup). Scale-out 6 months later: Node4 added in 30 minutes (hot-add, no downtime). No traditional NAS, no separate MediaAgent required.

Q111

How do you configure Commvault VSA proxy for VMware?

Answer:

VSA proxy configuration: (1) Install Commvault on a Windows Server (the proxy): include Virtual Server Agent component; (2) Proxy requirements: 8+ cores, 32GB RAM, 10GbE NIC; (3) Register proxy as client in CommCell; (4) In vCenter client properties: add proxy to access node list; (5) Proxy access: proxy must be able to reach vCenter (port 443) and ESXi hosts (port 902 – VDDK); (6) VSA proxy role: reads backup data from ESXi via VADP, sends to MediaAgent; (7) Multiple proxies: add multiple proxies for load balancing and redundancy; (8) Proxy VM: can be a VM on the same ESXi cluster (hot-add transport mode: efficient direct VMDK read).

Example:

VMware environment with 400 VMs: 2 VSA proxies configured (PROXY-01 and PROXY-02, both Windows 2022, 8 core, 32GB). Hot-add transport: proxies are VMs on the same vSphere cluster – VMDK attached directly (no network copy). Load balancing: Commvault alternates between proxies. 400 VMs backed up with 20 concurrent jobs (10 per proxy). Each proxy handles 10 VMs simultaneously. Backup window: 4 hours for all 400 VMs (incremental). Without 2 proxies: 8 hours (single proxy bottleneck).

Q112

How do you configure Commvault for Google Cloud Platform (GCP) backup?

Answer:

GCP backup configuration: (1) Create GCP service account with required roles (Compute Viewer, Storage Admin, Cloud SQL Viewer); (2) Download service account JSON key; (3) CommCell Console: New Client □ Virtual Server □ Google Cloud Platform; (4) Upload service account JSON credentials; (5) Commvault discovers GCP projects, regions, instances; (6) VSA subclient: select GCP VMs (by zone, label, instance name); (7) Storage: configure GCP Storage bucket as cloud library or use on-premises disk library; (8) Backup: GCP VMs backed up using persistent disk snapshots via GCP API; (9) Granular: file-level restore from GCP VM backups.

Example:

GCP backup for 80 Compute Engine VMs: GCP service account `svc-commvault@project.iam.gserviceaccount.com` with Compute Viewer + Storage Object Admin. JSON key uploaded to CommCell. Commvault discovers: 80 VMs across 3 regions (us-central1, europe-west1, asia-east1). VSA subclient: include all VMs with label 'backup:true'. Storage: GCP Storage bucket 'company-cv-backup' (Nearline for 30-day retention). Daily incremental backup: 2-hour window. Full backup Sunday. Monthly cost: ~\$800 GCP storage.

Q113

How do you install Commvault client agent using a silent push installation?

Answer:

Silent push installation from CommServe: (1) CommCell Console □ Client Computers □ Install Software □ Client; (2) Select clients by hostname or IP range; (3) Choose iDataAgents to install (File System, SQL, Oracle, etc.); (4) Provide credentials: domain admin for Windows; (5) Select installation path; (6) CommServe pushes installer over port 445 (SMB) and installs remotely; (7) Installation log available in CommCell Console; (8) Post-install: client registers with CommServe automatically; (9) For Linux clients: SSH credentials required for push; (10) Bulk install: install agents on 100+ clients simultaneously using client groups or IP ranges.

Example:

Bulk installation of 50 Windows clients: CommCell Console □ Install Software □ select 50 hostnames (imported from CSV) □ File System iDA + SQL Server iDA □ credentials: `svc-commvault-install@company.local` (domain admin) □ install path: `C:\Program Files\Commvault`. Installation initiated: parallel push to 50 clients simultaneously. 45 minutes later: 48 successful, 2 failed (WMI firewall blocked). Fix: enable WMI on 2 clients, retry. All 50 online. Default subclients created automatically.

Q114

What is the Commvault Cloud Gateway and how does it work?

Answer:

Cloud Gateway is a Commvault component that enables backup to/from cloud without requiring public IP or inbound connectivity. Functions: (1) Acts as a secure intermediary between on-premises CommCell and cloud; (2) Handles authentication and encrypted traffic to cloud providers; (3) Manages bandwidth throttling for cloud transfers; (4) Enables multicloud: single gateway managing AWS + Azure + GCP cloud libraries; (5) Deployment: can be on-premises or cloud-hosted; (6) No additional firewall holes required beyond outbound HTTPS (443); (7) Media streaming: optimises upload/download of large backup data to cloud. Available as virtual appliance or software component.

Example:

Hybrid cloud setup: On-premises MediaAgent □ Cloud Gateway (Windows Server, 8 core, 32GB, 10GbE) □ Internet □ AWS S3. Cloud Gateway handles: HTTPS encryption, connection retry on timeout, multipart upload for large backup files, bandwidth cap (100Mbps during business hours, uncapped overnight). Result: on-premises MediaAgent doesn't need direct internet access. All cloud traffic through dedicated Cloud Gateway server. Proxy whitelisted: *.amazonaws.com:443.

Q115

How do you configure Commvault data paths for optimised backup routing?

Answer:

Data paths in Commvault define how backup data flows from client to MediaAgent to storage. Configuration: (1) CommCell Console □ Storage □ MediaAgents □ Data Path; (2) Add data paths: specify which MediaAgent handles which clients (load balancing); (3) Preferred vs failover: configure primary MA and failover MA for each client group; (4) Network interfaces: configure specific NIC on MediaAgent for backup traffic (separate from management NIC); (5) Data path priority: ordered list of preferred routes; (6) Client-side: configure which network client uses for backup (if multi-homed); (7) LAN vs WAN: different data paths for local vs remote clients. Correct data path prevents backup traversing expensive or slow links.

Example:

Multi-site data path: DataCenter1 clients □ prefer MA-DC1 (LAN, 10GbE). Branch clients □ prefer MA-DC1 via WAN (throttled 50Mbps). If MA-DC1 fails □ failover to MA-DC2 (via inter-DC link). Client 'DB-SERVER-01' multi-homed (management: 10.1.1.x, backup: 10.2.1.x): configure Commvault to use 10.2.1.x for backup traffic (dedicated backup NIC). Result: backup traffic never traverses management network.

Q116

What are the steps to configure Commvault for hybrid cloud backup?

Answer:

Hybrid cloud backup setup: (1) On-premises: CommServe + MediaAgent + disk library (primary backup); (2) Cloud library: add AWS S3 or Azure Blob as cloud library (secondary copy); (3) Storage policy: 3-copy policy — primary (on-prem disk), secondary (cloud), optional tertiary (tape); (4) Auxiliary copy: schedule secondary copy to cloud (nightly or continuous); (5) Cloud bursting: if on-prem disk full □ overflow to cloud; (6) Retention tiers: 30 days on-prem, 1 year cloud standard, 7 years cloud archive; (7) Cloud-native workloads: Azure/AWS VMs backed up to cloud storage in same region; (8) M365: backed up directly to Azure Blob (cloud-to-cloud).

Example:

Hybrid architecture: CommServe on-prem. MediaAgent on-prem. Primary: NetApp NAS disk library (50TB, 30-day retention). Secondary: Azure Blob (Hot tier, 90-day retention, nightly auxiliary copy). Tertiary: Azure Blob Archive tier (7-year, monthly fulls, lifecycle policy after 90 days). M365 (5,000 users): backed up directly to Azure Blob (cloud-to-cloud, no on-prem traffic). On-prem VMs: backed up to NetApp, replicated to Azure. Cloud VMs: backed up directly to Azure Blob. Single CommServe manages all.

Q117

How do you configure Commvault for Oracle RAC backup?

Answer:

Oracle RAC backup configuration: (1) Install Oracle iDA on each RAC node or just one node (preferred node for backup); (2) Configure RMAN: Commvault interfaces with RMAN as a backup channel; (3) Block Change Tracking (BCT): enable on Oracle for efficient incrementals; (4) Subclient: specify Oracle RAC SID, select databases; (5) Preferred node: configure Commvault to back up from a specific RAC node (reduce impact on primary); (6) Archive log backup: configure frequency (e.g., every 15 minutes); (7) DG standby: configure backup from DataGuard standby (zero primary impact); (8) Test restore: restore individual tablespace or full RAC database to verify.

Example:

Oracle 19c RAC (4-node, 8TB): Oracle iDA on all 4 nodes. RMAN backup configured – backup runs on Node4 (least busy). BCT enabled: incremental reads only 300GB/day (4% change). Archive logs: Commvault backs up every 15 minutes. Schedule: daily incremental 11 PM, full Sunday 9 PM. DataGuard: backup preferred from standby (Commvault configured to prefer standby node) – zero impact on production 4-node RAC. Test: tablespace USERS restored from last night's backup + 4 archive log cycles □ point-in-time recovery confirmed.

Q118

How do you configure Commvault for Amazon AWS EC2 backup?

Answer:

AWS EC2 backup configuration: (1) Create IAM role/user with EC2 snapshot permissions (ec2:CreateSnapshot, ec2:DescribeInstances, etc.); (2) CommCell Console: New Client □ Virtual Server □ Amazon Web Services; (3) Enter IAM credentials (Access Key/Secret Key) or use IAM role (for EC2-hosted CommServe/MA); (4) Select AWS regions to protect; (5) Commvault discovers EC2 instances, EBS volumes; (6) VSA subclient: select instances by tag, region, VPC; (7) Backup method: EBS snapshot (via AWS API) + Backup Copy to S3 for long-term; (8) Network: Commvault MA in same AWS region for efficient data transfer (avoid egress costs).

Example:

AWS EC2 backup (200 instances, 3 regions): IAM policy with EC2+S3 permissions attached to Commvault role. CommCell: AWS client with IAM role credentials. Discovery: 200 EC2 instances discovered. VSA subclient: tag 'Backup: Required' (150 instances). MediaAgent: deployed as EC2 instance in each region (avoid cross-region data transfer costs). Backup: EBS snapshots created in seconds, backup copy to S3 Standard (30-day). S3 Glacier for 1-year retention (lifecycle policy). Monthly EC2 backup cost: \$450.

Q119

What is Commvault's ring topology and when is it used?

Answer:

Ring topology in Commvault: a configuration where backup data flows through a series of MediaAgents in a ring – MA1 backs up clients and copies to MA2, MA2 copies to MA3, MA3 may copy back or to a final repository. Use cases: (1) Multi-hop data replication (site A □ site B □ site C); (2) Data distribution: each site has a local copy + a copy at next site; (3) DR chain: each link in the ring provides local and remote copies; (4) WAN-efficient: each leg of the ring uses local WAN connection. Less common than traditional 3-copy policies – used in specific topologies like retail chains with regional distribution centres.

Example:

Retail chain ring topology: 500 stores □ Regional DCs (5 regions) □ Central HQ. Store backups: local disk (1-day retention). Regional DC: receives auxiliary copy from 100 stores, stores 7 days. Central HQ: receives auxiliary copy from 5 regional DCs, stores 90 days. Data flows: Store □ Regional DC □ HQ. Each link is WAN-efficient (short-haul). Restore: local file restore from store backup (<1 min), monthly historical restore from HQ archive. Ring ensures data always progresses toward central repository.

Q120

What is a Commvault Recover Point and how is it configured?

Answer:

Recovery points are specific backup versions available for restore. Configuration determines how many recovery points are retained: (1) Retention days: keep all backup versions for X days; (2) Retention cycles: keep last N full backup cycles; (3) Incremental versions: all incrementals between fulls are individual recovery points; (4) Snapshot recovery points: array snapshots provide additional short-term recovery points; (5) Extended retention: specific jobs (e.g., month-end) kept as long-term recovery points; (6) RPO calculation: recovery points determine achievable RPO (e.g., hourly incrementals = 1-hour RPO). More recovery points = more storage required but better RPO.

Example:

Recovery point strategy for ERP system: Daily full (1 recovery point/day). 4-hourly incremental (6 recovery points/day = 1-4 hour RPO). IntelliSnap hourly (24 recovery points/day = 1-hour RPO from array). Monthly full (extended retention, 7 years). Total recovery points at any time: $30 \times 6 + 30 \times 24 + 84$ monthly = 900+ recovery points. SLA: RPO = 1 hour (from IntelliSnap), RTO = 15 minutes (from IntelliSnap instant recovery).

Q121

What are the system requirements for installing a Commvault CommServe?

Answer:

CommServe system requirements: (1) OS: Windows Server 2016/2019/2022 (64-bit); (2) CPU: 8 cores minimum (16+ recommended for large environments); (3) RAM: 16GB minimum (32GB recommended for >500 clients); (4) Disk: 300GB minimum for CommServe software + SQL database (SSD strongly recommended for SQL data files); (5) SQL Server: SQL 2016/2017/2019 (Enterprise or Standard — Express not supported in production); (6) Network: static IP, DNS resolution to all clients and MediaAgents; (7) .NET Framework 4.8+; (8) Windows Firewall: ports 8400-8403, 443 open.

Example:

Production CommServe sizing for 800 clients: Windows Server 2019, 16 vCPU, 64GB RAM, 500GB SSD for SQL (separated from OS drive), SQL Server 2019 Standard. CommServe VM on high-availability cluster (vSphere HA). Monitoring: Commvault watchdog service restarts SQL if it fails. Backup: CommServe DR exported nightly to 3 locations.

Q122

What is the recommended Commvault deployment architecture for a medium enterprise?

Answer:

Medium enterprise (200-500 clients) architecture: (1) CommServe: dedicated Windows Server (physical or VM), SQL Server 2019 Standard; (2) MediaAgent(s): 2 MediaAgents for redundancy (Windows/Linux), each 8 core, 32GB RAM, 10GbE NIC; (3) Disk library: primary disk library on NAS (100TB+, SSD-backed), secondary on slower NAS or cloud; (4) VSA proxy: 2 Windows proxies for VMware backup (8 core, 32GB RAM each); (5) Index Server: dedicated Windows Server for content indexing; (6) DR: CommServe DR to remote location, auxiliary copy to cloud or remote disk.

Example:

250-client deployment: CommServe (VM, 8 core, 32GB, SSD). 2 MediaAgents (Windows, 8 core, 32GB, 10GbE). Primary disk library: NetApp NFS (80TB, 10:1 dedup = 800TB effective). Cloud library: Azure Blob (secondary copy, infinite scalability). 2 VSA proxies: VMware backup for 150 VMs. Index Server: 4 core, 16GB. Daily DR export to Azure Blob. All components connected on dedicated backup VLAN (10GbE).

Q123

What is the process for installing a Commvault MediaAgent?

Answer:

MediaAgent installation steps: (1) Download Commvault installer from software store or CommServe; (2) Run installer on target Windows/Linux server; (3) Select 'MediaAgent' role in component selection; (4) Enter CommServe hostname/IP and authentication credentials; (5) Select installation path and data paths for DDB (SSD recommended); (6) Configure network interface for backup traffic; (7) Complete installation — MediaAgent registers with CommServe automatically; (8) In CommCell Console: verify MediaAgent appears and is Online; (9) Create disk library on the MediaAgent and assign mount paths.

Example:

Linux MediaAgent installation: `./cvpkgadd -install -silent -cs 10.10.150 -cspassword Admin123 -role MA`. MediaAgent registers with CommServe. In CommCell Console: Storage ▢ MediaAgents ▢ MA-Linux-01 shows Online. Create disk library: Storage ▢ Libraries ▢ Add Disk Library ▢ mount path `/backup/disk1` (2TB NFS mount). Dedup enabled, DDB path: `/ddb/ddb01` (SSD-backed LUN). MediaAgent ready for use in 20 minutes.

Q124

How do you configure a storage policy in Commvault from scratch?

Answer:

Storage policy creation steps: (1) CommCell Console ▢ Policies ▢ Storage Policies ▢ New Storage Policy; (2) Name the policy and set type (Data Protection, Archival, or DR); (3) Select Primary copy: choose MediaAgent and disk library; (4) Configure deduplication: enable/disable, select DDB path; (5) Set primary copy retention (days AND cycles); (6) Optionally add Secondary copy: configure auxiliary copy MediaAgent, library, and retention; (7) Configure encryption per copy if required; (8) Click Finish. Assign storage policy to subclients via subclient properties.

Example:

Creating 'SQL-Production-SP': Name ▢ type: Data Protection. Primary copy: MediaAgent-01, DiskLib-Primary, dedup enabled (DDB on SSD), 30-day/2-cycle retention. Secondary copy: MediaAgent-02, DiskLib-Remote (off-site NAS), 90-day retention, auxiliary copy. Tertiary copy: Tape library, 7-year retention, monthly fulls only. Storage policy created in 10 minutes. Assigned to all SQL Server subclients. All SQL data now flows through this policy automatically.

Q125

How do you add a new client (Windows server) to Commvault?

Answer:

Adding a Windows client: (1) CommCell Console ▢ Client Computers ▢ Add Client; (2) Enter client hostname/IP; (3) Select OS type and iDataAgents to install (File System, SQL, etc.); (4) Provide credentials for push installation; (5) CommServe pushes Commvault software to the client; (6) Client registers with CommServe after install; (7) Configure default subclient content, storage policy, and schedule; (8) Run initial full backup. Alternative: pull installation — download installer on client and specify CommServe during install.

Example:

New Windows Server 2022 'SQLSRV-05' added: CommCell Console ▢ Add Client ▢ enter hostname SQLSRV-05 ▢ Windows ▢ select File System iDA + SQL Server iDA ▢ credentials: administrator ▢ push install. 15-minute installation. SQLSRV-05 appears in CommCell. SQL iDA: auto-discovers SQL instances ▢ creates default subclient ▢ assign 'SQL-Production-SP' ▢ schedule: daily full Sunday, incremental Mon-Sat. First backup Sunday night.

Q126

How do you configure VMware backup in Commvault?

Answer:

VMware backup configuration: (1) Add vCenter as a virtualisation client (Client ▢ New Client ▢ Virtual Server ▢ VMware); (2) Enter vCenter credentials (read/write access required); (3) Commvault discovers all ESXi hosts, clusters, VMs; (4) Add VSA proxy — Windows server with Commvault VSA proxy agent; (5) Create subclient: select VMs to protect (by VM, cluster, tag, datastore); (6) Assign storage policy and schedule; (7) Run discovery to pick up new VMs (automatic or manual). VSA proxy must have access to vCenter API and ESXi hosts.

Example:

VMware vSphere 7.0 configuration: New Virtual Server client ▢ vCenter URL: vcenter.company.local ▢ credentials: backup-svc@vsphere.local. 350 VMs discovered. VSA proxy: BACKUP-PROXY-01 (Windows 2019, 8 core, 32GB). Subclient 'All-Production-VMs': include VMs with tag 'env:production' (240 VMs). Storage policy: VM-Backup-SP. Schedule: daily incremental 10 PM, full Sunday 9 PM. Auto-discovery: every 4 hours (new VMs auto-included if tagged).

Q127

What is the Commvault software upgrade process?

Answer:

Commvault upgrade sequence (order matters): (1) Download latest service pack from Commvault's software store; (2) Upgrade CommServe first — run installer, select upgrade; (3) Verify CommServe is running post-upgrade; (4) Upgrade MediaAgents — via CommCell Console or push from CommServe; (5) Upgrade VSA proxies; (6) Upgrade iDataAgents on clients (can be done in waves); (7) Verify all components show correct version in CommCell Console; (8) Run test backup and restore post-upgrade. Never upgrade clients before CommServe — forward compatibility only.

Example:

Upgrading from SP26 to SP30: Saturday night maintenance. CommServe upgraded first (45 min). CommServe verified running (jobs resumed). MediaAgents upgraded (15 min each, 2 MAs = 30 min). VSA proxies upgraded (20 min each). Client upgrades: pushed in waves over 2 weeks (100 clients/night using CommCell Console push). All components on SP30 within 2 weeks. Zero backup failures during phased upgrade.

Q128

How do you configure deduplication in Commvault?

Answer:

Deduplication configuration: (1) Storage policy ▢ Primary copy properties ▢ Enable deduplication; (2) DDB location: select dedicated path on MediaAgent (must be fast SSD/NVMe — DDB is I/O intensive); (3) DDB type: Global or Local; (4) Source-side dedup: configure on subclient level (reduces network traffic); (5) MediaAgent-side: default — dedup happens at MediaAgent before writing; (6) DDB backup: configure automatic DDB backup (critical — DDB corruption without backup means data unrecoverable); (7) DDB move: move DDB to different path if storage changes needed.

Example:

DDB setup for new disk library: Storage Policy ▢ Primary copy ▢ Enable Deduplication. DDB path: /ddb/ssd01 on MediaAgent-01 (Samsung 1TB NVMe, dedicated). DDB partition size: 500GB (handles up to 250TB unique data). Source-side dedup: enabled on all VM subclients (reduces WAN traffic by 60%). DDB backup: daily to disk library (copy of DDB in case of failure). First week: dedup ratio builds from 1:1 to 8:1 as blocks accumulate.

Q129

How do you configure a cloud library in Commvault?

Answer:

Cloud library configuration: (1) CommCell Console ▢ Storage ▢ Libraries ▢ New ▢ Cloud Library; (2) Select cloud vendor (AWS S3, Azure Blob, GCP, etc.); (3) Enter credentials (access key/secret for AWS, storage account key for Azure); (4) Select or create a bucket/container; (5) Set storage class (S3 Standard, S3-IA, Glacier; Azure Hot, Cool, Archive); (6) Assign cloud library to a MediaAgent (MediaAgent acts as gateway); (7) Create storage policy copy targeting the cloud library; (8) Test with a small auxiliary copy job. Cloud library acts like a disk library – transparent to policies.

Example:

AWS S3 cloud library: New Cloud Library ▢ Amazon S3 ▢ Access Key/Secret Key entered ▢ bucket 'company-backup-aws' ▢ S3 Intelligent-Tiering storage class (automatic cost optimisation). Assigned to MediaAgent-02. Storage policy tertiary copy targets this library. Auxiliary copy (weekly, monthly fulls only) runs 2 AM Sunday. First copy: 500GB transferred to S3. Monthly cost: ~\$12 (S3 Intelligent-Tiering at \$0.023/GB). 7-year retention configured.

Q130

What is Commvault's CommCell Console versus Command Center – which to use for what?

Answer:

CommCell Console (Java-based rich client): Use for – (1) Advanced storage configuration (tape libraries, NDMP); (2) CommServe DR configuration; (3) DDB management; (4) Advanced client configuration; (5) Legacy features; (6) Detailed job-level troubleshooting. Command Center (web-based): Use for – (1) Day-to-day monitoring and dashboards; (2) Adding common workloads (VMs, databases, M365); (3) Plans and policy management (simplified); (4) Reports; (5) User and RBAC management; (6) Self-service restores; (7) Alert management. Both connect to same CommServe – changes in one reflected in the other.

Example:

Admin uses Command Center for: daily health check, new VM subclient creation, running a self-service restore for a user, reviewing SLA dashboard. Uses CommCell Console for: configuring new NDMP tape library, setting up DDB rebalancing, configuring CommServe DR backup path, troubleshooting a specific job's chunk-level errors. Both tools available simultaneously – no need to choose one permanently.

Q131

How do you configure a schedule policy in Commvault?

Answer:

Schedule policy creation: (1) CommCell Console ▢ Policies ▢ Schedule Policies ▢ New; (2) Name the policy; (3) Add schedules for each backup type: Full (specify day/time), Incremental (specify frequency/time), Differential, Log (for databases); (4) Specify run window (not before/not after times); (5) Configure missed job handling (run immediately when client available, skip, etc.); (6) Apply to subclients: subclient properties ▢ Schedule Policy ▢ select policy. Multiple subclients share the same schedule policy – change once, update all.

Example:

Schedule policy 'DB-Standard': Full backup: every Sunday 10 PM. Incremental: Mon-Sat 10 PM. Transaction logs: every 15 minutes (for SQL/Oracle subclients). Run window: 10 PM to 6 AM (won't start after 6 AM). Missed job: run immediately when client reconnects. Applied to 45 database subclients. Changing log frequency from 15 to 30 minutes: update policy ▢ all 45 subclients updated immediately.

Q132

How is the Commvault CommServe database backed up?

Answer:

CommServe database backup: (1) CommServe DR feature – automatic export of CommServe database (metadata, configs, job history) to specified paths; (2) Configured in CommCell Console ▢ CommServe ▢ DR Settings; (3) Export paths: local disk, UNC path (MediaAgent share), FTP, cloud; (4) Schedule: runs after every backup phase (every 4-8 hours typically); (5) Retention: keep last N DR backups; (6) SQL native backup: also configure SQL Server Agent job for CommServe DB backup as additional copy; (7) Commvault agent on CommServe: use Commvault to back up the CommServe server itself.

Example:

CommServe DR configuration: export to 3 locations – (1) Local C:\CommvaultDR\ (last 5 exports). (2) MediaAgent-01 UNC share \\MA01\CommvaultDR\ (last 10 exports). (3) Azure Blob via Commvault FTP export (last 30 exports). Runs after every backup phase (every 6 hours). SQL Agent job: full backup of CommVault DB daily to separate disk. CommServe server itself backed up by Commvault File System iDA. Triple redundancy for CommServe recovery.

Q133

How do you configure Commvault to integrate with Active Directory for authentication?

Answer:

AD integration steps: (1) CommCell Console □ Security □ Additional Settings □ Enable AD authentication; (2) Enter AD domain name and domain controller details; (3) Test AD connectivity; (4) Create external user groups: Security □ User Groups □ New External User Group □ enter AD group name; (5) Assign roles to the user group; (6) Users log into CommCell Console with domain\username or UPN; (7) Password policies enforced by AD; (8) User management: adding/removing from AD groups automatically changes Commvault access; (9) Service accounts for Commvault services: dedicated AD service account recommended.

Example:

AD integration for CORP.LOCAL domain: Domain controller 10.10.1.5. External group 'CORP\Backup-Admins' □ Commvault Admin role. External group 'CORP\App-Owners' □ Restore-Only role. External group 'CORP\Audit-Team' □ Report Viewer. Test: log in with domain admin □ success. Offboarding: disable AD user account □ Commvault access revoked immediately. No separate Commvault account deprovisioning required.

Q134

What are the steps to configure email alerts in Commvault?

Answer:

Email alert configuration: (1) CommCell Console □ Alerts □ Alert Settings □ Configure SMTP: enter SMTP server, port, authentication, sender address; (2) Create alert definitions: Alerts □ Add Alert □ select alert criteria (job failure, SLA miss, disk full, etc.); (3) Set notification threshold (e.g., alert on failure, or after 3 failures); (4) Set recipients (email addresses or distribution lists); (5) Enable/disable specific alert types per client or globally; (6) Test alert: Alerts □ Test Alert; (7) Command Center: Manage □ Notifications □ Email settings equivalent.

Example:

Alert configuration: SMTP: smtp.company.com:587, TLS, backup-noreply@company.com. Alerts created: Job failure □ backup-team@company.com + PagerDuty webhook. SLA miss □ it-manager@company.com. Disk library >80% □ storage-team@company.com. CommServe DB full □ on-call@company.com (SMS). License expiry 30 days □ admin@company.com. Test alert sent – received in 45 seconds. NOC integrated via SNMP for ticket auto-creation.

Q135

How do you configure encryption in a Commvault storage policy?

Answer:

Encryption configuration: (1) Storage policy copy □ Properties □ Encryption; (2) Enable encryption; (3) Select cipher: AES-256 (recommended), BlowFish, 3DES; (4) Key length: 256-bit for AES; (5) Re-encryption: option to re-encrypt existing data; (6) Key management: Commvault KMS (built-in) or external KMIP KMS; (7) Pass-phrase or certificate: set a strong pass-phrase (stored in CommServe DB – back this up!); (8) Network encryption: Client □ MediaAgent in-flight encryption via SSL. Encryption transparent to end users and restore operations.

Example:

Healthcare storage policy encryption: AES-256 enabled on primary and all secondary copies. Commvault KMS: pass-phrase '32-character random string' stored in CommServe + backed up to offline location. Network encryption: SSL enabled between all clients and MediaAgents (HIPAA: PHI encrypted in transit). Azure cloud copy: customer-managed encryption key (CMK) in Azure Key Vault. HIPAA audit: all backup data encrypted in transit and at rest – finding: compliant.

Q136

How do you configure a Commvault plan for simple backup management?

Answer:

Plan creation in Command Center: (1) Protect ▢ Plans ▢ Add Plan; (2) Select plan type: Server Plan (for VMs/servers), Laptop Plan (for endpoints), Database Plan; (3) Define backup frequency (daily, every 4 hours, etc.); (4) Set retention (30 days, 1 year, custom); (5) Select backup destination (storage pool/library); (6) Enable deduplication and encryption; (7) Define additional copies (snapshot copy, cloud copy); (8) Click Save. Assign plan to workloads: Command Center ▢ Protect ▢ [workload] ▢ select plan. Plan applies all settings automatically.

Example:

Server Plan 'Standard-VM-Plan': backup every day at 11 PM. Retention: 30 days. Backup destination: Primary-Disk-Pool (on-prem NAS). Additional copy: Cloud copy to Azure Blob (90-day retention, runs nightly). Dedup: enabled. Encryption: AES-256. Created in 5 minutes. Applied to 200 VMs: Command Center ▢ Virtual Machines ▢ select all ▢ Assign Plan ▢ Standard-VM-Plan. All 200 VMs now protected with consistent settings.

Q137

What is the process to configure NDMP backup for a NAS device?

Answer:

NDMP configuration: (1) Add NDMP host as a client: New Client ▢ NAS/NDMP ▢ enter NAS IP/hostname; (2) Enter NDMP credentials (configured on NAS device first); (3) Select NDMP agent type: V4 (most NAS) or V3 (older devices); (4) Commvault discovers NAS volumes; (5) Create subclient: select volumes to back up; (6) Configure storage policy with MediaAgent; (7) For 3-way copy: select tape library directly connected to NAS (data bypasses MediaAgent); (8) Run full backup. NDMP-specific settings: backup type, log backup, tape management.

Example:

NetApp ONTAP NDMP: New Client ▢ NAS ▢ IP: 10.10.1.100 ▢ NDMP credentials (ndmp_user). Volumes discovered: /vol/vol1 (20TB), /vol/vol2 (15TB), /vol/nas_home (10TB). Subclient: select all volumes. Storage policy: NDMP-SP (MediaAgent handles data, not 3-way). Schedule: nightly 11 PM. First backup: 45TB in 6 hours. Incremental: 2TB/night (4%). NetApp snapshot for hourly local recovery. Commvault NDMP for daily offsite copy.

Q138

How do you configure Commvault for Microsoft 365 backup?

Answer:

Microsoft 365 configuration: (1) Command Center ▢ Protect ▢ Applications ▢ Microsoft 365 ▢ Add; (2) Connect Microsoft 365 tenant: enter tenant domain, grant permissions (global admin + application consent required); (3) Commvault creates Azure AD application with required API permissions (Graph API, EWS); (4) Select workloads to protect: Exchange Online, SharePoint, OneDrive, Teams; (5) Select storage: disk library, cloud library, or HyperScale; (6) Configure retention per workload; (7) Run discovery (finds all mailboxes, sites, users); (8) First backup initiates. No agents required — API-based.

Example:

Microsoft 365 setup for 3,000 users: Command Center ▢ M365 ▢ Add ▢ tenant: company.onmicrosoft.com ▢ admin@company.com authenticated ▢ consent granted. Workloads: Exchange (3,000 mailboxes), SharePoint (500 sites), OneDrive (3,000 drives), Teams (200 teams). Storage: Azure Blob cloud library (same region as M365 tenant). Discovery: 3,500 objects found. First backup: 2.8TB Exchange, 1.5TB SharePoint, 4TB OneDrive. Daily incremental thereafter.

Q139

What is the Commvault software store and how is it used?

Answer:

The Software Store (or Download Center) is a Commvault-hosted repository for downloading: (1) Full installation packages; (2) Service packs and updates; (3) Client software packages (for push install to clients); (4) Documentation and release notes. CommServe can be configured to download updates from the Software Store automatically. Clients can be updated by: CommServe push (pulls from Software Store), direct download, or pre-staged package on internal network share. Software Store URL: cloud.commvault.com/webconsole/loginPage.do.

Example:

Software Store workflow: CommServe configured to sync with Commvault Software Store weekly. New service pack SP32 available: notification appears in Command Center. Admin downloads SP32 to CommServe (2GB download). Push upgrade initiated: CommServe upgrades itself, then pushes SP32 to all MediaAgents (via Software Store staging on CommServe). Client upgrades scheduled in maintenance windows. No manual downloads per client required.

Q140

How do you configure IntelliSnap for VMware in Commvault?

Answer:

IntelliSnap VMware configuration: (1) Add array: Storage □ Arrays □ Add □ select vendor (NetApp, Pure, EMC, etc.); (2) Enter array management IP and credentials; (3) Commvault discovers array volumes; (4) Enable IntelliSnap on subclient: subclient properties □ Snapshot □ Enable IntelliSnap; (5) Select snap engine (NetApp SnapCenter, Pure Storage, etc.); (6) Configure snapshot retention on array; (7) Configure Backup Copy job to copy snapshot to backup storage; (8) Test: run IntelliSnap backup and verify snapshot created on array. Requires VSA proxy with access to array management interface.

Example:

NetApp IntelliSnap for VMware: Array added □ NetApp ONTAP 9.10 □ cluster mgmt IP 10.10.1.200 □ credentials. VMware datastore on NetApp volume vol_vmdata01 mapped. IntelliSnap subclient 'VMware-IntelliSnap': 100 VMs on NetApp datastore. Snap engine: NetApp ONTAP. Snapshot created in 3 seconds (production zero-impact). Backup Copy: runs 2 AM, reads NetApp snapshot, writes to disk library (deduplicated). Array snapshots: 24 hours × 24 = 24 snapshots on array. Disk library: 30-day retention.

Q141

What is a Commvault subclient content configuration?

Answer:

Subclient content defines exactly what data is backed up. Configuration options: (1) Include — explicit paths, databases, VMs, mailboxes to include; (2) Exclude — patterns or paths to exclude (*.tmp, Recycle Bin, pagefile.sys); (3) Exceptions to exclude — within an excluded folder, re-include specific items; (4) Auto-discovery — automatically include new databases/VMs matching criteria; (5) Wildcard — use wildcards for dynamic content (D:\Data\Project_*); (6) VSA: include by VM name, cluster, datastore, tag, or power state. Content configuration determines backup scope and should be reviewed regularly.

Example:

*File server subclient content: Include: D:\Company_Data\. Exclude: *.tmp, *.log, Thumbs.db, D:\Company_Data\Archive\Temp\. Exception to exclude: D:\Company_Data\Archive\Temp\Important\ (re-include this subfolder). Result: 98% of D:\Company_Data\ backed up, temp files and logs excluded, one important temp subfolder preserved. Weekly review: new log folder found — added to exclusion list. Storage saved: 15% from exclusions.*

CommCell Administration

Q142

How do you create a custom report in Commvault?

Answer:

Custom report creation: (1) CommCell Console □ Reports □ Report Builder (or Report Writer); (2) Select data source: jobs, clients, media, storage; (3) Select columns: client name, job status, data transferred, duration, start time; (4) Apply filters: date range, client group, agent type; (5) Add groupings and sorting; (6) Preview report; (7) Save report with name; (8) Schedule: configure automatic run and email delivery; (9) Export formats: HTML, CSV, PDF; (10) Command Center: modern report builder with drag-and-drop interface and chart visualisation. Custom reports surface data not available in standard reports.

Example:

Custom report: 'Weekly SQL Backup Summary'. Columns: SQL Server name, database name, last backup date, backup size, duration, job status. Filter: SQL iDA clients only. Group by: server name. Schedule: every Monday 7 AM □ email to dba-team@company.com and it-manager@company.com as PDF. First delivery: 14 SQL servers, 85 databases, 99% success, 2.4TB backed up this week. DBA team uses it to verify all databases protected before Monday workday starts.

Q143

How do you configure network throttling in Commvault?

Answer:

Network throttling configuration: (1) CommCell Console □ Storage □ MediaAgent □ Properties □ Network Throttle; (2) Or subclient □ Properties □ Storage Device □ Throttle; (3) Time-based throttling: different limits at different times (business hours vs off-peak); (4) Per-client throttle: limit specific clients' backup bandwidth; (5) Per-network throttle: throttle WAN-connected clients more aggressively; (6) Units: MB/s or Mbps; (7) Throttle applies to: data transfer between client and MediaAgent; (8) Auxiliary copy: separate throttle for auxiliary copy jobs; (9) QoS integration: DSCP marking on Commvault traffic for network QoS policies.

Example:

WAN throttle setup: 20 branch offices, each 50Mbps WAN. Without throttle: backups at full speed saturate WAN, VoIP/video calls degrade. Throttle policy: 8 AM-6 PM: 10Mbps per branch (20% WAN). 6 PM-8 AM: 40Mbps per branch (80% WAN). Weekend: unlimited. Result: business hours – backup traffic 10Mbps (transparent to users). Night: 40Mbps – backups complete within SLA. Weekend: catch-up runs at full speed. All branch backups within 24-hour SLA.

Q144

What is a Commvault disaster recovery test (DRT) and how is it documented?

Answer:

Disaster Recovery Test documentation: (1) Test plan: define scope (which systems), objectives (RTO/RPO), test environment (isolated network), schedule; (2) Pre-test: verify backup versions current, DR site ready, runbook available; (3) Test execution: follow runbook, record actual times for each step; (4) Validation: application-level testing (login, run transaction, verify data integrity); (5) Issues log: document any failures or deviations; (6) Results: actual RTO vs target RTO, actual RPO vs target RPO; (7) Sign-off: business owner and IT manager sign; (8) Improvement actions: list changes to improve next test; (9) Frequency: most organisations test quarterly for critical systems.

Example:

DRT report for core banking: Test date: March 15, 2 AM (Saturday). Scope: 8 core banking VMs. Target RTO: 4 hours. Actual RTO: 3h 22m. Target RPO: 30 minutes. Actual RPO: 18 minutes (better than target). Issues: Step 3 (DB service start) failed once due to missing DNS entry □ manual fix 8 minutes. Action: add DNS entry to DR runbook pre-steps. Sign-off: CTO, Head of IT Operations. Next test: June 15. Result: PASS. Filed with compliance team as evidence.

Q145

How does Commvault manage media vault and tape rotation?

Answer:

Tape vault and rotation management: (1) CommCell Console □ Storage □ Libraries □ tape library □ Media □ Export/Import; (2) Export procedure: select tapes for export □ mark as 'Exported' □ physically remove from library; (3) Vault tracking: record which tapes are at which vault location; (4) Return schedule: import when needed for restore or rotation; (5) Rotation policy: define how often tapes rotate (weekly offsite cycle); (6) Commvault tracks: tape status (on-site/off-site), location, available from date; (7) Recall: locate specific tape in CommCell, request recall from vault; (8) Media destruction: when retention expires □ mark destroyed, Commvault removes from tracking.

Example:

Weekly tape rotation: Friday night – full backups complete on 10 tapes. Saturday morning: admin exports 10 tapes from library (CommCell marks: Exported, Location: Iron Mountain). 10 tapes from previous week: returned from Iron Mountain (iron-mountain delivers). Import: 10 returned tapes □ Commvault scans barcodes □ updates status: On-site, Available. Restore request: 'Need data from January 12.' CommCell: identifies tape CV20240112 at Iron Mountain. Recall request: tape arrives next business day. Restore: 4 hours after tape arrival.

Q146

What is Commvault's pre-scan process and why is it important?

Answer:

Pre-scan (also called file system scanning) identifies changed files before backup begins: (1) Accelerated File System Scanner (ACFS): scans file system in parallel using multiple threads before backup window; (2) Benefits: backup window reduced – scan complete before backup starts, so backup begins transferring immediately; (3) Change journal: uses OS change journal (USN journal on Windows, iNotify on Linux) to identify changes faster; (4) Impact: pre-scan runs at low priority (minimal production impact); (5) Large file servers: pre-scan especially valuable for servers with millions of files; (6) Configuration: subclient □ Advanced □ Enable Pre-scan. Separate pre-scan schedule from backup schedule.

Example:

50 million file server: Without pre-scan – backup starts 10 PM: scan phase 3 hours (10 PM-1 AM), transfer phase 4 hours (1 AM-5 AM). Total: 7 hours. With pre-scan – pre-scan starts 7 PM (3-hour scan, low priority, no production impact). Backup starts 10 PM: skip scan, begin transfer immediately. Transfer: 4 hours (10 PM-2 AM). Total backup window: 4 hours (saved 3 hours). Pre-scan overlaps with production hours but at low priority – no user impact.

Q147

How do you configure a Commvault subclient filter?

Answer:

Subclient filters exclude specific content from backup: (1) CommCell Console □ Subclient □ Properties □ Filters; (2) Wildcard filters: *.tmp, *.log, *.bak; (3) Path filters: specific directories excluded (C:\Windows\Temp); (4) Regular expression: advanced pattern matching; (5) System-generated exclusions: Commvault automatically excludes pagefile.sys, hibernation files, OS temp; (6) Exception to filter: within an excluded path, include a specific subfolder; (7) Global filters: set in CommCell properties to apply to all subclients; (8) Impact: excluded files not backed up, reducing backup size and time; (9) Best practice: review filters quarterly – new temp directories may accumulate.

Example:

Filter configuration for application server: Exclude: *.tmp, *.log (reduces backup by 15%), C:\Temp (all temp files), C:\Windows (OS files – covered by System State), D:\AppLogs (log files rotated separately). Exception: D:\AppLogs\Audit (audit logs must be backed up despite being in excluded path). After filter: backup size reduced from 800GB to 650GB (19% reduction). Backup window reduced from 4 hours to 3.5 hours. Audit logs still backed up via exception. Review: new path D:\AppData\Temp found – added to exclusion list.

Q148

What is the Commvault CommCell readiness check?

Answer:

CommCell readiness check (or CommCell diagnostic) validates the health of the CommCell environment: (1) CommServe health: SQL DB size, free space, connectivity; (2) MediaAgent check: all MAs online, DDB health, disk library space; (3) Client check: all clients communicating with CommServe; (4) Library check: disk libraries and tape libraries accessible; (5) Job check: any hung jobs, long-running jobs; (6) License check: license validity and usage; (7) Service check: all Commvault services running; (8) Access via: CommCell Console ▢ CommServe ▢ Diagnostic ▢ CommCell Readiness Report. Run daily as part of proactive monitoring.

Example:

Daily readiness check at 8 AM: CommServe SQL: 45GB, free space 200GB ▢. MediaAgents: MA-01 Online ▢, MA-02 Online ▢. Disk library: Primary 68% full (alert threshold 80% – OK), Secondary 45% ▢. Clients: 495/500 online (5 offline – weekend patching servers). License: 1,200/1,500 used ▢. Services: all running ▢. Hung jobs: 0 ▢. Issues: none critical. Action: check 5 offline clients (expected – patch schedule). Readiness report emailed to backup-team.

Q149

How do you manage Commvault CommCell SSL certificates?

Answer:

Commvault SSL certificate management: (1) Default: Commvault uses self-signed certificates for internal communication; (2) External access (Command Center): replace self-signed with CA-signed certificate (PFX format); (3) Import: CommCell Console ▢ CommServe ▢ Properties ▢ Security ▢ Certificate; (4) Renewal: certificates expire (typically 1-3 years) – renew before expiry to avoid service interruption; (5) CA-signed: use internal PKI (Active Directory Certificate Services) or public CA; (6) Certificate pinning: Commvault pins internal certificates – changing CommServe certificate requires coordinated update; (7) SAML: for SSO, upload identity provider certificate; (8) Alerts: Commvault alerts 30 days before certificate expiry.

Example:

Certificate management: Command Center HTTPS using self-signed cert ▢ browsers show 'Not Secure'. Fix: request certificate from internal CA (SubjectAltName: commserve.company.com). Export as PFX. CommCell Console ▢ CommServe ▢ Properties ▢ Security ▢ Upload PFX. Command Center HTTPS: now shows padlock (CA-signed). Expiry monitoring: calendar reminder 60 days before expiry. Renewal: request new cert, upload PFX, no service restart needed. Commvault internal SSL: managed automatically by Commvault.

Q150

How do you configure multi-tenancy in Commvault?

Answer:

Commvault multi-tenancy for MSPs or enterprise tenants: (1) CommCell Console ▢ Organisation ▢ Add Company/Tenant; (2) Each tenant/company gets isolated view: own clients, own operators, own storage; (3) Tenant admin: can manage only their clients and data; (4) Global admin: sees all tenants (MSP master view); (5) Storage isolation: separate storage policies and libraries per tenant; (6) Metering: per-tenant capacity reporting for billing; (7) Branding: tenant-specific portal branding (logo, URL); (8) Isolation: tenant A cannot see tenant B data; (9) Plans: shared infrastructure, separate logical isolation per tenant.

Example:

MSP managing 10 enterprise customers: Commvault multi-tenant setup. Tenant A (Finance Corp): 50 clients, dedicated storage policy, finance-admin@financecorp.com as tenant admin. Tenant B (Retail Inc): 30 clients, separate storage, retail-admin@retailinc.com. Tenant C (Healthcare): 20 clients, dedicated encrypted storage (HIPAA), separate MediaAgent. Global admin: MSP backup team sees all tenants. Monthly billing: Commvault report shows per-tenant TB consumed ▢ billing system integration. Tenant admins: manage their own workloads, restores, reports – no MSP involvement.

Q151

What is the Commvault CommCell synchronisation and when is it needed?

Answer:

CommCell sync operations: (1) CommServe DB sync: if CommServe DR restored from backup, re-sync jobs run since last export; (2) DDB sync: DDB reconciliation after DDB backup restore (validates DDB matches backup data); (3) Client sync: if client reinstalled, re-register and sync job history; (4) MediaAgent sync: if MA moved to new server, sync library associations; (5) Auxiliary copy sync: if aux copy fell behind, sync creates missing copies; (6) Commvault 'Data verification' is related: reads backup data from storage to verify it matches CommServe metadata. Sync operations usually run automatically but can be triggered manually.

Example:

DDB sync scenario: MediaAgent NVMe failed, DDB restored from backup (6-hour-old backup). Sync needed: 6 hours of backup data written after DDB backup. Commvault DDB sync runs: reads backup data from disk library for last 6 hours & updates DDB with any new unique blocks written in those 6 hours. Duration: 2 hours for 6-hour backlog. After sync: DDB fully consistent with backup data. All restore operations for that period now fully supported. Without sync: restores for that 6-hour window may fail with 'block not found' errors.

Q152

How do you configure and use Commvault's CommCell activity control?

Answer:

Activity control management: Levels — (1) CommCell level: all operations globally paused/resumed; (2) Client level: pause/resume for specific clients; (3) Agent level: pause specific iDA (pause SQL backup but keep File System backup running); (4) Operations controllable: Data Management (backup), Data Recovery (restore), Auxiliary Copy, Data Aging, Workflow; (5) Blackout window: schedule activity control to automatically pause/resume at specific times; (6) Override: specific jobs can be force-started despite activity control; (7) View: CommCell Console & Activity Control shows current state of all operations. Essential for maintenance windows and emergency pauses.

Example:

Scenario: storage array maintenance 2 AM-4 AM. Activity Control: (1) 1:55 AM — suspend all Data Management (backup) jobs globally. (2) Let running jobs complete. (3) 2 AM — array maintenance begins. (4) 4 AM — array maintenance complete. (5) Resume all Data Management. Jobs queued during maintenance now run. All within 5 minutes of activity control actions. Alternative: set automatic blackout window 2 AM-4 AM recurring Saturday — no manual action needed. Blackout windows are the preferred approach for regular maintenance.

Q153

What is the Commvault software cache and how is it managed?

Answer:

Software cache is a local repository of Commvault installation packages stored on the CommServe (or a dedicated cache server). Purpose: (1) Client push installations use packages from cache (faster than downloading per-client); (2) Cache stores: latest service pack installers for all supported OS/agents; (3) Management: CommCell Console & Administration & Software Cache & configure cache location and size; (4) Sync: cache syncs with Commvault Download Center automatically; (5) Distributed cache: for remote sites, configure a local cache server to reduce WAN usage for agent installations; (6) Cache validation: verify cache integrity before major upgrade rollout; (7) Space: cache can be large (100GB+) — plan storage accordingly.

Example:

Upgrade rollout using software cache: SP32 downloaded to CommServe cache (12GB). 500 clients to upgrade. Without distributed cache: each client downloads from CommServe via WAN (12GB × 500 = 6TB WAN traffic). With distributed cache: cache servers at each of 5 sites (Mumbai, Delhi, Bangalore, Hyderabad, Chennai). Cache synced to each site (12GB × 5 = 60GB WAN). Local clients download from local cache (LAN speed). WAN savings: 99%. Upgrade rollout: 1 night per 100 clients at LAN speed.

Q154

How do you configure Commvault for federated CommCell management?

Answer:

Federated CommCell management allows managing multiple CommCells from a single Command Center or CommCell Console:

- (1) Command Center: add remote CommCells (enter CommServe hostname and credentials);
- (2) Global search: find clients across all CommCells from one interface;
- (3) Cross-CommCell reporting: unified dashboard showing all CommCells;
- (4) Commvault Metrics Reporting: aggregates data from all CommCells for management reports;
- (5) Configuration: each CommCell remains independent — no central DB;
- (6) Delegation: specific operators manage only specific CommCells;
- (7) Cross-CommCell auxiliary copy: replicate data between CommCells for DR or global distribution.

Example:

Enterprise federation: 4 CommCells (Americas/EMEA/APAC/DR). Command Center federated view: global backup health dashboard shows 2,400 clients across all CommCells, 98.1% success rate this week. Alert: APAC CommCell MediaAgent offline (visible from Americas Command Center). Admin remotes to APAC CommServe and resolves. Without federation: must log into each Command Center separately (4 logins, 4 dashboards). With federation: 1 dashboard, global visibility. Quarterly capacity report: MRS generates consolidated 4-CommCell report for management.

Q155

What is the Commvault smart client groups feature?

Answer:

Smart client groups (dynamic client groups) automatically include clients based on rules rather than manual assignment:

- (1) Rule types: OS type (Windows Server, Linux), hostname pattern (SQLSRV*), application type (SQL Server installed), Active Directory OU, geographic location, tags;
- (2) Auto-membership: new clients matching rules auto-added to group;
- (3) Policies applied to group automatically apply to new members;
- (4) Examples: 'All-Windows-Servers' group (OS = Windows Server) auto-includes new Windows VMs. 'Production-SQL' group (hostname matches SQLPRD*);
- (5) Reduces admin work: no manual group management after initial rule setup.

Example:

Smart group 'All-Production-DBs': Rule — any client with SQL Server or Oracle iDA installed + hostname not matching *DEV* or *TEST*. Result: 45 SQL servers and 12 Oracle servers auto-included. New server 'SQLPRD-46' provisioned with SQL iDA installed □ automatically added to group □ inherits SQL-Production storage policy and schedule. Admin action: zero. Without smart groups: admin manually adds SQLPRD-46 to group (easy to forget □ unprotected server). Smart groups eliminate backup coverage gaps from human error.

Q156

How do you roll back a Commvault upgrade if issues are encountered?

Answer:

Upgrade rollback process:

- (1) CommServe rollback: Commvault supports rollback to previous service pack within 10 days (using pre-upgrade snapshot or DB backup);
- (2) Process: stop Commvault services, restore SQL DB from pre-upgrade backup, reinstall previous service pack version;
- (3) Client rollback: uninstall current version, install previous version;
- (4) MediaAgent rollback: similar to CommServe;
- (5) Staged upgrades: upgrade 10% of clients first — if issues, rollback only those 10%;
- (6) Snapshot: take VM snapshot of CommServe before upgrade for instant rollback;
- (7) Test: always test upgrade in non-production CommCell first;
- (8) Commvault Support: engage support before rollback if unclear.

Example:

SP32 upgrade issue: CommServe upgraded to SP32, then discovered all VMware VSA backups failing (known bug in SP32 — vendor patch pending). Rollback decision: revert CommServe to SP30. Process: (1) Stop Commvault services. (2) Restore SQL DB from pre-SP32 backup (CommServe DR export taken before upgrade). (3) Reinstall SP30 on CommServe (45 min). (4) Verify CommServe operational. (5) VMware backups: immediate success. SP32 MediaAgent upgrade: cancelled pending vendor bugfix. SP30 CommServe with SP32 clients: works (forward compatibility is one version).

Q157

What is Commvault job throttle and how does it differ from network throttle?

Answer:

Commvault has two distinct throttle mechanisms: (1) Job throttle (stream/resource throttle): limits the NUMBER of concurrent jobs or streams per MediaAgent or client — a resource count limit; configures how many jobs run simultaneously. (2) Network throttle: limits the DATA TRANSFER RATE (MB/s or Mbps) for backup traffic — a bandwidth limit. Job throttle: prevents resource overload (too many jobs = CPU/memory exhaustion). Network throttle: prevents bandwidth saturation (backup consuming all available WAN/LAN). Both can be configured simultaneously and work independently.

Example:

Scenario requiring both throttles: Remote office with 10Mbps WAN and 30 clients. Network throttle: total WAN backup capped at 8Mbps (80%). Job throttle: maximum 3 concurrent backup jobs from this office (prevent I/O spikes on NAS at branch). Without network throttle: backups consume 10Mbps □ VoIP breaks. Without job throttle: 30 clients start simultaneously □ NAS I/O bottleneck despite WAN limit. With both: 3 jobs run at a time, consuming max 8Mbps combined. All 30 clients backed up within the off-peak window.

Q158

How do you monitor Commvault backup jobs in real time?

Answer:

Real-time job monitoring: (1) CommCell Console □ Job Controller — shows all running, queued, pending jobs with status, progress %, start time, client name; (2) Command Center □ Jobs — web-based job monitoring dashboard; (3) Job details: click job to view events, errors, data transferred, throughput; (4) Active monitoring: set refresh rate (30 seconds auto-refresh); (5) Filter: filter by status (running, queued, failed), client, media agent; (6) API: GET /webconsole/api/Job for programmatic monitoring; (7) Alerts: configure real-time alerts for failures. Job Controller is the primary operational tool for daily backup management.

Example:

Morning health check routine: Job Controller opened — 485 jobs completed overnight (green). 3 jobs failed (red) — click each: Job 1: SQL VSS writer error (fix: restart SQL Server VSS writer service). Job 2: Client offline (server being patched). Job 3: Media full (add capacity to disk library). All 3 issues identified and resolved by 9 AM. Afternoon: 2 jobs currently running (large file servers), 10 minutes remaining each. All healthy.

Q159

What is CommCell Console's Administration Console and what can it configure?

Answer:

CommCell Console Administration Console provides configuration for: (1) CommServe settings — performance, logging, timeout values; (2) User and security management — users, roles, capabilities; (3) License management — view installed licenses, export license info; (4) Activity control — globally pause/resume backups, restores, auxiliary copies; (5) Network topology — CommCell routing for multi-network environments; (6) Commvault Firewall configuration — for environments with strict network segmentation; (7) Event Viewer — detailed system events and errors; (8) CommCell migration — export/import CommCell configurations. Accessible via CommCell Console □ right-click CommServe □ Properties.

Example:

CommServe administration tasks performed via Administration Console: (1) Activity Control: paused all backups for Saturday maintenance window (one click). (2) License: verified 1,250 VM licenses active, 50 remaining. (3) User: created new backup operator account for new team member. (4) Firewall: configured pre-defined routes for new branch office subnet. (5) Events: reviewed last 24h events — found 50 'low disk space' warnings on MediaAgent 02 (added capacity). All in one interface.

Q160

What is the Commvault Data Aging job and how do you configure it?

Answer:

Data Aging configuration: (1) CommCell Console □ Schedules □ Data Aging □ configure run time (default: daily at 3 AM); (2) Set aging frequency; (3) Data Aging evaluates all backup jobs, compares to retention settings; (4) Jobs past retention: media marked reusable; (5) Disable aging temporarily: Activity Control □ Disable Data Aging (for compliance holds); (6) Extended retention: mark specific jobs as 'do not age' (infinite retention); (7) Verify aging: run pruning report to confirm which jobs will be aged; (8) Never prune: jobs under legal hold or with compliance flags. Data Aging is the housekeeping function — essential for storage management.

Example:

Data Aging configuration: scheduled daily 3:30 AM (after nightly backup window). Extended retention: all monthly full backups flagged 'Extended Retention' (7 years) — Data Aging ignores these. Legal hold: 12 client backups under legal hold — never aged. Last night's aging: freed 2TB from 35-day-old jobs (past 30-day retention). Tape media: 3 tapes had all jobs aged □ marked scratch, returned to tape pool. Morning report: 2TB freed, 3 tapes available.

Q161

How do you troubleshoot a Commvault backup failure?

Answer:

Backup failure troubleshooting steps: (1) Job Controller: open failed job □ View Events (detailed error messages with codes); (2) Common error codes: check Commvault documentation or KB articles for specific code; (3) Client connectivity: verify CommServe can reach client (ping, telnet port 8400); (4) Service check: on client, verify Commvault services running (CvFWD, CVD services); (5) VSS check: for VMs/Windows servers, verify VSS writers healthy (vssadmin list writers); (6) Disk space: verify client, MediaAgent, and disk library have space; (7) Logs: client logs in installation directory, MediaAgent logs on MA server; (8) KB search: search Commvault Knowledge Base for error code.

Example:

SQL backup failure: error code 19:118. Step 1: View Events □ 'VSS backup failed, error 0x80042301'. Step 2: VSS check on SQL server □ vssadmin list writers □ 'SQL Server Writer: waiting for completion (timeout)'. Step 3: Restart SQL Server VSS Writer service. Step 4: Retry backup. Success. Root cause: SQL maintenance job held a lock on VSS writer. Fix: stagger SQL maintenance job to not conflict with backup window. Documented in change log.

Q162

What are Commvault's log files and where are they located?

Answer:

Key Commvault log file locations: (1) CommServe logs: C:\Program Files\Commvault\ContentStore\Log Files\ (Windows) — EvMgrS.log (event manager), JobManager.log (job scheduling), AppMgrS.log (application manager); (2) MediaAgent logs: same path on MediaAgent — CVMA.log (media agent service), DedupDB.log (dedup operations); (3) Client logs: C:\Program Files\Commvault\ContentStore\Log Files\ — ClMgrS.log, FileScan.log, DataClassifier.log; (4) Linux logs: /var/log/commvault/Log_Files/; (5) Log collection: CommCell Console □ Diagnostics □ Collect Log Files (zips all relevant logs for support).

Example:

Troubleshooting slow backup: MediaAgent CVMA.log shows repeated 'write timeout' errors to disk library. Disk library NAS: NetApp. NetApp logs: storage controller at 95% CPU during backup window. Root cause: NAS performance bottleneck, not Commvault issue. Fix: NetApp storage QoS policy added to prevent backup from consuming all NAS I/O. CVMA.log confirmed — timeout errors gone after QoS policy. Commvault log collection: CommCell Console □ Diagnostics □ Collect Logs □ sent to Commvault Support for additional analysis.

Q163

How do you configure Commvault activity control?

Answer:

Activity Control manages operational states at various levels: (1) Global: CommCell Console □ Activity Control □ disable/enable all backup, restore, auxiliary copy CommCell-wide; (2) Client level: right-click client □ Properties □ Activity Control □ disable backup or restore for specific client; (3) Agent level: disable specific iDataAgent operations; (4) Use cases: disable backups during major maintenance, enable/disable specific clients during migrations, temporarily disable restores during storage maintenance; (5) Override: individual job-level override available; (6) Status visible: disabled clients show with indicator in CommCell Console. All activity control changes logged in audit trail.

Example:

Maintenance window: Saturday 10 PM storage array maintenance (4-hour window). Activity Control: CommCell-wide backup disabled at 10 PM. Auxiliary copies: disabled. Restores: left enabled (users may need file recovery during maintenance). Array maintenance completes at 2 AM. Activity Control: re-enable all operations. Automated jobs that were queued during maintenance: now run immediately. All within 10-minute window. Zero manual job management required.

Q164

What is the Commvault Audit Trail and what does it record?

Answer:

Commvault's audit trail records all administrative actions for compliance and security: (1) Login events — user logins, failed attempts, logout; (2) Configuration changes — policy modifications, new clients added, deleted; (3) Restore operations — who initiated, what data, from where, to where; (4) Job operations — kill, suspend, resume; (5) Security changes — user creation, role assignment changes; (6) Media operations — tape export, import; (7) License events. Audit trail accessible: CommCell Console □ Audit Trail. Exportable to CSV for SIEM integration. Retention of audit data configurable.

Example:

Compliance audit requirement: 'Prove no unauthorised data access for last 6 months.' Commvault Audit Trail export: filtered for restore operations. Report shows: 452 restore operations, all by authorised users (backup-team group), all within working hours except 3 after-hours restores (all approved by manager per change records). No unauthorised restores detected. GDPR auditor satisfied. Audit trail records: user, timestamp, client, restored path, target location for every restore.

Q165

How do you configure Commvault disaster recovery (DR) runbook testing?

Answer:

DR testing in Commvault: (1) Commvault Disaster Recovery Orchestration — define recovery plans with VM boot order, dependencies, scripts; (2) Test failover: isolated environment (no production impact), VMs recovered to test network; (3) Replication validation: verify Live Sync data is current before test; (4) Recovery time measurement: Commvault records RTO for each phase; (5) Documentation: recovery report generated automatically; (6) Use IntelliSnap for fast test recovery from array snapshots; (7) Cleanup: delete test VMs after validation; (8) Rollback: always available during test (test VMs deleted, production unchanged).

Example:

Quarterly DR test for CRM application (20 VMs): Recovery plan: Boot Group 1 — Oracle DB (2 VMs), 5-min delay. Group 2 — App servers (8 VMs), 3-min delay. Group 3 — Web tier (10 VMs). Test environment: isolated VLAN 999 (no production routing). Commvault restores from latest backup to test VMs. Phase timing: DB up in 8 min, App in 5 min, Web in 4 min. Total RTO: 17 minutes (SLA: 30 min). CRM application validated. Recovery report auto-generated. Test VMs deleted. Production unaffected.

Q166

What is the CommCell Console Event Viewer?**Answer:**

CommCell Console Event Viewer shows system events generated by all Commvault components: (1) Job events — backup start/completion, errors within jobs; (2) System events — services starting/stopping, connectivity issues; (3) Media events — tape errors, disk library warnings; (4) Security events — login attempts, permission changes; (5) Filter by: severity (critical, major, minor, information), time range, component, type; (6) Real-time refresh; (7) Export to CSV; (8) Link to related job from event; (9) Integrated with alerts — events trigger email/SNMP alerts. Event Viewer is the primary diagnostic tool for investigating issues beyond job-level errors.

Example:

Event Viewer investigation: Monday morning, 15 events marked 'Critical'. Filter: Critical severity, last 12 hours. Events: 12 × 'Cannot connect to MediaAgent MA-02' (from CommServe). Investigation: MA-02 server had network configuration change by server team — replication NIC IP changed without notifying backup team. Fix: update MediaAgent IP binding in CommCell. Event Viewer identified root cause in 5 minutes vs hours of guessing.

Q167

How do you manage tape media in Commvault?**Answer:**

Tape media management: (1) CommCell Console □ Storage □ Libraries □ tape library □ Media; (2) Media status: blank (new), full (in use), append (partially written), scratch (reusable), overwrite protected; (3) Barcode scanning: Commvault reads barcode labels to track media; (4) Media rotation: export media for offsite storage, import when needed; (5) Media aging: when all jobs on a tape are past retention, tape marked scratch; (6) Cleaning tapes: configure cleaning slot and frequency; (7) Media labelling: Commvault assigns internal labels matching barcodes; (8) Missing media: track tapes at offsite vault, request recall when needed for restore.

Example:

Monthly tape rotation: 20 tapes exported to offsite vault (Iron Mountain) after weekly full backups complete. Commvault media tracking: these 20 tapes marked 'exported'. Vault location recorded in CommCell. 90-day-old tapes: retention passed, Commvault marks them 'eligible for recycle'. Tape recall: audit department needs data from 2-year-old tape. Commvault: tape barcode CV2024001 requested from Iron Mountain. Tape arrives next day, imported into library. Restore initiated from that tape. Process: 2 days total.

Q168

What is the Commvault Resource Manager and what does it control?**Answer:**

Resource Manager controls resource allocation in the CommCell to prevent overload: (1) Stream limits — maximum concurrent backup streams per MediaAgent or client; (2) Drive limits — tape drive allocation; (3) Network throttling — bandwidth limits per client, group, or time schedule; (4) Job priority — high/medium/low for scheduling; (5) DDB load balancing — distribute dedup I/O across multiple DDB partitions; (6) Throttle schedules — different limits during business hours vs off-peak; (7) Application-aware throttling — VMs, databases with different resource profiles. Prevents backup from saturating production networks or storage.

Example:

Resource Manager for WAN environment: 20 remote offices each with 200Mbps WAN links. Without throttle: backups consume 100% WAN = production network failure. Resource Manager: per-office WAN throttle 8AM-6PM: 20Mbps (10%). 6PM-8AM: 150Mbps (75%). Per-client stream limit: 2 streams max per remote client. Result: backups complete overnight within SLA, production WAN unaffected during business hours.

Q169

How do you perform a CommServe migration to new hardware?

Answer:

CommServe migration steps: (1) Build new Windows Server with SQL Server (matching version); (2) Install Commvault CommServe software on new server; (3) Stop Commvault services on old CommServe; (4) Back up the CommServe database (SQL backup); (5) Restore SQL database to new CommServe's SQL instance; (6) Update CommServe hostname in Commvault: client properties for all clients must reflect new CommServe if hostname changes; (7) Update IP/DNS if CommServe IP changes; (8) Start Commvault services on new CommServe; (9) Verify all components connect; (10) Test backup and restore. IP/hostname change requires client updates.

Example:

CommServe hardware refresh: New server: Windows 2022, SQL 2019, 32GB RAM, NVMe SSD. Process: Friday 10 PM – stop old CommServe services. SQL backup (15 min). Restore SQL to new server (20 min). Commvault installed on new server, DB attached. DNS updated: commserve.company.com → new IP. All clients use FQDN (hostname resolution updates automatically). Commvault services started on new server. Saturday 1 AM: all 500 clients reconnected. Test backup: success. Old CommServe decommissioned Monday. Zero backup downtime over the weekend.

Q170

How do you configure CommCell Console for multi-CommCell management?

Answer:

Multi-CommCell management: (1) CommCell Console → Add CommCell → enter remote CommServe hostname and credentials; (2) Each additional CommCell appears as a separate tree in CommCell Console; (3) Switch between CommCells without opening separate consoles; (4) Metrics Reporting Server: centralised cross-CommCell reporting (separate setup); (5) Global search: search for clients across all CommCells; (6) Operations: each CommCell administered independently but accessible from one console; (7) CommCell Sync: share license/settings between CommCells; (8) Cross-CommCell auxiliary copy: replicate data between CommCells.

Example:

Enterprise with 4 CommCells (Americas, EMEA, APAC, DR): CommCell Console configured with all 4. Admin switches between Americas and EMEA in seconds. Morning health check: view Job Controller for all 4 CommCells sequentially. Report: total 2,100 clients across all CommCells, 98.5% backup success globally. Alert from EMEA CommCell: disk library at 88% – added NAS capacity. All managed from single CommCell Console session.

Q171

What is the Commvault SLA report and how is it used?

Answer:

SLA (Service Level Agreement) reports show which clients met or missed their backup SLA: (1) Protected: clients backed up within SLA window; (2) Unprotected: clients that missed backup (never backed up, backup failed, or outside SLA window); (3) SLA window: configurable per client or globally (e.g., must be backed up in last 24 hours); (4) Categories: Met SLA, Missed SLA, Not configured; (5) Used for: daily health review, management reporting, compliance evidence; (6) Accessed: Command Center → Reports → SLA; (7) Scheduled: email SLA report daily to IT manager; (8) Drill-down: click missed SLA client to see last backup attempt and failure reason.

Example:

Daily SLA report review (Monday 8 AM): 492 clients total. Met SLA: 475 (96.5%). Missed SLA: 17. Analysis: 10 missed due to weekend patching (servers rebooting during backup window – reschedule patching). 5 failed (VSS issues – 3 fixed). 2 not configured (new servers provisioned Friday – assign plans today). Manager report: 96.5% SLA this week vs 98% target. Action plan: patching schedule change, 2 new server plans. Next week target: 98%+.

Q172

How do you upgrade Commvault MediaAgent without interrupting backup?

Answer:

MediaAgent upgrade without interruption: (1) Drain jobs: prevent new jobs from being assigned to the MA being upgraded (activity control); (2) Wait for running jobs to complete; (3) Run upgrade on MA; (4) MA service restarts automatically; (5) Re-enable for new jobs. For 2+ MediaAgents: (1) Upgrade MA1 while MA2 handles all jobs; (2) After MA1 completes, upgrade MA2; (3) Load balancing ensures no backup gap. Single MediaAgent: schedule upgrade during low-activity window; brief interruption for running jobs (they complete before upgrade, or suspend/resume). Upgrade time: 15-30 minutes per MA.

Example:

2 MediaAgent environment (MA1 and MA2). Upgrade plan: Saturday midnight. Step 1: Activity Control □ disable new jobs on MA1. Running jobs on MA1: 5 jobs, wait 20 minutes to complete. MA1 upgrade: 20 minutes. MA2 handles all jobs during MA1 upgrade. MA1 back online. Disable MA2, wait for MA2 jobs to complete (15 min). MA2 upgrade: 20 minutes. Both MAs upgraded within 75 minutes. Zero backup failures. Jobs ran on MA2 while MA1 was upgrading.

Q173

What is the Commvault capacity management and how do you plan storage?

Answer:

Storage capacity planning in Commvault: (1) Metrics: track disk library utilisation over time; (2) Growth projections: weekly data growth rate (% increase in protected data); (3) Dedup efficiency: track dedup ratio — higher ratio = more effective storage; (4) Retention impact: longer retention = more storage required; (5) Reports: Command Center □ Reports □ Storage □ Disk Library Usage; (6) Alerts: configure alert when disk library >80%; (7) Capacity planning formula: (Protected Data GB × Retention Days × Change Rate%) ÷ Dedup Ratio = Required Storage; (8) Cloud tier: overflow to cloud for capacity flexibility.

Example:

Capacity planning exercise: 500 clients, 10TB total data protected, 5% daily change rate, 30-day retention, 8:1 dedup ratio. Calculation: $(10TB \times 0.05 \times 30 \text{ days}) \div 8 = 1.875TB$ disk required for incrementals. Plus one full copy $(10TB \div 8 = 1.25TB)$. Total: ~3TB with dedup. Actual allocation: 5TB (safety margin). Cloud overflow: Azure Blob configured — if disk library >80%, auxiliary copy immediately goes to cloud. No unplanned capacity crisis.

Q174

How do you configure Commvault for high availability (HA)?

Answer:

Commvault HA configuration: (1) CommServe HA: Commvault native failover (active/passive), Windows Server Failover Cluster (WSFC), or vSphere HA for the CommServe VM; (2) CommServe Live Sync: real-time database replication to standby CommServe; (3) MediaAgent redundancy: 2+ MediaAgents so one can fail without stopping backups; (4) Disk library redundancy: RAID storage or mirrored NAS; (5) VSA proxy HA: 2+ VSA proxies for VMware backup; (6) Network redundancy: NIC teaming on all Commvault components; (7) DDB backup: regular DDB backup ensures dedup data recoverable.

Example:

HA architecture: CommServe on vSphere HA cluster (auto-restart on host failure, 2-min RTO). CommServe Live Sync to passive CommServe at DR site. 2 MediaAgents (active-active). Disk library: NetApp with RAID-DP (dual parity). 2 VSA proxies (round-robin load balancing). All components on 10GbE with NIC teaming. DDB backed up daily to disk library copy. Last test: primary CommServe VM simulated failure — vSphere HA restarted in 90 seconds, all pending jobs queued, resumed within 5 minutes. No backup gap.

Q175

What is Commvault's media management and what are media statuses?

Answer:

Commvault media management tracks physical and virtual media (tapes, disk volumes). Media statuses: (1) Blank — new, never written; (2) Active — currently being written (append mode); (3) Full — no more space; (4) Scratch — aged and reusable; (5) Overwrite Protected — locked, cannot overwrite; (6) Exported — sent offsite; (7) Bad — media errors, unusable; (8) Recycled — scratch media re-labelled for reuse; (9) Cleaning tape — tape cleaning cartridge. Disk media statuses similar: disk volume online/offline, full/available. MediaAgent manages all media transitions automatically.

Example:

Tape library media management: 100 LTO-8 tapes. Status breakdown: 20 active (being written tonight), 45 full (data within retention), 30 scratch (reusable — all jobs aged), 5 bad (hardware errors — removed from library), 0 blank (all tapes used at least once). 5 bad tapes: Commvault moved their data to auxiliary copies before marking bad. Physical inspection: 5 damaged cartridges ordered as replacements. Data integrity maintained.

Q176

How do you use Commvault's CommCell Console reports?

Answer:

CommCell Console reports provide operational and management information: (1) Job Summary Report — success/failure overview for a time period; (2) Media Report — tape/disk usage, status, availability; (3) Subclient Report — which subclients have/haven't been backed up; (4) License Report — license usage and available licenses; (5) SLA Report — clients meeting/missing SLA; (6) Capacity Report — storage usage by client, library, media type; (7) Data Transfer Report — throughput per job; (8) Dedup Savings Report — storage saved by deduplication; (9) Custom reports — build custom reports via report builder. Reports can be scheduled for automatic email delivery.

Example:

Weekly management report package (auto-emailed every Monday): (1) SLA Report: 97.8% this week. (2) Job Summary: 2,100 jobs, 98.2% success. (3) Capacity: disk library at 67% (comfortable). (4) Dedup Savings: 12.3:1 ratio saved \$45,000 in storage costs this month vs no-dedup scenario. (5) License: 1,240/1,300 licenses used (60 available, procure 100 more Q2). Management receives concise 5-report package without any admin effort after initial setup.

Q177

How do you configure CommCell notification for SLA breach?

Answer:

SLA breach notification configuration: (1) Command Center □ Manage □ Notifications □ SLA notifications; (2) Or CommCell Console □ Alerts □ Add Alert □ select 'SLA violation' alert type; (3) Configure threshold: missed SLA for X hours; (4) Set recipients: email, SNMP, webhook; (5) Frequency: alert immediately on miss, or batch daily report; (6) Per-client alerts: create alert scoped to specific critical clients (immediate notification for Tier-1 clients, daily digest for Tier-2); (7) Integration: send SLA miss to ServiceNow via webhook (auto-create incident); (8) Resolution: alert clears when client successfully backed up.

Example:

Tiered SLA alerting: Tier-1 (50 critical clients): immediate email to backup-oncall@company.com and PagerDuty page on SLA miss. Tier-2 (200 standard clients): daily SLA miss digest email to backup-team@company.com at 8 AM. Tier-3 (100 dev clients): weekly SLA report only. Alert integration: Tier-1 SLA miss □ ServiceNow incident auto-created with client name, last backup time, failure reason. On-call resolves within 30 minutes. SLA metrics tracked in ServiceNow for management review.

Q178

How do you configure email alerts in Commvault CommCell Console?

Answer:

Email alerts are configured via the Alert Manager under the CommCell node. Administrators define alert types (job failure, media error, license expiry), thresholds, recipient email addresses, and SMTP server settings. Alerts can be scoped to specific clients, agents, or the entire CommCell.

Example:

An alert is configured to send an email to storage-team@company.com whenever a backup job fails or runs more than 20% over the expected time window. The CommServe SMTP settings point to the corporate mail relay at smtp.company.com.

Q179

What is the role of the Workflow Engine in Commvault?

Answer:

The Workflow Engine allows automation of complex, multi-step data management tasks using pre-built or custom workflows. Workflows can trigger backups, approvals, notifications, scripts, and restores in sequence — enabling end-to-end automation of common operations without manual intervention.

Example:

A 'New Client Onboarding' workflow automatically installs the file system agent on a new server, creates a subclient with standard policies, runs an initial full backup, and sends a completion email to the backup team — all triggered by a single request.

Q180

What is a User Group in Commvault and how is it used for access control?

Answer:

A User Group in Commvault aggregates users and assigns them a common set of capabilities and permissions. Permissions can be scoped to CommCell components (clients, storage policies, reports) and operation types (browse, restore, schedule). This enables role-based access control (RBAC) across the CommCell.

Example:

A 'Backup Operators' user group has permissions to monitor and restart backup jobs for assigned clients but cannot modify storage policies or delete data. Junior admins are added to this group without receiving full CommCell admin rights.

Q181

How does Commvault's capacity licensing model work?

Answer:

Commvault's capacity-based license is measured against the front-end terabytes (FETB) of protected data — the size of the source data before backup, deduplication, or compression. The license covers all Commvault features for that data capacity, regardless of how many copies, agents, or servers are used.

Example:

A company with 50TB of source data (VMs, file servers, databases) purchases a 50 FETB license. Even if Commvault stores 200TB total (with multiple copies and no deduplication credit applied), the license is still calculated on the 50TB source footprint.

Troubleshooting

Q182

How do you resolve 'DDB is offline' errors in Commvault?

Answer:

DDB offline resolution: (1) Identify which MediaAgent's DDB is offline: CommCell Console □ Storage □ MediaAgents □ DDB status; (2) Root causes: DDB storage path unmounted/offline, DDB disk full, DDB files corrupted, MediaAgent service stopped; (3) Check DDB disk: verify disk path exists and has space on MediaAgent server; (4) Remount: if NFS/disk unmounted, remount it; (5) DDB bring online: CommCell Console □ MediaAgent □ DDB □ right-click □ Bring Online; (6) DDB verification: run DDB verification to check integrity; (7) DDB restore: if corrupted, restore from DDB backup; (8) During offline: new backup jobs using that DDB will fail. Restores from non-dedup copies still work.

Example:

DDB offline alert at 2 AM. Investigation: DDB path /ddb/nvme1 on MediaAgent-02 shows 'offline'. Server check: /dev/nvme1 partition not mounted (auto-mount failed after reboot). Fix: mount /dev/nvme1 /ddb/nvme1 □ verify DDB files present. CommCell Console □ MA-02 □ DDB □ Bring Online. DDB online in 3 minutes. All queued backup jobs resume. Root cause: /etc/fstab missing nofail option for NVMe mount. Fix: add nofail to fstab entry. Monitor: no further occurrences after fstab correction.

Q183

What causes backup data verification failures and how do you resolve them?

Answer:

Data verification failure causes: (1) Storage hardware failure: disk errors on NAS, tape read errors, cloud retrieval errors; (2) Data corruption: backup data blocks unreadable (bad blocks on storage); (3) Partial media: incomplete backup data on tape/disk; (4) Network interruption: cloud retrieval fails due to network issue; (5) Mediaagent issue: MA cannot mount or read from library; (6) Resolution: (a) Quick DV failure → run Full DV to confirm; (b) Confirm whether alternate copies exist (secondary copy); (c) If secondary copy clean: auxiliary copy to rebuild primary; (d) If all copies corrupt: data is lost — investigate storage hardware; (e) Alert immediately as data unrecoverable if all copies corrupt.

Example:

Full DV on monthly tape backup fails: 50 blocks unreadable on tape CV20240115. Investigation: tape cartridge shows physical damage (edge nick from vault handling). All backup data on this tape potentially unrecoverable. Secondary check: same data exists on disk library (primary copy — 90-day retention). Primary copy DV: PASS. Recovery plan: re-run auxiliary copy from primary disk to new tape (CV20240115_R). Secondary tape now clean. Physical tape CV20240115: marked 'Bad' in CommCell, quarantined. Prevention: tape handling training for vault staff.

Q184

How do you troubleshoot slow restore operations in Commvault?

Answer:

Slow restore performance troubleshooting: (1) Identify bottleneck: network, storage read speed, MediaAgent CPU, client write speed; (2) Network: check bandwidth between MediaAgent and client during restore; (3) Storage: disk library read speed — NAS I/O during restore (competing jobs?); (4) Tape: tape restore inherently sequential — cannot parallelize individual tape reads; (5) DDB for deduplicated restores: DDB I/O impacts restore throughput; (6) Multiple streams: increase restore streams for parallel recovery; (7) MediaAgent load: if MA handling many jobs, restore competes for resources; (8) Client disk: destination disk write speed limiting restore; (9) Cloud restore: egress bandwidth and latency from cloud storage.

Example:

VM restore running at 50MB/s (expected: 500MB/s). Investigation: (1) Network: backup LAN showing 100% utilization — 20 backup jobs running simultaneously. (2) MediaAgent: 95% CPU (dedup operations for backup jobs consuming resources). Root cause: restore competing with peak-hour backups. Fix: (1) Suspend 20 backup jobs (non-critical). (2) Restore now has full MediaAgent resources. (3) Restore speed: 480MB/s. Lesson: critical restores should pre-empt backup jobs. Configure restore job priority: Critical → automatically suspends lower-priority backup jobs.

Q185

What is a Commvault 'aged data but not pruned' issue and how to resolve it?

Answer:

Aged but not pruned: Data Aging has determined data is eligible for pruning but it hasn't been removed from storage. Causes: (1) Auxiliary copy pending: primary data not pruned until secondary copy confirmed complete (copy precedence); (2) Dependent jobs: archived jobs pointing to this data; (3) Legal hold: data flagged with hold cannot be pruned; (4) Data Aging job not running: scheduled job disabled or failing; (5) DDB mark for pruning queue backed up; (6) Resolution: (a) Run Auxiliary copy to satisfy copy precedence; (b) Verify legal hold not applied incorrectly; (c) Run Data Aging manually (right-click on storage policy → Run Data Aging).

Example:

Disk library at 95% despite 30-day retention policy — expected usage: 60%. Investigation: Data Aging report shows 500GB 'aged but not pruned'. Root cause: secondary copy (tape) not completed for these jobs (copy precedence configured: primary waits for secondary before pruning). Tape library: 2 drives offline for maintenance. Fix: bring tape drives online → auxiliary copy completes → Data Aging re-runs → 500GB pruned from disk library. Disk: drops to 72%. Prevention: tape library maintenance requires activity control suspension and auxiliary copy catchup plan.

Q186

What are the steps to troubleshoot a Commvault MediaAgent that won't come online?

Answer:

MediaAgent offline troubleshooting: (1) Server status: verify the server itself is powered on and accessible (ping, RDP); (2) Services: check Windows Services or Linux systemctl for Commvault MediaManager service; (3) Start service: if stopped, start manually; (4) Service errors: view Windows Event Log for Commvault service start failures; (5) Port: verify CommServe can reach MA on port 8400; (6) Certificate: if certificate mismatch, re-register MA; (7) DDB path: if DDB path missing, MA may refuse to start; (8) Disk space: MA server disk full can prevent services from starting; (9) Logs: review CVMA.log on MediaAgent for startup errors; (10) Re-register: if all else fails, deconfigure and re-add MA to CommCell.

Example:

MediaAgent offline after server reboot. Steps: (1) Server: accessible via RDP □. (2) Services: Commvault MediaManager service stopped (not set to auto-start after reboot). (3) Start service: service starts. (4) MA status in CommCell: Online □. (5) Root cause: automatic service recovery not configured. Fix: set all Commvault services to Auto (Delayed Start) in Windows Services. Also: DDB NFS mount not restored after reboot (missing from fstab). Fix: add to fstab. Lesson: always test server reboot after installing Commvault – verify services and mounts auto-recover.

Q187

How do you recover from an accidental deletion of backup data in Commvault?

Answer:

Recovery from accidental backup data deletion: (1) Prevention: no accidental deletion should be possible – DDB manages pruning, not manual deletion; (2) If disk library files deleted manually: backup data may be unrecoverable; (3) Alternate copy: check if secondary or tertiary copy exists (auxiliary copy); (4) Re-cataloging: if data exists but CommServe doesn't know about it, re-catalogue from storage; (5) CommServe DB restore: if CommServe DB has records of deleted data, restore DB from before deletion; (6) Commvault support: Commvault Support has tools for partial recovery of CommServe metadata; (7) Legal implications: document accidental deletion if under compliance/legal hold.

Example:

Scenario: admin accidentally deleted 50 backup jobs' data from disk library mount path (ran 'rm -rf' on wrong directory). Assessment: (1) Are secondary copies available? YES – tape library has secondary copy for all 50 jobs. (2) Disk library: those 50 jobs' blocks gone from disk. (3) Tape: fully intact. (4) Action: mark primary copy jobs as 'disabled' (don't re-use those disk blocks – they're gone). Restore operations: CommCell configured to use tape (secondary) for these jobs. Users can still restore from tape. No data loss. Lesson: restrict file system access to disk library mount paths.

Q188

What is Commvault's simpana.log and when is it used for troubleshooting?

Answer:

Simpana.log (legacy name, now often EvMgrC.log or CVD.log depending on version): primary client-side Commvault log file that records: (1) Connection events to CommServe; (2) Job start/stop events received from CommServe; (3) Service startup and errors; (4) Tunnel communication events; (5) Certificate events. Used for: diagnosing communication issues between client and CommServe, service startup failures, job receipt problems. Location: C:\Program Files\Commvault\ContentStore\Log Files\ (Windows) or /var/log/commvault/Log_Files/ (Linux). In current versions: primary logs are CMgrS.log (client manager service) and CVD.log (Commvault daemon).

Example:

Client not receiving backup jobs from CommServe. CVD.log on client: 'Failed to authenticate CommServe certificate – certificate mismatch'. Root cause: CommServe SSL certificate renewed but client still has old certificate cached. Fix: CommCell Console □ client □ right-click □ Push Certificates. New certificate pushed to client. CVD.log: 'Certificate updated successfully'. Next backup job: received and executed normally. Log analysis confirmed issue in 3 minutes vs 30 minutes of guessing.

Q189

What causes Commvault auxiliary copy lag and how do you resolve it?

Answer:

Auxiliary copy lag: auxiliary copy falling behind the primary backup (data on primary not yet copied to secondary). Causes: (1) Network bandwidth insufficient for copy rate; (2) Target MediaAgent or library bottleneck; (3) Source MediaAgent reading slower than backup rate; (4) Auxiliary copy schedule too infrequent; (5) Throttle too aggressive on aux copy; (6) Large backup jobs creating a large backlog. Resolution: (1) Increase auxiliary copy frequency or run continuously; (2) Remove throttle temporarily to catch up; (3) Add MediaAgent at target site; (4) Increase network bandwidth between sites; (5) Prioritise auxiliary copy in resource manager; (6) Run manual auxiliary copy to catch up.

Example:

Auxiliary copy 5 days behind (disk-to-tape). Discovery: Monday morning, 5TB data on disk not yet on tape. Root cause: weekend had unusually large full backups (100TB annual report data). Auxiliary copy running at 500GB/night – only moving 2.5% of backlog/night. Emergency resolution: (1) Suspend non-critical backups. (2) Remove auxiliary copy throttle (50Mbps → unlimited). (3) Add second tape drive to library (was using 1 of 2 drives). (4) Manual auxiliary copy run all Monday: catches up 4TB. (5) Tuesday: 1TB remaining, completes by Tuesday night. Normal operations resume.

Q190

How do you troubleshoot Commvault Command Center login issues?

Answer:

Command Center login troubleshooting: (1) URL: verify correct URL (<https://CommServeHostname/webconsole>); (2) Browser: test in different browser (Chrome, Firefox); (3) SSL certificate: if self-signed, browser may block – add exception; (4) Services: verify Commvault Web Server service running on CommServe; (5) Port: port 443 open on CommServe firewall (Windows Firewall + network firewall); (6) AD authentication: verify AD connectivity from CommServe (can CommServe reach domain controller); (7) Credentials: verify username format (DOMAIN\username or username@domain.com); (8) Account lockout: check if AD account locked; (9) Logs: IIS logs on CommServe for HTTP errors.

Example:

Command Center login fails with 'Unable to connect'. Investigation: (1) URL correct. (2) Browser: Chrome + Firefox both fail. (3) From CommServe itself: <https://localhost/webconsole> → loads. (4) External access: Port 443 blocked. Check: Windows Firewall on CommServe has no inbound rule for 443. Fix: add Windows Firewall inbound rule TCP 443. Test: external browser access works. Lesson: Commvault installer doesn't automatically add Windows Firewall rules for Command Center HTTPS – must be added manually post-installation.

Q191

What are causes of Commvault backup failing with 'Access Denied' on Windows?

Answer:

Access Denied on Windows backup causes: (1) Service account permissions: Commvault service account lacks read access to the backed-up path; (2) NTFS permissions: specific folders with restricted permissions; (3) Share permissions: if backing up via UNC path; (4) Volume Shadow Copy: VSS access denied (service account needs local admin); (5) SQL permissions: SQL service account not granted backup permissions; (6) UAC: User Account Control blocking Commvault agent; (7) Resolution: (a) Add service account to Local Admins (recommended for backup agents); (b) Grant specific NTFS permissions; (c) Verify VSS operator membership; (d) Check specific error in job events for exact path denied.

Example:

Incremental backup: 500 files backed up, 50 files 'Access Denied'. Job Events: 'Cannot access D:\Finance\Restricted\Payroll\' – access denied. Investigation: D:\Finance\Restricted\Payroll\ has custom NTFS permissions – only HR-Admin group has access. Commvault service account svc-commvault: not in HR-Admin. Fix options: (1) Add svc-commvault to HR-Admin group (simpler). (2) Grant svc-commvault 'Read' on that specific folder (principle of least privilege). Option 2 chosen – security team approves. Access Denied: resolved for Payroll folder.

Q192

How do you handle a Commvault job that is stuck in 'Waiting' state?

Answer:

Job stuck in Waiting state causes: (1) Waiting for resource: MediaAgent at stream limit, tape drive busy, disk library full; (2) Waiting for commit: previous backup job phase not committed in CommServe DB; (3) Waiting for network: client unreachable, network flap; (4) Waiting for mount: tape mount request pending (tape offline or ejection requested); (5) Priority conflict: lower-priority job waiting for high-priority to release resources; (6) Resolution: (a) Check resource availability (streams, library space, tape drives); (b) View job events for specific wait reason; (c) Kill stuck job and restart; (d) Check CommServe services if all jobs waiting (may indicate CommServe DB issue).

Example:

12 jobs stuck in Waiting for 4 hours. Job Events: 'Waiting for Drive – Library TapeLib01'. Tape library check: both LTO drives showing 'Busy' for 4 hours. Investigation: a previous auxiliary copy job started a tape mount 4 hours ago and the mount never completed (tape loader robot arm fault). Fix: resolve robot arm fault via tape library management console □ drive freed. Auxiliary copy completes mount □ frees drives. All 12 waiting jobs begin within 2 minutes. Robot arm: scheduled for maintenance.

Q193

How do you diagnose and fix CommServe database performance issues?

Answer:

CommServe DB performance issues: (1) Symptoms: jobs slow to start, CommCell Console unresponsive, queries timeout; (2) SQL performance: check SQL Server resource usage (CPU, memory, I/O); (3) Database fragmentation: CommServe DB index fragmentation degrades performance over time – run index maintenance; (4) DB growth: CommServe DB grows as job history accumulates – if disk full, SQL cannot write; (5) SQL maintenance: Commvault has built-in CommServe DB maintenance jobs (run automatically); (6) TempDB: SQL TempDB on slow storage degrades performance; (7) Memory: SQL needs sufficient RAM for buffer pool (CommServe SQL should have 16-32GB); (8) Antivirus: AV scanning SQL data files in real time causes extreme performance degradation.

Example:

CommCell Console taking 3-5 minutes to load. SQL check: SQL CPU 80%, wait type: PAGEIOLATCH_SH (disk I/O waits). CommServe SQL DB and TempDB on same HDD. Fix: move TempDB to NVMe SSD. SQL performance check: index fragmentation on JobHistory table: 87% fragmented. Fix: run index rebuild on JobHistory (2 hours). After both fixes: CommCell Console load: 15 seconds. Job scheduling latency: reduced from 5 minutes to 30 seconds. Preventive: weekly SQL index maintenance job scheduled via SQL Agent.

Q194

How do you use Commvault's diagnostic tool CVPing for connectivity testing?

Answer:

CVPing is a Commvault-specific connectivity testing tool that tests communication using Commvault protocols (not just ICMP ping). Usage: (1) Location: CommServe/MediaAgent/client installation directory □ CVPing.exe (Windows) or cvping (Linux); (2) Test CommServe-to-client: CVPing -cs [CommServe_IP] -cl [client_hostname]; (3) Test client-to-CommServe: run CVPing on client pointing to CommServe; (4) Results: shows if Commvault tunnel (port 8400) connection established, authentication status, latency; (5) More accurate than ping for diagnosing Commvault communication issues (includes authentication testing, not just network reachability).

Example:

Client backup failing with connection errors. Regular ping from CommServe to client: SUCCESS (10ms). CVPing from CommServe to client: FAILED – 'Authentication failed – certificate mismatch'. Regular ping showed network was fine. CVPing revealed the real issue: certificate authentication failure (not a network problem). Fix: push new certificate to client. CVPing after fix: SUCCESS – 12ms, authenticated. Backup retry: success. CVPing saved 30 minutes of network troubleshooting by immediately identifying the authentication layer as the problem.

Q195

What is Commvault's client deconfiguration and when is it used?

Answer:

Client deconfiguration removes a client from CommCell management without deleting its backup data. Use cases: (1) Server decommissioned but backup data still needed for compliance; (2) Re-registering a client with new hostname or IP; (3) Troubleshooting: deconfigure and re-add to resolve communication issues; (4) Migration: deconfigure from old CommCell before adding to new CommCell. Process: (1) CommCell Console ▢ right-click client ▢ Deconfigure; (2) Options: delete backup data (permanent — use carefully), or retain data; (3) Retain data: client removed from management but all backed-up data accessible for restore; (4) Re-register: reinstall Commvault on client, point to CommServe — re-registers automatically.

Example:

Server SQLSRV-05 being decommissioned. Compliance requirement: keep 7 years of backup data. Action: Deconfigure client SQLSRV-05 with 'Retain backup data'. Client no longer appears in CommCell as active client. Backup data: still on disk library and tape under storage policy retention. 3 years later: legal department requests data from SQLSRV-05 from 2021. Restore: CommCell Console ▢ Retired Clients ▢ SQLSRV-05 ▢ Browse ▢ restore files from 2021. Data accessible despite server decommissioned 3 years ago.

Q196

How do you resolve a Commvault backup failure with error 'Failed to connect to CommServe'?

Answer:

Resolution steps: (1) Verify CommServe hostname/IP is correct on the client (client properties ▢ CommServe details); (2) Network connectivity: ping CommServe from client; (3) Port check: telnet CommServe-IP 8400 (Commvault tunnel port — must be open); (4) Commvault services: on CommServe, verify all Commvault services running (CommVault Communications Manager, Galaxy Event Manager, etc.); (5) Firewall: check Windows Firewall and network firewalls for port 8400; (6) DNS: verify client can resolve CommServe FQDN; (7) Re-register client: if communication broken, re-register client with CommServe.

Example:

Branch office backup failing with 'Cannot connect to CommServe'. Investigation: ping CommServe IP from branch — success. Telnet port 8400: connection refused. Firewall rule review: branch firewall had TCP 8400 blocked after recent rule update (firewall refresh removed the Commvault rule). Fix: re-add TCP 8400 permit rule. Backup retry: success. Root cause: change management process failure — firewall change not reviewed for impact on Commvault. Prevention: add Commvault ports to firewall change impact checklist.

Q197

What causes 'Media not available' errors in Commvault?

Answer:

Media not available error causes: (1) Tape library: tape drive offline, tape ejected, library door open, robot arm fault; (2) Disk library: mount path offline (NAS unreachable, drive disconnected), permissions changed; (3) Cloud library: cloud storage credentials expired, bucket deleted, network access blocked; (4) MediaAgent: MediaAgent service stopped, MediaAgent VM powered off; (5) Library configuration: library not assigned to MediaAgent, mount path permissions incorrect; (6) Media status: all media full or in 'overwrite protected' state (need more tapes or disk space).

Example:

Backup jobs failing with 'Media not available' for disk library. Investigation: disk library mount path \\NAS01\Backup. Test: try to access from MediaAgent server ▢ 'Access denied'. Root cause: NAS01 service account password expired (90-day rotation policy). Fix: update service account password on NAS01, update Commvault credentials for mount path. Backup retry: success. Prevention: Commvault service account password rotation scheduled in team calendar with 2-week advance alert.

Q198

How do you recover from a corrupted Deduplication Database (DDB) in Commvault?

Answer:

DDB corruption recovery: (1) Commvault DDB verification: CommCell Console □ MediaAgent □ DDB □ Verify; (2) Minor corruption: Commvault can self-repair using DDB verification and repair process; (3) DDB backup: if DDB backup exists □ restore DDB from last good backup; (4) DDB resync: if no backup, Commvault can rebuild DDB from source data (time-consuming but possible); (5) DDB move: if storage failure, move DDB to healthy storage; (6) Full reconstruction: worst case – reconstruct DDB from all backup data (extremely time-consuming for large environments); (7) Prevention: always configure DDB backup, use SSD for DDB storage.

Example:

DDB corruption after NVMe drive failure on MediaAgent. Situation: DDB backup exists from last night. Recovery steps: (1) Replace failed NVMe with new NVMe. (2) Restore DDB from backup (45-minute restore for 500GB DDB). (3) Run DDB verification – 3 minor inconsistencies found. (4) DDB repair run – inconsistencies resolved. (5) Backup jobs resumed. (6) DDB resync for last 18 hours (since DDB backup): Commvault reconciles 18 hours of backup data with restored DDB (4-hour process). Total recovery time: 6 hours. All backup data intact – no data loss.

Q199

What does 'Backup completed with errors' mean in Commvault?

Answer:

'Completed with errors' (also 'Completed with one or more errors'): job ran but some items failed while others succeeded. Not a complete failure. Common causes: (1) File in use – file locked by application, couldn't be backed up; (2) Access denied – permissions issue on specific files/directories; (3) File not found – file deleted between scan and backup; (4) VSS partial failure – some VSS components failed; (5) Network timeout – specific file transfer timed out but others succeeded. Investigation: View Events on the job □ look for specific items that failed. Determine if the failed items are critical (important files missed = treat as failure).

Example:

SQL backup 'Completed with errors': View Events □ 'Cannot back up file: C:\Windows\Temp\~df4567.tmp – file in use'. Assessment: temp file – not critical. This is expected and acceptable. Add '*.tmp' to subclient exclusion list – no more errors. Second scenario: backup 'Completed with errors' □ Events show 'Finance_DB backup failed – VSS writer timeout'. This IS critical – entire database missed. Action: investigate VSS writer, escalate. Distinguish between ignorable file-in-use errors and critical application backup failures.

Q200

How do you troubleshoot slow backup performance in Commvault?

Answer:

Slow backup performance diagnostic: (1) Job history: compare throughput (MB/s) to previous successful jobs; (2) Network: check bandwidth utilisation on backup network (saturation?); (3) Storage I/O: check disk library disk latency (should be <10ms for disk library NAS); (4) Client CPU/IO: monitor source client during backup (is it the bottleneck?); (5) MediaAgent CPU/RAM: high CPU on dedup or MediaAgent process; (6) DDB performance: DDB on slow HDD instead of SSD dramatically reduces throughput; (7) Stream count: too many or too few streams; (8) Network throttle: check if throttle policy is limiting speed; (9) Antivirus: AV scanning backup traffic or MediaAgent data paths.

Example:

Backup throughput dropped from 500MB/s to 50MB/s overnight. Investigation: (1) Network: backup VLAN utilisation 15% – not saturated. (2) Storage: disk library latency 85ms (was 5ms). (3) NAS performance: NetApp ONTAP workload analysis – NAS CPU 98% due to antivirus scanning enabled on backup volume overnight. Fix: exclude backup volume from AV scanning on NAS. After fix: latency 4ms, throughput 520MB/s. Lesson: AV scan exclusions for backup volumes are critical.

Q201

What causes a Commvault job to remain in 'Queued' state indefinitely?

Answer:

Job stuck in Queued state causes: (1) MediaAgent offline – no MediaAgent available to process job; (2) Stream limit reached – all streams in use, job waiting for free stream; (3) Drive limit – all tape drives in use; (4) Disk library full – no space to write data; (5) Activity control disabled – backups globally disabled; (6) Schedule conflict – jobs queuing faster than completing; (7) Client network issue – CommServe cannot communicate with client to start job; (8) Resource manager throttle – at stream limit; (9) Job priority – low-priority job waiting for high-priority jobs to complete.

Example:

40 jobs stuck in Queue for 2 hours. Investigation: (1) MediaAgents: both MediaAgents online. (2) Stream limits: MediaAgent-01 at 20/20 streams, MA-02 at 20/20. Root cause: large file server backup started at 10 PM and is consuming 15 streams on MA-01 (plus 5 more from other large jobs). Other 40 jobs queuing waiting for streams. Fix: increase MA-01 stream limit from 20 to 30 (host has capacity). Alternative: add third MediaAgent. Jobs begin processing immediately after stream limit increase.

Q202

How do you recover when CommServe SQL database is corrupt?

Answer:

CommServe SQL database corruption recovery: (1) Assess severity: SQL cannot start = severe; SQL running but CommVault reports DB errors = moderate; (2) Restore from CommServe DR backup: stop Commvault services □ restore SQL DB from most recent DR export □ start services □ verify CommServe functional; (3) SQL repair: if minor corruption, SQL DBCC CHECKDB with REPAIR options; (4) Partial corruption: restore specific tables from DR backup; (5) No DR backup: Commvault Support may assist with partial recovery from SQL logs; (6) After recovery: run full CommServe DB validation; (7) Prevention: always configure CommServe DR with multiple export locations.

Example:

CommServe DB corruption after unexpected power loss (UPS failure). CommVault services fail to start (SQL database errors). Recovery: (1) Restore SQL DB from last night's CommServe DR export (from MediaAgent UNC share). (2) Restore takes 20 minutes (12GB CommServe DB). (3) Start Commvault services – CommServe operational. (4) Job history: some jobs from last 8 hours missing (between last DR export and failure). (5) These 8 hours' worth of backup data exists on disk library but isn't catalogued. (6) Commvault Support assists with re-cataloguing. Lesson: increase CommServe DR frequency from nightly to every 4 hours.

Q203

What are common VMware backup failures in Commvault and how to resolve them?

Answer:

Common VMware backup failures: (1) CBT reset required: 'CBT error' – reset CBT on VM (right-click VM □ Reset Changed Block Tracking); (2) Snapshot removal failure: orphaned snapshot – manually remove in vCenter; (3) vCenter API error: vCenter unavailable, credentials expired, API timeout; (4) VSA proxy issue: proxy VM CPU/memory exhausted, proxy network issue; (5) Quiesce failure: VMware Tools not installed/running – switch to crash-consistent; (6) Datastore lock: another operation holds lock – wait or force-remove lock; (7) VDDK error: VMware VDDK library version mismatch – update VDDK.

Example:

VMware backup failure: 'An error occurred while taking a snapshot: The operation is not allowed in the current state.' Root cause: VM already had 3 stale snapshots from a previous failed backup (CommVault failed to delete them). vCenter: VM □ Manage Snapshots □ delete all 3 orphaned snapshots. Commvault retry: backup completes. Prevention: Commvault snapshot management alerts enabled. Orphaned snapshot alert fires if snapshot exists >4 hours after backup job completion. Automated cleanup script checks nightly.

Q204

What is the Commvault log viewer and how do you use it?

Answer:

Commvault Log Viewer (CvLogDumper or Log Manager): (1) Access via CommCell Console □ Diagnostics □ Log Files □ select component; (2) View and filter logs from CommServe, MediaAgents, and clients; (3) Filter by: time range, severity, keyword, process; (4) Real-time tail: monitor logs as events happen; (5) Download logs for offline analysis; (6) Automated collection: CommCell Console □ Diagnostics □ Collect Log Files (bundles all logs for support cases); (7) Remote log collection: collect logs from clients without direct SSH/RDP access; (8) Log levels: adjustable from informational to debug (higher = more detail, more volume).

Example:

Troubleshooting intermittent job failure: client side ClMgrS.log collected. Filter: 'Error' severity, last 24 hours. Found: 'CEPH: Timeout connecting to mount path \\storage01\backup – 30 seconds'. Pattern: failures only between 10 PM and 11 PM. Cross-reference: Storage01 runs nightly backup tasks 10–11 PM consuming NAS I/O. Solution: stagger client backup start time to 11:30 PM (after storage01 backup window). Log analysis identified root cause in 20 minutes.

Q205

How do you diagnose and fix high MediaAgent CPU usage?

Answer:

High MediaAgent CPU diagnostic: (1) Identify process consuming CPU: Task Manager □ sort by CPU; (2) Common high-CPU processes: MediaAgent (data processing), DedupEngine (DDB operations), Scanner (file scanning on large jobs); (3) Root causes: (a) Too many concurrent streams – stream limit too high for MA hardware; (b) DDB on slow storage – DDB I/O waits cause CPU stall; (c) No DDB partitioning – large DDB single partition; (d) Encryption enabled – CPU overhead for AES operations; (e) Network 100% – CPU processing retransmits; (4) Resolution: reduce streams, move DDB to SSD, enable hardware encryption if available, add DDB partitions.

Example:

MediaAgent CPU 95% sustained. Profiling: DedupEngine.exe consuming 65% CPU. Investigation: DDB located on spinning HDD (7200 RPM) with 45ms average I/O latency. DDB size: 800GB – very I/O intensive. Fix: migrate DDB to NVMe SSD partition (new NVMe installed). After migration: DedupEngine CPU drops to 12%. MediaAgent overall CPU: 25%. Backup throughput improves from 180MB/s to 650MB/s. Cost of NVMe: \$200. Performance improvement: 3.5× throughput. ROI immediate.

Q206

What happens when a Commvault license expires?

Answer:

License expiry behavior: (1) 30-day warning: Commvault alerts 30 days before expiry; (2) Expiry day: backup jobs continue running (grace period); (3) Grace period (30 days typically): all functions continue; (4) Post-grace period: new backup jobs cannot start; existing backup data still accessible for restore; (5) Restore: restores continue even after license expiry; (6) Management: CommCell Console accessible; (7) Renewal: purchase new license, apply via CommCell Console □ License □ Update; (8) Emergency: Commvault Support can provide temporary extension license. Monitoring: License report shows expiry date prominently.

Example:

License expiry scenario: Friday 5 PM – license expires. All jobs running complete. Saturday morning: no new backup jobs start (license expired). On-call admin notified via alert. Emergency: call Commvault Support □ 48-hour temporary extension license issued (email). Apply extension: CommCell Console □ License □ Update □ apply emergency license. Monday morning: procurement purchases annual renewal. New license applied. All backups running by Monday 9 AM. 24-hour gap in backups – risk documented, management informed.

Q207

How do you resolve Commvault auxiliary copy failures?

Answer:

Auxiliary copy failure troubleshooting: (1) Job events: view auxiliary copy job events for specific error; (2) Source data: verify source copy is healthy (primary backup data readable); (3) Target media/library: verify target library and MediaAgent are online; (4) Network: check connectivity between MediaAgents (source and target); (5) Stream conflict: auxiliary copy may be queued behind backups (resource contention); (6) Data verification failure: if DV runs before auxiliary copy, bad blocks prevent copy; (7) Throttle: if WAN throttle is too low, auxiliary copy falls behind □ increase throttle or run manually; (8) Restart: kill failed auxiliary copy job and restart manually.

Example:

Nightly auxiliary copy (disk to tape) failing. Events: 'Failed to mount drive – no drives available'. Tape library: both LTO drives show 'busy'. Investigation: tape library had a firmware update earlier that put drives in maintenance mode. Fix: complete drive maintenance mode via tape library management interface. Drives come online. Auxiliary copy restart: success. Tape library firmware update procedure updated: must verify drives available after firmware update before Commvault jobs resume.

Q208

What is Commvault's self-healing capability?

Answer:

Commvault self-healing features: (1) Automatic job retry – failed jobs retry automatically (configurable attempts and delay); (2) Service watchdog – Commvault watchdog process monitors and restarts failed services; (3) MediaAgent failover – if one MA fails, jobs reassigned to available MA; (4) DDB self-repair – DDB verification and auto-repair for minor inconsistencies; (5) CBT auto-reset – automatically resets CBT for VMware VMs when CBT errors detected; (6) Orphaned job cleanup – CommServe detects stuck jobs at startup and marks them complete; (7) Mount path failover – if one NAS mount path fails, jobs fail over to alternate mount path.

Example:

MediaAgent MA-01 crashes at 2 AM (hardware fault). Self-healing: (1) CommServe detects MA-01 offline (5-minute heartbeat timeout). (2) All 15 running jobs on MA-01 automatically reassigned to MA-02. (3) Jobs resume on MA-02 (slight delay while MA-02 gets job instructions). (4) Watchdog on MA-01 server attempts service restart (fails – hardware issue). (5) Alert fires to on-call. (6) Morning: MA-01 hardware replaced, MA restored. (7) 12 jobs completed on MA-02. 3 required restart (were mid-stream at failover). Net result: minimal backup impact despite MA failure.

Q209

How do you troubleshoot Exchange iDA backup failures in Commvault?

Answer:

Exchange iDA backup failure troubleshooting: (1) VSS check: vssadmin list writers □ verify Exchange Writer 'Stable, No error'; (2) Exchange services: all Exchange services running (Exchange Management Service, System Attendant); (3) Commvault Exchange iDA service: verify CvFWD, CVD services on Exchange server; (4) DAG backup: for DAG environments, verify passive copy node is accessible; (5) Transaction log backup: check log backup is truncating logs (or log drive full); (6) Circular logging: if enabled, Commvault cannot take incremental backups (disable circular logging); (7) Database status: Exchange DB mounted and healthy (dismounted DB cannot be backed up).

Example:

Exchange incremental backup failing: 'Exchange Writer failed'. Exchange VSS writer check: 'Exchange Writer: Waiting for completion – 0x800423f4'. This means Exchange database has issues. Exchange shell: Get-MailboxDatabase -Status □ Database 'MailboxDB01' shows 'Mounted: False'. Exchange DB dismounted! Root cause: disk space issue caused auto-dismount. Fix: free disk space on Exchange database drive □ mount database (Mount-Database MailboxDB01). Exchange Writer returns to 'Stable'. Backup retry: success. Commvault alert: disk space alert added for Exchange DB drive.

Q210

How do you perform a test restore in Commvault?

Answer:

Test restore best practice: (1) Select a non-critical client and backup version; (2) Restore to alternate location (never restore over production to test); (3) Restore a cross-section: a file, a database, a VM — covers different iDA types; (4) Verify restored data: open files, query database, boot VM; (5) Measure RTO: time the restore operation; (6) Document results: success/failure, time taken, issues encountered; (7) Schedule regular tests (monthly for critical clients); (8) Application-level validation: don't just check files restored — verify application functionality. Test restores validate both backup integrity and team readiness.

Example:

Monthly test restore schedule: Week 1: file server restore test (restore 100 random files to staging server — verify file integrity via checksum comparison). Week 2: SQL restore test (restore Finance DB to SQL-TEST-01, run integrity checks). Week 3: VM restore test (restore HR-Server VM to isolated VLAN, boot, verify services). Week 4: Exchange granular restore (restore 50 random emails to test mailbox, verify content). All tests documented with pass/fail, duration. Q3 result: all tests passed. Average RTO: file 8 min, SQL 22 min, VM 35 min, Exchange granular 6 min.

Q211

What is Commvault's job priority and how does it affect scheduling?

Answer:

Commvault job priority settings: (1) High — job gets resources ahead of medium and low priority jobs; (2) Medium — default for most jobs; (3) Low — deferred when resources needed for higher priority; (4) Configuration: subclient properties □ General □ Backup Priority; (5) Use cases: Tier-1 critical systems (High), standard servers (Medium), dev/test environments (Low); (6) Priority does not guarantee immediate execution — just resource preference when queued; (7) Real-time priority change: Job Controller □ right-click running or queued job □ Change Priority; (8) Resource pools: combine priority with resource pools for granular control.

Example:

Priority conflict scenario: 200 jobs queued at 10 PM (all scheduled at same time). MediaAgent handles 20 concurrent. Without priority: FIFO — critical systems may wait hours. With priority: Tier-1 systems (15 jobs, High) start immediately. Tier-2 (85 jobs, Medium) start as Tier-1 complete. Tier-3 dev (100 jobs, Low) start as Medium complete. All Tier-1 complete by midnight (within RTO SLA). Tier-2 complete by 4 AM. Tier-3 complete by 8 AM. SLA compliance: 100% for Tier-1.

Q212

What causes Commvault deduplication ratios to decline?

Answer:

Declining dedup ratios causes: (1) New data types: encrypted data, already-compressed data (ZIP, JPG, video) — deduplication ineffective; (2) Changed backup content: new applications with unique data; (3) DDB fragmentation: DDB needs rebalancing over time; (4) Increased change rate: higher data change rate = more unique blocks; (5) Source data encryption: encrypted source data cannot be deduplicated; (6) New clients with unique data added; (7) DDB corruption: if DDB is rebuilt, ratio temporarily drops as it re-learns blocks; (8) Backup of already-deduplicated storage (VDI datastores with thin-provisioned VMs). Resolution: audit backup content for new data types, review DDB health.

Example:

Dedup ratio dropped from 9:1 to 4:1 over 2 months. Investigation: new workload added — video surveillance archive files (MP4). These are pre-compressed — dedup is ineffective. Also: new application using AES-256 encrypted databases (unique blocks per backup). Fix: exclude video surveillance files from dedup-enabled backup (move to separate non-dedup storage policy — cheaper storage appropriate for video). After exclusion: ratio returns to 8.5:1 for remaining data. Cost optimisation achieved.

Q213

How do you troubleshoot a backup job stuck in 'Waiting' state?

Answer:

A job in Waiting state means it is queued but cannot start due to resource constraints. Check: (1) MediaAgent availability and stream limits, (2) library drive/mount path availability, (3) network connectivity between client and MediaAgent, (4) other jobs consuming all available streams. Use Job Details > Resource Allocation to see which resource is limiting.

Example:

A backup job waits for 2 hours. Checking Resource Allocation shows 'No streams available on MediaAgent-01'. The disk library is configured for 10 max streams but 10 concurrent jobs are already running. Increasing the stream limit or waiting for running jobs to complete resolves the issue.

Q214

What steps do you take when a restore job fails with 'Media Not Available' error?

Answer:

Media Not Available indicates the backup media required for restore is offline, exported, or unavailable. Steps: (1) Identify which media (tape or library path) is needed using the restore job details, (2) Check media status in the Library Manager, (3) If tape – bring it back to the library and run an inventory, (4) If disk – verify mount path accessibility, (5) If cloud – verify connectivity and credentials.

Example:

A restore fails with 'Media Not Available'. The Library Manager shows the required tape is marked 'Exported' (sent offsite). The admin contacts the offsite vault, retrieves the tape, loads it in the library, runs inventory, and the restore succeeds on the second attempt.

Q215

How do you investigate high backup job duration – backup taking longer than expected?

Answer:

Investigate by: (1) Checking throughput in Job Details – compare MB/s to baseline, (2) Reviewing Phase timings (scan, backup, verify) to identify the slow phase, (3) Checking network bandwidth between client and MediaAgent, (4) Reviewing deduplication hit ratio – low hit ratio increases data written, (5) Checking client resource usage (CPU, disk I/O) during backup window.

Example:

A VM backup that normally takes 1 hour now takes 4 hours. Phase analysis shows the backup phase is slow. Network monitoring reveals the 1Gbps link between the ESXi host and MediaAgent is saturated by concurrent backups. Staggering backup schedules reduces contention and restores normal duration.

Q216

What is the CommServe Event Viewer and how is it used for troubleshooting?

Answer:

The Event Viewer in CommCell Console displays system-wide events generated by CommServe, MediaAgents, clients, and agents. Events are categorized by severity (Information, Warning, Critical) and include event codes, descriptions, and timestamps. It is the first place to look for root cause when jobs fail unexpectedly.

Example:

After several backup failures, the admin opens Event Viewer and finds Critical event 19:261 – 'MediaAgent disk space below threshold'. The disk library's drive is 95% full. Pruning aged jobs and expanding the library resolves all subsequent backup failures.

Security & Compliance

Q217

What is Commvault's approach to Zero Trust security?

Answer:

Commvault Zero Trust alignment: (1) Least-privilege access: RBAC scoped to minimum required permissions; (2) MFA: multi-factor authentication for admin access (SAML/SSO with MFA); (3) Network segmentation: backup infrastructure on isolated VLAN; (4) Verify explicitly: CommServe validates every client request (certificate-based); (5) Assume breach: immutable copies assume production compromised – backup copy independent of production credentials; (6) Air gap: ultimate Zero Trust for backup – physically/logically isolated copy; (7) Privileged access management: PAM integration (CyberArk) for service account credentials; (8) Continuous monitoring: anomaly detection, audit log SIEM integration.

Example:

Zero Trust implementation for backup: (1) MFA: all backup admin accounts require Azure AD MFA for Command Center login. (2) Separate credentials: cloud backup storage uses IAM role with no production system access. (3) Network: CommServe on backup VLAN – production VMs cannot reach CommServe directly. (4) Immutable: S3 Object Lock on all cloud copies – cannot be deleted even by compromised admin account. (5) PAM: CyberArk manages CommServe SQL sa password, rotated every 30 days. (6) Anomaly: Commvault alerts if backup change rate exceeds 200% baseline. Result: ransomware audit: backup infrastructure hardened against all attack vectors.

Q218

How does Commvault handle data sovereignty and residency requirements?

Answer:

Data sovereignty with Commvault: (1) Geographic restriction: configure backup storage in specific regions (e.g., EU-only for GDPR); (2) Cloud region selection: AWS EU-West, Azure West Europe, GCP europe-west1; (3) Separate storage policies: EU clients □ EU storage policy □ EU cloud library; non-EU clients □ US storage policy □ US cloud library; (4) MediaAgent placement: deploy MediaAgents in required data regions; (5) Data residency report: confirm all EU citizen data backed up to EU storage only; (6) Cross-border transfer control: prevent auxiliary copy from sending EU data to non-EU storage; (7) Data classification: tag data by sovereignty requirement using Commvault Activate.

Example:

GDPR data residency: EU subsidiary data must stay in EU. Configuration: (1) MediaAgent deployed in Frankfurt (EU-West). (2) Azure Blob storage: West Europe region only. (3) EU Storage Policy: primary copy □ Frankfurt MediaAgent □ Azure West Europe. (4) Non-EU clients use US storage policy (US MediaAgent, US-East S3). (5) Enforce: RBAC prevents EU storage policy from being assigned to non-EU clients. (6) Audit: quarterly report verifying EU client data stays in West Europe Azure region. No EU data ever leaves EU – GDPR Article 44 compliant.

Q219

What is Commvault Activate's sensitive data discovery?

Answer:

Commvault Activate sensitive data discovery: (1) Pattern matching: uses regex patterns to identify PII (SSNs, credit card numbers, email addresses, passport numbers, NHS numbers); (2) Entity extraction: identifies people, organisations, locations in unstructured content; (3) File type analysis: identifies sensitive content by file type (medical records, legal documents); (4) ML classification: machine learning classifies content beyond simple patterns; (5) Source coverage: scans backed-up file servers, email archives, SharePoint, databases; (6) Risk scoring: each data store gets a risk score based on: sensitivity of content, access controls, staleness; (7) Reporting: compliance dashboard for GDPR, HIPAA, PCI-DSS, CCPA.

Example:

PCI-DSS scope assessment using Activate: Financial organisation scans 2TB of backed-up file server data for cardholder data. Activate findings: (1) 4,500 Excel files containing credit card numbers across 80 file shares. (2) 12 shares with public access. (3) Risk score: Critical for 12 public shares. Action: (1) Remove public access from 12 shares. (2) Encrypt 4,500 files or move to PCI-controlled storage. (3) Reduce PCI scope: delete unnecessary card data (per PCI requirement). PCI QSA: scope reduced by 40%. Activate scans ongoing monthly.

Q220

What is Commvault's WORM storage support and why is it important?

Answer:

WORM (Write Once, Read Many) storage: data written once cannot be modified or deleted. Commvault WORM support: (1) Cloud WORM: AWS S3 Object Lock, Azure Blob Immutable Storage, GCP Bucket Lock; (2) Tape WORM: LTO-M WORM tape cartridges (hardware enforced); (3) Disk WORM: NetApp SnapLock, EMC DataDomain Retention Lock, HDS WORM; (4) Commvault compliance copy: marks copies as WORM – cannot be aged, deleted, or overwritten by Commvault or any external process; (5) Air gap + WORM: ultimate ransomware protection – air-gapped AND immutable copy. Importance: ransomware cannot encrypt/delete WORM data; regulators require immutable copies for financial and healthcare data.

Example:

Financial services WORM requirement (SEC Rule 17a-4): all trade records must be retained for 7 years on WORM storage. Commvault: (1) Monthly full backup copied to AWS S3 Object Lock (Compliance mode, 7-year retention). (2) Object Lock Compliance mode: even AWS cannot delete objects before retention period. (3) Commvault marks these copies as 'Compliance Retention – no aging'. SEC audit: 'Demonstrate 7-year immutable storage.' Evidence: S3 bucket policy + Object Lock configuration + Commvault retention report. SEC auditor: satisfied.

Q221

How does Commvault protect against insider threats?

Answer:

Insider threat protection: (1) RBAC: operators can only access their scope – cannot access other departments' data; (2) Two-person rule: some critical operations (large data export, delete) require approval from second user; (3) Audit trail: every operation logged with user, timestamp, action – tamper-evident; (4) SIEM integration: audit events forwarded to Splunk – anomalous restore patterns trigger alerts; (5) Data classification: identify which backups contain sensitive data (Activate) – enhanced monitoring on sensitive restores; (6) Immutable audit: audit logs stored in WORM storage (cannot be deleted by insider); (7) Read-only credentials: backup service accounts write-only to cloud storage (cannot read backup data).

Example:

Insider threat scenario: suspicious restore activity detected. Commvault audit trail + Splunk alert: user jdoe@company.com performed 15 restore operations of HR department files to personal laptop between 11 PM and 2 AM (outside working hours). Normal pattern: 0-2 restores/week during business hours. Alert: Splunk correlation rule fires. SOC investigates. Finding: jdoe is HR system administrator, accessed payroll data before resignation. HR team informed. RBAC review: jdoe's access revoked immediately. Audit trail admissible for HR investigation.

Q222

What is Commvault's approach to key management for encryption?

Answer:

Commvault key management options: (1) CommServe Key Management Server (KMS): built-in, keys stored in CommServe SQL DB; simple, no external dependency; risk: lose CommServe = lose keys (mitigated by CommServe DR); (2) External KMIP KMS: connect to external key server (Thales, Vormetric, HashiCorp Vault, AWS KMS) via KMIP protocol; enterprise-grade; keys separated from data; (3) Pass-phrase: simpler per-copy pass-phrase (not recommended for large environments); (4) HSM integration: Hardware Security Module via external KMIP for highest security; (5) Key rotation: external KMS allows key rotation without re-encrypting data (envelope encryption). Best practice: external KMIP for regulated industries.

Example:

Healthcare key management: PHI backup encrypted with AES-256. Keys managed by Thales CipherTrust (KMIP-compatible HSM). Integration: CommCell = KMIP = Thales. Keys never stored on CommServe. Key rotation: quarterly – Thales rotates master key, Commvault data keys re-wrapped (no re-encryption of backup data). Key access: requires MFA to Thales console. Key audit: Thales logs every key access – cross-referenced with Commvault restore audit. Key backup: Thales HSM clustered (2 nodes). HIPAA encryption requirement: satisfied with HSM-grade key management.

Q223

How does Commvault support ISO 27001 compliance?

Answer:

Commvault ISO 27001 support: (1) A.12.3 (Backup): documented backup procedures, tested recovery, retention policies – Commvault provides all; (2) A.10.1 (Cryptography): encryption at rest and in transit – Commvault AES-256 + TLS; (3) A.9 (Access Control): RBAC with least privilege – Commvault RBAC; (4) A.12.4 (Logging and Monitoring): audit trail and SIEM integration; (5) A.17 (Business Continuity): DR testing, documented RTOs/RPOs, tested recovery; (6) A.12.1 (Operational Procedures): documented backup runbooks from Commvault jobs; (7) Evidence: Commvault reports and audit trails serve as ISO audit evidence.

Example:

ISO 27001 audit preparation (Annex A.12.3): auditor requests evidence of: (1) Backup policy documented □ (Commvault storage policies + schedule). (2) Regular testing of restores □ (monthly test restore reports). (3) Backup stored offsite □ (cloud secondary copies). (4) Access controls on backup □ (RBAC report). (5) Encryption □ (AES-256 configuration screenshot). (6) Monitoring and alerting □ (SLA reports, alert configuration). ISO 27001 auditor: 'Annex A.12.3 – No findings.' Commvault evidence package prepared in 2 hours.

Q224

What is Commvault's anomaly detection and how does it identify ransomware?

Answer:

Commvault anomaly detection monitors backup patterns for unusual activity: (1) Baseline: Commvault establishes normal daily change rate per client (e.g., 3% daily change rate for a file server); (2) Anomaly threshold: alert if change rate exceeds X% of baseline (e.g., >200%); (3) Ransomware signature: encrypted files have specific file extension patterns (.locked, .encrypted, .crypted) – Commvault detects new extension proliferation; (4) Honeypot files: Commvault creates hidden decoy files – if modified, immediate alert; (5) Alert: fires before ransomware completes encryption – enabling faster response; (6) Integration: alerts forwarded to SIEM/SOAR for automated response (isolate infected VMs).

Example:

Ransomware detection in action: File server normal change rate: 2% (20GB/day out of 1TB). Tuesday 2:47 AM: Commvault anomaly alert fires. Reason: change rate jumped to 95% (950GB changed in 30 minutes). Also: 12,000 new files with .locked extension detected across all shares. Commvault alert: email to SOC + PagerDuty page. SOC response 2:52 AM: isolates file server from network. Investigation: ransomware infection began 2:40 AM. Detection: 7 minutes after infection. Clean backup: 2:30 AM backup (before infection). Recovery: from 2:30 AM backup. Data loss: 10 minutes.

Q225

What is Commvault's support for multi-factor authentication (MFA)?

Answer:

Commvault MFA support: (1) SAML 2.0 SSO: integrate with Azure AD, Okta, PingFederate, ADFS – MFA enforced at identity provider level; (2) Direct MFA: Commvault v11 SP32+ supports direct TOTP (time-based OTP) for local Commvault accounts; (3) Adaptive MFA: if identity provider supports it (Azure AD Conditional Access) – require MFA for unusual login locations; (4) Configuration: CommCell Console □ Security □ Identity Management □ Add SAML Identity Provider; (5) Scope: MFA for Command Center logins – CommCell Console logins (Java) also support SAML; (6) Bypass: emergency break-glass account without MFA (stored offline, used only for CommServe failure scenarios).

Example:

MFA implementation: Azure AD configured as SAML identity provider for Commvault Command Center. All admin logins redirected to Azure AD sign-in page. Azure AD Conditional Access policy: 'If signing into Commvault from any location: require MFA (Authenticator app)'. Result: backup admin login □ redirected to Microsoft login □ MFA challenge (push notification to phone) □ approved □ Commvault Command Center opens. Phishing attack on backup admin: attacker gets password but cannot satisfy MFA □ access denied. MFA blocked 100% of credential-only attacks.

Q226

How does Commvault support SOC 2 compliance requirements?

Answer:

Commvault SOC 2 support (Trust Service Criteria): (1) Availability (A1): backup SLA ensures data availability, DR capability; (2) Confidentiality (C1): encrypted backups, RBAC on access, WORM copies; (3) Processing Integrity (PI1): data verification confirms backup accuracy; (4) Security (CC): audit trails, MFA, network segmentation, RBAC; (5) Privacy (P): sensitive data discovery, retention management, right to erasure; (6) Evidence for Type 2 audit: 12 months of SLA reports, audit trails, test restore records, encryption configuration; (7) Change management: Commvault audit trail shows all configuration changes (CC7 - change management control).

Example:

SOC 2 Type 2 audit support: auditor requests 12 months of evidence for CC6.1 (Access Controls) and CC7.2 (Change Management).

Commvault evidence: (1) CC6.1: RBAC configuration report, user access review logs (quarterly reviews documented), MFA enabled screenshot. (2) CC7.2: Audit trail export showing all CommCell configuration changes for 12 months (who changed what, when). Total evidence: 50-page report generated in 2 hours from Commvault audit trail. Auditor: 'CC6.1 and CC7.2 – operating effectively. No exceptions noted.'

Q227

What is Commvault's privacy wipe capability?

Answer:

Privacy wipe (data subject erasure): Commvault can remove specific personal data from backup copies to support GDPR Article 17 (Right to Erasure): (1) Identify: use Activate to find all backup copies containing the specific person's data; (2) Wipe: for file/email backups – specific files or emails containing the individual's PII removed from backup copies; (3) VM backups: more complex – cannot remove individual records from a VM backup without restoring, editing, re-backing up; (4) Content Store archive: individual email items can be deleted from archive; (5) Audit: privacy wipe actions logged for regulatory evidence; (6) Limitation: granular wipe from block-level VM backups is technically complex – often handled at application level instead.

Example:

GDPR erasure request for departed employee: (1) Activate search: identify all backup copies containing employee's personal data (email, files, HR records). Found: 90 days of email backups, 3 file server backups. (2) Content Store archive: individual emails deleted (privacy wipe). (3) File server backups: specific files wiped from last 90 days. (4) VM backups: employee's application records deleted at source (production) + any new backup captures no PII. Old VM backups: noted as 'may contain historical data' – business retention applied for 30 more days then aged. (5) Audit log: privacy wipe documented for DPA review.

Q228

How does Commvault handle backup security in a DMZ environment?

Answer:

DMZ backup challenges: DMZ servers have restricted connectivity – cannot freely communicate with internal backup infrastructure. Solutions: (1) CommCell Firewall: client-initiated tunnel from DMZ server to CommServe (single outbound port 8400); (2) Internal proxy: deploy Commvault proxy in DMZ – DMZ servers communicate with proxy, proxy communicates to internal MediaAgent; (3) Restricted data path: DMZ client sends backup data through the Commvault tunnel (not a separate data port); (4) Firewall rules: minimum: DMZ → internal CommServe: TCP 8400 outbound only; (5) Data isolation: DMZ backup data stored in separate storage policy (prevent cross-contamination with internal backup data).

Example:

PCI-DSS CDE (Cardholder Data Environment) in DMZ: (1) Commvault firewall enabled. (2) CDE servers: outbound TCP 8400 to internal CommServe only. (3) Internal proxy: DMZ proxy deployed – CDE servers talk to proxy, proxy forwards to internal MA. (4) No inbound connections from internal to DMZ (PCI requirement). (5) Separate storage policy: CDE-Backup-SP with AES-256 encryption (PCI requirement). PCI QSA: 'Backup network traffic from CDE: single outbound port 8400. No inbound connections. Data encrypted. Backup data isolated.' PCI compliant backup design.

Q229

How does Commvault support GDPR compliance for data protection?

Answer:

Commvault GDPR support: (1) Data discovery: Activate module scans backups for EU citizen PII; (2) Right to erasure: identify and delete PII from backup data (privacy wipe); (3) Data residency: restrict backup data to specific geographic regions (EU-only cloud regions); (4) Access controls: RBAC limits who can access personal data in backups; (5) Audit logging: all data access and restore operations logged; (6) Retention management: automatic deletion of personal data after configured retention period; (7) Legal hold exemption: legal hold overrides retention for compliance investigations; (8) Consent tracking: integration with data governance platforms.

Example:

GDPR subject access request (SAR): Individual requests all their personal data. Commvault Activate: search backup data for this person's email, name, ID number. Finds data in 15 mailbox backups, 3 file server backups. Report generated: 'All backup copies containing PII for Data Subject ID: 45892'. GDPR team reviews. Data subject also requests erasure: Commvault privacy wipe removes specific PII items from backup data. Audit trail documents every action for DPA (Data Protection Authority) review.

Q230

What is Commvault's ransomware recovery process?

Answer:

Commvault ransomware recovery: (1) Detection: anomaly detection/SIEM identifies ransomware; (2) Isolation: infected systems isolated from network; (3) Assessment: identify ransomware timeline (when encryption started); (4) Threat scan: Commvault scans backup versions to find clean recovery point; (5) Isolated recovery environment: provision clean VMs in isolated network from backup; (6) Validation: verify recovered systems are clean (additional threat scan); (7) Production cutover: connect validated clean systems to production after confirmation; (8) Root cause analysis: determine and fix infection vector before reconnecting.

Example:

Ransomware at 2:30 AM. Detection: SIEM alert from Commvault anomaly detection (data change rate 1,200% of normal). Response timeline: 2:35 AM – incident declared, 50 VMs isolated. 3:00 AM – Commvault threat scan on 20 backup versions (hourly). Clean point identified: 1:15 AM backup (before ransomware). 3:30 AM – isolated recovery VMs provisioned from 1:15 AM backup. 4:00 AM – application validation in isolated environment. 5:00 AM – production cutover. 2.5-hour data loss (1:15 AM to 3:45 AM) vs days/weeks with ransom payment alternative.

Q231

What are Commvault's RBAC best practices?

Answer:

Commvault RBAC best practices: (1) Principle of least privilege: users get minimum permissions needed; (2) Separate roles: backup admin ≠ restore operator ≠ report viewer; (3) AD integration: manage users via AD groups (no local Commvault accounts for users); (4) Service accounts: dedicated AD service accounts for Commvault services (not user accounts); (5) Regular access reviews: quarterly review of user roles (remove departed employees immediately); (6) Scope limiting: application owners can only restore their own applications; (7) Audit trail: enable and retain audit trail for all operations; (8) No shared admin: individual admin accounts (traceable to person).

Example:

RBAC implementation for 500-client environment: Roles defined: Backup Admin (5 people: full CommCell access), Restore Operator (10 people: restore only, scoped to their business unit), DB Admin (3 people: SQL/Oracle operations only), Auditor (2 people: report viewer only). All via AD groups. Quarterly review: 2 former employees found in Backup Admin group (missed offboarding) – removed. Service accounts: svc-commvault-agent, svc-commvault-sql – dedicated AD accounts, 90-day password rotation via CyberArk.

Q232

How does Commvault protect backup data from being deleted or modified by ransomware?

Answer:

Ransomware protection of backup data: (1) Air-gapped copies: tape offsite (physical air gap), or cloud with Object Lock (logical air gap); (2) Immutable storage: AWS S3 Object Lock, Azure Blob Immutable Storage – backup data cannot be deleted even by admin during retention period; (3) Separate credentials: cloud backup storage uses credentials not accessible from production environment; (4) Network segmentation: MediaAgents not accessible from production VLANs; (5) Read-only cloud: cloud IAM policy allows only write (no delete); (6) Honeypot files: detect ransomware accessing backup storage early; (7) WORM tape: hardware-enforced write-once tapes.

Example:

Immutable backup setup: AWS S3 bucket 'company-immutable-backup' with Object Lock (Compliance mode, 30-day retention). IAM role for MediaAgent: s3:PutObject, s3:GetObject, s3:ListBucket – NO s3:DeleteObject. Even if ransomware compromises the MediaAgent server, it cannot delete S3 objects (no delete permission + Object Lock protection). Additionally: MediaAgent on isolated VLAN (no production access). Tape library: LTO-M WORM cartridges for monthly backups. Three-layer protection: VLAN isolation + S3 immutability + WORM tape.

Q233

What audit trail and reporting does Commvault provide for compliance?

Answer:

Compliance audit capabilities: (1) Backup audit trail: every backup job recorded (client, time, data amount, success/failure, operator); (2) Restore audit: every restore documented (who restored, what data, source backup, target location, timestamp); (3) Access audit: all admin operations logged (configuration changes, user additions, policy changes); (4) Legal hold tracking: all legal hold operations documented; (5) Data Aging log: all pruning/deletion operations recorded; (6) SIEM integration: events forwarded to Splunk/QRadar via syslog; (7) Immutable audit: Commvault audit log can be exported to immutable storage for tamper-proof compliance evidence.

Example:

PCI-DSS audit requirement: 'Document all access to cardholder data for last 12 months.' Commvault audit trail query: filter restore operations for PCI-scope clients. Result: 8 restore operations in 12 months – all by authorised backup operators with documented change requests. No unauthorised access. Commvault also shows: audit log not modified (integrity verified via hash). PCI QSA: satisfied. Audit trail submitted as evidence for requirement 10.2 (audit log completeness).

Q234

How does Commvault handle legal hold for eDiscovery?

Answer:

Legal hold in Commvault: (1) Mark specific backup data as legal hold: it will not be aged/deleted regardless of retention policy; (2) Apply at: client level, subclient level, or specific backup job level; (3) Legal hold applies to: backup data (backups), archive data (Content Store), and email archives; (4) Legal hold manager: tracks all active holds, owners, and scope; (5) Remove hold: only authorised users can remove (RBAC); (6) Hold audit: all hold placement and removal actions logged with user, timestamp, reason; (7) eDiscovery: search held data for specific content without restoring everything.

Example:

Legal hold for investigation: 'Executive VP John Smith – email preservation order received.' Action: Commvault legal hold applied to VP's Exchange mailbox backup and Content Store archive. Hold reason: 'Litigation Case #2024-0045'. From this point: VP's mailbox backup data not deleted even when retention period expires. Hold duration: indefinite until removed by Legal team. eDiscovery: search held archive for 'Project Alpha' – 2,300 emails found, exported to PST for attorneys. Hold active for 18 months until case resolution.

Q235

How does Commvault meet HIPAA requirements for healthcare organisations?

Answer:

HIPAA compliance with Commvault: (1) Encryption: AES-256 at rest, TLS in transit — PHI encrypted throughout backup lifecycle; (2) Access controls: RBAC limits who can access patient data in backups; (3) Audit controls: all access to PHI backup data logged (required by HIPAA §164.312(b)); (4) Backup integrity: data verification ensures PHI recoverable (HIPAA §164.310(d)(2)(iv)); (5) DR capability: documented recovery procedures (HIPAA contingency plan); (6) Business Associate Agreement: Commvault can sign BAA for cloud services; (7) Geographic restriction: PHI backups stored only in US regions (cloud); (8) Retention: HIPAA requires 6-year record retention — Commvault LTR supports this.

Example:

Hospital HIPAA audit: EHR (Epic) system backed up by Commvault. Evidence provided: (1) Encryption certificates (AES-256 at rest, TLS 1.2 in transit). (2) RBAC report: only 3 backup admins have access to Epic client. (3) Audit trail: all restores of Epic data over last 2 years documented. (4) DR test report: quarterly tests showing Epic recovery in <2 hours. (5) Retention policy: 6-year minimum via LTR. HIPAA auditor: 'Security Rule Technical Safeguards — compliant.' No findings on backup-related controls.

Q236

What is Commvault's Threat Scan and how does it integrate with antivirus?

Answer:

Threat Scan integration: (1) Commvault mounts a specific backup version on the MediaAgent; (2) Integrated AV engine (ClamAV) or external AV (symlinked via ICAP) scans mounted data; (3) Scan report: clean, suspicious, or infected files identified; (4) Recovery action: only clean backup versions recommended for recovery; (5) Quarantine: infected backup versions flagged, not used for recovery; (6) Automated scanning: configure Commvault to scan all new backups automatically; (7) Integration with SIEM: scan results forwarded to security tools; (8) Reduces risk of restoring infected backup data after ransomware attack.

Example:

Post-ransomware recovery: 8 hourly backup versions available (24 hours of backups before attack). Threat Scan run on each: Versions 1-5 (last 5 hours): infected — ransomware files detected. Version 6 (6 hours ago): CLEAN. Versions 7-8 (older): also clean. Commvault recommends recovery from Version 6. Admin confirms — 6-hour data loss accepted. Recovery from Version 6 initiated. Secondary validation: recovered VMs run CrowdStrike scan before connecting to production — no threats detected. Confidence in clean recovery: high.

Q237

How does Commvault secure the backup infrastructure?

Answer:

Commvault infrastructure security hardening: (1) OS hardening: CIS benchmark on CommServe, MediaAgent, and proxy servers; (2) Network segmentation: backup components on dedicated VLAN (not accessible from end-user VLANs); (3) Encrypted communications: all inter-component traffic encrypted (SSL/TLS); (4) Certificate management: replace self-signed certificates with CA-signed; (5) Privileged access: MFA for admin accounts (SAML/SSO integration); (6) PAM integration: CyberArk manages Commvault service account passwords; (7) Vulnerability management: patch Commvault and OS regularly; (8) Immutable audit: audit logs protected from tampering.

Example:

Commvault security assessment findings and remediation: Finding 1: self-signed certificates on CommServe web UI. Fix: CA-signed certificate installed. Finding 2: admin accounts no MFA. Fix: Azure AD SSO with MFA configured. Finding 3: CommServe on shared VLAN with production servers. Fix: CommServe moved to dedicated backup VLAN. Finding 4: Commvault service account password never rotated (3 years). Fix: CyberArk vault, 90-day rotation. Post-remediation: penetration test passed. Security posture improved from Medium to Low risk.

Q238

What is Commvault's Privacy Risk Analysis and how is it used?

Answer:

Privacy Risk Analysis (part of Commvault Activate): analyses backup and archive data to identify privacy risks: (1) PII discovery: locates personal data (names, SSNs, credit cards, emails) across all backed-up data; (2) Risk scoring: assigns risk score to each data store based on sensitivity and access controls; (3) Data owner identification: maps data to responsible owners; (4) Stale data identification: old sensitive data that should be deleted; (5) Regulatory mapping: maps discovered data to GDPR, CCPA, HIPAA requirements; (6) Reporting: dashboard for compliance officers; (7) Remediation: identifies data to delete, restrict, or encrypt.

Example:

CCPA compliance program: Activate Privacy Risk Analysis run on 200 file servers. Findings: 1.2 million files with California resident PII (names, addresses, SSNs). High-risk: 15,000 files with SSNs on file shares accessible by 500+ employees. Risk score: Critical. Remediation: restrict access to high-risk shares (15 people need access, not 500), encrypt files with PII. Track: 3-month remediation plan, progress tracked in Activate dashboard. CCPA compliance improved from 45% to 92% measured posture.

Q239

How does Commvault integrate with SIEM tools for security monitoring?

Answer:

SIEM integration options: (1) Syslog forwarding: CommServe forwards events via syslog (UDP/TCP 514) to SIEM (Splunk, QRadar, ArcSight); (2) SNMP traps: backup events sent as SNMP traps to NMS/SIEM; (3) REST API: SIEM pulls events via Commvault REST API; (4) CEF/LEEF format: Common Event Format for Splunk, Log Enhanced Event Format for QRadar; (5) Commvault Splunk App: pre-built Splunk dashboards for Commvault data; (6) Event types forwarded: job failures, SLA misses, login events, configuration changes, media errors; (7) Anomaly detection: Commvault anomaly alerts forwarded to SIEM for correlation.

Example:

Splunk integration: CommServe syslog output configured to Splunk Heavy Forwarder. Splunk Commvault App installed: dashboards created for – backup success rate trend, SLA compliance history, failed job analysis, media alerts. Security use case: Splunk correlation rule – 'If Commvault logs 5+ failed login attempts to CommServe in 1 hour then create Splunk alert then SOC investigates potential brute-force attack.' First month: 3 alerts – 2 false positives (admin mistyped password), 1 real – retired contractor account attempted access. Account removed.

Q240

What is Commvault's Air Gap Protect solution?

Answer:

Air Gap Protect is Commvault's solution for creating an isolated, immutable copy of backup data separated from production. Features: (1) Isolated network: backup data vault not connected to production or backup network (true air gap); (2) Automated transfer: data moves to air-gap storage via controlled, one-way data flow; (3) Immutability: WORM storage in the vault – data cannot be modified; (4) Ransomware resilience: ransomware cannot encrypt or delete air-gapped copies; (5) Recovery testing: test recovery from air-gap storage without connecting vault to production; (6) Appliance-based: Commvault Air Gap Protect appliance or certified 3rd-party storage.

Example:

Financial institution air gap: Air Gap Protect appliance in separate secured room (no network connection). Weekly data transfer via automated, unidirectional data diode protocol. Air gap vault: 3 years of weekly full backups on WORM disks. Ransomware attack hits production and backup infrastructure (both on same network). Recovery: air gap appliance brought online in isolated environment then systems recovered from 7-day-old air-gapped copy. Cost: 7 days of re-entry vs paying \$5M ransom. Air gap was the last line of defense – and it worked.

Q241

What is Data Encryption in Commvault and at what levels can it be applied?

Answer:

Commvault supports AES-256 encryption at three levels: (1) Client-side — data is encrypted before leaving the client, (2) MediaAgent-side — data is encrypted as it is written to the storage library, (3) Network — data in transit between client and MediaAgent is encrypted. Encryption keys can be managed by Commvault or integrated with external key managers.

Example:

For a healthcare client, client-side encryption is enabled on all subclients to ensure HIPAA compliance. Data leaves the hospital server already encrypted, travels encrypted over the network, and is stored encrypted on the disk library — never exposed in plaintext outside the source server.

Q242

What is the Commvault Security IQ and what compliance features does it provide?

Answer:

Security IQ (formerly Risk Analysis) is a Commvault module that scans backup data for sensitive content — PII, PCI data, credentials — and provides risk reports, compliance dashboards, and data classification. It identifies over-exposure of sensitive files without requiring a separate DLP tool.

Example:

Security IQ scans a file server backup and identifies 15,000 files containing Social Security Numbers in unprotected shared folders. The compliance report is used to remediate access controls before a PCI audit.

Q243

How does Commvault support GDPR compliance?

Answer:

Commvault supports GDPR via: (1) Data discovery — locating personal data across backup sets, (2) Right to erasure — identifying and deleting specific user data from backups, (3) Audit trails — complete job and access logs, (4) Encryption — protecting personal data at rest and in transit, (5) Retention management — enforcing data retention and deletion policies aligned with GDPR timelines.

Example:

A user submits a GDPR erasure request. The admin uses Commvault's eDiscovery to locate the user's data across all backup copies, marks it for deletion, and generates an audit report documenting the erasure — satisfying the 30-day GDPR compliance requirement.

Q244

What is Ransomware Protection in Commvault and how does it work?

Answer:

Commvault's ransomware protection includes: (1) Immutable backup copies using WORM (Write Once Read Many) storage or object lock on cloud, (2) Anomaly detection that flags unusual backup behaviour (sudden size changes, high encryption rates), (3) Air-gap copies to isolated media, (4) Honeypot file monitoring. Together, these ensure clean recovery points exist even if production data is encrypted by ransomware.

Example:

Commvault detects that a backup job's data size increased by 800% overnight and flags it as a potential ransomware anomaly. The admin investigates, confirms a ransomware attack, and restores from the last clean backup point — which is protected on immutable object storage and cannot be deleted by the attacker.

Q245

What is the purpose of Audit Trail in Commvault?

Answer:

The Audit Trail records all administrative actions taken in the CommCell — logins, policy changes, deletions, restores, and configuration modifications — with timestamps and user identity. It provides forensic traceability for compliance audits and helps detect unauthorized activity.

Example:

After a storage policy is unexpectedly modified, the Audit Trail shows that user 'jsmith' changed the retention period from 30 days to 7 days at 2:14 AM — outside business hours. This triggers an internal security investigation.

Q246

How does Commvault support multi-tenancy and data isolation in MSP environments?

Answer:

Commvault's Managed Service Provider (MSP) mode uses CommCell Console companies and tenant-level RBAC to provide complete data and administrative isolation between tenants. Each tenant sees only their own clients, jobs, and data. Billing reports are generated per tenant. Data encryption keys can be tenant-specific.

Example:

An MSP manages Commvault for 20 different customers on a shared CommCell. Customer A's admin logs in and sees only their 50 clients with no visibility into Customer B's data. Separate MediaAgents and storage libraries can be dedicated per customer for physical isolation.

CCA Exam Preparation

Q247

What is the Commvault data path and how does data flow during backup?

Answer:

Commvault data flow during backup: (1) CommServe reads schedule □ starts job at scheduled time; (2) CommServe instructs client iDA to start backup; (3) Client iDA reads source data (files, DB, VM); (4) Data sent from client to MediaAgent (via network); (5) MediaAgent receives data, performs deduplication (compares blocks against DDB); (6) Unique blocks: compressed, optionally encrypted, written to disk library; (7) Duplicate blocks: only pointer stored in DDB; (8) CommServe records job results, media location, data size in SQL DB; (9) Job complete: CommServe sends completion status. Note: CommServe NEVER handles backup data — only metadata.

Example:

Exam data flow question: 'During backup, which component writes data to the disk library?' Answer: MediaAgent. Data flow: Client (iDA reads data) □ network □ MediaAgent (dedup, compress, encrypt) □ Disk Library. CommServe does NOT handle backup data — it handles only job metadata (instructions to start, records of completion). Trick: if asked 'CommServe writes to disk library' □ WRONG. Always: MediaAgent writes data. CommServe writes metadata to SQL DB.

Q248

What is the difference between a Commvault client and an agent?

Answer:

Client vs Agent distinction: (1) Client = a logical entity in CommCell representing a physical/virtual machine, application, or cloud resource. It's the container that holds agents; (2) Agent (iDataAgent/iDA) = the specific backup software for a particular workload installed on or associated with a client. A single client can have multiple agents. Example: 'SQL-Server-01' is the CLIENT. It has: File System iDA + SQL Server iDA = 2 AGENTS on one client. CCA exam: client is the registration (the server in CommCell); agents determine WHAT can be backed up on that client.

Example:

Exam: 'A Windows server runs SQL Server and IIS. What clients and agents are needed in Commvault?' Answer: ONE client: Windows-Web-SQL-Server. THREE agents on that client: (1) File System iDA (backs up IIS web files, OS state). (2) SQL Server iDA (backs up SQL databases with VSS, log backups). (3) System State agent (backs up Windows registry, boot files). Three protection capabilities from one client registration. Adding Exchange to the same server: just add Exchange iDA to the existing client (no new client needed).

Q249

What is Commvault's changed block tracking and which hypervisors support it?

Answer:

Changed Block Tracking (CBT) is a hypervisor feature tracking which disk blocks changed since last backup — enables efficient incremental backups. Commvault support: (1) VMware: CBT (VADP Changed Block Tracking) — natively supported since vSphere 4.0; (2) Microsoft Hyper-V: RCT (Resilient Change Tracking) — Hyper-V 2016+; (3) Nutanix AHV: CRT (Changed Region Tracking) — Nutanix native; (4) AWS: EBS changed block tracking (AWS-native incremental snapshots); (5) Azure: Azure incremental snapshots. Without CBT/CRT: each incremental reads entire VM disk (inefficient). With CBT: reads only changed blocks (10× faster for typical change rates). CCA: know which abbreviation belongs to which hypervisor.

Example:

Exam: 'What feature does Commvault use to identify changed blocks in VMware VMs for efficient incremental backup?' Answer: CBT — Changed Block Tracking (VMware terminology). Commvault accesses CBT via VADP API. If the question mentions Hyper-V: answer is RCT. If Nutanix: CRT. Each hypervisor has its own mechanism but the result is the same: only changed blocks sent during incremental backup. CBT disabled = full disk read every backup (slow). CBT enabled = only delta (fast).

Q250

What is a Commvault storage pool and how is it different from a library?

Answer:

Storage Pool (used in plans-based management): a logical grouping of storage resources (disk libraries, cloud libraries) that can be assigned to plans as a destination. Library: a specific storage resource (physical disk, NAS mount, tape, cloud bucket). Relationship: Storage Pool = abstract collection of libraries. Assign plans to storage pool → plan picks best available library from pool. Benefits: (1) Flexibility: plan doesn't need to know which specific library — just the pool; (2) Load balancing: multiple libraries in pool for distribution; (3) Failover: if one library full, pool uses next. In traditional Commvault: storage policy directly references a library. Plans reference storage pools.

Example:

Exam: 'In Commvault Plans, what is assigned as the backup destination — a specific library or a storage pool?' Answer: Storage pool. Admin creates Storage Pool 'Primary-Storage' containing: 3 disk libraries (NetApp-Vol1, NetApp-Vol2, AWS-S3). Server Plan references 'Primary-Storage'. Commvault picks the best available library from the pool for each job. If NetApp-Vol1 is full: jobs go to NetApp-Vol2 automatically. This pool abstraction prevents 'library full' job failures — always has an available target within the pool.

Q251

What are Commvault's SLA Protection Levels and how are they assigned?

Answer:

Commvault SLA Protection Levels define RPO expectations for clients: (1) Gold (Tier 1): most critical — RPO <1 hour, continuous protection (CDR) or hourly snapshot; (2) Silver (Tier 2): standard business — RPO 4-8 hours, 4-hourly backup; (3) Bronze (Tier 3): non-critical — RPO 24 hours, daily backup; (4) Unprotected: no backup — monitor and alert for coverage gaps; (5) Levels defined by: business value, data volatility, recovery requirements; (6) Assignment: assign protection level when onboarding client; (7) Reporting: SLA dashboard shows distribution — % clients at each level, compliance rate per level. CCA exam: match protection level to appropriate backup strategy.

Example:

Exam scenario: 'Hospital EHR system (RPO requirement: 1 hour). File server for non-critical documents (RPO: 24 hours). Dev environment (RPO: weekly). Assign protection levels.' Answer: EHR → Gold (hourly IntelliSnap/CDR). File server → Silver (4-hourly incremental). Dev → Bronze (daily or even weekly). Or unprotected with documented acceptance. SLA levels align with business criticality and determine backup frequency, retention, and storage tier.

Q252

What are the Commvault backup job phases and what happens in each?

Answer:

Commvault backup job phases: (1) Initialization: CommServe allocates resources, contacts client, prepares job parameters; (2) Scan phase: client scans file system or database to identify items for backup (for incremental: compare with last backup to identify changes); (3) Backup phase: data read from client, sent to MediaAgent, deduplicated, stored; (4) Commit phase: job metadata committed to CommServe SQL DB (media location, items backed up, success); (5) Finalization: CommServe updates job status, sends alerts if needed. For VM backup: snapshot creation before Scan phase; snapshot deletion after Backup phase. Understanding phases helps identify where a slow or failed job is stuck.

Example:

Exam: 'A backup job is stuck in Scan phase for 6 hours. What is the most likely cause?' Answer: The file system scan is taking too long – common for very large file servers with millions of files. Solutions: (1) Enable Accelerated File System Scanner (pre-scan). (2) Add filters to exclude temp/log files (reducing scan scope). (3) Enable USN journal tracking (only scan recently changed files). (4) Check if file system has errors causing scan to hang. A job stuck in Scan hasn't sent any data yet – disk library and MediaAgent are not the bottleneck.

Q253

What is Commvault's extended retention and how is it configured?

Answer:

Extended retention keeps specific backup jobs beyond the standard retention period for compliance or business requirements. Configuration: (1) Storage policy copy □ Properties □ Extended Retention; (2) Define rules: keep every weekly full for 1 month, every monthly full for 1 year, every yearly full for 7 years; (3) Rules applied at backup time: when a full backup job runs on a Sunday that is the last Sunday of the month □ automatically flagged as monthly extended retention; (4) Manual flag: right-click a specific job □ Do Not Age (infinite retention); (5) Extended retention jobs: exempt from normal Data Aging (never pruned by standard aging). Common use: SOX (7-year), HIPAA (6-year), legal hold.

Example:

Exam: 'A company must retain month-end financial backups for 7 years but all other backups for only 30 days. How is this configured?' Answer: Extended retention rules on the storage policy copy. Standard retention: 30 days, 2 cycles (for daily backups). Extended retention rule: 'For every monthly full backup (run on last day of month) □ retain 7 years.' Commvault automatically applies 7-year retention to month-end full backups and 30-day retention to all other backups. No manual intervention required after rule is set.

Q254

What happens to a Commvault backup job when the source client restarts mid-backup?

Answer:

Client restart mid-backup behavior: (1) If client reboots: backup job detects connection loss; (2) Job enters 'Waiting' or 'Interrupted' state; (3) After client comes back online and Commvault services restart: CommServe detects client reconnected; (4) Job behavior depends on configuration: (a) Restart from checkpoint: for file backups, restart from last completed checkpoint (not from beginning); (b) Database backups: may need full restart (consistency concern); (5) If job cannot restart: it fails and is retried per retry policy; (6) Data backed up before restart: valid and usable even if job ultimately fails (partial backup = partial recovery possible). CCA: checkpoint restart is key for large file server efficiency.

Example:

Exam: 'A 10-hour file server backup is 7 hours in when the client reboots due to Windows Update. What happens to the backup job?' Answer: Job enters Waiting state. After client comes back online (30 minutes), Commvault resumes from last checkpoint (around the 7-hour mark). Job completes with remaining 3 hours of work – total backup time: 10.5 hours. Without checkpoint restart: entire 10 hours would restart from scratch – total 17+ hours. Checkpoint restart is critical for large file servers with maintenance windows.

Q255

What is a Commvault proxy and what are the types?

Answer:

Commvault proxies are intermediary servers that perform data-intensive operations on behalf of clients: (1) VSA proxy: performs VM backup data movement — reads from vCenter/ESXi/Hyper-V; (2) Index proxy: performs content indexing for search and eDiscovery; (3) Client proxy: acts as intermediary for clients that cannot communicate directly with CommServe (DMZ, NAT); (4) MediaAgent as proxy: MediaAgent itself acts as data proxy; (5) Cloud proxy: facilitates connection to cloud services. Proxies reduce load on production systems (database servers, VMs) by offloading data processing to the proxy server.

Example:

Exam: 'In a VMware environment, which server actually reads the VM backup data from the ESXi host?' Answer: VSA proxy (Virtual Server Agent proxy). The VSA proxy connects to vCenter, initiates VM snapshot, reads CBT-identified blocks from ESXi, and sends data to MediaAgent. The ESXi hosts and VMs themselves do NOT run Commvault software — proxy handles everything remotely. Without VSA proxy: no agentless VMware backup. Each production VM would need a separate Commvault agent (much higher overhead).

Q256

How does Commvault determine which backup copy to use for a restore?

Answer:

Restore copy selection in Commvault: (1) Copy precedence: configured in storage policy — which copy is preferred for restore (typically Primary, then Secondary); (2) Fastest available: Commvault may prefer disk over tape for faster restore; (3) Availability: if primary copy media offline, falls back to secondary copy automatically; (4) User selection: admin can specify which copy to restore from (override auto-selection); (5) Encrypted copies: if restoring encrypted copy, requires matching key/pass-phrase; (6) Geographic selection: for DR, prefer copy stored at DR site; (7) Latest version: for point-in-time restore, selects the appropriate copy holding the requested time's data.

Example:

Exam: 'Primary disk library is offline. Admin initiates a restore. What happens?' Answer: Commvault checks copy precedence. Primary copy (disk) inaccessible. Falls back to Secondary copy (tape library). Restore initiated from tape. Tape restore is slower but functional. If tape also unavailable: restore fails with 'No valid copy available'. Lesson: copy precedence and multiple copies ensure restore availability even when primary storage fails. This is the operational value of the 3-2-1 rule.

Q257

What is Commvault's snap revert capability?

Answer:

Snap revert instantly reverts a production volume or VM to a previous IntelliSnap snapshot state — without restoring data from backup. Process: (1) IntelliSnap creates array snapshots (e.g., every 2 hours); (2) Snap revert: storage array reverts the volume to a specific snapshot state; (3) Commvault orchestrates the revert via array API; (4) Speed: near-instant (seconds to minutes) vs hours for traditional restore; (5) Use case: logical corruption, accidental deletion, failed update; (6) Limitation: requires IntelliSnap-compatible storage array; snapshot must still be available on array (within array snapshot retention); (7) Vs backup restore: snap revert is faster but limited by array snapshot retention; backup restore covers longer history.

Example:

Exam: 'Production Oracle database has logical corruption after a bad stored procedure at 10 AM. IntelliSnap hourly snapshots available on NetApp array. What is the fastest recovery option?' Answer: Snap revert to 9 AM snapshot (1 hour before corruption). Process: Commvault → NetApp API → revert Oracle volume to 9 AM snapshot. Database reverts in 45 seconds. Oracle services restarted. Application online in 5 minutes. Data loss: 1 hour. Traditional backup restore of 10TB Oracle: 3 hours minimum. Snap revert: 5 minutes. When IntelliSnap available, always prefer snap revert for speed.

Q258

What are the common Commvault backup transport modes for VMware?

Answer:

VMware backup transport modes (how data moves from ESXi to VSA proxy): (1) SAN transport: VSA proxy directly reads ESXi VMDK from shared SAN – fastest, requires proxy on same SAN; (2) Hot-add (virtual disk): proxy is a VM on same ESXi/cluster – VMDK temporarily attached to proxy, read directly – very fast, no network required; (3) NBD (Network Block Device): data transferred over network from ESXi to proxy – standard network transport; (4) NBDSSL: encrypted NBD (slower due to encryption overhead); (5) Automatic: Commvault selects best available mode (SAN > Hot-add > NBD). CCA: hot-add is most common for virtual environments; SAN for physical proxies with FC access.

Example:

Exam: 'VSA proxy is a VM running on the same VMware cluster as the backed-up VMs. Which transport mode provides the best performance?' Answer: Hot-add. VMDK temporarily attached to proxy VM (direct disk read, no network copy). Performance: close to raw disk read speed. SAN transport would require physical proxy with SAN access. NBD uses network (slower). Commvault selects hot-add automatically if proxy is on same cluster. Configuration: ensure proxy VM has 'Change disk mode' rights in vCenter – required for VMDK hot-add operation.

Q259

How do Commvault jobs interact with SQL Always On Availability Groups?

Answer:

SQL Always On Availability Groups (AG) with Commvault: (1) Preferred backup replica: configure Commvault to back up from the preferred secondary replica (reduce primary impact); (2) AG-aware backup: Commvault understands AG topology, automatically connects to preferred replica; (3) Log shipping: transaction log backups from secondary do not break the log chain; (4) AG failover resilience: if AG fails over, Commvault redirects backup to new primary automatically; (5) Instance-level subclient: configure subclient at AG listener level (not per-node) for seamless failover; (6) VSS integration: Commvault coordinates with SQL AG VSS writer for consistent backups across all replicas.

Example:

Exam: 'SQL Always On AG with primary and 2 secondaries. How should Commvault be configured to minimize impact on primary while ensuring consistent backups?' Answer: Configure Commvault SQL iDA subclient using the AG listener name (not a specific node). Set preferred backup replica: Secondary preferred. Commvault reads from secondary – zero primary IO impact. Log backups: from secondary also (log chain maintained). AG failover: Commvault automatically uses new primary if all secondaries unavailable. Result: consistent backup, zero primary impact, AG-aware failover.

Q260

What does the Commvault 'Data Verification' job validate?

Answer:

Data Verification validates three things: (1) Readability: backup data can actually be read from storage (disk, tape, or cloud); (2) Integrity: data checksums match what was stored (no silent corruption); (3) Completeness: all expected blocks are present (no missing chunks). Types: Quick DV = verifies headers and a sample of checksums (fast, not exhaustive). Full DV = reads every block and verifies all checksums (thorough). CCA exam: DV does NOT test whether the data can be successfully restored (it only tests readability) – actual restore test is the only true recovery validation. Both are needed: DV (frequent) + restore test (periodic).

Example:

Exam distinction: 'Is data verification the same as a test restore?' Answer: No. Data Verification: confirms backup data is readable and uncorrupted on storage media. It does NOT restore data to the source or a test system – no functional validation. Test restore: actually restores data to a test location and validates the application starts, data is accessible, and the restore process works end-to-end. Best practice: DV monthly (validates storage integrity), test restore quarterly (validates recovery capability). Both are required for a complete backup validation program.

Q261

What is the Commvault CommCell ID and why is it important?

Answer:

CommCell ID is a unique numeric identifier assigned to a CommCell at installation. Importance: (1) Identifies the CommCell uniquely in all backup metadata and media labels; (2) Media written by CommCell 12345 can only be restored using that same CommCell (or after importing to another CommCell with this ID); (3) DR recovery: when rebuilding CommServe on new hardware, the CommCell ID must be restored correctly (from CommServe DR backup) to access existing backup media; (4) Multi-CommCell: CommCell ID distinguishes each CommCell when viewing in a federated console; (5) Support: CommCell ID used to identify customer environments in Commvault Support cases.

Example:

Exam: 'CommServe hardware fails. Admin installs CommServe on new hardware and creates a NEW installation (new CommCell ID). What happens to existing backup data?' Answer: Existing backup data is inaccessible – media was written with the OLD CommCell ID. The new CommServe (different CommCell ID) cannot read or restore from media written by the old CommCell. Correct approach: restore CommServe DR backup to new hardware – this restores the original CommCell ID, all configurations, and job history, allowing access to all existing backup media.

Q262

What are the Commvault data protection levels (backup, replication, snapshot) and when is each used?

Answer:

Three data protection tiers: (1) Backup: copy of data to separate storage, managed retention, multiple versions, long-term (days to years). Best for: recovery from deletion, corruption, compliance. RPO: hours to days. RTO: minutes to hours. (2) Replication (Live Sync/CDR): continuous or near-continuous copy of changes to secondary location. Best for: DR failover, business continuity. RPO: seconds to minutes. RTO: minutes to hours. (3) Snapshot (IntelliSnap): near-instant point-in-time array snapshot. Best for: fast local recovery, backup window elimination. RPO: minutes (snapshot frequency). RTO: seconds to minutes. Use all three in a tiered protection strategy.

Example:

Exam: 'Critical financial trading application – RPO: 15 seconds, RTO: 1 minute. Which protection tier?' Answer: Snapshot + Replication. Snapshot (IntelliSnap every 15 minutes for operational recovery) + CDR Continuous Replication (seconds RPO for DR). Backup in addition for long-term retention. Backup alone cannot achieve 15-second RPO. Replication alone doesn't provide long-term recovery points. Three-tier strategy: IntelliSnap (local fast recovery) + CDR (DR) + Backup (long-term, compliance). Cost: highest – justified by criticality of trading application.

Q263

How does Commvault handle the backup of encrypted virtual machines?

Answer:

Encrypted VM backup challenges: (1) vSphere-encrypted VMs: VMware encryption applied at VM storage layer – Commvault reads encrypted data (cannot deduplicate – encrypted data appears unique to dedup engine); (2) Backup data: stored encrypted on backup storage (good for security, bad for dedup efficiency); (3) NBD encrypted mode: for vSphere encrypted VMs, NBD transport mode required (hot-add not allowed for encrypted VMs in some versions); (4) Key management: ensure KMS key available during backup; (5) Restore: restore encrypted VM, encryption maintained through restore; (6) Dedup: configure separate storage policy without dedup for encrypted VMs (dedup is ineffective for encrypted data anyway).

Example:

Exam: 'Company encrypts all VMs using vSphere VM Encryption (KMS integrated). Admin notices backup dedup ratio dropped from 8:1 to 1:1:1 after enabling VM encryption. Why?' Answer: vSphere VM encryption encrypts at the VMDK layer. Commvault reads encrypted blocks – each block appears unique to the dedup engine (encryption randomises blocks). Dedup is ineffective on encrypted source data. Solution: separate storage policy without dedup for encrypted VMs (saves DDB I/O and storage for dedup overhead). Encrypted VMs stored as-is on backup media. Note: if Commvault encrypts AFTER dedup, dedup still works – the order matters.

Q264

What is the importance of NTP in a Commvault environment?

Answer:

NTP (Network Time Protocol) importance: (1) Job scheduling: CommServe schedules jobs based on system time — time skew causes jobs to run at wrong times; (2) Certificate validation: Commvault SSL certificates are time-sensitive — time skew >5 minutes causes certificate validation failures □ communication breaks; (3) Log correlation: troubleshooting requires correlating logs from CommServe, MediaAgent, and clients — time skew makes this impossible; (4) Job history: backup timestamps incorrect with time skew — compliance evidence unreliable; (5) Recommended: all CommCell components (CommServe, MAs, clients) synced to same NTP source. Time drift: 5+ minutes causes Commvault communication failures.

Example:

Exam: 'After migrating CommServe to a new data centre, all backup jobs fail with authentication errors. No network changes were made. What is the first thing to check?' Answer: NTP/time synchronization. Data centre migration may have changed the NTP source or caused a time drift. If CommServe and clients have >5 minutes time difference: SSL certificate validation fails □ all communications fail. Verification: compare system time on CommServe vs a client server. Fix: configure all systems to same NTP server (internal NTP relay or public time.windows.com). Authentication errors resolve within minutes of time sync.

Q265

What Commvault operation would you use to find all backup versions of a specific file?

Answer:

Finding all versions of a specific file: (1) CommCell Console or Command Center □ Find (search feature); (2) Enter filename or use wildcards (report*.xlsx); (3) Select scope: specific client or all clients; (4) Select time range: last 30 days, last year, all time; (5) Results: list of all backup versions with date, size, and backup job; (6) Action from search: restore directly from search results; (7) Index Server required: file content search requires Index Server; (8) Hash-based search: search by file hash to find all copies of identical files (useful for dedup auditing); (9) Alternatively: Browse specific client □ navigate to file path □ view versions.

Example:

Exam: 'An auditor needs all backup versions of FinanceReport.xlsx from the last 2 years across all file servers. What Commvault feature would you use?' Answer: Find (search). Command Center □ Find □ 'FinanceReport.xlsx' □ scope: All File Server Clients □ time range: last 2 years. Results: 85 versions across 8 file servers (different file servers had different copies). From results: auditor can view each version, restore specific versions for comparison, or export list of all versions with timestamps. Without Find: would need to browse each file server individually — hours of work.

Q266

How does Commvault's retention interact with legal hold?

Answer:

Retention vs Legal Hold interaction: Normal retention: Data Aging ages and prunes data when both retention criteria (days AND cycles) are met. Legal hold: overrides all retention rules — data under legal hold is NEVER aged or pruned regardless of retention expiry. The hold is the absolute override. Interaction rules: (1) Data within retention + legal hold: both protect it (normal retention + hold = doubly protected); (2) Data past retention + legal hold: Data Aging would normally prune — but hold prevents pruning (hold wins); (3) Hold released within retention: normal retention now applies again; (4) Hold released past retention: Data Aging prunes at next run (no further protection without hold). CCA: hold > retention in all cases.

Example:

Exam: 'Email backup retention: 90 days. Legal hold applied to a specific mailbox on day 1. Day 95: Data Aging runs. What happens?' Answer: Day 90: retention expires for this mailbox backup. Data Aging evaluates — normally would prune. But legal hold active □ hold overrides □ data NOT pruned. Data Aging skips this job. Day 95: same — hold still active, data retained. Resolution: legal counsel releases hold after 2 years. Day 1: Data Aging at next run evaluates the data — it's 2 years old, way past 90-day retention □ pruned at next Data Aging cycle.

Q267

What are Commvault's main licensing categories and what do they cover?

Answer:

Commvault licensing categories: (1) Commvault Complete (capacity-based): all features included – backup, replication, archive, analytics, cloud management – priced per TB protected; (2) Commvault Backup and Recovery: traditional backup and recovery features – VM, physical, cloud, SaaS; (3) Commvault Disaster Recovery: replication and DR orchestration; (4) Commvault Archive: long-term archiving, Content Store; (5) Commvault Activate: data intelligence, compliance, eDiscovery (separate add-on); (6) Per-workload: VM licenses, server licenses, SaaS user licenses; (7) CommCell license: CommServe license for each CommCell; (8) MediaAgent license: included with CommServe in most bundles. CCA: capacity licensing is the modern standard.

Example:

Exam: 'A company has 500 VMs, 100 physical servers, 3,000 Microsoft 365 users, and needs GDPR compliance scanning. Which Commvault license components are needed?' Answer: (1) Commvault Complete (covers VMs + physical + M365 backup and replication). (2) Commvault Activate add-on (GDPR sensitive data discovery). OR: if not needing DR replication: Commvault Backup and Recovery + Activate. Capacity: measure total deduplicated protected data in TB for capacity license. SaaS: M365 3,000 users may be per-user licensed separately in some bundles.

Q268

What topics are covered in the Commvault Certified Associate (CCA) exam?

Answer:

CCA exam topics: (1) Commvault architecture – CommServe, MediaAgent, iDA, Storage Policy, SubClient; (2) Backup concepts – backup types, retention, deduplication; (3) VMware backup – VSA, VADP, CBT; (4) Restore operations – full VM, granular, point-in-time, cross-platform; (5) Storage management – disk libraries, tape, cloud, auxiliary copy; (6) Administration – job monitoring, troubleshooting, alerts; (7) Security – RBAC, encryption, ransomware protection; (8) Cloud integration – Microsoft 365, AWS, Azure backup; (9) Policies – storage policy, schedule policy, plans; (10) Compliance – retention, legal hold, eDiscovery. Hands-on product knowledge essential.

Example:

CCA study plan: Week 1: Architecture + components (CommServe, MediaAgent, iDA hierarchy). Week 2: Backup types + storage policies + schedules. Week 3: VMware backup + restore operations. Week 4: Troubleshooting + administration. Week 5: Cloud + compliance + security. Week 6: Practice exams. Resources: Commvault University (free training), Commvault documentation, hands-on trial environment. CCA validates practical deployment and management skills.

Q269

What is the CommServe and why is it the most critical component?

Answer:

CommServe is the central management and metadata server – the brain of the entire CommCell. It is most critical because: (1) All job scheduling and orchestration depends on it; (2) All backup metadata (what was backed up, when, where) stored in CommServe SQL DB; (3) Without CommServe: no new backup jobs can run; (4) Without CommServe: restore operations cannot be initiated (need metadata to find data); (5) VRAs/Agents continue existing operations briefly, but nothing new can start; (6) Recovery depends on CommServe DR backup being current. CCA exam: CommServe failure = most severe single-point-of-failure in CommCell – hence CommServe HA and DR are critical.

Example:

Exam question: 'Which single component failure would prevent ALL new backup jobs from starting?' Answer: CommServe. MediaAgent failure: reduces capacity but other MAs continue. Client failure: only that client affected. iDA failure: only that application's backup affected. CommServe failure: ALL jobs stop (job scheduling resides here). This underscores why CommServe DR backup, HA clustering, and Live Sync are standard practices in production deployments.

Q270

What is the difference between a Storage Policy and a Plan in Commvault?

Answer:

Storage Policy (traditional model): Complex, multi-object — storage policy + schedule policy + subclient configured separately. More granular control, requires more configuration steps. Plan (simplified model): Single object combining backup frequency, retention, and storage destination. Applied directly to workloads. Plans auto-create the underlying storage policies and schedules. CCA exam: Plans introduced in v11 to simplify management; Storage Policies still exist and are used for advanced configurations. Plans = simplified front-end; Storage Policies = the underlying mechanism. Know both — CCA tests both models.

Example:

Exam: 'What is the simplest way to protect a new VM with daily backup and 30-day retention?' Answer: Assign a Plan (e.g., Daily-30Day-Plan) directly to the VM in Command Center. Plans handle the underlying storage policy, schedule, and subclient creation automatically. Storage policy would require: create storage policy + schedule policy + configure subclient separately — more steps, more control. CCA tests: understanding when to use Plans (day-to-day), when to use Storage Policies (advanced custom configuration).

Q271

What is the role of the MediaAgent and what happens if it fails?

Answer:

MediaAgent role: data movement engine — reads from clients, deduplicates, compresses, writes to storage. MediaAgent failure impact: (1) Jobs assigned to that MA stop or fail; (2) Jobs can be reassigned to another MA (if one exists); (3) DDB on failed MA: unavailable until MA recovers (affects restore of deduped data from that MA); (4) Tape drives managed by that MA: inaccessible; (5) Single MA environment: all backups stop; (6) Recovery: restart MA service, or rebuild MA from backup. CCA exam: understand MA redundancy importance — always deploy 2+ MediaAgents in production.

Example:

Exam scenario: 'You have 2 MediaAgents. MA-01 fails at 3 AM during backup window. What happens?' Answer: Jobs running on MA-01 fail (or are suspended). CommServe detects MA-01 offline. If jobs configured for failover: reassigned to MA-02 automatically. If not: admin must manually reassign jobs to MA-02. Ongoing restores using MA-01 DDB: queued until MA-01 recovers. MA-02 DDB not affected. Lesson: configure MediaAgent failover and ensure MA-02 has capacity for all jobs if MA-01 fails.

Q272

What are the two retention criteria in Commvault and how do they work together?

Answer:

Commvault uses TWO retention criteria that BOTH must be satisfied before data is pruned: (1) Retention Days — data must be retained for at least X days; (2) Retention Cycles — data must be retained for at least N full backup cycles. Data is NOT deleted until BOTH criteria are met. This dual-criteria system prevents scenarios where: data is too old (days criterion met) but not enough backup cycles have occurred to ensure a valid recent backup exists. CCA exam favourite: this dual-criteria concept is frequently tested with scenarios requiring calculation of actual retention duration.

Example:

Exam question: 'Retention: 30 days AND 2 cycles. Full backup runs weekly. A full backup job runs on January 1. When can this job be pruned?' Answer: 30 days from January 1 = January 31 (days criterion met on Jan 31). 2 cycles = 2 weekly fulls needed after Jan 1 = January 8 (cycle 1) and January 15 (cycle 2) — cycles criterion met January 15. BOTH criteria met on January 31 (the later of the two). Data retained until January 31 minimum.

Q273

What is the difference between incremental and differential backup?

Answer:

Incremental backup: backs up data changed since the LAST backup (full or incremental). Smallest backup, fastest. Restore requires: last full + ALL incrementals since. Differential backup: backs up data changed since the LAST FULL backup. Larger than incremental but smaller than full. Restore requires: last full + ONLY the last differential. CCA exam comparison: Incremental = faster backup, more backups to restore. Differential = larger backup, simpler restore. Synthetic full = eliminates the restore complexity of incremental-forever by merging at MediaAgent.

Example:

Exam calculation: Monday full (100GB). Tuesday incremental (5GB changed). Wednesday incremental (4GB changed). To restore Wednesday state: need Monday full + Tuesday incremental + Wednesday incremental = 109GB total to restore. If differential instead: Wednesday differential = all changes since Monday = 9GB. Restore: Monday full + Wednesday differential = 109GB restored – same data, but only 2 pieces vs 3. Tradeoff: incremental = 5+4=9GB total backup vs differential = 9GB for just Wednesday.

Q274

What is a subclient and what are its key properties?

Answer:

A subclient is the fundamental backup unit that defines WHAT, HOW, and WHEN data is backed up. Key properties: (1) Content – files, databases, VMs included/excluded; (2) Storage policy – where data goes (links to MediaAgent and library); (3) Schedule policy – when backups run; (4) Backup type – full, incremental, differential; (5) Number of streams – parallel backup threads; (6) Pre/post scripts – actions before/after backup; (7) Encryption – enable/disable at subclient level; (8) Activity control – pause/resume this specific subclient. CCA exam: subclient is the most granular controllable backup object – every backup configuration element is set here or inherited from policy.

Example:

Exam question: 'An admin needs to back up SQL databases with different schedules: FinanceDB daily, HR-DB weekly. What should they configure?' Answer: Two subclients within the SQL iDA: Subclient-Finance (content: FinanceDB, schedule: daily incremental, 4-hour log backup). Subclient-HR (content: HR-DB, schedule: weekly full only). Each subclient has its own independent schedule. One storage policy can serve both, or different policies if retention differs. Granular control at subclient level is the key.

Q275

How does source-side deduplication differ from MediaAgent-side deduplication?

Answer:

Source-side deduplication: (1) Dedup fingerprinting done on the client before sending data to MediaAgent; (2) Only unique blocks sent over the network; (3) Benefit: significant WAN/LAN bandwidth reduction; (4) Cost: client CPU overhead for fingerprinting; (5) Best for: WAN-connected branch offices, slow networks. MediaAgent-side deduplication: (1) All data sent from client to MediaAgent first; (2) Dedup fingerprinting and block comparison done at MediaAgent; (3) Benefit: zero client CPU overhead; (4) Drawback: all data transferred regardless of uniqueness; (5) Best for: high-speed LAN connections where bandwidth is not a constraint.

Example:

Exam scenario: '20 branch offices each with 100Mbps WAN links back up 50GB/night to central MediaAgent. What type of dedup should be configured?' Answer: Source-side deduplication. Without: 20 offices × 50GB = 1TB WAN traffic/night = 93Mbps average (saturates 100Mbps links). With source-side dedup: only unique blocks sent. If 70% similarity across offices: 1TB × 30% = 300GB WAN traffic (28Mbps average). WAN well within limits. Client CPU overhead acceptable for this use case.

Q276

What is the CommCell and how many CommServers can it have?

Answer:

CommCell is the administrative domain managed by Commvault software. It consists of: CommServe (1 active, 1 optional standby) + all MediaAgents + all clients. CommCell always has exactly ONE active CommServe – this is fundamental to Commvault architecture. Multiple CommCells = multiple CommServe servers, each managing their own domain. A standby CommServe (for HA) is passive – it activates only when primary fails. CCA exam: a single CommCell = one active CommServe. Large enterprises may have multiple CommCells managed from one CommCell Console (multi-CommCell view).

Example:

Exam question: 'Can a CommCell have 2 active CommServe servers?' Answer: No – a CommCell has exactly one active CommServe. You can have: (1) One CommServe + one passive standby (HA). (2) Two CommCells each with one CommServe (independent management domains, viewable from one CommCell Console). Distinguish: CommCell (one active CommServe) vs multiple CommCells (multiple active CommServes, multiple management domains).

Q277

What is an auxiliary copy and what are its use cases?

Answer:

Auxiliary copy replicates backup data from one storage policy copy to another: Primary copy → Secondary/Tertiary copies. Use cases: (1) 3-2-1 backup rule: disk → tape and disk → cloud; (2) DR: local disk → remote site disk; (3) Compliance: local backup → long-term tape/cloud archive; (4) Copy management: create tape copies for offsite, cloud copies for long-term retention; (5) IntelliSnap Backup Copy: from array snapshot to backup storage; (6) Selective: copy only specific jobs (e.g., monthly fulls) to tape. Auxiliary copy is separate from backup – runs as its own job, can be scheduled independently.

Example:

Exam question: 'Which operation copies backup data from on-site disk to an off-site tape library for DR?' Answer: Auxiliary copy. The primary backup job writes to the local disk library (primary copy). A scheduled auxiliary copy job reads from the primary copy and writes to the tape library (secondary copy). These are two separate operations. Without auxiliary copy: only one copy exists (no 3-2-1). With auxiliary copy: disk (on-site) + tape (off-site) = two media types, two locations.

Q278

What is the Virtual Server Agent (VSA) and what does it protect?

Answer:

VSA (Virtual Server Agent) is the Commvault iDataAgent for virtualisation platforms. It provides agentless VM backup by communicating with hypervisor APIs (no agent needed inside the VM). Protects: (1) VMware vSphere VMs (via VADP); (2) Microsoft Hyper-V VMs; (3) Microsoft Azure VMs; (4) AWS EC2 instances; (5) Nutanix AHV VMs; (6) GCP instances. Key capabilities: image-level backup, CBT incremental, application-consistent quiescing, granular file restore, cross-hypervisor restore, Live Browse. CCA exam: VSA is deployed on a proxy server (not inside the VM), communicates with vCenter/hypervisor API.

Example:

Exam: 'How does Commvault back up VMs without installing an agent inside each VM?' Answer: Through the VSA (Virtual Server Agent) deployed on a separate proxy server. The VSA proxy communicates with vCenter (VMware) using VADP, creates VM snapshots, reads changed block data via CBT, and sends it to the MediaAgent – all without touching the VM itself. The guest VMs need VMware Tools for application-consistent (quiesced) backups but no Commvault agent.

Q279

What are common CCA exam scenarios for restore operations?

Answer:

Common CCA restore scenarios: (1) 'File accidentally deleted from file server': □ File System iDA □ Browse and Restore □ select file version from yesterday □ restore to original location. (2) 'VM corrupted, need to recover': □ VSA □ select VM backup □ restore to original or alternate vCenter/datastore. (3) 'SQL database needs point-in-time recovery': □ SQL iDA □ Point-in-time restore □ specify date/time. (4) 'VM disk full, need one file from 3 months ago': □ Live Browse □ mount backup □ copy file. (5) 'DR scenario, restore VMs to DR site': □ out-of-place restore □ alternate vCenter/ESXi/datastore. Each scenario has a specific restore path — know which iDA and which restore type applies.

Example:

Exam scenario: 'A user accidentally deleted their OneDrive folder containing 5,000 files 2 weeks ago. Microsoft's 30-day retention has not expired, but the company also uses Commvault. What options exist?' Answer: (1) Microsoft OneDrive recycle bin (if within 30 days). (2) Commvault M365 backup: browse to 2-week-old backup □ restore OneDrive folder directly to user's OneDrive. Commvault option preferred if Microsoft recycle bin unavailable or for granular item selection. Commvault can restore specific files rather than entire folder, and to alternate users.

Q280

What is IntelliSnap and why is it used instead of traditional backup?

Answer:

IntelliSnap coordinates array-level snapshots with Commvault. Used instead of traditional backup when: (1) Backup window too small for traditional backup (snapshot = seconds vs hours); (2) Production impact unacceptable (snapshot minimally impacts production); (3) Need for fast local recovery (restore from array snapshot in minutes); (4) Application-consistent snapshots at high frequency (e.g., hourly snapshots of large DB). IntelliSnap is not a replacement for backup — it's complementary. Backup Copy then moves snapshot data to long-term storage. CCA exam: IntelliSnap = fast snapshot, minimal impact; Backup Copy = traditional copy of snapshot data for long-term retention.

Example:

Exam: 'A 10TB Oracle database needs to be protected. The backup window is 2 hours but full backup takes 8 hours. What feature resolves this?' Answer: IntelliSnap. (1) IntelliSnap triggers array snapshot in seconds (backup window: ~2 seconds). (2) Oracle iDA quiesces Oracle before snapshot for app-consistency. (3) Backup Copy runs in background: reads snapshot, sends to MediaAgent storage. (4) Array snapshot available for instant recovery locally. (5) Backup Copy provides long-term retention copy. Backup window problem solved — production not impacted for hours.

Q281

What is the purpose of data verification in Commvault?

Answer:

Data verification (DV) validates that backed-up data is readable and intact: (1) Quick DV: checks data headers and checksums — fast; (2) Full DV: reads every block of backup data from storage — thorough; (3) Purpose: detect corrupt or unreadable backup data before it's needed for recovery; (4) Schedule: run DV regularly (weekly quick, monthly full) especially on tape; (5) Result: if DV fails — data is corrupt (know before a real DR); (6) Recovery from DV failure: restore data from alternate copy (secondary copy should be clean); (7) CCA exam: DV is a proactive validation step — never rely on backup data without periodic verification.

Example:

Exam question: 'What Commvault feature would you use to confirm that last month's full backup is still readable and recoverable?'

Answer: Data Verification. Run Full Data Verification on the specific backup job. DV reads all backup data from storage, verifies checksums. If verification passes: data is confirmed intact and recoverable. If verification fails: the backup is corrupt — initiate restore from secondary copy (auxiliary copy). Without DV, you wouldn't discover corruption until you need to recover — possibly too late.

Q282

What is Live Browse and when would you use it over a full restore?

Answer:

Live Browse: Commvault mounts a backup as a live file system – browse and copy specific files without a full restore. Use when: (1) Only a few specific files needed (not the entire backup); (2) Source system is running (don't want to restore over it); (3) Need to compare backup content with current files; (4) Quick ad-hoc file recovery; (5) Need immediate access – no time for full restore. Use full restore instead when: (1) Entire system recovery needed; (2) Multiple files across different directories; (3) VM needs to be fully recovered (not just files). CCA exam: Live Browse = faster for few files; full restore = needed for full system recovery.

Example:

Exam: 'An admin needs to retrieve a specific configuration file from a 200GB VM backup from 2 weeks ago. What is the most efficient method?' Answer: Live Browse. (1) Mount the 2-week-old VM backup. (2) Navigate to the configuration file path. (3) Copy the file. (4) Unmount. Total time: 5 minutes. Alternative (full restore): 200GB restore to staging server = 45 minutes + find file + copy. Live Browse 9× faster for single-file recovery. Full restore only needed if entire VM needs recovery.

Q283

How does Commvault's deduplication database (DDB) impact restore operations?

Answer:

DDB impact on restore: (1) Restore reads from backup storage, reassembles deduplicated blocks using DDB; (2) DDB must be available and healthy for restore of deduplicated data; (3) DDB offline: restore of dedup data from that MediaAgent fails; (4) DDB corruption: restore may fail or return incomplete data; (5) DDB on different MediaAgent: restore must use the MA that manages the DDB for that storage; (6) Migrating DDB: if DDB moved to new storage, restore still works (DDB updated with new path); (7) DDB backup importance: DDB backup allows DDB recovery – without it, dedup data unrecoverable even if backup data exists.

Example:

Exam: 'A MediaAgent with the DDB fails. The backup data is on a disk library (NAS) still accessible from other MAs. Can the restore proceed?' Answer: No – not immediately. While backup data is on the NAS (accessible), the DDB on the failed MA is needed to reassemble the deduplicated blocks. Without DDB: cannot restore dedup data. Recovery options: (1) Restore MA (quick restart). (2) Restore DDB from DDB backup to another MA. (3) Wait for MA to come online. Lesson: DDB is as critical as backup data itself.

Q284

What is the Commvault Disaster Recovery backup (CommServe DR) and what does it contain?

Answer:

CommServe DR backup exports the CommServe SQL database to protected locations. Contents: (1) All CommCell configuration (clients, MediaAgents, storage policies, schedules); (2) Complete job history (all backup jobs run, results, media locations); (3) Media tracking (which data is on which tape, disk, cloud); (4) User and security configuration; (5) Alert and notification settings; (6) License information. Without CommServe DR: if CommServe fails, cannot initiate restores (no metadata to know where data is). With CommServe DR: rebuild CommServe from DR backup, all metadata restored, all backups accessible for recovery. CCA exam: CommServe DR = most critical backup in the CommCell.

Example:

Exam: 'CommServe server hardware fails completely. There is no CommServe DR backup. What is the impact?' Answer: Catastrophic. (1) All backup data still exists on disk/tape/cloud (data is not lost). (2) But: cannot initiate restores (no metadata to know where data is). (3) Cannot run new backup jobs. (4) Recovery: Commvault Support may be able to assist with partial recovery, but it's extremely difficult and time-consuming. (5) Lesson: CommServe DR is not optional – it should be configured from day 1, exported to 3 locations, verified regularly.

Q285

What is the key difference between in-place restore and out-of-place restore?

Answer:

In-place restore: data restored to its original source location. Overwrites current data. Use when: source system failed and needs recovery, OR source data is lost/corrupt and you want to recover it to the same location. Out-of-place restore: data restored to a different location. Original source untouched. Use when: (1) Source still running (don't overwrite); (2) Testing recovery; (3) Forensic investigation; (4) Need comparison of backup version with current; (5) DR site recovery; (6) Migration. CCA exam: almost always prefer out-of-place for testing; in-place only when source is unavailable or deliberately overwriting.

Example:

Exam question: 'A user reports that a report.xlsx file has wrong data. They think yesterday's version was correct. How should the admin restore without risking the current version?' Answer: Out-of-place restore. Restore report.xlsx from yesterday's backup to a different folder (e.g., D:\Restore\report_yesterday.xlsx). User compares both versions. If yesterday's is correct: user copies it over the current version. If not: try the day-before version. Original current file preserved until confirmed correct version found. Never in-place restore when source data may still be valid.

Q286

What is the purpose of a schedule policy in Commvault?

Answer:

Schedule policy is a reusable template that defines WHEN backups run. It contains: (1) Backup frequency and type (full, incremental, differential); (2) Start time and time window (not before/not after); (3) Retry settings for missed jobs; (4) Synthetic full schedule (if used). Purpose: apply consistent schedule to many subclients – change once, update all. Separation from storage policy: schedule policy controls WHEN; storage policy controls WHERE and HOW LONG. CCA exam: understand the three policy types – storage policy (where/retention), schedule policy (when), and how they apply to subclients.

Example:

Exam: 'An admin wants 50 SQL servers to have the same backup schedule: full Sunday, incremental Mon-Sat, logs every 15 min. What is the most efficient approach?' Answer: Create one Schedule Policy 'SQL-Standard' with these schedules, apply to all 50 SQL Server subclients. When the log frequency needs to change from 15 to 30 minutes: update the one schedule policy □ all 50 servers updated automatically. Without schedule policy: update each of 50 subclients individually = 50× the work.

Q287

What are the CCA exam tips for Commvault storage policy concepts?

Answer:

CCA exam storage policy key concepts: (1) One active CommServe per CommCell; (2) Storage policy links subclients to MediaAgent and storage; (3) Primary copy always exists; secondary copies are optional; (4) Retention: BOTH days AND cycles must be met before data aged; (5) Auxiliary copy = data movement between copies within storage policy; (6) DDB is per-copy (primary copy has its own DDB); (7) Encryption per copy – can encrypt some copies differently; (8) Dedup: only disk libraries (not tape); (9) Cloud library = a disk library in CommVault terms; (10) Multiple subclients can share one storage policy.

Example:

Exam trap question: 'Can tape libraries be deduplicated in Commvault?' Answer: No – deduplication works on disk libraries (and cloud libraries which are disk-type). Tape media does not support Commvault DDB-based deduplication. However: tape hardware compression is available. Dedup happens on disk (primary copy), then compressed+encrypted data copied to tape via auxiliary copy. The tape copy is already benefitting from dedup on disk – just the tape itself doesn't do inline dedup.

Q288

How does Commvault handle application-consistent backups for databases?

Answer:

Application-consistent database backups: (1) SQL Server: VSS (Volume Shadow Copy Service) used to quiesce SQL, ensure consistent backup; (2) Oracle: RMAN (Oracle Recovery Manager) integration — backup taken via RMAN ensuring block-level consistency; (3) Exchange: Exchange VSS writer quiesces Exchange DB; (4) SAP HANA: HANA backup API used for consistent DB freeze; (5) Without app-consistency: crash-consistent backup (like pulling power plug) — DB may require recovery on restore; (6) Application-consistent: guaranteed clean DB state, no transaction replay needed. CCA exam: application-consistent = uses application's own API to quiesce; crash-consistent = hypervisor snapshot only.

Example:

Exam: 'Why is it important to use the SQL iDA instead of just backing up the SQL .mdf and .ldf files with the File System iDA?' Answer: File System iDA on SQL files — files are locked during backup (SQL writes to them constantly), resulting in a corrupt or inconsistent backup. SQL recovery requires significant replay. SQL iDA uses VSS: signals SQL Server to quiesce writes, creates consistent snapshot, then backup reads consistent state. Restored SQL database is 100% clean — no recovery needed. Always use application-aware iDA for database backups.

Q289

What happens when a backup exceeds its scheduled backup window in Commvault?

Answer:

When a backup exceeds its schedule window: (1) If 'Do not start after X time' configured: job won't start if it would run past the cutoff; (2) Running jobs: if job is already running and the cutoff is reached, configurable behavior: (a) continue to completion, (b) kill and restart next cycle; (3) Missed job policy: (a) run immediately when possible, (b) skip to next scheduled run, (c) run at specific catch-up time; (4) Impact on next scheduled job: if previous job still running when next is scheduled, new job may queue or run in parallel depending on stream limits; (5) Alert: configure alert for jobs running beyond expected duration. CCA exam: know the configurable options for handling exceeded windows.

Example:

Exam scenario: 'Backup scheduled 10 PM, must complete by 6 AM. Server is busy and backup runs until 8 AM. What would prevent this from happening in future?' Answer: Configure 'Run during maintenance window only' in schedule policy with 10 PM to 6 AM window, and 'Kill job if running past end of window'. OR: identify why backup is slow (investigation needed). For large file servers: enable source-side dedup, filter unneeded files, increase streams. Prevention better than kill — investigate root cause of slow backup.

Q290

What are CCA exam-level distinctions between Commvault backup types?

Answer:

CCA backup type exam summary: Full — reads all data, baseline, largest, slowest, fewest restore pieces. Incremental — reads changes since LAST backup (any type), smallest, fastest backup, most restore pieces needed. Differential — reads changes since LAST FULL, medium size, simple restore (full + one diff). Synthetic Full — created at MediaAgent from full + incrementals, no client read, combines benefits of full restore with incremental backup speed. Transaction Log — database logs only, enables point-in-time recovery. Selective Full — full backup at specific times (e.g., monthly). CCA: match each type to its best use case — incremental forever + synthetic full is modern standard.

Example:

Exam scenario: 'Client needs: (1) Fast daily backups, (2) Fast restores, (3) No large weekly full backup window impacting production. What strategy?' Answer: Incremental forever + weekly synthetic full. Daily: incremental (fast, low production impact). Weekly: synthetic full created at MediaAgent (zero client impact — client not read). Restore from: synthetic full (fast, like a regular full). This strategy meets all 3 requirements. Traditional weekly full would violate requirement 3 (large window, production impact).

Q291

What is the Commvault Metrics Reporting Server (MRS) and what does it provide?

Answer:

Metrics Reporting Server: centralised reporting across multiple CommCells. Provides: (1) Consolidated job reports (all CommCells in one view); (2) SLA compliance dashboard (% protection across all CommCells); (3) Capacity and storage reports; (4) License consumption tracking; (5) Scheduled report delivery; (6) Trend analysis over time. Who uses it: IT management, backup team leads, compliance officers. When needed: organisations with multiple CommCells (e.g., global enterprises, MSPs managing multiple customers). CCA exam: MRS = reporting only, not operational management (can't run jobs from MRS). Separate installation from CommServe.

Example:

Exam: 'A global company has 6 CommCells (one per region). Management wants a single monthly backup success report for all regions. What component provides this?' Answer: Metrics Reporting Server (MRS). MRS aggregates data from all 6 CommCells and generates a consolidated report: global backup success rate, SLA compliance per region, total data protected, and storage consumption. Without MRS: admin would need to log into each of 6 CommCells separately and compile reports manually. MRS automates this and delivers by email monthly.

Q292

How does Commvault's RBAC model provide granular security?

Answer:

Commvault RBAC granularity: (1) Roles: predefined (Admin, Operator, Report Viewer) or custom; (2) Capabilities: 100+ individual permissions within roles (e.g., 'Run Backup', 'Browse and Restore', 'Configure Policies'); (3) Scope: limit access to specific clients, client groups, storage policies; (4) User groups: assign roles to groups (AD-integrated); (5) Two-user rule: some operations (e.g., restore) can require two different users; (6) Tenant isolation: each tenant/business unit sees only their data; (7) Time-based access: restrict access to specific hours. CCA exam: RBAC is highly granular — users can restore their own files but not others, even with 'Restore' role.

Example:

Exam: 'Finance team should be able to restore their own server backups but not view or restore other departments' data. How is this configured?' Answer: Create user group 'Finance-Restore-Operators'. Role: Restore Operator. Scope: limited to 'Finance Client Group' (only Finance servers). Result: Finance operators log into CommCell, see only Finance servers, can restore only Finance data. IT, HR, and other departments' clients are not visible to Finance operators. Principle of least privilege enforced through RBAC scope configuration.

Q293

What is Commvault's global deduplication and why is it important for large environments?

Answer:

Global deduplication: all clients sharing the same DDB (Deduplication Database) benefit from cross-client deduplication. Importance for large environments: (1) Maximises storage savings — identical blocks across all clients stored once; (2) VDI: 1,000 desktops with same OS image = OS blocks stored once (massive savings); (3) SQL servers: transaction log blocks often repetitive across instances; (4) Scale: one DDB can serve hundreds of clients via multiple MediaAgents with DDB sync; (5) Reporting: global dedup ratio shows true storage efficiency; (6) Requirement: all clients must use the same storage policy copy/DDB to benefit. CCA exam: global dedup = cross-client, not just within one client.

Example:

Exam scenario: 'A company has 500 Windows 10 VDI desktops, each 40GB. Without dedup: $500 \times 40\text{GB} = 20\text{TB}$. With per-client dedup: 20TB (no cross-client benefit if done per-VM). With global dedup: Windows 10 OS (25GB, same across all desktops) stored once = 25GB. Unique user data (15GB average per VM) = 7.5TB. Total: ~7.5TB instead of 20TB. Global dedup ratio: 2.7:1 in this example. For environments with high data similarity, global dedup dramatically reduces storage needs.

Q294

What are the key Commvault components a CCA candidate should know by memory?

Answer:

CCA must-know components: (1) CommServe – management server, SQL DB, one per CommCell; (2) MediaAgent – data mover, dedup engine, can have multiple; (3) iDataAgent (iDA) – client-side agent, application-specific; (4) Subclient – fundamental backup unit (what, when, where); (5) Storage Policy – where data goes and retention; (6) Schedule Policy – when backups run; (7) DDB (Deduplication Database) – block hash database on MediaAgent; (8) CommCell – all components managed by one CommServe; (9) Disk Library – backup storage target; (10) Auxiliary Copy – data replication between copies; (11) VSA – virtual server backup (agentless VM backup); (12) IntelliSnap – hardware snapshot integration.

Example:

CCA exam mental map: Data flow – Client iDA □ (reads data) □ MediaAgent □ (dedup via DDB) □ writes to Disk Library. CommServe orchestrates: schedules job, tells MA which library, monitors progress, records in SQL DB. Restore: CommServe queries SQL DB (where is data?) □ tells MA which library/media □ MA reads data □ sends to client. Every component has a specific role. Understanding data flow answers most architecture questions. Memorise: CommServe = brain, MediaAgent = muscle, iDA = hands, Library = storage.

Q295

What are common exam traps in CCA certification questions?

Answer:

Common CCA exam traps: (1) Dedup on tape: WRONG – dedup only on disk/cloud libraries, not tape; (2) CommServe moves data: WRONG – MediaAgent moves data, CommServe only orchestrates; (3) One CommServe per CommCell: TRUE – multiple MAs allowed, one CommServe active; (4) Both retention criteria: data pruned only when BOTH days AND cycles met; (5) VSA needs agent inside VM: WRONG – VSA is agentless (proxy outside VM); (6) Incremental restores: needs full + ALL incrementals (not just last incremental); (7) Synthetic full is a real full backup: TRUE – same restore speed as traditional full; (8) Data Aging deletes within retention: WRONG – never deletes within retention period.

Example:

Practice exam traps: Q: 'Where does Commvault store backup metadata?' A: CommServe SQL database (not the backup storage). Q: 'Which component handles data deduplication?' A: MediaAgent (not CommServe, not client by default). Q: 'An admin removes a media from the library. Which retention criteria determines when data is pruned?' A: Both days AND cycles – tape removal doesn't affect retention policy. Q: 'Can deduplication be enabled on a tape library?' A: No – only disk/cloud libraries support Commvault DDB-based deduplication.

Q296

How do you explain CommCell architecture using a simple analogy for the CCA exam?

Answer:

CommCell components analogy: CommServe = the 'dispatcher' – knows all drivers, trucks, and warehouses; schedules routes; tracks deliveries; never touches packages. MediaAgent = the 'truck driver' – actually moves data (packages) from pickup (client) to storage (warehouse); handles loading/unloading; knows warehouse layout. iDataAgent = the 'warehouse picker' at client side – knows how to package specific items (SQL, Exchange, files) correctly for transport. Storage Library = the 'warehouse' – where packages are stored. Subclient = the 'shipping manifest' – defines what gets packed, when, and to which warehouse. CCA exam: if unsure which component does what, use this analogy to reason through answers.

Example:

Exam question reasoning: 'Which component is responsible for compressing backup data?' □ MediaAgent (the truck driver compresses the packages before loading). 'Which component decides WHEN a backup runs?' □ CommServe (the dispatcher controls the schedule). 'Which component knows how to back up an Oracle database?' □ Oracle iDataAgent (the specialist picker who knows Oracle packaging). 'Where is the list of what was backed up stored?' □ CommServe SQL DB (the dispatcher's delivery records). Analogy helps answer architecture questions without memorising every detail.

Q297

What are the key differences between CommCell Console and Command Center in Commvault?

Answer:

CommCell Console is the legacy Windows-based thick client GUI, offering deep configuration and advanced management. Command Center is the modern browser-based interface providing a simplified, role-appropriate view of operations, monitoring, and self-service. For CCA, both are covered — Command Center is the strategic direction for new deployments.

Example:

An operator uses Command Center for daily monitoring and job management. An architect uses CommCell Console for advanced storage policy configuration and agent-level tuning. In Commvault 11 and later, Command Center handles most common tasks without needing the thick client.

Q298

What is the Commvault HyperScale X platform?

Answer:

HyperScale X is Commvault's hyperconverged backup appliance platform that combines compute, storage, and software in a single scale-out node. Nodes are added to expand capacity and performance linearly. It eliminates the need for separate MediaAgent and storage hardware by integrating all components in a pre-configured, validated stack.

Example:

An enterprise starts with 3 HyperScale X nodes providing 500TB of backup capacity. As data grows, they add 2 more nodes — capacity scales to 830TB and throughput increases proportionally without reconfiguring existing infrastructure.

Q299

What is the significance of the Recovery Point Objective (RPO) in Commvault backup design?

Answer:

RPO defines the maximum acceptable data loss in time — how old the most recent backup can be during a recovery. Commvault backup schedules (full, incremental, differential, log backup frequency) are designed to meet the RPO. Lower RPO requires more frequent backups and often log-based or continuous protection.

Example:

A business requires RPO of 1 hour for a critical database. Commvault is configured with database transaction log backups every 30 minutes, ensuring that at most 30 minutes of data is lost in any failure — comfortably within the 1-hour RPO target.

Q300

For the CCA exam, what is the key architectural component that distinguishes Commvault from traditional backup tools?

Answer:

The key differentiator is the unified platform approach — a single CommServe manages all agents, storage, policies, and data types across physical, virtual, cloud, and SaaS environments. Traditional backup tools are often siloed by workload. Commvault's single metadata repository, policy engine, and reporting layer eliminate the need for multiple backup products.

Example:

In a hybrid environment with on-prem VMs, AWS workloads, Microsoft 365, and SAP HANA, Commvault manages all four from one CommServe with unified reporting and SLA tracking — whereas traditional tools would require four separate backup products for the same coverage.