

PVE

Proxmox VE Admin Certification

300 Interview Questions & Answers

KVM | LXC | Clustering | Storage | Networking | HA | Security

Prepared by TechShiksha

akhileshchaudhari.com

300

Questions

9

Sections

Expert

Level

For educational and interview preparation purposes | 2025

Table of Contents

Section 1: Proxmox VE Fundamentals	Q1–Q30
Section 2: KVM Virtual Machines	Q31–Q70
Section 3: LXC Containers	Q71–Q100
Section 4: Networking	Q101–Q140
Section 5: Storage Management	Q141–Q180
Section 6: High Availability & Clustering	Q181–Q220
Section 7: Security & Access Control	Q221–Q250
Section 8: Monitoring, Logging & Performance	Q251–Q280
Section 9: Advanced Topics & Troubleshooting	Q281–Q300

This guide covers comprehensive Proxmox VE topics for exam preparation and professional interviews. Questions range from fundamental concepts to advanced troubleshooting and best practices.

Section 1: Proxmox VE Fundamentals (Q1–Q30)

Q1. What is Proxmox VE?

Answer: Proxmox Virtual Environment (Proxmox VE) is an open-source server virtualization management platform that integrates KVM hypervisor and LXC containers, software-defined storage, and networking into a single solution, managed via a web-based interface.

Q2. What are the two main virtualization technologies in Proxmox VE?

Answer: Proxmox VE supports KVM (Kernel-based Virtual Machine) for full hardware virtualization of virtual machines, and LXC (Linux Containers) for lightweight OS-level virtualization.

Q3. What Linux distribution is Proxmox VE based on?

Answer: Proxmox VE is based on Debian GNU/Linux. Each major Proxmox VE version aligns with a Debian release — for example, Proxmox VE 8.x is based on Debian 12 (Bookworm).

Q4. What is the default web interface port for Proxmox VE?

Answer: The Proxmox VE web management interface listens on port 8006 over HTTPS (TLS). Access it via `https://:8006` in a browser.

Q5. What is the difference between a VM and a Container in Proxmox VE?

Answer: A VM (KVM) runs a complete operating system with its own kernel using hardware virtualization, providing full isolation. A Container (LXC) shares the host kernel, making it lighter and faster but limited to Linux-compatible workloads.

Q6. What kernel does Proxmox VE ship with?

Answer: Proxmox VE ships with a custom hardened Linux kernel based on the Ubuntu/Debian kernel, patched with additional drivers, ZFS support, and optimizations for virtualization workloads.

Q7. What is a Proxmox VE Cluster?

Answer: A Proxmox VE Cluster is a group of Proxmox nodes that share configuration, storage, and network settings, enabling centralized management, live migration of VMs, and high availability through a unified web interface.

Q8. What is the minimum number of nodes required for a Proxmox HA cluster?

Answer: A minimum of three nodes is recommended for a High Availability cluster in Proxmox VE. This is because the quorum mechanism requires a majority vote; three nodes tolerate one node failure while maintaining quorum.

Q9. What is Corosync in Proxmox VE?

Answer: Corosync is the cluster communication framework used by Proxmox VE. It handles cluster membership, quorum, and heartbeat messaging between nodes to ensure consistency and detect node failures.

Q10. What is the role of pvecm?

Answer: pvecm (Proxmox VE Cluster Manager) is the command-line utility used to create, join, and manage Proxmox VE clusters. Commands include pvecm create, pvecm add, pvecm status, and pvecm nodes.

Q11. What file systems are supported by Proxmox VE storage?

Answer: Proxmox VE supports multiple storage backends including LVM, LVM-Thin, ZFS, Directory (ext4/xfs), NFS, CIFS/SMB, Ceph RBD, iSCSI, Gluster, and ZFS over iSCSI.

Q12. What is ZFS in the context of Proxmox VE?

Answer: ZFS (Zettabyte File System) is an advanced file system and volume manager supported natively in Proxmox VE. It offers features like snapshots, data integrity checking (checksums), RAID-Z, compression, and deduplication.

Q13. What is Ceph in Proxmox VE?

Answer: Ceph is a distributed software-defined storage platform that can be deployed directly on Proxmox VE nodes. It provides highly available block storage (RBD), object storage, and file storage with no single point of failure.

Q14. What is the difference between LVM and LVM-Thin in Proxmox VE?

Answer: LVM provides thick-provisioned volumes where storage is fully allocated immediately. LVM-Thin uses thin provisioning, allocating storage on demand, enabling overprovisioning and snapshot support.

Q15. What is a storage pool in Proxmox VE?

Answer: A storage pool is a configured storage resource in Proxmox VE that can be shared across nodes in a cluster. It defines where VM disk images, ISO files, container templates, and backups are stored.

Q16. How does Proxmox VE handle VM backups?

Answer: Proxmox VE provides built-in backup functionality using the `vzdump` utility. Backups can be scheduled via the web GUI, stored locally or on NFS/CIFS shares, and support snapshot, suspend, and stop modes.

Q17. What are the three backup modes in Proxmox VE?

Answer: The three backup modes are: Snapshot (live backup with minimal downtime using snapshots), Suspend (VM is briefly suspended during backup), and Stop (VM is fully stopped during backup – most consistent but causes downtime).

Q18. What is `vzdump`?

Answer: `vzdump` is the command-line tool in Proxmox VE used to create backups of VMs and containers. It supports all backup modes and can compress backups in various formats (gzip, lzo, zstd).

Q19. What is a Proxmox VE node?

Answer: A Proxmox VE node is a physical or virtual server running Proxmox VE software. Each node can host multiple VMs and containers. Nodes can be grouped into a cluster for unified management.

Q20. What is VMID in Proxmox VE?

Answer: VMID (Virtual Machine ID) is a unique numeric identifier assigned to each VM or container in Proxmox VE. IDs 100 and above are used for user VMs/containers; IDs below 100 are reserved for system use.

Q21. What is the default storage created during Proxmox VE installation?

Answer: During installation, Proxmox VE creates a local storage (directory-based at `/var/lib/vz`) and a local-lvm storage (LVM-Thin pool on the root volume group) by default.

Q22. What is the Proxmox VE subscription model?

Answer: Proxmox VE is free and open-source. Subscriptions are available (Basic, Standard, Premium, Community) that provide access to the enterprise repository, stable updates, and commercial support from Proxmox Server Solutions GmbH.

Q23. What is the pve-no-subscription repository?

Answer: The pve-no-subscription repository is the free/community repository for Proxmox VE updates, usable without a paid subscription. It may contain less-tested packages compared to the enterprise repository.

Q24. How do you update Proxmox VE from the command line?

Answer: Run: `apt update && apt dist-upgrade`. It is recommended to use `apt dist-upgrade` rather than `apt upgrade` to properly handle package transitions in Proxmox VE.

Q25. What is the Proxmox VE API?

Answer: Proxmox VE provides a comprehensive RESTful API that allows programmatic management of all resources — nodes, VMs, containers, storage, networks, and cluster settings. The API is used by the web UI, CLI tools, and third-party integrations.

Q26. What authentication methods does Proxmox VE support?

Answer: Proxmox VE supports PAM (Linux users), PVE realm (built-in user database), LDAP/Active Directory integration, OpenID Connect (OIDC), and two-factor authentication (TOTP, WebAuthn/FIDO2).

Q27. What is the default administrator account in Proxmox VE?

Answer: The default administrator account is `root@pam`, which uses PAM (Linux system) authentication. This account has full superuser privileges on the Proxmox VE node.

Q28. What is role-based access control (RBAC) in Proxmox VE?

Answer: Proxmox VE uses RBAC to control user access. Permissions are assigned by combining Users/Groups + Roles + Paths. Built-in roles include Administrator, PVEAdmin, PVEVMAdmin, PVEAuditor, and others.

Q29. What is a resource pool in Proxmox VE?

Answer: A resource pool is a logical grouping of VMs, containers, and storage in Proxmox VE. Pools simplify permission management by allowing administrators to grant users access to a collection of resources at once.

Q30. What are Proxmox VE tasks?

Answer: Tasks are background operations tracked by Proxmox VE, such as VM start/stop, backup, migration, or snapshot creation. Task logs are viewable in the web GUI under the Tasks panel and stored in `/var/log/pve/tasks/`.

Section 2: KVM Virtual Machines (Q31–Q70)

Q31. What is KVM?

Answer: KVM (Kernel-based Virtual Machine) is a full virtualization technology built into the Linux kernel. It turns the host into a hypervisor, allowing guest VMs to run isolated operating systems with near-native hardware performance.

Q32. What CPU features are required for KVM in Proxmox VE?

Answer: KVM requires hardware virtualization extensions: Intel VT-x (vmx flag) for Intel CPUs or AMD-V (svm flag) for AMD CPUs. These must be enabled in the system BIOS/UEFI.

Q33. How do you check if KVM is supported on a Proxmox host?

Answer: Run: `egrep '(vmx|svm)' /proc/cpuinfo`. If the command returns output, the CPU supports hardware virtualization. You can also run: `kvm-ok` or check `dmesg` for KVM module messages.

Q34. What is the difference between CPU types in Proxmox VE VM settings?

Answer: Proxmox VE allows selecting the CPU model exposed to VMs. 'host' passes through the actual CPU model and flags for best performance. 'kvm64' offers a generic baseline compatible with migration. Custom models can be configured for specific workloads.

Q35. What is QEMU in relation to KVM?

Answer: QEMU (Quick EMUlator) is the userspace emulator that Proxmox VE uses alongside KVM. QEMU provides device emulation (network, disk, USB, etc.) while KVM handles the actual CPU/memory virtualization in kernel space.

Q36. What is a VM snapshot in Proxmox VE?

Answer: A VM snapshot captures the entire state of a VM (disk, memory, and configuration) at a specific point in time. Snapshots allow rolling back to a previous state instantly without restoring from a full backup.

Q37. What storage types support VM snapshots in Proxmox VE?

Answer: VM snapshots require storage that supports snapshots, such as LVM-Thin, ZFS, Ceph RBD, or Proxmox Backup Server storage. Traditional LVM and NFS directory storage do not support live VM snapshots.

Q38. What is the difference between a snapshot and a clone in Proxmox VE?

Answer: A snapshot is a point-in-time state save of a VM used for rollback. A clone creates an independent copy of a VM. Linked clones share base storage with the original (space-efficient), while full clones are completely independent.

Q39. How do you create a VM template in Proxmox VE?

Answer: Configure a VM to the desired state, then right-click the VM in the web GUI and select 'Convert to Template', or run: `qm template`. Templates are read-only and used to quickly deploy new VMs via cloning.

Q40. What is Cloud-Init support in Proxmox VE?

Answer: Cloud-Init is a standard for configuring cloud VMs on first boot. Proxmox VE supports Cloud-Init by attaching a special Cloud-Init drive to a VM, allowing automated configuration of hostname, IP, DNS, SSH keys, and user accounts.

Q41. What is the Proxmox VE VM configuration file format?

Answer: VM configurations are stored as plain-text files at `/etc/pve/nodes//qemu-server/.conf`. They use a key=value format and contain all VM settings including CPU, memory, network, disk definitions, and boot order.

Q42. What is virtio in Proxmox VE?

Answer: VirtIO is a paravirtualized device driver framework that provides high-performance network (virtio-net), block (virtio-blk/virtio-scsi), and other devices to guest VMs. It avoids full hardware emulation overhead for significantly better I/O performance.

Q43. What is the recommended disk controller for VMs in Proxmox VE?

Answer: VirtIO SCSI (virtio-scsi-pci) is the recommended disk controller for Linux VMs in Proxmox VE as it provides the best performance. For Windows VMs, you may need to install VirtIO drivers first before switching from emulated IDE/SATA.

Q44. How do you pass through a physical GPU to a VM in Proxmox VE?

Answer: GPU passthrough (PCIe passthrough) requires enabling IOMMU in the BIOS (Intel VT-d or AMD-Vi), adding iommu kernel parameters, binding the GPU to the `vfio-pci` driver, and adding the PCIe device to the VM configuration.

Q45. What is live migration in Proxmox VE?

Answer: Live migration moves a running VM from one cluster node to another without downtime or service interruption. It requires shared storage (so disk data is accessible from both nodes) and sufficient network bandwidth on the migration network.

Q46. What is offline migration in Proxmox VE?

Answer: Offline migration moves a stopped VM from one node to another. Unlike live migration, it does not require shared storage – the VM disk image is physically transferred between nodes. The VM experiences downtime during the migration.

Q47. What is the qm command used for?

Answer: qm (QEMU Manager) is the Proxmox VE command-line tool for managing KVM virtual machines. Common commands: qm list, qm start , qm stop , qm shutdown , qm migrate, qm snapshot, qm config .

Q48. What is UEFI/OVMF in Proxmox VE VMs?

Answer: OVMF (Open Virtual Machine Firmware) is the UEFI firmware implementation for KVM VMs in Proxmox VE. It enables UEFI boot (required for Secure Boot, GPT disks, and some modern OSes). It stores boot settings in a separate EFI disk image.

Q49. What is Secure Boot in Proxmox VE?

Answer: Secure Boot is a UEFI feature that validates OS bootloader signatures against trusted keys. Proxmox VE supports Secure Boot for VMs using OVMF firmware and the Machine type q35. It requires enrolling the appropriate keys.

Q50. What is the q35 machine type in Proxmox VE?

Answer: q35 is a modern QEMU machine type that emulates Intel ICH9 chipset with PCIe support, required for features like UEFI/Secure Boot, NVMe emulation, and some USB 3.0 support. The alternative is i440FX, an older compatible machine type.

Q51. How do you resize a VM disk in Proxmox VE?

Answer: In the web GUI: VM □ Hardware □ select disk □ Disk Action □ Resize Disk. Via CLI: qm resize +G. Note: this only increases the disk size; the partition/filesystem inside the guest must be resized separately.

Q52. What is VirtIO Balloon driver?

Answer: The VirtIO Balloon driver allows dynamic memory adjustment of a running VM. The host can instruct the guest to release memory back to the host or reclaim it, helping balance memory usage across multiple VMs without rebooting.

Q53. What is CPU pinning in Proxmox VE?

Answer: CPU pinning (affinity) assigns specific host CPU cores to a VM's vCPUs using the `cpuaffinity` option. This can improve performance for latency-sensitive workloads by reducing CPU cache misses and scheduler interference.

Q54. What are NUMA settings in Proxmox VE VMs?

Answer: NUMA (Non-Uniform Memory Access) settings allow Proxmox VE to expose virtual NUMA topology to VMs. Enabling NUMA for large VMs on multi-socket hosts can improve performance by aligning VM memory and CPU access patterns.

Q55. How do you configure hugepages for VMs in Proxmox VE?

Answer: Hugepages allocate larger memory pages (2MB or 1GB) for VMs, reducing TLB (Translation Lookaside Buffer) misses and improving memory performance. Configure via `/etc/sysctl.conf` (`vm.nr_hugepages`) and enable in VM settings.

Q56. What is the Proxmox VE firewall?

Answer: Proxmox VE has a built-in firewall based on `iptables`/`nftables`. It supports datacenter-level, node-level, and per-VM/container rules. Rules can filter traffic by IP, port, protocol, and direction (in/out). It also supports security groups and IP sets.

Q57. What is a VM template vs a VM clone?

Answer: A template is a read-only master image used to create multiple identical VMs. A clone is a new, writable VM created from a template or another VM — either linked (shares template storage) or full (independent copy).

Q58. How do you enable nested virtualization in Proxmox VE?

Answer: Add `nested=1` to the KVM module options: `echo 'options kvm-intel nested=1' > /etc/modprobe.d/kvm-intel.conf` (Intel) or `kvm-amd` for AMD. Then set the VM CPU type to 'host' or include the `vmx`/`svm` flag.

Q59. What is the purpose of the qemu-guest-agent?

Answer: The `qemu-guest-agent` is a daemon running inside the VM guest. It enables the host to communicate with the guest for graceful shutdown, filesystem freeze before snapshots, IP address reporting, and file system trimming.

Q60. What are PCI device passthrough requirements in Proxmox VE?

Answer: PCIe passthrough requires: IOMMU enabled in BIOS (VT-d/AMD-Vi), IOMMU kernel parameter (`intel_iommu=on` or `amd_iommu=on`), device bound to `vfio-pci` driver, and ideally the device in its own IOMMU group to avoid ACS issues.

Q61. What is 'noVNC' in Proxmox VE?

Answer: noVNC is the browser-based VNC console used in Proxmox VE's web GUI. It allows accessing a VM's graphical console directly in the browser without installing additional client software. Proxmox VE also supports SPICE protocol for enhanced display.

Q62. What is SPICE in Proxmox VE?

Answer: SPICE (Simple Protocol for Independent Computing Environments) is a remote display protocol available for Proxmox VE VMs. It supports higher resolutions, audio, USB redirection, and clipboard sharing. Requires the SPICE client (`virt-viewer`) on the client side.

Q63. How do you import a disk image into a Proxmox VE VM?

Answer: Use the `qm importdisk` command: `qm importdisk [--format ].` After importing, attach the disk to the VM via the Hardware tab in the GUI or using `qm set`.

Q64. What formats are supported for VM disk import in Proxmox VE?

Answer: Proxmox VE supports importing disk images in raw, qcow2, vmdk, vhd, and vhdx formats via `qm importdisk`. The images are converted to the target storage format during import.

Q65. How do you convert a VMware VM for use in Proxmox VE?

Answer: Export the VMware VM as OVF/OVA. Then import using `qm importovf` or convert the VMDK disk with `qemu-img convert -f vmdk -O qcow2 source.vmdk dest.qcow2`, and import the converted image using `qm importdisk`.

Q66. What is the difference between qcow2 and raw disk formats?

Answer: Raw is a simple flat image with no overhead, offering maximum performance but no built-in features. qcow2 (QEMU Copy-On-Write v2) supports snapshots, compression, encryption, and sparse allocation but with slightly more overhead.

Q67. What is storage migration in Proxmox VE?

Answer: Storage migration moves a VM's disk images between different storage backends. It can be done online (for running VMs with compatible storage) or offline. Use the web GUI (Migrate) or: `qm move_disk`.

Q68. How do you configure a serial console for a VM?

Answer: Add a serial device to the VM (Hardware ▢ Add ▢ Serial Port), then configure the guest OS to use the serial device as console (e.g., add `console=ttyS0` to kernel boot parameters in Linux). Access via: `qm terminal`.

Q69. What is the CPU limit setting in Proxmox VE VMs?

Answer: The CPU limit caps how much host CPU time a VM can consume. A value of 1 means one vCPU's worth of CPU time maximum, regardless of how many vCPUs the VM has. Useful for preventing noisy-neighbor problems.

Q70. What is 'CPU units' in Proxmox VE?

Answer: CPU units set the relative CPU scheduling priority for a VM using the Linux cgroups `cpu.shares` mechanism. Higher values get more CPU time during contention. Default is 1024. Range is 8–262144.

Section 3: LXC Containers (Q71–Q100)

Q71. What is LXC?

Answer: LXC (Linux Containers) is an OS-level virtualization method that uses Linux kernel features (namespaces, cgroups) to run isolated Linux environments. LXC containers share the host kernel, making them faster and more resource-efficient than VMs.

Q72. What is the `pct` command used for?

Answer: `pct` (Proxmox Container Toolkit) is the command-line utility for managing LXC containers in Proxmox VE. Commands include: `pct list`, `pct create`, `pct start`, `pct stop`, `pct exec --`, `pct snapshot`.

Q73. Where are LXC container configurations stored?

Answer: LXC container configurations are stored as text files at `/etc/pve/lxc/.conf` on the Proxmox VE node. The format is similar to VM configurations and contains network, storage, resource, and security settings.

Q74. What is a container template in Proxmox VE?

Answer: A container template is a compressed OS root filesystem (tar.gz or tar.xz) used as the base for creating new LXC containers. Templates for popular distributions are available via the Proxmox template repository (pveam command).

Q75. How do you download container templates in Proxmox VE?

Answer: Use the pveam command: `pveam update` (refresh list), `pveam available` (list available templates), `pveam download` (download). Templates can also be downloaded directly via the web GUI under Storage ▢ Content ▢ Templates.

Q76. What are the key differences between privileged and unprivileged containers?

Answer: Privileged containers run as root inside with UIDs mapped 1:1 to host UIDs — easier to configure but less secure. Unprivileged containers use UID/GID mapping, shifting container UIDs to unprivileged ranges on the host, providing better isolation.

Q77. What is UID/GID mapping in unprivileged LXC containers?

Answer: UID/GID mapping shifts container user IDs to high ranges on the host (e.g., container UID 0 becomes host UID 100000). This means container root has no special privileges on the host, greatly improving security.

Q78. Can LXC containers run Docker?

Answer: Yes, but with caveats. Docker in LXC requires enabling nesting (features: nesting=1 in container config) and may require keyctl. Unprivileged containers need additional capabilities. Using a VM is more secure and reliable for Docker workloads.

Q79. What is a bind mount in LXC containers?

Answer: A bind mount exposes a directory from the Proxmox host inside the container. Configured via `mp: /host/path,mp=/container/path` in the container config. Useful for sharing data but reduces isolation (especially in privileged containers).

Q80. How do you limit CPU usage for an LXC container?

Answer: Set `cpulimit` (e.g., `cpulimit: 2`) to limit CPU cores, `cpuunits` for relative weight, and `cores` to set the number of visible CPU cores. These use Linux cgroups under the hood.

Q81. How is memory limited for LXC containers?

Answer: Set memory: for RAM limit and swap: for swap limit in the container configuration. These are enforced via cgroups memory subsystem.

Q82. What are container features in Proxmox VE?

Answer: Container features enable special kernel capabilities: nesting (containers within containers), keyctl (kernel key management), fuse (FUSE filesystem), mknod, mount (specific fs types), and SMB mounts. Configured in features: key=val in container config.

Q83. How do you take a snapshot of an LXC container?

Answer: Via GUI: right-click container □ Snapshot □ Take Snapshot. Via CLI: `pct snapshot --description 'note'`. Snapshots require storage that supports them (ZFS, LVM-Thin, Ceph RBD).

Q84. Can LXC containers be live migrated?

Answer: LXC containers support live migration (CRIU-based checkpoint/restore) in Proxmox VE, but it has limitations and requires shared storage. Offline migration (stop, move, start) is more reliable and commonly used.

Q85. What network interface types are available for LXC containers?

Answer: LXC containers support veth (virtual Ethernet) interfaces — the most common type. Each container veth is connected to a host-side bridge (e.g., `vmbr0`) using a pair of virtual Ethernet devices.

Q86. How do you access the console of an LXC container?

Answer: Via GUI: click the container □ Console. Via CLI: `pct enter` (opens a shell inside the container) or `pct console` (attaches to the container console). You can also use `lxc-attach -n`.

Q87. What is the difference between `pct enter` and `pct exec`?

Answer: `pct enter` starts an interactive shell session inside the container. `pct exec --` runs a specific command inside the container and returns the output, useful for scripting.

Q88. How do you resize an LXC container's rootfs?

Answer: Use: `pct resize rootfs G`. The rootfs partition is resized, and the filesystem inside (usually `ext4`) is automatically expanded. No downtime is required for most storage backends.

Q89. What is `vzdump`'s behavior with LXC containers?

Answer: `vzdump` backs up LXC containers by archiving the root filesystem. It supports snapshot mode (if storage supports it), suspend mode, and stop mode. Container backups include all files in the rootfs and the container configuration.

Q90. What AppArmor profile is used for LXC containers?

Answer: Proxmox VE applies the `lxc-container-default` AppArmor profile to containers by default, which provides MAC (Mandatory Access Control) security restrictions. Privileged containers may need `apparmor=unconfined` for certain workloads.

Q91. Can Windows run in an LXC container?

Answer: No. LXC containers share the Linux host kernel, so only Linux-based operating systems can run as containers. For Windows, you must use a KVM virtual machine which provides full hardware virtualization.

Q92. What is the `/dev/lxc` directory?

Answer: The `/dev/lxc` directory on the Proxmox VE host contains the control interfaces for running LXC containers, including console and monitor sockets. These are managed by the LXC runtime and should not be modified manually.

Q93. How do you clone an LXC container?

Answer: Via GUI: right-click container ▢ Clone. Via CLI: `pct clone [--full] [--storage ]`. Without `--full`, a linked clone is created (requires snapshot-capable storage). `--full` creates an independent copy.

Q94. What happens to container IPs during migration?

Answer: IP addresses configured in the container are migrated along with the container. However, network connectivity depends on having the same bridge and network configuration on the target node. Static IPs should work seamlessly.

Q95. How do you run a container in a resource pool?

Answer: When creating a container, assign it to a resource pool via `--pool` option in `pct create` or in the GUI. Resource pools group containers and VMs for permission management and organization.

Q96. What is the `/etc/network/interfaces` file in an LXC container?

Answer: Inside the container, `/etc/network/interfaces` (Debian/Ubuntu) or equivalent network config manages the container's own network settings. However, in Proxmox VE, the network interface assignment (bridge, VLAN, MAC) is controlled by the host-side container config.

Q97. How do you enable VLAN tagging for an LXC container network?

Answer: Add `tag=` to the network interface entry in the container configuration: `net0: name=eth0,bridge=vmbr0,tag=100`. The container will receive packets tagged with VLAN 100 on the bridge.

Q98. What is the impact of `nesting=1` on container security?

Answer: Enabling nesting allows running containers within containers (or Docker inside LXC). It grants additional capabilities (`sys_admin`) and partially relaxes AppArmor profiles, reducing isolation. Use only in trusted environments.

Q99. How do you migrate a container to another node?

Answer: Via GUI: right-click container ▢ Migrate. Via CLI: `pct migrate [--online]`. For offline migration, the container is stopped, disk data is transferred via `rsync/vzdump`, and the container is started on the target.

Q100. What storage types support LXC containers?

Answer: LXC containers can reside on directory storage (`ext4/xfs`), LVM, LVM-Thin (recommended for snapshots), ZFS, and Ceph RBD. NFS can be used for backup storage but is not recommended as primary container storage due to filesystem limitations.

Section 4: Networking (Q101–Q140)

Q101. What is a Linux bridge in Proxmox VE?

Answer: A Linux bridge (e.g., `vmbr0`) is a virtual network switch on the host that connects VMs/containers to the physical network or to each other. Proxmox VE creates `vmbr0` during installation and bridges it to the physical NIC.

Q102. What is the difference between `vmbr0` and a physical NIC?

Answer: `vmbr0` is a virtual bridge interface managed by the Linux kernel. The physical NIC (e.g., `eno1`) is the actual hardware interface. `vmbr0` acts as a switch — VMs connect to it and packets are forwarded to/from the physical NIC via bridging.

Q103. What is a VLAN-aware bridge in Proxmox VE?

Answer: A VLAN-aware bridge is a single bridge configured with `vlan_aware: yes` that supports multiple VLANs. VMs and containers can be tagged with a VLAN ID, allowing trunk-port behavior on a single bridge without creating separate bridges per VLAN.

Q104. What is Open vSwitch (OVS) in Proxmox VE?

Answer: Open vSwitch is an advanced software-defined networking switch supported in Proxmox VE. OVS supports VLAN trunking, bonding, VXLAN, OpenFlow, and more. It requires installing the `openvswitch` packages and is an alternative to Linux bridges.

Q105. What is network bonding in Proxmox VE?

Answer: Network bonding (link aggregation) combines multiple physical NICs into a single logical interface for increased bandwidth or redundancy. Modes include LACP (802.3ad), active-backup, balance-rr, and others. Configured in `/etc/network/interfaces`.

Q106. What is the network configuration file in Proxmox VE?

Answer: Proxmox VE network configuration is stored in `/etc/network/interfaces` on each node. In a cluster, a pending changes file is maintained at `/etc/pve/nodes//config`. The GUI shows pending changes that must be applied (revert or apply button).

Q107. What is a VXLAN in Proxmox VE?

Answer: VXLAN (Virtual Extensible LAN) is a network overlay technology that encapsulates Layer 2 frames in UDP packets, allowing VMs on different physical hosts/subnets to appear on the same L2 network. Used in SDN configurations in Proxmox VE.

Q108. What is Proxmox VE SDN (Software Defined Networking)?

Answer: Proxmox VE SDN is a built-in network management framework introduced in Proxmox VE 7. It provides zones, VNets, subnets, and DHCP/DNS management for defining virtual networks, supporting VXLAN, EVPN, and Simple zones.

Q109. What is the difference between SDN Simple Zone and VXLAN Zone?

Answer: Simple Zone uses existing bridges/VLANs and just organizes network configurations. VXLAN Zone creates overlay networks using VXLAN encapsulation, enabling L2 connectivity between VMs on different nodes without a shared physical L2 network.

Q110. What is EVPN in the context of Proxmox SDN?

Answer: EVPN (Ethernet VPN) with BGP is an advanced SDN zone type in Proxmox VE for multi-node L2 overlay networks with distributed routing. It uses FRRouting (frr) daemon for BGP-based MAC/IP route distribution.

Q111. What is a static route in Proxmox VE?

Answer: A static route is a manually configured routing table entry that directs traffic for a specific network via a defined gateway. Configured in `/etc/network/interfaces` with `up ip route add ...` or via the `iproute2` post-up commands.

Q112. What MTU setting is recommended for VXLAN in Proxmox?

Answer: VXLAN adds 50 bytes of encapsulation overhead. If the physical network MTU is 1500, VXLAN interfaces should be configured with MTU 1450 to avoid fragmentation. The underlying physical NICs should support jumbo frames (9000 MTU) in production.

Q113. How do you configure a static IP for a Proxmox VE host?

Answer: Edit `/etc/network/interfaces` and set: `iface vmbr0 inet static, address , gateway`. Then apply with: `ifreload -a` or `systemctl restart networking`. Alternatively, configure during installation via the installer GUI.

Q114. What is a trunk port vs access port in Proxmox VE networking?

Answer: A trunk port carries multiple VLANs (tagged traffic). An access port carries a single VLAN (untagged). In Proxmox VE, VM network interfaces can be set with trunks (comma-separated VLAN IDs) or a single tag for access port behavior.

Q115. What is the purpose of vmbr0 having no IP when used as a pure bridge?

Answer: When `vmbr0` is used purely as a switch bridge for VMs, it needs no IP. The host can use a separate management IP on a different bridge or interface. This isolates host management traffic from VM traffic.

Q116. What is MAC address management in Proxmox VE?

Answer: Proxmox VE auto-assigns a unique MAC address to each VM/container network interface in the range `52:54:00:XX:XX:XX` (QEMU) or `BC:24:11:XX:XX:XX` (Proxmox). MACs can be manually specified. They persist in the VM/CT config file.

Q117. What is IP aliasing in Proxmox VE network config?

Answer: IP aliasing adds multiple IP addresses to a single network interface using sub-interfaces (e.g., `vmbr0:0`). Common in Proxmox for assigning additional IPs to the host for services or routing without creating additional bridges.

Q118. How do you add a second network interface to a running VM?

Answer: Via GUI: VM □ Hardware □ Add □ Network Device. Via CLI: `qm set --net virtio,bridge=vmbr0`. The interface is hotplugged if the VM is running and the guest OS supports virtio hotplug.

Q119. What is firewall zone management in Proxmox VE?

Answer: Proxmox VE firewall supports three levels: Datacenter (global rules/defaults), Node (host-level rules), and VM/CT (per-instance rules). Rules at each level are merged, with higher levels applying to all lower levels.

Q120. What are IPSets in the Proxmox VE firewall?

Answer: IPSets are named groups of IP addresses/CIDR ranges used in firewall rules. Instead of specifying individual IPs, you reference the IPSet name (e.g., `+management`), making firewall management more organized and scalable.

Q121. What is a security group in Proxmox VE firewall?

Answer: A security group is a named set of firewall rules that can be applied to multiple VMs/containers. This allows reusing common rule sets (e.g., 'web-server' group with HTTP/HTTPS rules) across many instances efficiently.

Q122. What is the firewall policy in Proxmox VE?

Answer: The firewall policy defines default actions for INPUT, OUTPUT, and FORWARD traffic. Proxmox VE defaults to DROP for INPUT and OUTPUT, and ACCEPT for FORWARD at the datacenter level when the firewall is enabled for a VM/CT.

Q123. What is an OVS IntPort in Proxmox VE?

Answer: An OVS IntPort (Internal Port) is a virtual network interface on an Open vSwitch bridge that allows the Proxmox host itself to participate in OVS networking. It behaves like a regular Linux interface connected to the OVS bridge.

Q124. How do you configure LACP bonding with Open vSwitch?

Answer: Create an OVS Bond interface in `/etc/network/interfaces` with `ovs_type OVSBond`, `bond_mode active-backup` or `balance-tcp` (for LACP), `ovs_options bond_mode=balance-tcp lacp=active`. The bond members are physical NICs.

Q125. What is the Proxmox VE firewall log?

Answer: Proxmox VE firewall logs dropped/rejected packets to `/var/log/kern.log` and via syslog. Logging verbosity can be configured per rule (`log=nolog`, `alert`, `crit`, `err`, `warning`, `notice`, `info`, `debug`) in the firewall rule configuration.

Q126. What is Proxmox VE's approach to IPv6?

Answer: Proxmox VE fully supports IPv6 for host management, VMs, and containers. Configure IPv6 addresses on bridges in `/etc/network/interfaces` using `inet6` stanzas. Guest VMs can receive IPv6 via SLAAC, DHCPv6, or static assignment.

Q127. What is bridge-filter vs ebtables in Proxmox VE?

Answer: Proxmox VE uses ebtables (bridge-level firewalling) as part of its firewall implementation to filter traffic at Layer 2 before it reaches the iptables/nftables Layer 3 rules. This allows filtering by MAC address and VLAN tags.

Q128. What is the purpose of the cluster network (corosync ring) in Proxmox?

Answer: The corosync cluster network carries heartbeat and cluster communication traffic between nodes. It should be a dedicated low-latency, reliable network — ideally a direct link or separate physical network from the management and storage networks.

Q129. What is the storage/migration network in Proxmox VE?

Answer: A dedicated storage/migration network is used for VM live migration traffic and storage replication. Separating this traffic from management and VM networks prevents saturation of production interfaces during migrations.

Q130. How many NICs are recommended for a production Proxmox VE deployment?

Answer: A best-practice production deployment uses at least 4 NICs: 1 for management (host access), 1 for VM/client traffic, 1 for cluster/corosync heartbeat, and 1 for storage/migration. Bonded pairs double each for redundancy.

Q131. What is promiscuous mode in Proxmox VE networking?

Answer: Promiscuous mode on a bridge or NIC allows it to receive all packets regardless of destination MAC. Required for some nested virtualization scenarios, VLAN setups, or network monitoring VMs. Enabled via: `ip link set promisc on`.

Q132. What is the VLAN native/pvid setting in Proxmox SDN?

Answer: The PVID (Port VLAN ID) sets the default VLAN for untagged traffic on a VLAN-aware bridge port. Traffic arriving untagged is assigned to the PVID. Leaving traffic tagged with the PVID is stripped of the tag before delivery.

Q133. How do you apply pending network changes in Proxmox VE?

Answer: Pending network changes (shown with a yellow asterisk in the GUI) are applied by clicking Apply Configuration in the Network panel, or via CLI: `ifreload -a`. This applies changes without requiring a full system reboot.

Q134. What is trunks option in VM network config?

Answer: The trunks parameter in a VM's network interface configuration (`net0: ...,trunks=100;200`) allows the VM to receive and send VLAN-tagged frames for the specified VLAN IDs. The VM's OS must handle VLAN tagging itself.

Q135. What is DHCP relay in Proxmox SDN?

Answer: Proxmox SDN supports DHCP relay (via `dnsmasq`), forwarding DHCP requests from VMs in SDN subnets to an external DHCP server. This allows centralized DHCP management without running a server in every subnet.

Q136. What is a Proxmox VE bond slave?

Answer: A bond slave is a physical NIC that is a member of a bonding interface. Slave NICs are configured without IPs; the bond master interface carries the IP. Slave health is monitored using MII or ARP monitoring.

Q137. What is ARP monitoring in NIC bonding?

Answer: ARP monitoring sends ARP requests to specified IP targets (`arp_ip_target`) to verify connectivity of bond slaves. If an ARP reply is not received, the slave is considered failed. An alternative to MII link-state monitoring.

Q138. What is MII monitoring in NIC bonding?

Answer: MII (Media Independent Interface) monitoring checks the physical link state of bond slave NICs at a configured interval (miimon in milliseconds). If the link goes down (cable unplugged), the slave is removed from the bond.

Q139. How do you add a VLAN interface to the Proxmox host?

Answer: Create a VLAN sub-interface in `/etc/network/interfaces`: `iface eno1.100 inet static with address/gateway`. Or attach the VLAN to a bridge: `iface vmbr100 inet static with bridge-ports eno1.100`. Apply with `ifreload -a`.

Q140. What is Proxmox VE's default behavior for VM network traffic filtering?

Answer: By default (without firewall rules), Proxmox VE passes all network traffic for VMs through the bridge without filtering. When the firewall is enabled for a VM, default policies apply. MAC/IP spoofing protection can be enabled per interface.

Section 5: Storage Management (Q141–Q180)

Q141. What is thin provisioning in Proxmox VE?

Answer: Thin provisioning allocates disk space on demand rather than pre-allocating the full disk size. This allows VMs to be provisioned with large disks while only consuming actual used space on the storage backend, enabling storage overcommitment.

Q142. What is the difference between qcow2 and raw on LVM storage?

Answer: On LVM storage, raw format is used (logical volumes are raw block devices). qcow2 is used on directory/file-based storage. LVM provides better performance for VMs but lacks snapshot capabilities without LVM-Thin.

Q143. What is a Proxmox Backup Server (PBS)?

Answer: Proxmox Backup Server is a dedicated enterprise backup solution tightly integrated with Proxmox VE. It supports incremental backups (only changed blocks are transferred), deduplication, encryption, compression, and efficient restoration.

Q144. How does Proxmox Backup Server incremental backup work?

Answer: PBS tracks changed blocks since the last backup using dirty bitmaps (KVM feature). On subsequent backups, only changed blocks are sent to PBS, dramatically reducing backup duration and network/storage usage.

Q145. What is deduplication in Proxmox Backup Server?

Answer: Deduplication in PBS stores identical data chunks only once across all backups. When multiple VMs or snapshots share identical data blocks, PBS stores the block once and references it multiple times, saving significant storage space.

Q146. What is a ZFS pool in Proxmox VE?

Answer: A ZFS pool (zpool) is the top-level storage container in ZFS. It consists of one or more vdevs (virtual devices) which can be single disks, mirrors, or RAID-Z arrays. The pool provides the storage namespace for ZFS datasets and volumes.

Q147. What is RAID-Z in ZFS?

Answer: RAID-Z is ZFS's software RAID implementation. RAID-Z1 (single parity, like RAID-5), RAID-Z2 (double parity, like RAID-6), and RAID-Z3 (triple parity) provide redundancy. ZFS RAID-Z avoids the RAID-5 write hole problem through copy-on-write.

Q148. What is ZFS ARC?

Answer: ARC (Adaptive Replacement Cache) is ZFS's main memory cache for frequently accessed data. It uses available RAM to cache disk reads, significantly improving read performance. ARC size is tunable via `zfs_arc_max` in `/etc/modprobe.d/zfs.conf`.

Q149. What is ZFS L2ARC?

Answer: L2ARC (Level 2 ARC) extends the ZFS read cache to fast SSDs. Data evicted from RAM ARC is stored in L2ARC for subsequent reads. L2ARC is a second-level read cache and does not provide redundancy.

Q150. What is ZFS SLOG/ZIL?

Answer: ZIL (ZFS Intent Log) records synchronous write transactions for crash recovery. SLOG (Separate Intent Log) moves the ZIL to a dedicated fast SSD, dramatically improving synchronous write performance (important for NFS, iSCSI, databases).

Q151. What is ZFS compression in Proxmox VE?

Answer: ZFS compression (typically lz4 or zstd) transparently compresses data before writing to disk. lz4 is extremely fast with minimal CPU overhead. Compression often improves I/O performance by reducing the amount of data written to slower disks.

Q152. What is ZFS deduplication and its trade-offs?

Answer: ZFS deduplication stores identical data blocks only once, saving storage for redundant data. Trade-off: requires very large amounts of RAM for the dedup table (DDT) — ~5GB RAM per TB of deduplicated data. CPU overhead is also significant.

Q153. What is Ceph OSD?

Answer: A Ceph OSD (Object Storage Daemon) manages a single storage device (disk) in the Ceph cluster. Each OSD stores data, handles replication, recovery, and backfilling. A typical Ceph cluster has one OSD per physical disk.

Q154. What is a Ceph Monitor (MON)?

Answer: Ceph Monitors maintain the cluster map (CRUSH map, OSD map, PG map) and handle quorum for cluster decisions. A minimum of 3 monitors is required for a production HA Ceph cluster.

Q155. What is a Ceph MDS?

Answer: Ceph MDS (Metadata Server) handles the filesystem metadata for CephFS (the Ceph distributed filesystem). MDS is not required for Ceph block storage (RBD) or object storage (RGW) deployments.

Q156. What is the Ceph CRUSH map?

Answer: The CRUSH (Controlled Replication Under Scalable Hashing) map defines how Ceph distributes and replicates data across OSDs based on the physical topology (hosts, racks, DCs). It ensures data is spread across failure domains.

Q157. What is a Ceph placement group (PG)?

Answer: A Placement Group (PG) is a logical partition of a Ceph pool. Data (objects) are first mapped to a PG, then the PG is mapped to OSDs via CRUSH. PG count affects performance and rebalancing behavior.

Q158. What is a Ceph pool replication size?

Answer: Replication size defines how many copies of each data object are stored in a Ceph pool. A size of 3 means 3 copies across 3 different OSDs. `min_size` defines the minimum replicas required for writes to succeed (usually 2).

Q159. What is Ceph RBD?

Answer: Ceph RBD (RADOS Block Device) provides distributed block storage for VMs in Proxmox VE. VM disk images are stored as RBD images, striped across the Ceph cluster. RBD supports snapshots, cloning, and thin provisioning.

Q160. How do you add Ceph storage to Proxmox VE?

Answer: Install Ceph on Proxmox nodes (Datacenter ▢ Ceph ▢ Install), create monitors, create OSDs from disks, create a pool, then add the pool as storage in Proxmox (Datacenter ▢ Storage ▢ Add ▢ RBD).

Q161. What is NFS storage in Proxmox VE?

Answer: NFS (Network File System) storage mounts a remote NFS share on Proxmox VE nodes for storing VM disk images, ISOs, backups, and container templates. NFS is easy to set up but has limitations for VM disk images (no native snapshot support).

Q162. What is CIFS/SMB storage in Proxmox VE?

Answer: Proxmox VE can mount Windows CIFS/SMB shares for backup and ISO storage. It's configured similarly to NFS with server, share, credentials, and domain options. SMB storage is generally used for backup targets, not primary VM storage.

Q163. What is iSCSI storage in Proxmox VE?

Answer: iSCSI (Internet Small Computer System Interface) presents remote storage as block devices over a network. Proxmox VE supports using iSCSI LUNs as VM disks (directly or via LVM over iSCSI). It provides block-level access with good performance.

Q164. What is Proxmox VE's storage content types?

Answer: Storage content types define what can be stored: images (VM disks), rootdir (container rootfs), vztmpl (container templates), iso (ISO images), backup (vzdump backups), snippets (config files, hook scripts). Not all storage types support all content types.

Q165. What is the recommended storage for VM disks in production?

Answer: For production, Ceph RBD (distributed, HA, no single point of failure) or ZFS with RAID-Z (local high-performance) are recommended. LVM-Thin is also widely used for local storage. NFS is suitable for backups and shared ISOs but not ideal for VM disks.

Q166. What is storage replication in Proxmox VE?

Answer: Storage replication (for ZFS storage) synchronizes VM disk data between cluster nodes at configurable intervals (minimum 1 minute) using ZFS send/receive. It provides a near-real-time copy on another node, enabling fast failover without shared storage.

Q167. What is the difference between shared and local storage in Proxmox clusters?

Answer: Shared storage (NFS, Ceph, iSCSI) is accessible from all cluster nodes simultaneously, enabling live migration and HA. Local storage (LVM, ZFS, directory) is only accessible on a single node, requiring offline migration or storage replication.

Q168. What is a storage content migration?

Answer: Storage content migration moves VM disk images between storage backends. For running VMs on shared storage, this can be done online (storage migration). For local storage, the VM typically needs to be stopped (or use replication).

Q169. How do you check ZFS pool health in Proxmox VE?

Answer: Run: `zpool status` (shows pool health, disk status, and errors) or `zpool list` (shows pool sizes and usage). In the GUI: Storage ▢ select ZFS storage ▢ Content shows usage. Errors show as DEGRADED or FAULTED status.

Q170. What is ZFS scrub?

Answer: A ZFS scrub verifies all data in the pool against stored checksums, detecting and correcting silent data corruption (bit rot). Schedule regular scrubs (monthly/quarterly) via a cron job: `zpool scrub`. Scrubs can run while the pool is in use.

Q171. What is the autoexpand feature in ZFS?

Answer: ZFS autoexpand (autoexpand=on) automatically expands a ZFS pool when the underlying disk/partition is enlarged. Useful when expanding virtual disks or replacing physical disks with larger ones in a mirror or RAID-Z.

Q172. What is Proxmox VE's backup retention policy?

Answer: Backup retention (configured per backup job) defines how many backups to keep: last N backups, keep daily/weekly/monthly/yearly backups. Older backups are automatically pruned based on the retention policy when new backups complete.

Q173. What compression algorithm is recommended for vzdump backups?

Answer: *zstd (Zstandard) is the recommended compression algorithm for vzdump backups. It offers excellent compression ratios with much faster compression/decompression speeds compared to gzip or lzo.*

Q174. What is sparse file support in Proxmox VE storage?

Answer: *Sparse files store only written (non-zero) data, with unwritten areas consuming no actual disk space. Directory-based storage (ext4, XFS) supports sparse files for VM disk images in raw and qcow2 format, saving disk space.*

Q175. How do you identify storage used by a specific VM?

Answer: *Use: qm config to see disk references, then check storage capacity with pvesm status. In the GUI: VM □ Hardware □ Disk shows storage location. You can also use: pvesm list to list all content on a storage.*

Q176. What is the 'discard' option for VM disks?

Answer: *Enabling discard (TRIM/UNMAP) allows the guest OS to inform the hypervisor when disk blocks are no longer in use. On thin-provisioned storage (LVM-Thin, ZFS), this reclaims unused space. Requires virtio-scsi with iothread or VirtIO BLK.*

Q177. What is IO throttling in Proxmox VE?

Answer: *IO throttling limits the maximum read/write IOPS and bandwidth for VM disk devices. Configured per disk via mbps_rd, mbps_wr, iops_rd, iops_wr settings in the VM configuration. Prevents VMs from monopolizing storage I/O.*

Q178. What is a Proxmox VE storage plugin?

Answer: *Proxmox VE uses a plugin-based storage architecture. Each storage type (dir, lvm, zfspool, rbd, nfs, cifs, iscsi, etc.) is a plugin. Custom storage plugins can be written and integrated, extending Proxmox VE to support additional storage backends.*

Q179. What is the /var/lib/vz directory in Proxmox VE?

Answer: *The /var/lib/vz directory is the root of the default local storage (directory-type). It contains subdirectories: images/ (VM disk images), dump/ (vzdump backups), template/ (ISOs and CT templates). It resides on the system disk.*

Q180. What is the pvesm command?

Answer: *pvesm* (Proxmox VE Storage Manager) is the CLI tool for managing storage.

Commands: *pvesm status* (list storage), *pvesm list* (content), *pvesm add/remove* (manage storage configs), *pvesm path* (resolve volume path).

Section 6: High Availability & Clustering (Q181–Q220)

Q181. What is Proxmox VE High Availability (HA)?

Answer: Proxmox VE HA automatically monitors VMs and containers on cluster nodes and restarts them on other nodes if their host fails. It uses the Corosync cluster stack, quorum, and watchdog fencing to ensure reliability.

Q182. What is quorum in a Proxmox VE cluster?

Answer: Quorum is the minimum number of cluster nodes that must agree (vote) for the cluster to function. With N nodes, quorum requires more than $N/2$ votes. Quorum prevents split-brain scenarios where two partitioned groups each think they're the active cluster.

Q183. What is a split-brain scenario?

Answer: Split-brain occurs when a cluster partition causes two isolated groups of nodes to each believe they are the only active part of the cluster. This can lead to data corruption. Proxmox VE prevents split-brain via quorum and fencing.

Q184. What is fencing in Proxmox VE HA?

Answer: Fencing forcibly powers off or isolates a failed node to ensure it cannot corrupt data before HA restarts its VMs on other nodes. Proxmox VE uses watchdog-based fencing — if a node loses quorum, the watchdog triggers a system reset.

Q185. What is the HA watchdog in Proxmox VE?

Answer: The Proxmox VE HA watchdog is a kernel watchdog device (e.g., `/dev/watchdog`). It must be regularly fed by the HA manager. If a node loses cluster connectivity, the watchdog is not fed and triggers a hardware reset, fencing the node.

Q186. What is a QDevice in Proxmox VE clusters?

Answer: A QDevice is an external tie-breaker that provides an additional vote to a Proxmox cluster. Essential for 2-node clusters where one node failure would lose quorum. QDevice runs on a small external server (Raspberry Pi, VM, etc.) using `corosync-qdevice`.

Q187. What is the minimum node count for HA in Proxmox VE?

Answer: Three nodes is the recommended minimum for a proper HA cluster. With three nodes, losing one node still maintains quorum (2/3 votes). A 2-node cluster requires a QDevice to maintain quorum when one node fails.

Q188. What HA states can a resource be in?

Answer: HA resources cycle through states: started, stopped, disabled, ignored, error, freeze, recovery. The HA manager reconciles the current state with the requested state (started/stopped) and takes actions accordingly.

Q189. What is the HA group in Proxmox VE?

Answer: An HA group is a set of cluster nodes where a specific VM/container is allowed to run, with optional priority settings. Groups let administrators control which nodes host specific workloads and in what preferred order for HA failover.

Q190. What is the HA manager daemon?

Answer: The Proxmox VE HA manager runs as `pve-ha-lrm` (Local Resource Manager) on each node and `pve-ha-crm` (Cluster Resource Manager) globally. LRM manages local resources; CRM coordinates cluster-wide HA decisions and failover.

Q191. How long does HA failover typically take in Proxmox VE?

Answer: HA failover typically takes 1–2 minutes: ~30 seconds for the watchdog to trigger, then the CRM detects the failure, and finally VMs are started on surviving nodes. Actual time depends on VM boot time and cluster health.

Q192. What is migration type in HA configuration?

Answer: HA migration type can be live (online migration, no downtime, requires shared storage) or relocate (offline migration, requires VM stop). The type affects how the HA manager moves resources during maintenance or failure scenarios.

Q193. What happens to HA VMs during a cluster partition?

Answer: In a cluster partition, the minority partition (without quorum) will fence its nodes (watchdog reset). VMs on the minority side are restarted on the quorate (majority) partition to ensure service continuity.

Q194. What is Proxmox VE datacenter HA configuration file?

Answer: HA configuration is stored in the Proxmox cluster filesystem at `/etc/pve/ha/`. Key files: `groups.cfg` (HA groups), `resources.cfg` (HA-managed VMs/CTs with requested states), `manager_status` (current HA manager status).

Q195. What is node maintenance mode in Proxmox HA?

Answer: Maintenance mode (set via the GUI or `ha-manager set node --state maintenance`) gracefully migrates HA resources away from a node before maintenance. HA won't restart resources on the node while in maintenance mode.

Q196. What is the `corosync.conf` file?

Answer: `corosync.conf` (at `/etc/pve/corosync.conf` in a cluster) defines cluster configuration: node list with addresses, cluster name, quorum settings, and network transports (unicast `knet` or multicast). It is managed by Proxmox and should not be edited manually.

Q197. What ports does Corosync use?

Answer: Corosync uses UDP port 5405 (primary) by default for cluster communication. Additional ports (5406, 5407) may be used for redundant rings or `knet` links. These ports must be open between all cluster nodes.

Q198. What is the `pvecm nodes` command?

Answer: `pvecm nodes` lists all nodes in the Proxmox VE cluster, showing their node ID, name, and quorum votes. Use `pvecm status` for overall cluster health including quorum status and expected vs actual votes.

Q199. How do you safely remove a node from a Proxmox cluster?

Answer: Migrate all VMs/containers from the node, then run `pvecm delnode` from another cluster member. The node should be offline or reset before deletion to avoid data inconsistency. Remove residual config files from `/etc/pve/nodes/`.

Q200. What is storage replication vs HA in Proxmox VE?

Answer: Storage replication synchronizes disk data between nodes (ZFS send/receive) providing a local copy for fast failover. HA handles automatic VM restart on node failure. They complement each other: replication provides the data, HA handles the failover logic.

Q201. What is the cluster filesystem (pmxcfs) in Proxmox VE?

Answer: pmxcfs (Proxmox Cluster File System) is a database-driven filesystem stored in RAM and synchronized across all cluster nodes via Corosync. It stores configuration files (/etc/pve/) and requires quorum for write operations.

Q202. What is /etc/pve in Proxmox VE?

Answer: /etc/pve is the Proxmox cluster configuration directory mounted via pmxcfs. It contains VM/CT configs, storage configs, HA configs, firewall rules, cluster settings, and user data. Changes are automatically synchronized to all cluster nodes.

Q203. What happens to /etc/pve if a node loses quorum?

Answer: Without quorum, pmxcfs enters read-only mode on the minority node – configuration changes cannot be made. Running VMs continue operating, but no new administrative changes (start, stop, create) are possible until quorum is restored.

Q204. What is the expected votes setting in Corosync?

Answer: Expected votes (two_node: 1 in corosync.conf for 2-node clusters) defines the quorum expectation. Incorrect expected_votes can cause false quorum loss or false quorum state. Proxmox manages this automatically via pvecm.

Q205. How do you troubleshoot HA issues in Proxmox VE?

Answer: Check: `systemctl status pve-ha-lrm pve-ha-crm corosync`, `pvecm status` (quorum/nodes), `ha-manager status`, logs in `/var/log/syslog` and `/var/log/pve/tasks/`. Verify watchdog operation: `cat /dev/watchdog` (should trigger reset if not managed).

Q206. What is Proxmox VE cluster join?

Answer: Cluster join adds a new node to an existing cluster: `pvecm add`. The new node receives the cluster configuration, joins the Corosync ring, and its /etc/pve is synchronized. All cluster nodes must be reachable.

Q207. What is the cluster link in Proxmox VE?

Answer: Cluster links are the network interfaces used for Corosync cluster communication. Proxmox VE supports up to 8 redundant cluster links for resilience. Configured via `pvecm` or in `corosync.conf` as `ring0_addr`, `ring1_addr` etc.

Q208. What is the difference between HA priority and HA group order?

Answer: Priority within an HA group defines which node is preferred to run the resource (higher priority = preferred). Multiple nodes with the same priority create equal failover targets. Order is used for sequential service dependency management.

Q209. What is Proxmox VE replication schedule?

Answer: Replication schedules define how frequently ZFS-based storage replication runs between nodes. The minimum interval is 1 minute (*/* in cron notation). More frequent replication = smaller recovery point objective (RPO) but more I/O overhead.

Q210. What is the Proxmox VE HA 'freeze' state?

Answer: The freeze state prevents the HA manager from taking automatic actions (start/stop/migrate) on a resource. Used during maintenance operations when temporary unmanagement is needed without disabling HA entirely for the resource.

Q211. What is migration bandwidth limiting in Proxmox VE?

Answer: Migration bandwidth limiting (migration: type=insecure,bandwidth=100 in datacenter.cfg) caps the network bandwidth used for VM migrations, preventing migration traffic from saturating network links used by production VMs.

Q212. What is the fence_pve agent?

Answer: fence_pve is a fencing agent for Proxmox VE that can be integrated with external cluster managers. It uses the Proxmox API to forcibly stop/reset a VM or node. Used in complex multi-cluster setups requiring API-based fencing.

Q213. How do you check Corosync ring status?

Answer: Run: corosync-cfgtool -s (shows ring status and errors), corosync-quorumtool -v (verbose quorum info), pvecm status (Proxmox-level cluster status). Check /var/log/syslog for corosync errors and ring communication problems.

Q214. What is the Proxmox VE cluster token?

Answer: The cluster token is a shared secret used for cluster node authentication (SSH fingerprints and corosync crypto keys). When joining a cluster, the new node is authenticated and receives the token. Stored securely in /etc/pve/.

Q215. What is the HA simulator in Proxmox VE?

Answer: Proxmox VE includes a built-in HA simulator (`ha-manager sim`) for testing and understanding HA behavior without a real cluster. It simulates node failures, quorum events, and HA state transitions in a controlled environment.

Q216. What is node evacuation in Proxmox VE?

Answer: Node evacuation migrates all VMs and containers from a node to other cluster nodes. Initiated via GUI (Node ▢ Migrate All) or CLI (for VMs: iterate with `qm migrate`). Used before planned maintenance to prevent service disruption.

Q217. What IPMI/BMC role does Proxmox VE HA use?

Answer: While Proxmox VE's primary fencing is watchdog-based, integrating IPMI/BMC (Baseboard Management Controller) allows out-of-band power control. External IPMI fencing provides an additional reliability layer for node isolation.

Q218. What is the meaning of HA resource state 'error'?

Answer: The error state means the HA manager failed to achieve the requested state (e.g., failed to start a VM repeatedly). The manager stops retrying and waits for operator intervention. The resource must be manually inspected and re-enabled.

Q219. How do you force-stop an HA resource?

Answer: Use: `ha-manager set vm: --state stopped`, or disable it: `ha-manager set vm: --state disabled`. You can also use the GUI: HA ▢ Resources ▢ select resource ▢ set state. Never use `qm stop` directly on HA-managed VMs.

Q220. What is Proxmox VE cluster-wide VM lock?

Answer: Proxmox VE uses VM locks (lock: migrate, backup, snapshot, clone) stored in `/etc/pve` VM configuration files to prevent concurrent conflicting operations on a VM. Locks are automatically set and released, visible in `qm config` output.

Section 7: Security & Access Control (Q221–Q250)

Q221. What authentication realms does Proxmox VE support?

Answer: Proxmox VE supports: `pam` (Linux PAM system authentication), `pve` (built-in Proxmox user database), LDAP/Active Directory (external directory integration), and OpenID Connect (OAuth2/OIDC for SSO).

Q222. How do you create a user in Proxmox VE?

Answer: Via GUI: Datacenter □ Permissions □ Users □ Add. Via CLI: `pveum user add @ --password --comment 'description'`. Users need permissions assigned separately via ACL entries.

Q223. What is the difference between pam and pve realm?

Answer: pam realm authenticates against Linux system users (`/etc/passwd`, `/etc/shadow`). pve realm uses Proxmox VE's own internal user database, managed via pveum. pam users can SSH to the host; pve users are Proxmox-only.

Q224. How do you configure LDAP authentication in Proxmox VE?

Answer: Datacenter □ Permissions □ Realms □ Add □ LDAP. Configure server, port, base DN, user attribute (uid), and optionally TLS. Test the connection before saving. Users from LDAP still need Proxmox permission assignments.

Q225. What is Two-Factor Authentication (2FA) in Proxmox VE?

Answer: Proxmox VE supports TOTP (Time-based One-Time Password, compatible with Google Authenticator/Authy) and WebAuthn/FIDO2 hardware keys. 2FA can be required per-user or enforced via realm settings.

Q226. What are the built-in roles in Proxmox VE?

Answer: Built-in roles include: Administrator (full access), NoAccess (explicitly deny), PVEAdmin (most admin tasks, no system), PVEAuditor (read-only), PVEDatastoreAdmin, PVEDatastoreUser, PVEVMAdmin, PVEVMUser, PVESysAdmin, and PVEPoolAdmin.

Q227. How do you create a custom role in Proxmox VE?

Answer: Via CLI: `pveum role add --privs`. Via GUI: Datacenter □ Permissions □ Roles □ Create. Custom roles combine specific privileges from Proxmox VE's privilege list for fine-grained access control.

Q228. What is an ACL (Access Control List) in Proxmox VE?

Answer: An ACL entry in Proxmox VE maps a User/Group + Role to a Path. For example: `/vms/100 + PVEVMUser □ user@pve` grants that user VM user access to VM 100 only. ACLs define the RBAC permissions in Proxmox VE.

Q229. What is the permission propagation option in Proxmox ACLs?

Answer: When propagate is enabled for an ACL entry, the permission applies to the path and all sub-paths (e.g., /nodes applies to /nodes/node1, /nodes/node2, etc.). Disabling propagation applies the permission only to the exact path specified.

Q230. What privileges are included in the PVEVMAdmin role?

Answer: PVEVMAdmin includes: VM.Allocate, VM.Audit, VM.Backup, VM.Clone, VM.Config.*, VM.Console, VM.Migrate, VM.Monitor, VM.PowerMgmt, VM.Snapshot, VM.Snapshot.Rollback. It allows full VM management but not node-level or storage administration.

Q231. How do you reset the root password in Proxmox VE?

Answer: Boot from Proxmox VE installer or Debian rescue media, mount the system partition, chroot, and run passwd root. Alternatively, use the GRUB recovery mode: append init=/bin/bash to kernel parameters, mount -o remount,rw /, run passwd.

Q232. What is the Proxmox VE firewall default-in-policy?

Answer: The default INPUT policy when the firewall is enabled is DROP for VMs/containers. At the datacenter level, the default can be set to ACCEPT, REJECT, or DROP. Node-level default for the host itself can be configured separately.

Q233. What is MAC spoofing protection in Proxmox VE?

Answer: MAC spoofing protection (enabled per network interface) prevents a VM from changing its MAC address to one different from what is configured in Proxmox. Enforced via ebtables rules on the bridge port. Prevents MAC-based network attacks.

Q234. What is IP spoofing protection in Proxmox VE?

Answer: IP spoofing protection (enabled per VM interface) blocks traffic from a VM that has a source IP not matching its configured IP address. Enforced via ebtables ARP filter and iptables. Requires correct IP configuration in the VM network settings.

Q235. How do you audit user actions in Proxmox VE?

Answer: Proxmox VE logs all API calls (and thus all user actions) to /var/log/pve/api2/ with request details, user, and results. Task logs are in /var/log/pve/tasks/. The GUI shows task history per-node. Syslog also captures authentication events.

Q236. What is the token-based API authentication in Proxmox VE?

Answer: Proxmox VE supports API tokens as an alternative to user/password authentication for API access. Tokens are created per user, can have their own permissions, and can be revoked independently. Ideal for automation and third-party integrations.

Q237. What is the difference between API token full-privilege and separated privilege?

Answer: A full-privilege API token inherits all permissions of its owner user. A token with privilege separation has explicitly defined permissions (a subset of the owner's permissions), allowing safer, least-privilege automation access.

Q238. What is Proxmox VE's approach to SSH security?

Answer: Proxmox VE uses SSH for node-to-node cluster communication (VM migration, configuration sync). It generates cluster SSH keys during cluster creation. Best practice: disable password SSH login, use key-based auth for admin access, and restrict SSH to management networks.

Q239. What TLS certificates does Proxmox VE use?

Answer: Proxmox VE uses self-signed certificates by default for the web interface. ACME integration (Let's Encrypt) allows automatic valid TLS certificate issuance. Custom certificates can be installed via GUI (Node ▢ Certificates) or CLI.

Q240. What is ACME/Let's Encrypt integration in Proxmox VE?

Answer: Proxmox VE has built-in ACME client integration for automatic TLS certificate management via Let's Encrypt or other ACME providers. Supports HTTP-01 and DNS-01 challenges. Configured under Node ▢ Certificates ▢ ACME.

Q241. What are Proxmox VE user groups?

Answer: User groups allow assigning permissions to multiple users simultaneously. Create a group (pveum group add), add users to it (pveum user modify --groups), then assign ACLs to the group path instead of individual users.

Q242. How do you restrict a user to only see their own VMs?

Answer: Create a resource pool, assign the user PVEVMUser (or custom) role on /pool/, and add their VMs to the pool. The user only sees and interacts with VMs in their pool. Optionally assign storage permissions for ISO/backup access.

Q243. What is the pvestatd service?

Answer: `pvestatd` is the Proxmox VE statistics daemon. It collects and stores performance metrics (CPU, memory, network, disk I/O) for nodes, VMs, and containers. Data is stored in RRD files and displayed in the GUI performance graphs.

Q244. What is the pvedaemon service?

Answer: `pvedaemon` is the main Proxmox VE API server daemon. It handles all API requests from the web GUI, CLI tools, and external API clients. It runs as root and manages authentication, permission checks, and action dispatch.

Q245. What is TOTP and how is it configured in Proxmox VE?

Answer: TOTP (Time-based One-Time Password) generates 6-digit codes that refresh every 30 seconds using a shared secret. In Proxmox VE, users configure TOTP via Datacenter □ Permissions □ Users □ Two Factor, scanning a QR code with an authenticator app.

Q246. What is the pam_oath module in Proxmox VE?

Answer: Proxmox VE integrates with `libpam-oath` for TOTP-based 2FA for PAM realm users. The OATH credentials are stored in `/etc/pve/priv/tfa.cfg` (encrypted). When enabled, users must enter both password and TOTP code.

Q247. What is realm-specific 2FA enforcement?

Answer: Proxmox VE allows enforcing 2FA at the realm level – all users in the realm must have 2FA configured to log in. Set via GUI (Realm □ Edit □ TFA required) or via `pveum realm modify --tfa type=totp`.

Q248. How do you limit API access to specific IP addresses?

Answer: Proxmox VE's built-in firewall can restrict access to port 8006 (web GUI/API) to specific IP ranges. At the node firewall level, create INPUT rules allowing 8006 only from management IPs and set default INPUT policy to DROP.

Q249. What is the Proxmox VE audit log location?

Answer: Detailed audit logs are at `/var/log/pve/api2/` (API access logs), `/var/log/auth.log` (PAM auth), and `/var/log/syslog` (general system events). Task-specific logs are in `/var/log/pve/tasks/` organized by task ID.

Q250. What is a Proxmox VE permission token?

Answer: Not to be confused with API tokens, permission tokens in Proxmox VE refer to the authentication ticket (cookie) issued after successful login. Tickets expire after 2 hours and must be refreshed. They are used for subsequent API calls within a session.

Q251. What is the CSRFPreventionToken in Proxmox VE API?

Answer: The CSRFPreventionToken is returned alongside the authentication ticket on login. It must be included as a header (CSRFPreventionToken:) in all POST/PUT/DELETE API requests to prevent Cross-Site Request Forgery attacks.

Section 8: Monitoring, Logging & Performance (Q251–Q280)

Q252. What monitoring tools are available in Proxmox VE?

Answer: Proxmox VE includes built-in monitoring via the web GUI (RRD graphs for CPU, memory, network, disk). External integration via InfluxDB (push metrics), Graphite, and Prometheus exporters (node_exporter, pve-exporter) is supported.

Q253. How do you configure InfluxDB metric export in Proxmox VE?

Answer: Datacenter □ Metric Server □ Add □ InfluxDB. Specify server, port (8086 for HTTP, 8089 for UDP), and optionally authentication token and bucket/org for InfluxDB v2. All nodes send metrics automatically after configuration.

Q254. What RRD data is stored by Proxmox VE?

Answer: Proxmox VE stores RRD (Round Robin Database) data for: node CPU/memory/network/disk I/O, VM/CT CPU/memory/disk read-write/network I/O. Data is retained at multiple resolutions (hour, day, week, month, year) with automatic aggregation.

Q255. What is the pveperf command?

Answer: pveperf is a Proxmox VE benchmarking utility that tests disk I/O performance (sequential read/write and random IOPS) on the local system. Useful for baseline performance measurements and storage comparison.

Q256. How do you check VM performance graphs in Proxmox VE?

Answer: In the web GUI: click a VM □ Summary tab shows CPU/memory usage. More Graphs link shows detailed RRD graphs for CPU, memory, network I/O, and disk I/O over various time ranges (hour, day, week, month, year).

Q257. What logs are important for troubleshooting in Proxmox VE?

Answer: Key logs: `/var/log/syslog` (general), `/var/log/pve/tasks/` (task logs), `/var/log/pve/api2/` (API logs), `/var/log/auth.log` (auth events), `dmesg` (kernel messages). For VMs: check guest logs. For Ceph: `/var/log/ceph/`.

Q258. How do you check task logs in Proxmox VE?

Answer: Via GUI: any task listed in the Tasks panel has a log icon. Click it to view detailed log output. Via CLI: logs are stored in `/var/log/pve/tasks/` where UPID is the task ID. Use: `cat /var/log/pve/tasks/active` for recent tasks.

Q259. What is the `journalctl` command's role in Proxmox troubleshooting?

Answer: `journalctl` provides access to the `systemd` journal, containing structured logs from all system services. Use: `journalctl -u pve-cluster` (cluster logs), `journalctl -u qemu-server@` (VM-specific), `journalctl -f` (follow live), `journalctl --since '1 hour ago'`.

Q260. How do you monitor Ceph cluster health?

Answer: Use: `ceph status` (overall health), `ceph health detail` (detailed warnings/errors), `ceph osd status` (per-OSD status), `ceph df` (space usage). In Proxmox GUI: Ceph □ Status/OSDs/Pools shows a visual dashboard.

Q261. What is the IOPS metric and why does it matter?

Answer: IOPS (Input/Output Operations Per Second) measures how many read/write operations a storage system can handle per second. Critical for database VMs and latency-sensitive workloads. HDD typically offers 100-200 IOPS; NVMe SSDs can exceed 1 million IOPS.

Q262. What is CPU steal time in Proxmox VE VMs?

Answer: CPU steal time (`%st`) appears in guest OS metrics when the hypervisor cannot provide requested CPU cycles to the VM because the host CPU is overcommitted. High steal time indicates the host needs more CPU resources.

Q263. How do you check memory balloon status for a VM?

Answer: Use: `qm monitor`, then type: `info balloon` (shows current and target balloon size). Also check via the web GUI VM Summary page. High balloon deflation (host reclaiming memory) may indicate memory pressure on the host.

Q264. What is the purpose of the Proxmox VE Status plugin?

Answer: The Status plugin (introduced in Proxmox VE 7) allows custom integration of metrics into the Proxmox VE GUI via external metric collectors. It provides a framework for extending the built-in monitoring with custom data sources.

Q265. How do you check disk I/O wait on a Proxmox host?

Answer: Use: `iostat -xz 1` (detailed I/O stats with utilization), `iotop` (per-process I/O usage), `dstat` (combined system stats), or check `/proc/diskstats`. High `%iowait` in `top/htop` indicates storage bottlenecks.

Q266. What is the `pveceph` command?

Answer: `pveceph` is the Proxmox VE CLI for managing Ceph clusters deployed on Proxmox nodes. Commands: `pveceph install`, `pveceph init`, `pveceph mon create`, `pveceph osd create`, `pveceph pool create`, `pveceph status`.

Q267. How do you identify memory-hungry VMs in Proxmox VE?

Answer: Check the web GUI Summary/More Graphs per VM for memory usage over time. Via CLI: use `qm monitor` then `info kvm` (shows various stats). On the host: use `top`, `htop`, or `systemd-cgtop` to see cgroup-level memory usage per VM.

Q268. What is the `pve-firewall` service?

Answer: The `pve-firewall` service applies and manages Proxmox VE firewall rules (`iptables/nftables/ebtables`). It reads rules from `/etc/pve/firewall/` and applies them. Restart with: `systemctl restart pve-firewall`. Check status: `pve-firewall status`.

Q269. How do you enable verbose logging for Proxmox VE API?

Answer: Edit `/etc/default/pvedaemon` and add `--loglevel debug` option. Or set via syslog: configure `rsyslog` to capture `pvedaemon` messages at debug level. For temporary verbose output: use the `pvesh` command (`pve shell`) to trace API calls.

Q270. What is the `pvesh` command?

Answer: `pvesh` (Proxmox VE Shell) is a CLI API client that can call any Proxmox VE REST API endpoint directly from the command line. Example: `pvesh get /nodes`, `pvesh create /nodes//qemu --vmid 100 ...`. Useful for scripting and debugging.

Q271. What is SMART monitoring for disks in Proxmox VE?

Answer: Proxmox VE can display SMART (Self-Monitoring, Analysis and Reporting Technology) data for physical disks via the Node ▢ Disks ▢ select disk ▢ SMART Data panel. Install `smartmontools` and enable it for proactive disk health monitoring.

Q272. How do you check network bandwidth usage per VM in Proxmox?

Answer: Via GUI: VM ▢ More Graphs ▢ Network. Via CLI: check `/proc/net/dev` on the host for the tap/veth interface corresponding to the VM. Use `iftop` or `nethogs` on the host to see real-time per-interface bandwidth.

Q273. What is the meaning of 'load average' on a Proxmox host?

Answer: Load average represents the average number of processes in the run queue over 1, 5, and 15 minutes. On a Proxmox host, it reflects combined load from all VMs and the hypervisor. Load > number of CPU cores indicates CPU contention.

Q274. What is balloon driver memory reporting?

Answer: The VirtIO balloon driver inside the guest reports actual used memory to the hypervisor. This allows Proxmox to display accurate guest memory usage statistics rather than just the configured maximum memory of the VM.

Q275. What is the sysstat package in Proxmox VE?

Answer: `sysstat` provides `sar` (System Activity Reporter), `iostat`, and `mpstat` utilities for historical performance data collection and analysis. Install with: `apt install sysstat`. Configure data collection interval in `/etc/sysstat/sysstat` and enable the `sysstat` service.

Q276. How do you monitor Corosync cluster communication health?

Answer: Use: `corosync-cfgtool -s` (ring status), `corosync-quorumtool -l` (node list with votes), `pvecm status` (cluster status). Check `syslog` for Corosync messages. High packet loss or ring faults appear as `FAULTY` in `corosync-cfgtool` output.

Q277. What is the pve-exporter for Prometheus?

Answer: `pve-exporter` (`prometheus-pve-exporter`) is a community Prometheus exporter that exposes Proxmox VE metrics (nodes, VMs, storage, cluster) in Prometheus format. Combined with Grafana dashboards, it provides comprehensive cluster monitoring.

Q278. How do you analyze CPU topology in a Proxmox VM?

Answer: Use: `lscpu` inside the VM to see vCPU topology. On the host: `qm config` shows CPU socket/cores/numa settings. Use `numactl --hardware` to check NUMA topology. `lstopo` (hwloc package) shows graphical CPU topology.

Q279. What is the proxmox-backup-debug tool?

Answer: Proxmox Backup Server provides `proxmox-backup-debug` as a CLI tool for diagnosing backup issues, inspecting backup indices, verifying chunk stores, and analyzing backup catalog entries without full restore.

Q280. How do you capture network packets for a VM in Proxmox?

Answer: Use `tcpdump` on the host on the VM's tap interface: `tcpdump -i tapi -w /tmp/capture.pcap`. Find the tap interface name in: `ip link show | grep tap`. This captures all traffic going to/from the VM at the hypervisor level.

Q281. What is the Proxmox VE guest OS memory usage vs configured memory?

Answer: Configured memory is the maximum a VM can use. Actual guest memory usage (visible in the GUI) is reported by the balloon driver. Without balloon driver, Proxmox assumes full memory allocation. Guest OS free memory is not automatically returned to the host without balloon.

Q282. How do you check cluster-wide resource usage in Proxmox VE?

Answer: In the GUI: Datacenter □ Summary shows cluster-wide CPU, memory, and storage usage. Via CLI: `pvesh get /cluster/resources` returns JSON with all nodes, VMs, and storage usage. Use `pvesh get /nodes//status` for per-node stats.

Section 9: Advanced Topics & Troubleshooting (Q281–Q300)

Q283. How do you recover a Proxmox VE node that lost its cluster configuration?

Answer: If `/etc/pve` is empty (pmxcfs issue): restore the `corosync.conf` from a backup or another node, restart `pve-cluster` service. If cluster is inaccessible, use `pvecm expected 1` to reduce expected votes and restore quorum for emergency single-node operation.

Q284. What is the `pvecm expected` command?

Answer: `pvecm expected` temporarily reduces the expected vote count, allowing a cluster to achieve quorum with fewer nodes. Used for emergency single-node maintenance. Must be followed by proper cluster repair; never leave in reduced-quorum state permanently.

Q285. How do you restore a VM from a Proxmox Backup Server backup?

Answer: Via GUI: select the PBS storage ▢ Backups ▢ select backup ▢ Restore. Via CLI: `qmrestore : .` For PBS: `proxmox-backup-client restore` or via the PBS web interface ▢ Datastore ▢ Backups ▢ Restore.

Q286. What is the `qm` terminal command used for?

Answer: `qm` terminal connects to a VM's serial console (if configured). This allows console access to VMs even when the graphical/VNC console is unavailable, such as during network issues or for headless server OS debugging.

Q287. How do you fix a 'lock' stuck on a VM?

Answer: If a VM has a stuck lock (e.g., from a crashed backup), remove it with: `qm unlock` or manually edit `/etc/pve/nodes//qemu-server/.conf` and remove the lock: line. Always investigate why the lock occurred before removing it.

Q288. What is the `vzdump --dumpdir` option?

Answer: `--dumpdir` specifies the directory where `vzdump` saves backup files. Example: `vzdump 100 --mode snapshot --dumpdir /mnt/backup/proxmox`. This overrides the storage setting and saves to a local path.

Q289. How do you migrate a Proxmox VE installation to new hardware?

Answer: Backup all VMs/containers. Install fresh Proxmox VE on new hardware. Restore backups, or if using shared/Ceph storage, just recreate VM configs pointing to existing disk images. For ZFS, `zpool export/import` can move the pool to new hardware.

Q290. What causes 'task already running' errors in Proxmox VE?

Answer: This error occurs when trying to start a VM/CT that already has a running process, or when a previous task didn't clean up properly. Check: `qm status`, `ps aux | grep kvm`. Kill stale processes and remove PID files in `/var/run/qemu-server/` if necessary.

Q291. What is the `rbd bench` command and when is it useful?

Answer: `rbd bench --io-type write` tests the I/O performance of a Ceph RBD image. Used to benchmark Ceph storage performance for VM workloads. Also use: `rados bench`, `fio` on a mounted volume for comprehensive Ceph performance analysis.

Q292. How do you change the cluster IP address of a Proxmox node?

Answer: Edit `/etc/pve/corosync.conf` to update the node's `ring0_addr` (and `ring1_addr` if used). Restart corosync: `systemctl restart corosync`. Update `/etc/hosts` on all nodes. For the management IP, update `/etc/network/interfaces` and apply.

Q293. What is the proxmox-ve metapackage?

Answer: The `proxmox-ve` package is a Debian metapackage that depends on all Proxmox VE core components (`pve-kernel`, `pve-manager`, `corosync`, etc.). Installing or upgrading this package ensures all components are at compatible versions.

Q294. How do you enable kernel modules for Proxmox VE?

Answer: Add module names to `/etc/modules-load.d/.conf` for persistent loading at boot. For immediate loading: `modprobe`. For module options: create `/etc/modprobe.d/.conf` with options `=`. Update initramfs: `update-initramfs -u -k all`.

Q295. What is the KSM (Kernel Same-page Merging) feature in Proxmox?

Answer: KSM is a Linux kernel feature that identifies and merges identical memory pages across VMs, reducing total host memory consumption. Enabled in Proxmox VE by default. Configured via `/sys/kernel/mm/ksm/`. Useful for VMs with similar OS workloads.

Q296. How do you troubleshoot a VM that won't start?

Answer: Check: `qm status` (current state), task log in GUI or `/var/log/pve/tasks/`, `dmesg` (kernel errors), `/var/log/syslog` (qemu errors). Common causes: insufficient resources (CPU/memory/storage), configuration errors, locked storage, or missing CD-ROM ISO.

Q297. What is the pve-ha-lrm vs pve-ha-crm service?

Answer: `pve-ha-lrm` (Local Resource Manager) runs on each node, managing resources locally and executing commands from the CRM. `pve-ha-crm` (Cluster Resource Manager) runs as one active instance cluster-wide, making global HA decisions and coordinating failover.

Q298. How do you rebuild the initramfs in Proxmox VE?

Answer: Run: `update-initramfs -u -k all` to rebuild for all installed kernels, or `update-initramfs -u -k $(uname -r)` for the running kernel. Required after adding kernel modules, changing module parameters, or updating storage drivers (e.g., for ZFS).

Q299. What is the Proxmox VE upgrade path between major versions?

Answer: Follow official upgrade guides on the Proxmox wiki. Typically: update all packages in current version, address any deprecations, then run: `sed -i '/bullseye/bookworm/g' /etc/apt/sources.list` files and `apt dist-upgrade`. Always backup before major upgrades.

Q300. How do you diagnose high CPU steal in Proxmox VMs?

Answer: On the host: check CPU overcommitment with `top/htop` (total vCPU vs physical cores), `perf stat` or `turbostat` for CPU frequency/utilization. Reduce VM CPU allocation, enable CPU limits, or migrate VMs to less-loaded nodes to reduce steal.

Q301. What is the Proxmox VE community forum and documentation?

Answer: Proxmox VE documentation is available at pve.proxmox.com/pve-docs/. The community forum at forum.proxmox.com provides peer support. Official paid support is available to subscription holders. The wiki at pve.proxmox.com/wiki/ contains How-Tos and guides.

Q302. What are best practices for a production Proxmox VE deployment?

Answer: Key practices: use 3+ node clusters for HA, dedicated cluster/storage networks, enterprise-grade storage (Ceph or ZFS with RAID), regular backups with tested restore procedures, subscription for enterprise repos, hardware IPMI for OOB management, monitoring with Prometheus/Grafana, and documented runbooks for common failure scenarios.
