SD-WAN

ZERO TO HERO

Deep Dive: Scenarios, Projects & Daily Practical Knowledge

Gautam Shirsath

LinkedIn: linkedin.com/in/gautam-s2024

10+	50+	20+	A to Z
Chapters	Scenarios	Projects	Coverage

Table of Contents

Chapter 11: Real-World Project Scenarios A to Z

Chapter 12: Career, Certifications & Next Steps

CHAPTER 1

SD-WAN Foundations

What it is, why it exists, and how the whole thing hangs together

Before you can master SD-WAN, you must understand the pain it was built to solve. This chapter takes you from the old world — expensive, rigid, slow MPLS networks — all the way to the modern SD-WAN architecture. By the end, you'll understand exactly how the three planes work together and why every enterprise in the world is migrating.

1.1 The Problem with Traditional WAN

Imagine you are the IT manager at a company with 50 branch offices across India. Each office needs to connect to the main data center in Mumbai. In the traditional model, you buy MPLS circuits from a telecom provider — essentially renting a private road on the internet backbone. Here is what your daily life looks like:

- **Cost:** Each MPLS circuit costs Rs. 80,000–3,00,000/month per branch. For 50 branches, that's potentially crores every month just on connectivity.
- **Speed of change:** Need a new branch office in Pune? Call the telecom. Wait 4–12 weeks for provisioning. The branch cannot operate until the circuit arrives.
- **Rigid routing:** Traffic always goes through the data center. Even if an employee is opening Gmail or YouTube — it bounces all the way to the DC and back. This wastes bandwidth and adds 100+ ms of unnecessary latency.
- **No visibility:** You know if the link is up or down. You do NOT know which application is consuming bandwidth, how well it is performing, or where the bottleneck is.
- **No redundancy:** If the MPLS circuit goes down, the branch is dead. Adding a backup link means adding another MPLS circuit — more cost, more complexity.

■ REAL-WORLD SCENARIO: The Frustrated Bank IT Manager

Priya is IT manager at a private bank with 80 branches. Every Monday morning, her phone rings — some branch has a slow MPLS circuit causing core banking to lag. She opens a ticket with the telecom. They tell her the SLA window is 4 hours. Customers are angry. Tellers are stuck. Priya manually reroutes traffic by logging into each router individually — a process that takes 45 minutes per site. When she finally fixes Branch A, Branch C starts failing. She has zero centralized visibility, zero automation, and zero backup paths. This is the world SD-WAN was built to replace. With SD-WAN: Priya logs into ONE dashboard. She sees ALL 80 branches on one screen. The system has already automatically switched Branch A to broadband backup — no phone call, no ticket, no manual work. She gets an alert on her phone before any customer complains.

1.2 What SD-WAN Actually Is

SD-WAN (Software-Defined Wide Area Network) is a technology that separates the **intelligence** of a network (the control plane — deciding where traffic goes) from the **hardware** (the data plane — actually moving packets). By centralizing intelligence in software, SD-WAN achieves what was previously impossible:

- **Any transport:** Use MPLS, broadband, LTE, 5G, satellite — simultaneously. The software decides the best path in real time based on actual measured quality.
- **Application awareness:** SD-WAN deep-inspects packets (DPI) to identify applications — not just IP addresses. It knows the difference between a Zoom call and a file backup.
- **Centralized management:** All 80 branches managed from a single GUI. One policy change propagates everywhere in seconds.
- **Automated failover:** If MPLS degrades, traffic moves to broadband in sub-second time — before users notice. No human intervention required.
- **Zero-touch deployment:** Ship a router to a branch. Plug it in. It calls home, downloads its config, and comes online automatically.

SD-WAN is a software layer that sits above your physica
real-time decisions about which path each application sh
centrally managed. It turns a dumb pipe into a smart, self-h

1.3 Core Architecture: Underlay, Overlay & The Three Planes

Every SD-WAN architect must be crystal clear on these foundational concepts. They are the vocabulary of every SD-WAN conversation.

The Underlay Network

The **underlay** is the physical, real-world network — the actual cables, fiber, radio towers, and ISP infrastructure that carries bits from A to B. The underlay doesn't know anything about applications. It just moves packets. Examples:

- MPLS circuit from Tata Communications — private, reliable, expensive
- Broadband connection from local ISP — cheap, public, variable quality
- 4G LTE from Airtel — mobile, flexible, suitable for backup
- SD-WAN connects to ALL of these simultaneously

The Overlay Network

The **overlay** is a logical network built on TOP of the underlay using encrypted tunnels. SD-WAN creates IPsec tunnels between all edge devices. From the perspective of the overlay, all sites are directly connected — even if the physical paths are completely different ISPs. The overlay hides the complexity of the underlay. It also provides:

- Encryption: All traffic in the overlay tunnels is AES-256 encrypted

- Path independence: If underlay path changes, overlay stays up
- QoS: Traffic can be prioritized within overlay regardless of underlay provider

The Three Planes

Plane	Function	Component	Analogy
Management Plane	Configure policies, monitor health, provision devices	Orchestrator / vManage	The CEO — sets strategy and mor
Control Plane	Compute routes, distribute policies, authenticate devices	Controller / vSmart	The CFO — makes the intelligent
Data Plane	Forward actual user packets at wire speed per policy	Edge Device / vEdge / CPE	The employee — executes the ord

These three planes communicate constantly. The Orchestrator tells the Controller what policies to apply. The Controller tells the Edge devices which routes are best and how to forward traffic. The Edge devices report back health statistics to the Controller and Orchestrator. This feedback loop happens every few seconds — enabling truly dynamic networks.

■ REAL-WORLD SCENARIO: Coffee Chain — Nationwide SD-WAN Rollout

BrewCo has 200 coffee shops across India. Each shop needs: POS (payment terminal) connectivity, surveillance cameras, customer Wi-Fi, and back-office SAP access. With MPLS, provisioning all 200 shops would take 12+ months and cost crores per year. With SD-WAN: Each shop gets a dual broadband (primary Jio + backup Airtel) setup. The edge device auto-configures via ZTP. POS traffic goes over the best link with highest priority. Cameras use the backup link. Customer Wi-Fi is isolated in a separate VRF. SAP goes encrypted over the overlay to the DC. Total provisioning time per shop: 2 hours. Total IT headcount required on-site: Zero. Monthly savings vs MPLS: 70%.

■■ PRACTICE PROJECT: Draw Your First SD-WAN Topology

TASK: On paper or draw.io (free), sketch a simple SD-WAN for a company with: 1 Data Center (Mumbai), 3 Branch Offices (Delhi, Pune, Bangalore), and 1 Cloud (AWS). Requirements: (1) Each branch has 2 ISP links. (2) DC has MPLS + Broadband. (3) Show where the Orchestrator, Controller, and Edge devices live. (4) Draw overlay tunnels between all sites. (5) Label which traffic goes which path: Voice = MPLS preferred, Internet = Local breakout, SAP = MPLS primary + Broadband backup. VALIDATION: Can you answer these? Where does a Zoom call from Delhi go? What happens if the Mumbai MPLS fails? How does a new branch in Chennai get provisioned?

CHAPTER 2

Transports, Links & Underlay Networks

Understanding every connection type and when to use each

The power of SD-WAN comes from its ability to use ANY transport simultaneously. But not all transports are equal. A daily SD-WAN engineer must deeply understand each option — its strengths, weaknesses, cost structure, and ideal use cases.

2.1 MPLS — The Private Highway

Multiprotocol Label Switching (MPLS) is the gold standard of enterprise WAN. It works by assigning 'labels' to packets at the edge of the provider network, allowing them to be forwarded along pre-defined paths (Label Switched Paths / LSPs) without needing to examine IP headers at every hop. Think of it as a highway with dedicated lanes — guaranteed speed, guaranteed priority, guaranteed SLA.

- **Latency:** Extremely low and consistent. Typically 5-20ms within a country.
- **Jitter:** Near-zero. Packets arrive in order, consistently spaced.
- **Packet Loss:** Near-zero under normal conditions. Provider guarantees SLA.
- **Security:** Private network — traffic doesn't traverse public internet. BUT it's not encrypted by default! SD-WAN adds encryption on top.
- **Cost:** High. 3-10x more expensive than broadband for same bandwidth.
- **Provisioning:** Slow. New circuits take 4-12 weeks.
- **Ideal for:** Voice, video, real-time trading systems, core banking. Any application where latency and jitter must be guaranteed.

■ **PRO TIP:** In SD-WAN, MPLS is your premium lane. Use it only for traffic that truly needs SLA guarantees. Route everything else over broadband to save cost.

2.2 Broadband / DIA — The Public Road

Broadband (or **Dedicated Internet Access / DIA**) is regular internet connectivity — the same kind that runs to your home. DIA is dedicated (not shared), making it more reliable than consumer broadband, but it still uses the public internet backbone with no SLA on path quality. SD-WAN transforms broadband from 'unreliable backup' into a fully usable production link by constantly measuring its quality and routing traffic appropriately.

- **Latency:** Variable. Typically 20-80ms domestic, higher international.
- **Cost:** Low. Often 70-80% cheaper than MPLS per Mbps.

- **Bandwidth:** High and increasing. 100Mbps–10Gbps easily available.
- **Provisioning:** Fast. 3-10 days in urban areas.
- **Ideal for:** Internet breakout, SaaS apps (O365, Salesforce), bulk file transfers, backup paths, guest Wi-Fi.

2.3 LTE/5G — The Wireless Safety Net

LTE and 5G cellular connections are increasingly used in SD-WAN as backup or even primary links. They are invaluable for remote locations where fiber isn't available, for mobile deployments (retail kiosks, pop-up stores), and as last-resort backup when both wired links fail.

- **Latency:** LTE: 30-100ms. 5G: 5-30ms. Much improved with 5G.
- **Cost:** Per-GB pricing. Can get expensive with high traffic. Best used for backup or low-volume traffic.
- **Reliability:** Excellent in urban areas. Poor in basements, remote areas.
- **Ideal for:** Backup link, remote/rural sites, mobile deployments, disaster recovery.

Feature	MPLS	Broadband/DIA	LTE/5G	Satellite
Latency	5-20ms	20-80ms	30-100ms	500-700ms
Jitter	Very Low	Variable	Moderate	High
Packet Loss	< 0.1%	0.1-2%	0.5-3%	1-5%
Cost (per Mbps)	High $$$	Low $	Medium $$	Very High $$$$
Provisioning	4-12 weeks	3-10 days	Same day	2-4 weeks
Max Bandwidth	10Gbps	10Gbps	1Gbps (5G)	100Mbps
SLA Available	Yes (99.99%)	Limited	No	Limited
Best For	Voice/Video/Banking	SaaS/Internet	Backup/Remote	Remote/Offshore

■ REAL-WORLD SCENARIO: Retail Bank with 3 Link Types

PeopleFirst Bank has 120 branches. The IT team builds this link strategy: Tier-1 Branches (large urban) get MPLS (10Mbps) + Broadband DIA (100Mbps). Tier-2 Branches (mid-size) get Broadband DIA (50Mbps) + LTE backup. Tier-3 Branches (rural/small) get LTE primary + LTE backup (different carriers). SD-WAN Policy: Core banking traffic (TCP port 8443 to DC IP range) = MPLS primary, Broadband failover, threshold: latency >50ms or loss >0.5% triggers switch. ATM monitoring traffic = Broadband primary, LTE failover. Employee web browsing = Local breakout via Broadband/LTE. Security cameras = LTE (separate VRF). Result: Bank saves 45% on WAN costs, achieves 99.98% uptime across all 120 branches.

■■ PRACTICE PROJECT: Link Cost-Benefit Analysis

TASK: Design the optimal link strategy for these 3 scenarios: Scenario A: E-commerce company, 1 warehouse (Pune), needs 500Mbps consistently for order processing. Scenario B: Insurance company, 300 field agent offices across India, each needing <10Mbps, mostly browsing + VoIP. Scenario C: Oil rig offshore, needs 5Mbps for SCADA monitoring, must be 24/7 available. For each: Choose link types, justify the choice, estimate monthly cost (use typical ISP rates), calculate savings vs pure MPLS, and define what SLA metrics you would measure.

CHAPTER 3

SD-WAN Control Plane & Routing

How SD-WAN makes routing decisions and distributes intelligence

The control plane is the brain of SD-WAN. Understanding it deeply separates a junior engineer who can follow guides from a senior engineer who can diagnose complex routing issues and design sophisticated policies. This chapter goes deep into OMP, BGP integration, and route policy design.

3.1 How the Controller Works

The SD-WAN Controller (Cisco calls it **vSmart**) is a centralized routing engine. Instead of each router figuring out routes independently (like traditional OSPF/BGP), the controller computes routes centrally and distributes them to all edge devices. This is a radical departure from traditional networking.

Here is exactly what the controller does, step by step:

- **Step 1 — Authentication:** When an edge device boots up, it authenticates with the controller using mutual TLS certificates. No certificate = no connection. This prevents rogue devices from joining.
- **Step 2 — Route Advertisement:** Edge devices send their local network prefixes (subnets) to the controller via OMP.
- **Step 3 — Policy Application:** The controller applies route policies (what routes to accept, modify, or filter) and computes which routes should be distributed to which sites.
- **Step 4 — Route Distribution:** The controller pushes the computed routes back to all relevant edge devices. Each edge device updates its routing table.
- **Step 5 — Continuous Update:** This process repeats continuously. Link failures, new sites, policy changes — all propagate within seconds.

3.2 OMP — Overlay Management Protocol

OMP (Overlay Management Protocol) is Cisco's proprietary routing protocol used between edge devices and the vSmart controller. It is similar to BGP in concept (path-vector, policy-based) but designed specifically for SD-WAN overlay networks. OMP carries three types of routes:

- **OMP Routes:** Regular IP prefixes (like BGP prefixes) — 'I can reach 10.1.0.0/24'
- **TLOC Routes (Transport Locator):** Information about transport endpoints — 'I am reachable at IP 203.0.113.1 via color mpls or color biz-internet'. TLOC is the unique 3-tuple: System IP + Color + Encapsulation (IPsec/GRE).

- **Service Routes:** Advertisements of services like firewall, IDS, NAT that exist at a site. Enables service chaining.

TLOC (Transport Locator) is unique to Cisco SD-WAN. E
device gets a 'color' — a label identifying which transpo
'biz-internet', 'lte', 'public-internet', 'metro-ethernet'. When y
over MPLS', you're actually saying 'prefer the TLOC with c
transport-aware policies without hard-coding IP addresses

3.3 BGP, OSPF Integration

SD-WAN doesn't replace your LAN routing protocols — it integrates with them. Here is how the routing stack works end to end:

- **LAN Side:** Branches typically run OSPF or EIGRP to exchange routes with internal subnets (VLANs, servers). The edge device participates in LAN routing.
- **Redistribution:** LAN routes are redistributed into OMP and advertised to the controller, which distributes them to other sites.
- **DC Side:** Data centers often run BGP with the SD-WAN edge to exchange routes with internal infrastructure and cloud connections.
- **Internet Side:** If using BGP with ISP (for dynamic routing), SD-WAN can import BGP prefixes into OMP for overlay-aware internet routing.

■ REAL-WORLD SCENARIO: Hospital Failover Routing

City General Hospital has 3 sites: Main Hospital (HQ), West Wing (branch), ICU Building (branch). Traffic types: EMR (Electronic Medical Records) on TCP 443, VoIP (SIP on UDP 5060), Medical imaging (DICOM, large files), and general internet browsing. Routing design: EMR traffic → MPLS primary (SLA: latency <30ms, loss <0.1%). If MPLS degrades, OMP policy triggers immediate failover to broadband — healthcare cannot wait. VoIP → MPLS primary, LTE secondary. Imaging → Broadband (high bandwidth, latency-tolerant). Internet → Local breakout directly, never goes to HQ. The OMP controller runs in active-active pair for redundancy. If a controller fails, edge devices continue forwarding on last known good routes for up to 12 hours — called 'Data Plane Resilience'. This is critical for hospitals where a controller outage cannot mean a network outage.

■■ PRACTICE PROJECT: Configure OMP Route Policies

TASK: Write the logic (in plain English or pseudo-config) for these OMP policies: Policy 1 — Hub preference: Make all branches prefer routing through the Mumbai HQ for DC-bound traffic, but allow direct branch-to-branch for VoIP. Policy 2 — Route leaking between VRFs: Branch A is in VRF 'corp'. Branch B has a shared printer server at 192.168.50.0/24 in VRF 'mgmt'. Allow Branch A users to reach the printer. Policy 3 — Topology restriction: The satellite site should only be able to reach the DC, not other branches (reduce expensive satellite hops). For each policy, describe: which OMP route type is involved, where the policy is applied (controller/edge), and what the failure scenario looks like.

CHAPTER 4

Application-Aware Routing (AAR)

Making the network serve applications, not the other way around

Application-Aware Routing is arguably the most powerful feature of SD-WAN. It allows the network to make forwarding decisions based on what application is running — not just where the destination IP is. Combined with real-time link quality measurement, it creates a self-optimizing network.

4.1 How SD-WAN Identifies Applications

SD-WAN edge devices use **Deep Packet Inspection (DPI)** to identify applications. Unlike a basic firewall that only looks at ports, DPI examines the actual content of packets to determine what application is communicating. Identification methods:

- **Port-based:** TCP 443 = HTTPS, UDP 5060 = SIP. Simple but often wrong (everything uses 443 now).
- **Signature-based:** Match known byte patterns in packet payloads. Cisco has 3000+ application signatures (NBAR2). Identifies Zoom, Salesforce, SAP, etc.
- **URL/FQDN-based:** DNS queries reveal destination. *.zoom.us → Zoom. *.office.com → Office 365.
- **Machine Learning:** Advanced SD-WAN can identify apps from behavioral patterns even when encrypted (flows, timing, packet sizes).

DPI needs to see packets to identify an app. But by the tim
already been forwarded. SD-WAN solves this with 'first-pa
for the first few packets while identification runs, then r
vendors use IP/port as a 'hint' for the first packet and
operational consideration — very short flows (< 5 packets)

4.2 SLA Policies & Thresholds

An **SLA Policy** in SD-WAN defines the quality requirements for a specific application or traffic class. The SD-WAN system continuously measures link quality and compares it against SLA thresholds. When a link breaches a threshold, traffic moves to the next-best link.

The three core metrics measured per-link, per-second:

- **Latency:** Round-trip delay. Measured using BFD (Bidirectional Forwarding Detection) probe packets sent every 1 second. Voice requires < 150ms. Interactive apps < 200ms.
- **Jitter:** Variation in delay. A call with 20ms latency but 50ms jitter will sound broken. Voice requires < 30ms jitter.

- **Packet Loss:** Percentage of packets not received. Even 1% loss is audible on voice. Video stutters at 0.5%.

Application	Max Latency	Max Jitter	Max Loss	Priority
VoIP / Unified Comm	150ms	30ms	0.5%	Critical
Video Conferencing	200ms	50ms	1%	High
Core Banking / ERP	200ms	100ms	0.5%	High
Interactive Web Apps	300ms	150ms	2%	Medium
Email / Office 365	500ms	N/A	5%	Normal
Software Updates	No SLA	N/A	10%	Low
Backup / Replication	No SLA	N/A	Retransmit	Best Effort

4.3 Dynamic Path Selection — The Heart of SD-WAN

Dynamic Path Selection (DPS) is the real-time engine that evaluates links against SLA policies and routes traffic accordingly. Here is exactly how it works:

- **BFD Probes:** SD-WAN edge devices send BFD (Bidirectional Forwarding Detection) Hello packets across each WAN link every 1 second. These small probes measure latency, jitter, and loss in real time.
- **Color-based tunnels:** Each WAN link has a color (mpls, internet, lte). BFD probes run inside IPsec tunnels between edge devices on each color.
- **SLA evaluation:** Every second, for each active application flow, the system checks: Is the current link meeting SLA? If yes, keep. If no, pick next best.
- **Path switching:** When a path switch is needed, the system moves new flows immediately and can gracefully migrate existing flows (to avoid reset).

■ REAL-WORLD SCENARIO: Law Firm — Zoom vs Contract Database

Kapoor & Associates law firm has 3 offices: Mumbai HQ, Delhi branch, Pune branch. Links: MPLS 50Mbps + Broadband 200Mbps at each site. SCENARIO: Friday 3pm — 40 lawyers simultaneously on Zoom depositions. Also, the Mumbai server is syncing a 50GB contract archive to Delhi. WITHOUT SD-WAN: Both Zoom and the file sync compete for bandwidth. Zoom calls break up. Lawyers are frustrated during critical depositions. WITH SD-WAN AAR: The system identifies Zoom (UDP 8801, Zoom signature) as high-priority. SLA policy: Zoom → MPLS preferred, latency threshold 150ms. File sync → Broadband only. Even if broadband is saturated, Zoom is untouched on MPLS. The system automatically shapes the file sync bandwidth to avoid congesting the broadband link's DSCP queues. Zoom calls are crystal clear. File sync takes 30 minutes longer — which nobody cares about.

■■ PRACTICE PROJECT: Build a 5-Application AAR Policy

TASK: Design a complete Application-Aware Routing policy for an e-commerce company with 2 links: MPLS (50Mbps, guaranteed) and Broadband (200Mbps, best-effort). Applications to cover: 1. Payment gateway traffic (TCP to IP 10.50.0.0/24, port 8443) 2. Warehouse video surveillance (RTSP, multicast) 3. Employee Zoom/Teams calls 4. SAP ERP (proprietary ports, internal DC) 5. General internet browsing For each app: Define the identification method, SLA requirements, primary path, fallback path, what metric breach triggers failover, and what happens if ALL paths are down. Present as a table. Then write the policy in Cisco vManage pseudo-config format.

CHAPTER 5

Security in SD-WAN

Encryption, segmentation, zero trust, and next-gen firewall integration

Security is not an afterthought in SD-WAN — it is embedded at every layer. Because SD-WAN routes corporate traffic over the public internet, robust security is non-negotiable. This chapter covers everything from the cryptographic foundations to real-world segmentation designs.

5.1 IPsec Tunnels — How Data is Actually Protected

Every SD-WAN tunnel is an **IPsec tunnel**. IPsec operates at Layer 3 and provides authentication + encryption for all traffic. Here is exactly what happens when a packet enters an SD-WAN edge device:

- **Step 1 — Classify:** Packet arrives from LAN. Edge device checks policy to determine destination and required security.
- **Step 2 — Encrypt:** Packet is encrypted using AES-256-GCM. The original packet becomes the IPsec payload. Encryption happens in hardware on modern devices — no performance penalty.
- **Step 3 — Authenticate:** HMAC-SHA-256 signature is added. The receiver can verify the packet was not tampered with in transit.
- **Step 4 — Encapsulate:** An outer IP header is added with the WAN IP addresses of the two edge devices. This is what travels across the internet.
- **Step 5 — Transmit:** The encrypted, authenticated, encapsulated packet goes out the WAN interface.
- **Decapsulation at receiver:** Reverse process. Outer header stripped, HMAC verified, packet decrypted, original packet delivered to LAN.

Before IPsec can encrypt traffic, the two endpoints m
SD-WAN uses IKEv2 (Internet Key Exchange version 2)
exchange — a mathematical method that lets two parties
transmitting the key itself. Even if an attacker intercepts
cannot derive the key. Keys are automatically rotated ever
potential compromise.

5.2 PKI — Certificate-Based Authentication

Every device in an SD-WAN fabric has a **digital certificate** — a cryptographic identity document issued by the SD-WAN controller acting as a Certificate Authority (CA). This is how the network knows you're a legitimate device:

- **On provisioning:** The Orchestrator generates a unique certificate for each device containing: Device Serial Number, Organization ID, Signing CA, Public Key, and Validity Period.
- **On connect:** When a device tries to connect to the controller, it presents its certificate. The controller verifies: Is this cert signed by our CA? Is it expired? Has it been revoked?
- **Mutual authentication:** Both sides authenticate each other. A rogue device with no cert or a wrong cert is rejected before any data flows.
- **Certificate Revocation:** If a device is stolen or compromised, its certificate is revoked. Immediately, that device cannot join any SD-WAN tunnel.

5.3 Segmentation with VRF

VRF (Virtual Routing and Forwarding) creates completely separate routing tables within a single edge device. Traffic in different VRFs cannot communicate unless explicitly permitted (via route leaking or a firewall). This is the foundation of network segmentation in SD-WAN.

- **Corp VRF:** Employee workstations, corporate apps. Full access to DC.
- **Guest VRF:** Visitor Wi-Fi. Internet only. Zero access to Corp.
- **IoT VRF:** Printers, cameras, sensors. Restricted to device management only.
- **OT VRF:** Industrial control systems. Completely isolated, may not even have internet access.

■ REAL-WORLD SCENARIO: Manufacturing Plant — OT/IT Network Isolation

AutoParts Ltd. has a factory floor with 200 industrial robots (OT systems), SCADA controllers, and 500 office employees (IT users), all in the same building. RISK: If a ransomware attack hits the office network, it must NOT be able to reach SCADA systems and shut down production lines. This happened to multiple factories in 2023. SD-WAN DESIGN: 3 VRFs deployed: VRF 'corp' — All employee PCs, laptops. Access to internet + DC apps. VRF 'scada' — All OT/SCADA devices. NO internet access. Reaches only OT management server in DC via dedicated encrypted tunnel. Completely isolated. VRF 'mgmt' — IT management tools, monitoring. Limited cross-VRF visibility for ops team. SD-WAN edge device has hardware-enforced VRF separation. Even with the SD-WAN device compromised, kernel-level VRF tables prevent cross-contamination. Result: Factory survived a major ransomware attack in 2024 — OT was untouched, production continued, only office systems were affected.

■■ PRACTICE PROJECT: Design a Segmented SD-WAN Security Architecture

TASK: A hospital network needs the following segments: 1. Clinical staff (EMR access, video consultations) 2. Administrative staff (email, Office 365, HR systems) 3. Medical IoT devices (infusion pumps, monitors — NEVER internet-connected) 4. Patient entertainment/guest Wi-Fi 5. Building management (HVAC, access control) Design: Create a VRF architecture. Define which VRFs can route to each other. Define what traffic from each VRF can reach the internet. Define IPsec policy per VRF (which ones need encryption even on internal LAN). Define certificate policy — which devices get certificates, who is the CA. Draw the architecture. Write a 1-page security policy document for this hospital.

CHAPTER 6

Zero-Touch Provisioning (ZTP) & Day-0 Operations

How to deploy hundreds of sites without sending technicians

ZTP is the feature that makes SD-WAN deployments at scale actually feasible. Without ZTP, deploying 500 retail stores would require sending a network engineer to each site — astronomical cost. With ZTP, a non-technical person plugs in the device and the network configures itself automatically.

6.1 What Happens at Boot — The ZTP Dance

When a factory-fresh SD-WAN edge device is powered on, a precise sequence of events unfolds. Understanding this sequence is critical for troubleshooting failed ZTP deployments — which happen regularly in production.

- **Step 1 — DHCP Request:** Device sends DHCP broadcast on WAN interface. Gets an IP address from the ISP. Internet connectivity is now available.
- **Step 2 — DNS Lookup:** Device looks up 'devicehelper.viptela.com' (Cisco) or vendor-specific DNS. This resolves to the cloud ZTP service.
- **Step 3 — ZTP Service Contact:** Device sends its Serial Number to the ZTP service. The cloud service checks: Is this serial pre-registered by a customer?
- **Step 4 — Redirect to Orchestrator:** ZTP service returns the address of the customer's Orchestrator (vBond in Cisco terms). Device is directed home.
- **Step 5 — Orchestrator Registration:** Device contacts Orchestrator, presents serial number. Orchestrator verifies against allowed device list. Issues certificate. Device is now authenticated.
- **Step 6 — Controller Discovery:** Orchestrator tells device where the Controllers (vSmart) are. Device establishes OMP session.
- **Step 7 — Configuration Download:** Controller pushes templates and policies from vManage. Device applies configuration. Site is live.

For ZTP to work, three things must happen BEFORE the d be added to the Orchestrator's allowed list. (2) A device t associated with this serial. (3) The device must have th address pre-loaded (this happens at factory for new device these is the #1 reason ZTP fails in production.

■ **REAL-WORLD SCENARIO: Deploying 100 Retail Stores in 30 Days**

FashionRetail Ltd. acquired 100 new store locations and needs them all live in 30 days. Traditional approach: Send network engineer to each store (100 trips). Cost estimate: Rs. 15,000 per site visit × 100 = Rs. 15,00,000. Plus 3+ months time. SD-WAN ZTP Approach: Week 1: IT team creates one 'standard store' template in vManage. Variables defined: site name, store ID, regional MPLS gateway IP. Week 2: 100 edge devices ordered from Cisco. Serial numbers loaded into vManage. Each serial linked to the standard template with site-specific variables. Week 3: Devices shipped directly from Cisco to stores. Store manager (non-technical) receives device with instructions: 'Plug cable 1 into ISP1, Cable 2 into ISP2. Press power button. Wait 15 minutes.' Week 4: All 100 stores auto-configured. IT team monitors provisioning status in vManage. Zero site visits required. Total cost: ZERO travel. Time saved: 3 months. Risk: Near-zero (fully templated, tested).

■■ PRACTICE PROJECT: ZTP Lab Simulation

TASK: Using Cisco DevNet SD-WAN Sandbox (free at developer.cisco.com), perform: 1. Create a device template for a standard branch (2 WAN interfaces, 1 LAN, OSPF on LAN) 2. Add a simulated device serial to the allowed devices list 3. Attach the template to the device with variables: hostname=branch-test-01, system-ip=1.1.1.100, site-id=100 4. Verify: Template pushes successfully. Check vManage shows device in 'in-sync' state. 5. Simulate ZTP failure: Remove device from allowed list. Observe what happens. 6. Bonus: Write a Python script using Cisco SD-WAN REST API to list all devices and their sync status.

CHAPTER 7

Cloud Connectivity & Direct Internet Access

Connecting branches to cloud natively — no backhauling

The cloud is now where enterprise workloads live. Office 365, Salesforce, SAP HANA Cloud, AWS, Azure — if your SD-WAN can't connect to cloud efficiently, it is useless. This chapter covers the complete cloud connectivity strategy.

7.1 Local Internet Breakout — The Game Changer

Traditional WAN forces ALL internet traffic through the data center — even cloud apps. This creates a huge problem as SaaS usage grows. With SD-WAN, each branch can **break out directly to the internet** for cloud apps, while still sending DC-bound traffic through the overlay.

- **Before:** Branch user opens Outlook → traffic goes branch→MPLS→DC→Internet→Microsoft→DC→MPLS→branch. Total hops: 6. Latency added: 80-150ms. DC bandwidth wasted.
- **After:** Branch user opens Outlook → SD-WAN identifies Office 365 → routes directly to internet → Microsoft. Total hops: 2. Latency: 10-30ms. DC bandwidth saved.

With local breakout, you typically split traffic into two cate
overlay, full security scanning) and UNTRUSTED (internet/
cloud security like Zscaler). The split is done in SD-WAN
IP, FQDN lists (Microsoft publishes its IP ranges), and app
Zoom all publish their IP/domain lists specifically for this pu

7.2 Cloud OnRamp for SaaS

Even with local breakout, there is a problem: not all internet paths to a SaaS provider are equal. The internet is full of congestion. Cloud OnRamp for SaaS solves this by continuously probing multiple paths and selecting the best one.

Here is how it works for Office 365:

- SD-WAN edge at each branch sends probes to *.office.com from each WAN interface (MPLS, broadband, LTE)
- Probes measure latency and loss to the Microsoft edge PoP
- System picks the WAN interface providing best path to Microsoft
- This can change hour by hour as internet conditions vary
- Result: 25-40% improvement in SaaS app performance with no hardware changes

7.3 Cloud OnRamp for IaaS — AWS, Azure, GCP

When your applications run in cloud VPCs (not just SaaS), you need SD-WAN tunnels INTO the cloud. Cloud OnRamp for IaaS creates a virtual SD-WAN edge inside your cloud VPC — extending the overlay fabric all the way to the cloud.

- **AWS:** Deploy Cisco CSR1000v or vEdge in AWS VPC. Establish IPsec tunnels back to physical SD-WAN edges. Branch traffic to AWS goes optimally — direct internet or overlay depending on policy.
- **Azure:** Same concept using Azure Virtual WAN or a virtual router in Azure.
- **Multi-region:** Deploy virtual edges in AWS Mumbai + AWS Singapore. Branches pick the nearest cloud region automatically.

■ REAL-WORLD SCENARIO: SaaS-Heavy Startup Goes SD-WAN

TechStartup.io has 200 employees across Mumbai, Bangalore, Hyderabad. Their entire stack is SaaS: Office 365, Slack, Jira, Salesforce, AWS-hosted product. Problem: Old MPLS setup routes ALL traffic through Mumbai DC. Bangalore users accessing Slack: 180ms latency (Bangalore→MPLS→Mumbai→Internet→Slack). With local broadband: 35ms (Bangalore→Internet→Slack). SD-WAN Solution: Primary: Broadband DIA at each office (300Mbps, cheap) Backup: LTE (for failover only) SaaS breakout: Enabled for Office 365, Slack, Salesforce using FQDN lists AWS workload: Cloud OnRamp to AWS Mumbai ap-south-1 region Security: All internet traffic through Zscaler cloud proxy (integrated with SD-WAN) Result: Latency to SaaS dropped 80%. Monthly WAN cost dropped 60% (no MPLS). The startup scaled from 200 to 500 employees with zero new infrastructure.

■■ PRACTICE PROJECT: Design Cloud OnRamp Architecture for AWS

TASK: A company runs their core application in AWS (Mumbai region, ap-south-1) with a 3-tier architecture: Load Balancer, Application Servers, RDS Database. They have 5 branches (Chennai, Delhi, Pune, Kolkata, Ahmedabad) each with Broadband (100Mbps) + LTE backup. Design the Cloud OnRamp for IaaS setup: 1. Where to place the virtual SD-WAN edge in AWS (VPC subnet, AZ) 2. IPsec tunnel design (how many tunnels, redundancy) 3. Routing: Which traffic goes through overlay to AWS vs. direct internet 4. Security: VPC security groups for the virtual edge, what ports to allow 5. Cost estimate: EC2 instance type, data transfer costs 6. Failover: What happens if AWS ap-south-1 is unavailable? Define the DR path.

CHAPTER 8

SASE — The Evolution of SD-WAN

When networking and security converge in the cloud

SASE (Secure Access Service Edge, pronounced 'sassy') is the convergence of SD-WAN with cloud-delivered security. Defined by Gartner in 2019, SASE has become the dominant enterprise network strategy. If SD-WAN is the car, SASE is the entire transportation system — including roads, traffic laws, and police.

8.1 The 5 Components of SASE

Component	Full Name	What It Does	Example Vendors
SD-WAN	Software-Defined WAN	Smart routing, link optimization, application steering	Cisco, VMware, Fortinet
SWG	Secure Web Gateway	Web filtering, malware scanning, URL categorization for all web traffic	Zscaler, Netskope, Symantec
CASB	Cloud Access Security Broker	Controls which cloud apps employees use, stops data leakage	Netskope, McAfee, Zscaler
ZTNA	Zero Trust Network Access	Identity-based access to apps — replaces VPN	Zscaler, Cloudflare, Okta
FWaaS	Firewall as a Service	Cloud-delivered enterprise firewall with IPS/IDS, full L7 inspection	Palo Alto, Fortinet, Check Point

8.2 Zero Trust Network Access (ZTNA) — The VPN Replacement

Traditional VPN: User connects to VPN → gets full network access to company network. If user's laptop is compromised, attacker has full network access. This is catastrophic. ZTNA flips this model completely: **Assume Breach. Verify Everything. Grant Minimum Access.**

- **Identity verification:** Every access request requires the user to prove identity (MFA, certificate, device posture check).
- **Device posture:** Is the device encrypted? Is antivirus updated? Is it a company-managed device? Fail any check → deny access.
- **Application-level access:** User gets access to SPECIFIC apps, not the whole network. Salesforce user? Can reach Salesforce. Nothing else.
- **Continuous verification:** Access is re-verified constantly. If something changes (anomalous behavior, new location), access is revoked immediately.

■ **REAL-WORLD SCENARIO: Remote Workforce Transformation**

COVID-2020: 10,000-employee company sends everyone home overnight. They had 500 VPN licenses. VPN infrastructure collapses. Everyone is locked out. SASE SOLUTION deployed in 72 hours: SD-WAN: Employees get home router SD-WAN devices (or software clients) on their home broadband. ZTNA: Replaces VPN. Every employee authenticates with MFA. Gets access only to their needed apps. SWG: All internet traffic from home devices goes through cloud SWG — same web filtering as office. CASB: Shadow IT blocked. Employees can't upload company files to personal Dropbox. FWaaS: Home network is inspected. Compromised family member device can't infect work laptop. 10,000 employees fully operational within 72 hours. Security posture IMPROVED vs VPN era. IT team went from fighting fires to monitoring dashboards.

■■ PRACTICE PROJECT: SASE Architecture Design

TASK: Design a full SASE architecture for a 500-person financial services firm with: - 3 offices (Mumbai HQ, Delhi, Pune) - 200 full-time remote employees - 100 third-party contractors (untrusted devices) - Applications: Core banking (on-prem DC), Office 365, Salesforce, custom trading app (AWS) Design deliverables: 1. SASE platform choice (single vendor like Fortinet or multi-vendor like Cisco+Zscaler) 2. ZTNA policy for each user group (office, remote employee, contractor) 3. Security service chaining diagram (which traffic hits which security service) 4. CASB policy: List 5 specific controls for Office 365 and Salesforce 5. SWG URL policy: Define 10 categories, allow/block decision with business justification 6. Incident response: What happens when a remote employee's device is detected as compromised?

CHAPTER 9

Day-2 Operations — Monitoring, Troubleshooting & Optimization

The daily life of an SD-WAN engineer

Most SD-WAN courses teach you how to deploy. Few teach you what to do at 2am when the CFO's Zoom call drops during a board meeting. This chapter is about real operational knowledge — the stuff that makes the difference between a junior and a senior SD-WAN engineer.

9.1 Key Metrics to Monitor Daily

Metric	What It Means	Normal Range	Alert Threshold	Action
WAN Link Loss	Packets dropped on WAN	0 - 0.1%	> 0.5%	Check ISP, activate backup
WAN Latency	RTT per link	5-80ms	> 150ms	Check routing, ISP issue
Tunnel State	IPsec tunnel up/down	Up	Down	Check IKE, certificates, NAT
CPU Utilization	Edge device CPU	< 60%	> 85%	Check traffic volume, DPI loa
OMP Session	Controller connectivity	Up	Down	Check control plane connectiv
Certificate Expiry	Days until cert expires	> 90 days	< 30 days	Renew certificates ASAP
BFD Session State	Per-tunnel probing	Up	Down/Flap	Investigate WAN instability
App Performance	SLA breaches per app	0	> 0	Review AAR policy, link qualit

9.2 Troubleshooting Playbooks

Playbook 1: Application Performance Complaint

User says: 'My Zoom calls are terrible from the Bangalore office'

Step	Action	Command / Tool	What to Look For
1	Verify the complaint is real	vManage → Monitor → Applications → Zoom	High latency, jitter > 50ms, loss > 1%
2	Check which WAN link Zoom is using	vManage → Monitor → Real-time → App Route	Is it on MPLS or broadband? Correct per policy?
3	Check link quality at time of complaint	vManage → Alarms → Historical	Was there a link quality event during the call?
4	Verify SLA policy is applied correctly	vManage → Configuration → Policies	Is Zoom in the SLA policy? Correct thresholds?

5	Check edge device health	SSH to edge → show bfd sessions	BFD sessions up? Probe loss rates?
6	Contact ISP if link quality degraded	ISP ticket portal	Get ISP to run diagnostics on circuit ID

Playbook 2: Site Connectivity Lost

Alert: 'Branch site offline — no overlay connectivity'

- **Immediate:** Check vManage dashboard. Is control connection (OMP) up? Is data plane (BFD) up? Are any tunnels up?
- **If OMP down but BFD up:** Controller connectivity issue. Check Orchestrator/Controller health. Check firewall rules — TCP 12346 must be open.
- **If both OMP and BFD down:** Site has no WAN. Check physical connections. Check ISP. Check edge device power/status LEDs.
- **If primary down, backup should auto-kick in:** If backup didn't activate, check if LTE SIM has data. Check backup interface configuration.
- **If single tunnel down but others up:** IPsec negotiation failure. Check NAT-T settings. Check certificate validity. Bounce the tunnel.

■ REAL-WORLD SCENARIO: Outage at Call Center — 2am Crisis

3am. 500 call center agents cannot reach the CRM. The SD-WAN engineer gets the call. INVESTIGATION: vManage shows: Pune call center site — OMP session DOWN, BFD sessions DOWN on both links. Edge device unreachable via out-of-band console. This means the device itself has an issue. ROOT CAUSE FOUND: Power spike corrupted the edge device flash memory. Device is boot-looping. SHORT-TERM FIX: IT team at the call center manually connects an LTE router as emergency backup. SD-WAN is bypassed. CRM traffic restored in 45 minutes. Call center resumes operations. LONG-TERM FIX: Replacement edge device overnighted. ZTP re-provisions automatically next morning. Normal SD-WAN restored by 9am. LESSON: Always have an out-of-band management path (OOB console, separate management cellular). Always stock 2-3 spare edge devices per 50 production sites. Define an emergency bypass procedure in your runbook.

■■ PRACTICE PROJECT: Build a Network Operations Runbook

TASK: Create a professional SD-WAN Operations Runbook document covering: Section 1: Daily health check procedure (what to check every morning, checklist format) Section 2: Alert response matrix (for each alert type, define: severity, response time, first responder, escalation path) Section 3: Top 10 troubleshooting playbooks (site down, app slow, tunnel flapping, certificate expiry, ZTP failed, BGP neighbor down, VRF leak, high CPU, policy not applied, backup link not activating) Section 4: Change management process (how to push a policy change safely, rollback plan) Section 5: Maintenance window procedures (how to update firmware on 100 edges safely) This runbook should be something a Level 1 NOC engineer can use at 2am with no escalation needed.

CHAPTER 10

Cisco SD-WAN Deep Dive

vManage, vSmart, vBond, vEdge — the dominant enterprise platform

Cisco SD-WAN (originally Viptela, acquired 2017, now 'Cisco Catalyst SD-WAN') is the market leader with 35%+ market share. Understanding it deeply is the single most valuable thing you can do for your SD-WAN career.

10.1 The Four Core Components

Component	Old Name	Role	Deployment	Criticality
SD-WAN Manager	vManage	Management & Orchestration GUI. All configuration, monitoring, reporting	On-prem VM or Cisco cloud	Critical — but degraded ops
SD-WAN Controller	vSmart	Control plane. Distributes routes and policies via OMP. Authenticates devices	On-prem or cloud (cluster)	Critical — deploy 2+ always
SD-WAN Validator	vBond	Orchestrates initial device onboarding. First contact for ZTP.	DMZ or Cisco cloud. Must be Internet reachable	Critical — ZTP fails without
SD-WAN Edge	vEdge / CSR / ISR / Cat8K	Data plane. Forwards traffic. Runs AAR/DPI/IPsec.	Branch/DC/Cloud. Physical or virtual	Critical — this is the network

10.2 Template Architecture

Templates are how Cisco SD-WAN manages configuration at scale. Instead of configuring each device individually, you define templates that are applied to many devices. Templates have two levels:

- **Feature Templates:** Configure one specific feature. Examples: AAA template, BGP template, VPN interface template, OSPF template. One feature template can be reused across many device templates.
- **Device Templates:** Combine multiple feature templates into a complete device configuration. One device template applied to 50 branches = 50 identically configured branches. Variables (site name, IP, etc.) fill in the blanks.

Device templates use variables ({{variable_name}}) for
template to a device, vManage asks for variable value
hostname, tunnel-interface-ip, ospf-neighbor-ip. You car
crucial for deploying 100+ sites. Best practice: Define AL
hard-code them in templates. This allows one template to s

10.3 Cisco SD-WAN Policy Framework

Cisco SD-WAN has a powerful but complex policy framework with three main policy types:

- **Localized Policies:** Applied directly on the edge device. Controls local traffic (ACLs, QoS, route policies for LAN-side traffic). Think of it as the device's own local rules.
- **Centralized Policies:** Computed centrally on vSmart, distributed to edges. Controls overlay behavior — topology (who can talk to whom), traffic (AAR, data policies), and VPN segmentation. This is where the magic happens.
- **Security Policies:** Firewall, IPS, URL filtering rules applied at edge. Configured in vManage and pushed to edges that support security features.

■ REAL-WORLD SCENARIO: Enterprise Migration to Cisco SD-WAN

LargeBank Corp: 250 branches, 15 regional offices, 2 data centers, 50,000 users. Migrating from pure MPLS to Cisco Catalyst SD-WAN with hybrid MPLS+broadband. MIGRATION STRATEGY — 5 phases: Phase 1 (Month 1-2): Deploy vManage, vSmart (clustered x3), vBond. Build and test templates. Deploy at 5 pilot branches. Validate all policies. Phase 2 (Month 3-4): Deploy at 25 regional offices. These are complex sites with BGP, multiple VRFs, and high traffic. Extra testing time required. Phase 3 (Month 5-8): Mass deploy remaining 225 branches using ZTP. 100 sites per month cadence. Support team monitors each site as it comes online. Phase 4 (Month 9-10): Enable advanced features: AAR, Cloud OnRamp, SASE integration. Phase 5 (Month 11-12): Decommission MPLS circuits where broadband is proven stable. Target: 40% MPLS cost reduction.

■■ PRACTICE PROJECT: Full Cisco SD-WAN Lab Build

TASK: Using Cisco DevNet SD-WAN Sandbox or CML (Cisco Modeling Labs — has free trial): Build 1: 1 vManage + 1 vSmart + 1 vBond + 2 vEdge devices Build 2: Create device templates for Hub (DC) and Branch profiles Build 3: Configure OMP between vSmart and both vEdges Build 4: Create a centralized policy: Branch1 and Branch2 can talk directly for VoIP (hub-bypass), but all other traffic must go through hub Build 5: Configure AAR — Zoom goes MPLS (color=mpls), bulk traffic goes internet (color=biz-internet) Build 6: Simulate MPLS link failure — verify Zoom traffic moves to internet fallback Build 7: Add a new 'Branch3' using ZTP from scratch Document all configs. Screenshots of vManage showing healthy state.

CHAPTER 11

Real-World Project Scenarios A to Z

Five complete industry deployments with all design decisions explained

This chapter presents five complete, real-world SD-WAN deployment scenarios. Each scenario is a complete design exercise — the kind you would encounter in an actual enterprise project. Work through each one as if you are the lead architect.

11.1 Scenario: Indian Private Bank — MPLS + SD-WAN Hybrid

Client: FirstWealth Bank. 180 branches, 2 DCs (Mumbai primary, Delhi DR), 45,000 users.

Existing infrastructure: Pure MPLS (BSNL), 2-10Mbps per branch. Monthly WAN bill: Rs. 2.4 crore.

Requirements: Reduce WAN cost by 50%. Improve branch connectivity to 50+ Mbps. Maintain RBI compliance (data must stay in India). Core banking must have 99.99% uptime. Complete deployment within 9 months.

Design Decisions:

- **Topology:** Hub-and-spoke. All branches connect to Mumbai DC hub. Direct branch-to-branch only for VoIP (to reduce latency).
- **Transport:** Tier-1 branches (>200 users): Keep MPLS (10Mbps) + Add Broadband (100Mbps). Tier-2 branches (<200 users): Replace MPLS with dual Broadband. Use LTE as tertiary backup.
- **Encryption:** AES-256 on all tunnels. Even MPLS traffic encrypted — RBI considers MPLS 'private' but bank policy requires encryption end-to-end.
- **Segmentation:** VRF 'banking' for core banking. VRF 'employee' for general apps. VRF 'atm' for ATM network. No inter-VRF routing at branches — only at DC via firewall.
- **AAR:** Core banking (TCP 5555) → MPLS primary, SLA: loss <0.1%, latency <50ms. VoIP → MPLS primary (T1 branches) or best broadband path. Internet → Local breakout. Never goes to DC.
- **Result:** WAN cost reduced to Rs. 1.2 crore/month (50% savings). Branch bandwidth increased 5-10x. Zero core banking outages in first year.

11.2 Scenario: Retail Chain — 500 Stores, ZTP Rollout

Client: QuickMart Retail. 500 stores across India, opening 5 new stores per month.

Requirements: POS must never fail (revenue loss). Guest Wi-Fi isolated. Surveillance cameras on separate segment. New stores live within 2 hours of opening.

Design Decisions:

- **Standard Store Template:** 1 single template for all 500 stores. Variables: store-id, city, primary-ISP-IP, secondary-ISP-IP.
- **Links:** Dual broadband (different ISPs per store). LTE as emergency backup.
- **VRFs:** VRF 'pos' — POS terminals only, highest priority, NEVER goes to guest ISP. VRF 'cctv' — cameras, local storage, no internet. VRF 'guest' — customer Wi-Fi, internet breakout, isolated.
- **AAR:** POS traffic (TCP 8443 to payment gateway IP) → primary broadband, failover to LTE, SLA: loss <0.5%. Guest traffic → secondary broadband, no failover.
- **ZTP process:** Store electrician unpacks device, plugs 2 WAN cables + LAN switch cable + power. 15-minute auto-provision. NOC team confirms green status in vManage.

11.3 Scenario: Healthcare — HIPAA-Compliant SD-WAN

Client: MedCare Hospitals. 25 hospitals, 150 clinics, 500,000 patient records.

Requirements: HIPAA compliance (encryption, audit trails, access controls). EHR (Epic) must be available 99.99%. Telemedicine must have HD video quality. Medical IoT devices (infusion pumps, monitors) must be network-isolated.

HIPAA-Specific Design Decisions:

- **Encryption everywhere:** AES-256 on all WAN tunnels including MPLS. HIPAA requires encryption of PHI in transit. Document this for compliance audit.
- **Audit logging:** All traffic flows logged. 90-day retention minimum. vManage logs sent to SIEM for security monitoring.
- **Access control:** Each VRF has explicit firewall policy. PHI-containing systems in isolated VRF. Only clinician devices have access.
- **Medical IoT:** Separate VRF with RFC 1918 space. Can only communicate to medical device management server. Cannot reach internet or other VRFs.
- **Telemedicine:** Identified by application signature. Routed to MPLS with strict SLA policy: latency <100ms, jitter <20ms.

11.4 Scenario: Manufacturing — OT/IT Convergence

Client: AutoSteel Ltd. 8 factories, 200 robots per factory, SCADA systems controlling production. Recent ransomware attack cost $5M.

Critical Design Decisions:

- **Air-gap simulation via VRF:** OT VRF has zero default route to internet. OT systems can ONLY reach OT management server in DC. This is enforced at policy level.
- **Micro-segmentation:** Each production line is its own /28 subnet within OT VRF. A compromised robot cannot reach other production lines.
- **Anomaly detection:** SD-WAN logs all OT traffic. Any new connection type from OT device triggers immediate alert.
- **Out-of-band management:** SD-WAN management traffic runs on separate cellular connection — if production network is attacked, management is unaffected.

11.5 Scenario: School District — SD-WAN + SASE

Client: Delhi Public Schools. 80 schools, 50,000 students, federal funding requires content filtering (CIPA compliance).

Design:

- **Low-cost links:** Each school gets two cheap broadband connections. No MPLS (budget constraints). SD-WAN provides resilience.
- **SASE integration:** All student web traffic goes through cloud SWG (Zscaler or similar). Content filtering for CIPA compliance. No explicit proxy needed at schools.
- **Segmentation:** VRF 'staff' — teacher/admin computers. VRF 'student' — Chromebooks, tablets. VRF 'byod' — personal devices, internet only.
- **E-rate funding:** Design documented for FCC E-rate subsidy application. SD-WAN qualifies under Category 2 services.

CHAPTER 12

Career, Certifications & Your 90-Day Plan

How to become a sought-after SD-WAN professional

12.1 Job Roles & Salary Landscape

Role	Experience	Key Skills	India Salary	USA Salary
Network Engineer (SD-WAN)	0-2 years	CCNA, basic SD-WAN ops, monitoring	Rs. 4-8 LPA	$70K-90K
SD-WAN Implementation Engineer	2-4 years	CCNP, deployment, templates, troubleshooting	Rs. 8-15 LPA	$90K-115K
SD-WAN Solutions Architect	4-7 years	Design, vendor-agnostic, cloud, SASE, presales	Rs. 15-25 LPA	$120K-150K
Network Architect (SD-WAN/SASE)	7+ years	Multi-vendor, cloud architecture, leadership	Rs. 25-40 LPA	$150K-180K
SD-WAN Product Manager	5+ years + MBA	Product strategy, vendor mgmt, roadmap	Rs. 30-50 LPA	$160K-200K

12.2 Certification Roadmap

Certification	Vendor	Level	Cost (USD)	Validity	Why Get It
CCNA (200-301)	Cisco	Foundation	$330	3 years	Essential prerequisite for ever
CCNP ENSDWI (300-415)	Cisco	Professional	$400	3 years	The top SD-WAN cert. Most j
NSE 4 / NSE 7 SD-WAN	Fortinet	Professional	Free-$400	2 years	Best for Fortinet shops, free N
VCP-NV (VeloCloud)	VMware	Professional	$250	2 years	Good for VMware/Broadcom
PCNSE (Prisma)	Palo Alto	Professional	$350	2 years	Best SASE/Prisma cert
AWS ANS-C01	Amazon	Specialty	$300	3 years	Excellent for cloud+SD-WAN
JNCIP-ENT	Juniper	Professional	$400	3 years	Juniper SD-WAN / Session S

12.3 Top Interview Questions & Model Answers

Q: What is the difference between the underlay and overlay in SD-WAN?

A: The underlay is the physical transport infrastructure — MPLS circuits, broadband connections, LTE links — the actual cables and ISP networks that carry bits. The overlay is a logical, encrypted virtual network built on top of the underlay using IPsec tunnels. From the overlay's perspective, all sites appear directly connected. The overlay abstracts away the physical complexity and provides consistent encryption and QoS regardless of which underlay is used.

Q: How does Application-Aware Routing work?

A: AAR uses Deep Packet Inspection (DPI) to identify applications from packet signatures, ports, and URLs. SLA policies define quality requirements per application — for example, Zoom requires latency <150ms and loss <1%. BFD probes measure real-time quality of each WAN link every second. When a link's quality degrades below the SLA threshold for an application, that traffic is automatically moved to the next-best link that meets the SLA. This happens within seconds, transparently to users.

Q: What happens if the SD-WAN controller goes down?

A: The control plane and data plane are separated in SD-WAN. If the controller (vSmart) fails, edge devices continue forwarding traffic using their last-known routing table — this is called Data Plane Resilience or Forward Data Plane. Traffic continues flowing for hours. However, routing changes won't propagate (new routes, policy changes won't work). This is why controllers should always be deployed in active-active pairs minimum. In Cisco SD-WAN, deploy minimum 2 vSmart controllers for production.

Q: Explain the ZTP process for a new branch deployment.

A: ZTP (Zero-Touch Provisioning) works in these steps: (1) Pre-stage: Add device serial to Orchestrator, create and attach device template with variables. (2) On-site: Non-technical staff plug in device and power. (3) Boot: Device gets DHCP IP, does DNS lookup for vendor cloud ZTP service. (4) ZTP service redirects to customer's Orchestrator. (5) Orchestrator verifies serial, issues certificate, tells device where controllers are. (6) Device contacts controller, downloads template from vManage, applies configuration. Site is live.

Q: How would you troubleshoot a site where Zoom calls are dropping?

A: Step 1: Confirm in vManage → Monitor → Applications that Zoom is showing high latency/loss. Step 2: Check which WAN link Zoom is currently using vs. the configured SLA policy. Step 3: Check real-time BFD probe statistics for all WAN links at that site. Step 4: Review historical alarms — was there a link quality event during call drops? Step 5: SSH to edge device, run 'show app-route statistics' for Zoom flows. Step 6: If link quality poor, open ISP ticket. If routing unexpected, verify SLA policy is pushed correctly in vManage. Step 7: Temporarily change SLA thresholds or manually steer traffic to test hypothesis.

12.4 Your 90-Day Zero-to-Hero Study Plan

Days	Focus	Activities	Milestone
1-15	Foundations & Concepts	Read Chapters 1-3 of this guide. Watch Cisco CCNA [illegible]	[illegible]
16-30	Hands-On Lab Setup	Create Cisco DevNet account (free). Access SD-WAN Sandbox. [illegible]	[illegible]

31-50	Advanced Features	Complete Chapters 7-9. Implement AAR policies in la[illegible]	Full AAR policy working. [illegible]
51-70	Certification Prep	Register for Cisco CCNP ENSDWI (300-415). Study [illegible]	Practice exam score consistently > 75% [illegible]
71-85	Projects & Portfolio	Complete 3 project designs from this guide. Docume[illegible]	LinkedIn profile updated. [illegible] portfolio [illegible]
86-90	Job Search Launch	Apply to SD-WAN implementation engineer roles. Pr[illegible]	Active applications submitted. [illegible]

ath

SD-WAN is not just a technology — it's the foundation
operates. Every company with more than 10 offices is eith
migrate. The engineers who truly understand it — not jus
the routing protocols, the security model, the cloud in
well-compensated. The projects and scenarios in this gui
through every single one of them. Don't just read — build
how networking is truly learned. Connect with me on Link
ideas, and career advice: linkedin.com/in/gautam-s2024

Appendix: Master Glossary

Term	Full Form	One-Line Meaning
AAR	Application-Aware Routing	Routing decisions based on application type and real-time link quality
BFD	Bidirectional Forwarding Detection	Lightweight probe protocol measuring link quality every second
CASB	Cloud Access Security Broker	Controls which cloud applications employees are allowed to use
CPE	Customer Premises Equipment	The physical edge device installed at a branch office
DIA	Dedicated Internet Access	A business-grade internet circuit (dedicated, not shared)
DPI	Deep Packet Inspection	Examining packet content to identify applications, not just ports
DSCP	Differentiated Services Code Point	A 6-bit marking in IP header indicating traffic priority
FWaaS	Firewall as a Service	Enterprise firewall delivered from the cloud, no hardware needed
IKEv2	Internet Key Exchange v2	Protocol for securely negotiating IPsec encryption keys
IPsec	IP Security	A suite of protocols for encrypting and authenticating IP traffic
LLQ	Low Latency Queuing	A QoS queue that always forwards VoIP before anything else
MPLS	Multiprotocol Label Switching	Private WAN technology using labels to forward traffic on provider netwo
OMP	Overlay Management Protocol	Cisco's proprietary routing protocol between vEdge and vSmart
PKI	Public Key Infrastructure	System for issuing and managing digital certificates
SASE	Secure Access Service Edge	Convergence of SD-WAN with cloud-delivered security services
SD-WAN	Software-Defined WAN	Software layer providing intelligent, policy-based WAN management
SLA	Service Level Agreement	Agreed quality metrics (latency, loss, jitter) for a service
TLOC	Transport Locator	Cisco SD-WAN identifier: System IP + Color + Encapsulation type
VRF	Virtual Routing and Forwarding	Separate routing table on one device, enabling network segmentation
vBond	vBond Orchestrator	Cisco SD-WAN component handling initial device onboarding and ZTP
vEdge	Virtual/Physical Edge	Cisco SD-WAN data plane device (physical or virtual)
vManage	vManage NMS	Cisco SD-WAN management and orchestration GUI
vSmart	vSmart Controller	Cisco SD-WAN control plane, distributes routes via OMP
WAN	Wide Area Network	Network connecting geographically dispersed sites
ZTNA	Zero Trust Network Access	Identity-based access model — verify user/device before every access

ZTP	Zero-Touch Provisioning	Automatic device configuration on first boot without manual intervention