

VMware Virtualization

Complete Course Guide

From Fundamentals to Advanced VMotion

Author

Gautam Shirsath

<https://www.linkedin.com/in/gautam-s2024>

Table of Contents

Chapter 1	The History of VMware	6
Chapter 2	Virtualization Fundamentals	14
Chapter 3	VMware vSphere Architecture	24
Chapter 4	Installing ESXi	34
Chapter 5	vCenter Server	44
Chapter 6	Virtual Machines (VMs) Deep Dive	54
Chapter 7	VMware Networking	66
Chapter 8	VMware Storage	80
Chapter 9	VMotion — Complete Guide	96
Chapter 10	High Availability (HA)	116
Chapter 11	Distributed Resource Scheduler (DRS)	128
Chapter 12	VMware Security	140
Chapter 13	Monitoring & Performance	154
Chapter 14	Backup & Disaster Recovery	166
Chapter 15	VMware Cloud & Modern Trends	178
Chapter 16	Hands-On Labs & Exercises	192

Appendix	Glossary & Quick Reference	204
----------	----------------------------	-----

Chapter 1: The History of VMware

Understanding where VMware came from helps you appreciate the revolution it brought to the world of computing. Before VMware, if you needed five servers, you bought five physical machines — each running one operating system, consuming space, electricity, and budget. VMware changed everything.

1.1 The Birth of an Idea — Stanford University, 1998

VMware was founded in 1998 by Mendel Rosenblum and Diane Greene, both from Stanford University, along with three colleagues: Scott Devine, Mohammad Rosenblum, and Edward Wang. The founding team was built on deep academic research into operating systems and computer architecture.

Mendel Rosenblum had been working on a research project at Stanford called 'Disco' — a system that could run multiple operating systems on the same hardware. The insight was groundbreaking: what if the OS could be abstracted from the hardware entirely? What if a piece of software could pretend to be hardware?

The Founding Moment

In a Stanford lab, Mendel Rosenblum asked a question that changed computing forever: 'What if software could simulate hardware so perfectly that an OS couldn't tell the difference?' This question became VMware. The company was incorporated in 1998, and within a year, they had a working product.

1.2 VMware Workstation — The First Product (1999)

In February 1999, VMware released VMware Workstation 1.0 — a product that allowed you to run multiple operating systems simultaneously on a single desktop PC. At the time, this seemed like magic. A developer could run Windows and Linux on the same machine at the same time, without rebooting.

Real-world example: Imagine a software developer in 1999 who needs to test her application on Windows 98, Windows NT, and Red Hat Linux. Before VMware, she needed three separate computers. After VMware Workstation, she only needed one. Her productivity skyrocketed and her hardware budget plummeted.

★ **Pro Tip:** VMware Workstation is still sold today and is used by developers worldwide for testing, development, and running legacy applications.

1.3 ESX Server — Bringing Virtualization to the Data Center (2001)

While Workstation was a desktop product, VMware set its sights on the data center. In 2001, VMware released ESX Server — a bare-metal hypervisor that ran directly on hardware without requiring a host operating system. This was a completely different architecture called 'Type 1 hypervisor.'

ESX (Elastic Sky X) was a technical marvel. Instead of sitting on top of Windows or Linux, it replaced the OS entirely. It owned the hardware directly, leading to much better performance, stability, and security. Data centers around the world began adopting it rapidly.

Before ESX Server	After ESX Server
1 physical server = 1 application	1 physical server = 10+ virtual machines
Server utilization: 5-15%	Server utilization: 70-80%
New server = 6-week procurement	New server = Minutes to provision
Hardware failure = Application downtime	Hardware failure = VM migrated automatically
Data center full of blinking lights	Same workloads, far less hardware

1.4 EMC Acquisition — 2004

In 2004, storage giant EMC Corporation acquired VMware for approximately \$635 million. This was considered a bargain in hindsight. EMC recognized that storage and virtualization were deeply complementary — if VMs were going to proliferate, you needed smart storage to support them.

The acquisition actually gave VMware more resources and distribution channels. EMC let VMware operate largely independently, preserving its innovative culture. This turned out to be brilliant — by 2007, VMware's IPO was one of the largest in tech history.

1.5 The vSphere Revolution — 2009

In 2009, VMware released vSphere 4.0 — a complete platform that bundled ESXi (the lightweight successor to ESX), vCenter Server, and a suite of enterprise features including VMotion, High Availability, and Distributed Resource Scheduler. This release marked VMware's transformation from a hypervisor company to a full cloud infrastructure platform.

The 'i' in ESXi stands for 'integrated' — the management console was integrated directly into the hypervisor, reducing the attack surface and eliminating the Service Console that existed in ESX. ESXi had a tiny 150MB footprint compared to the gigabytes required by ESX.

 **Note:** VMware retired the original ESX (without the 'i') in 2012. All modern VMware deployments use ESXi.

1.6 Dell-EMC and Broadcom Era

In 2016, Dell acquired EMC for \$67 billion — the largest tech acquisition in history at that time. VMware came along as part of the deal. Dell held VMware as a publicly traded subsidiary, giving it financial resources while maintaining operational independence.

In 2023, Broadcom completed its \$61 billion acquisition of VMware, marking yet another chapter. Broadcom has repositioned VMware's offerings, focusing heavily on enterprise licensing and subscription models. The core technology — ESXi, vCenter, VMotion — remains the bedrock of enterprise virtualization worldwide.

Year	VMware Milestone
1998	VMware founded at Stanford University
1999	VMware Workstation 1.0 released
2001	ESX Server launched — enters data centers
2003	vMotion introduced — live VM migration
2004	EMC acquires VMware for \$635M
2007	VMware IPO — one of the biggest in tech history
2009	vSphere 4.0 released — the modern platform begins
2012	ESX retired, ESXi becomes the only hypervisor
2016	Dell acquires EMC (and VMware) for \$67 billion
2019	vSphere 7.0 with Kubernetes integration
2023	Broadcom acquires VMware for \$61 billion

Chapter 2: Virtualization Fundamentals

Before diving into VMware features, we need to understand virtualization at its core. Think of virtualization as real estate: instead of one business owning an entire building and using only 10% of it, virtualization lets you divide that building into apartments and rent them out. Each tenant thinks they have their own building — but they're actually sharing resources efficiently.

2.1 What Is Virtualization?

Virtualization is the process of creating a software-based (virtual) representation of something — whether that's a server, storage, network, or even a desktop. In the context of VMware, we primarily talk about server virtualization: using software to make one physical computer act like many separate computers.

Real-world analogy: Think of a virtualization platform like a luxury apartment building. The building (physical server) has floors, plumbing, electricity, and structure. VMware is like the building management company that divides the building into individual apartments (virtual machines). Each tenant (guest OS) has their own private space — their own 'virtual apartment' — and can't see into other apartments.

Key Definition

A Virtual Machine (VM) is a software emulation of a physical computer. It has virtual CPU, virtual RAM, virtual hard disk, and virtual network adapters. The VM's operating system believes it's running on real hardware — because VMware makes sure of that illusion.

2.2 Types of Hypervisors

A hypervisor is the software layer that creates and manages virtual machines. There are two types, and understanding the difference is critical for any VMware professional.

Type 1 Hypervisor (Bare-Metal)

A Type 1 hypervisor runs directly on the physical hardware — there is no host operating system underneath it. It IS the operating system, in a sense. VMware ESXi is a Type 1 hypervisor. Other examples include Microsoft Hyper-V and Xen.

Performance: Excellent. Security: Superior. Stability: Industrial-grade. This is what runs in every serious data center on Earth.

Type 2 Hypervisor (Hosted)

A Type 2 hypervisor runs on top of an existing operating system, like an application. VMware Workstation and Fusion are Type 2. You install them just like you'd install Microsoft Word.

Performance: Good, but with overhead from the host OS. Best used for development, testing, and desktop scenarios.

Type 1 (ESXi)	Type 2 (Workstation)
Runs on bare metal	Runs on host OS (Windows/macOS/Linux)
Used in data centers	Used on developer workstations
Maximum performance	Some overhead from host OS
Requires dedicated hardware	Shares hardware with host OS
Managed via vCenter or web UI	Managed via desktop application
Enterprise-grade features	Developer/test features

2.3 Key Virtualization Concepts

Host vs. Guest

The 'host' is the physical machine and the hypervisor that runs on it. The 'guest' is the virtual machine running on top of the hypervisor. Your host might be a Dell PowerEdge server running ESXi 8.0. Your guests might be Windows Server 2022, Ubuntu 22.04, and RHEL 9 — all running simultaneously on the same host.

vCPU — Virtual CPU

When you create a VM, you assign it virtual CPUs (vCPUs). A vCPU doesn't correspond 1:1 with a physical CPU core — VMware schedules vCPUs across physical cores dynamically. You can over-provision vCPUs (assign more vCPUs to VMs than physical cores exist) because VMs rarely all need CPU at the same time.

Real-world example: A physical host has 16 physical CPU cores. You might run 10 VMs, each with 4 vCPUs = 40 total vCPUs. This is 2.5:1 over-provisioning. This is perfectly normal and works because the VMs take turns using the physical CPUs.

vRAM — Virtual Memory

Similar to vCPU, virtual machines are assigned RAM. VMware has memory management techniques like Transparent Page Sharing (TPS), which lets multiple VMs share identical memory pages, and memory ballooning, which reclaims unused memory from VMs when the host is under memory pressure.

Thin vs. Thick Disk Provisioning

When you create a virtual disk for a VM, you choose how to allocate space on the storage system. Thick Eager Zeroed: Space pre-allocated and zeroed immediately. Best

performance. Thick Lazy Zeroed: Space pre-allocated, zeroed on first write. Medium performance. Thin: Space allocated on demand as data is written. Best storage efficiency.

★ **Pro Tip:** For production databases like SQL Server or Oracle, always use Thick Eager Zeroed disks for maximum performance. For development VMs, thin provisioning saves enormous amounts of storage.

💡 **Note:** Thin provisioning can lead to 'storage over-commitment' — promising more space than physically exists. Always monitor storage utilization carefully when using thin provisioning.

2.4 Why Virtualization Matters — The Business Case

Let's talk numbers. A typical physical server in a non-virtualized environment runs at 5-15% CPU utilization on average. That means 85-95% of your expensive hardware sits idle. Virtualization fixes this with dramatic results.

Metric	Typical Result from Virtualization
Server count reduction	10:1 or better (10 physical servers → 1)
Power savings	40-60% reduction in electricity costs
Data center space	Up to 80% reduction in rack space needed
VM provisioning time	From weeks to minutes
Disaster recovery time	From hours to minutes
IT staff productivity	2-3x improvement in servers managed per admin
Hardware costs	50-70% reduction over 3-5 years

A mid-sized company with 200 physical servers might consolidate to 20 hosts running the same workloads. The electricity savings alone often pay for the VMware licensing in under 2 years.

Chapter 3: VMware vSphere Architecture

vSphere is VMware's flagship enterprise product — a complete platform for server virtualization. Think of vSphere like a city: ESXi hosts are the individual buildings, vCenter Server is the city hall that manages everything, and the VMs are the people and businesses living in those buildings.

3.1 vSphere Components Overview

vSphere is not a single product — it's a suite. Understanding each component and how they fit together is the foundation of VMware expertise.

ESXi — The Hypervisor

ESXi is the heart of vSphere. It's the bare-metal hypervisor that installs directly on physical servers and creates the virtualization layer. Every physical server in your VMware environment runs ESXi. It's extremely lightweight — the entire OS footprint is about 150MB — and it can run hundreds of VMs per host on modern hardware.

vCenter Server

vCenter Server is the centralized management platform. While ESXi manages individual hosts, vCenter manages entire environments — multiple hosts, clusters, networks, and storage. Without vCenter, features like VMotion, HA, and DRS don't work. vCenter is what transforms individual ESXi hosts into an intelligent, unified infrastructure.

vSphere Client

The vSphere Client is the web-based management interface. You connect to vCenter Server through a browser (Chrome, Firefox, Edge) and get a rich, graphical dashboard for managing your entire VMware environment. Gone are the days of the old Windows-only vSphere Client — the modern HTML5 client works from any browser.

Architecture Visual

Physical Hardware → ESXi Hypervisor → Virtual Machines (running inside ESXi) All ESXi Hosts → vCenter Server (manages them all) Administrator → vSphere Client (browser) → vCenter Server → ESXi Hosts

3.2 Clusters, Hosts, and Datacenters

Logical Hierarchy in vCenter

vCenter organizes your infrastructure in a logical hierarchy. Understanding this hierarchy is essential for managing large environments effectively.

1. vCenter Server: The top-level management system
2. Datacenter: A logical container representing a physical location. You might have 'New York DC' and 'Chicago DC' as two datacenters.
3. Cluster: A group of ESXi hosts that share resources. VMs in a cluster can move between hosts automatically. This is where HA and DRS live.
4. Host: An individual ESXi server. Hosts are the physical workhorses.
5. VM: The workloads running on hosts within clusters.

What Is a Cluster?

A cluster is a group of ESXi hosts that pool their resources — CPU, RAM, and sometimes storage — into a single shared pool. When VMs are in a cluster, VMware can automatically balance workloads across all hosts and restart VMs on healthy hosts if one fails.

Real-world example: You have 4 ESXi hosts, each with 32 cores and 256GB RAM. In a cluster, that becomes 128 total cores and 1TB of RAM that VMware manages as one pool. A VM needing 16 cores can be placed on whichever host has the most available capacity.

3.3 vSphere Editions

Edition	Key Features
vSphere Essentials	Basic HA, ideal for small businesses, limited to 3 hosts
vSphere Standard	VM encryption, vMotion, basic storage features
vSphere Enterprise Plus	Full DRS, advanced networking, Storage DRS, Auto Deploy
vSphere with Tanzu	Kubernetes integration for containerized workloads
VMware Cloud Foundation	Complete SDDC stack: compute, storage, networking, management

3.4 vSphere Networking Architecture

Networking in VMware is virtualized just like compute. Physical network adapters (NICs) on ESXi hosts are connected to virtual switches that create virtual networks for VMs.

Standard vSwitch (vSS)

A vSphere Standard Switch is configured individually on each ESXi host. If you want the same VLAN on 10 hosts, you configure it 10 times. It's simpler but doesn't scale well.

Distributed vSwitch (vDS)

A vSphere Distributed Switch spans multiple ESXi hosts and is managed centrally from vCenter. Configure a network once and it's instantly available on all hosts in the cluster. This is essential for large environments and is required for some advanced features.

☆ **Pro Tip:** Always use Distributed vSwitches in production environments. Standard Switches are fine for small setups or testing, but you'll thank yourself later when you avoid the pain of managing per-host network configs.

3.5 vSphere Storage Architecture

Storage in VMware is accessed through several protocols and technologies. Understanding your storage architecture is critical for performance and data protection.

Datastores

A datastore is where VMware stores virtual machine files. It's essentially a storage pool abstracted from the underlying hardware. Datastores can be on local disk, SAN, NAS, or vSAN.

VMFS — VMware File System

VMFS (Virtual Machine File System) is VMware's proprietary file system, optimized for storing virtual machine files. It supports concurrent access from multiple ESXi hosts, which is what makes live migration (VMotion) and HA possible — multiple hosts can read/write the same storage simultaneously.

Storage Type	Use Case
Local Storage (VMFS)	Small environments, single host, no shared storage needed
SAN (FC/iSCSI)	Enterprise environments, high performance, shared access
NAS (NFS)	File-based shared storage, simpler setup, good for backups
vSAN	Hyper-converged, uses host's local SSDs/HDDs as shared storage
VMware Cloud Storage	Cloud-based storage for hybrid/cloud workloads

Chapter 4: Installing and Configuring ESXi

Installing ESXi is straightforward — but configuring it correctly is where the skill lies. In this chapter, we'll walk through the complete process from hardware requirements to a production-ready ESXi host.

4.1 Hardware Requirements

ESXi doesn't run on just anything. VMware maintains a Hardware Compatibility List (HCL) that lists every tested and supported server, storage device, and network adapter. Always check the HCL before purchasing hardware for a VMware environment.

Component	Minimum / Recommended
CPU	Minimum: 64-bit x86 CPU, 2 cores. Recommended: Intel Xeon or AMD EPYC, 16+ cores
RAM	Minimum: 4GB (ESXi itself uses 4GB+). Recommended: 256GB+ for production
Storage (boot)	1GB minimum for ESXi. Recommended: 32GB+ USB, SD, or NVMe for boot
Storage (VM data)	Varies. SAN, NAS, or NVMe SSDs for datastores
Network	Minimum: 1 Gbps NIC. Recommended: 10/25/40 Gbps NICs, redundant
CPU Features	Hardware virtualization required: Intel VT-x or AMD-V

4.2 ESXi Installation — Step by Step

ESXi installation is done via ISO — you boot from the ESXi installer and follow a wizard. Here's the complete process:

- Download the ESXi ISO from the VMware Customer Connect portal (requires a license or eval account)
- Create a bootable USB drive using Rufus (Windows) or dd (Linux/macOS) — write the ISO to the USB
- Boot the physical server from the USB drive. Press F11 or F12 during POST to access boot menu

9. The ESXi installer loads (this takes 1-2 minutes) — you'll see the purple VMware screen
10. Press Enter to begin, then F11 to accept the End User License Agreement
11. Select the disk to install ESXi on — choose your boot device (SSD, USB, or SD card)
12. Choose your keyboard layout (US is default for most environments)
13. Set the root password — use a strong password and document it securely
14. Press F11 to confirm and begin installation — takes 3-10 minutes
15. Remove the USB installer and press Enter to reboot
16. ESXi boots and presents the DCUI (Direct Console User Interface) — a simple text-based management screen

Post-Installation First Steps

After installation, you must: 1) Set a static IP address (F2 > Configure Management Network). 2) Set hostname and DNS. 3) Enable SSH if needed for management. 4) Connect to vCenter Server. Do these before the host sees production traffic.

4.3 Configuring Management Network

The management network is how vCenter Server and administrators connect to the ESXi host. It must be stable, have a static IP, and be accessible from your management workstations.

Configure via DCUI

17. At the ESXi console, press F2 and enter root credentials
18. Navigate to 'Configure Management Network'
19. Select 'Network Adapters' and choose which physical NIC(s) to use for management
20. Select 'IPv4 Configuration' — choose 'Set static IPv4 address'
21. Enter IP address, subnet mask, and default gateway
22. Select 'DNS Configuration' — enter your DNS servers and hostname
23. Press Escape and confirm 'Y' to restart the network with new settings

☆ **Pro Tip:** Use a dedicated management VLAN for ESXi hosts. This separates management traffic from VM traffic and improves security — administrators can't accidentally route into VM networks.

4.4 ESXi Networking — vmkernel Ports

In ESXi, the host's own network traffic flows through 'vmkernel ports' (VMkernel adapters). These are separate from the network adapters used by VMs. Each vmkernel port has a specific purpose.

vmkernel Port Type	Purpose
Management	vCenter communication, SSH, DCUI access, host management
vMotion	Live VM migration traffic — dedicated port required for vMotion
vSAN	Storage traffic between vSAN nodes
Fault Tolerance Logging	FT logging traffic for VMware Fault Tolerance
Provisioning	VM cloning, snapshot migration, and OVF deployments
Replication	vSphere Replication traffic

 **Note:** Always use separate vmkernel ports for separate traffic types. Mixing vMotion and management traffic on the same adapter can cause VMs to pause during migrations when the network is saturated.

Chapter 5: vCenter Server — The Command Center

If ESXi hosts are individual soldiers, vCenter Server is the general. vCenter Server is the centralized management platform that turns a collection of individual ESXi hosts into a unified, intelligent infrastructure. Without vCenter, you're flying blind — you can manage hosts one at a time but lose all the enterprise features that make VMware worth the investment.

5.1 What vCenter Server Does

vCenter Server provides centralized management, monitoring, and control for your entire VMware environment. Here's what you get:

- Unified management of all ESXi hosts from a single web interface
- Enables VMotion — live migration of running VMs between hosts
- Enables High Availability — automatic VM restart on host failure
- Enables DRS — automatic workload balancing across cluster hosts
- Centralized alerting, performance monitoring, and capacity planning
- Role-based access control — granular permissions for different admin teams
- VM templates and content libraries for rapid VM deployment
- Lifecycle management — ESXi patching and updates from a single console

5.2 vCenter Deployment Options

vCenter Server Appliance (VCSA)

The vCenter Server Appliance is a pre-packaged Linux VM that you deploy on an ESXi host. It's the modern, recommended deployment method. VMware ships it as an OVA/ISO and you simply run the installer — it deploys itself as a VM.

VCSA comes in different sizes based on the number of hosts and VMs you'll manage:

Deployment Size	Hosts / VMs Supported
Tiny	Up to 10 hosts / 100 VMs
Small	Up to 100 hosts / 1,000 VMs
Medium	Up to 400 hosts / 4,000 VMs
Large	Up to 1,000 hosts / 10,000 VMs
X-Large	Up to 2,000 hosts / 35,000 VMs

☆ **Pro Tip:** Start with Small sizing even if your environment is tiny today — storage is cheap and resizing a vCenter appliance later is a hassle. Plan for 3-5 years of growth.

5.3 Deploying vCenter Server — Step by Step

24. Download the VCSA ISO from VMware Customer Connect
25. Mount the ISO on your workstation and run the installer (vcsa-ui-installer)
26. Stage 1 — Deployment: Enter the target ESXi host IP, credentials, and where to deploy the VCSA VM
27. Configure the VCSA VM name, root password, and select deployment size
28. Configure the VCSA network settings: IP address, gateway, DNS, hostname (use FQDN)
29. Stage 1 completes — VCSA VM is deployed and powered on
30. Stage 2 — Setup: Configure the vCenter Server settings
31. Set up the SSO (Single Sign-On) domain — typically 'vsphere.local'
32. Set the SSO administrator password
33. Choose whether to join VMware's CEIP (Customer Experience Improvement Program)
34. Click Finish — Stage 2 begins. Takes 10-20 minutes.
35. Access vCenter at <https://your-vcenter-fqdn/ui> using `admin@vsphere.local`

5.4 Adding ESXi Hosts to vCenter

Once vCenter is running, you add ESXi hosts to manage them. This is straightforward:

36. In vSphere Client, create a Datacenter object (right-click vCenter > New Datacenter)
37. Create a Cluster inside the Datacenter (right-click Datacenter > New Cluster)
38. Right-click the Cluster > Add Host
39. Enter the ESXi host's IP address or FQDN
40. Enter the root credentials for the ESXi host
41. Review the host's SSL thumbprint and accept
42. Assign a license key (or the host enters evaluation mode)
43. Place the host in lockdown mode if desired (prevents direct ESXi access)
44. Click Finish — the host appears in vCenter

💡 **Note:** Add all ESXi hosts to the same cluster if they'll share storage and you want HA/DRS/VMotion to work between them. Only hosts in the same cluster can use these features together.

5.5 vCenter Roles and Permissions

vCenter has a sophisticated role-based access control (RBAC) system. Different teams get different levels of access. A VM administrator might be able to power VMs on and off but can't modify networking. A storage admin can manage datastores but can't touch VMs. Security teams can monitor without changing anything.

Built-in Role	What They Can Do
Administrator	Full control of everything in vCenter
Read-only	View everything but change nothing
VM Power User	Power VMs on/off, take snapshots, configure VMs
VM User	Interact with VMs (console access, power on/off only)
Resource Pool Administrator	Manage resource pools and assign VMs to them
Datastore Consumer	Allocate space on datastores, browse file systems
Network Administrator	Configure virtual networking
No Access	Can't see or do anything in the object

Chapter 6: Virtual Machines — Deep Dive

The virtual machine is the fundamental unit of VMware infrastructure. Everything else — ESXi, vCenter, networking, storage — exists to create, run, protect, and manage virtual machines. In this chapter, we go deep on VM creation, configuration, and management.

6.1 Anatomy of a Virtual Machine

A virtual machine is a collection of files stored in a folder on a datastore. Each VM has a core set of files:

File Extension	Purpose
.vmx	VM configuration file — the 'brain' of the VM. Text file listing all settings.
.vmdk	Virtual disk descriptor — describes the virtual disk structure
-flat.vmdk	The actual virtual disk data — this is often the largest file
.nvram	VM's BIOS/UEFI settings — persisted between reboots
.vmsd	Snapshot metadata — records snapshot hierarchy
.vmsn	Snapshot data file — saves VM memory state during snapshot
.vmss	Suspend file — VM state when suspended
.log	VM activity log — useful for troubleshooting

6.2 Creating a Virtual Machine

Creating a VM in vCenter is done through a guided wizard. Here's the complete process:

45. Right-click a cluster, host, or folder in vCenter > New Virtual Machine
46. Select 'Create a new virtual machine' (vs. deploy from template or clone)
47. Name the VM and select a location in the inventory
48. Select the compute resource (which cluster or host to run on)
49. Select the datastore to store VM files on

50. Choose compatibility level — higher levels support more features but may not work on older hosts
51. Select the Guest OS type (Windows, Linux, Other) — this tells VMware which VMware Tools version to install
52. Customize hardware: vCPU count, vRAM, virtual disks, network adapters, etc.
53. Review and click Finish
54. Mount installation ISO and power on the VM to begin OS installation

VM Hardware Customization Options

- CPU: Number of vCPUs, cores per socket, CPU hot-add enabled
- Memory: RAM allocation, hot-add, memory reservation
- Hard Disk: Size, provisioning type (thin/thick), controller type (SCSI/NVMe/SATA)
- Network Adapter: Which virtual switch, adapter type (VMXNET3 recommended), MAC address
- CD/DVD Drive: Point to ISO datastore path or physical CD/DVD
- Video Card: VRAM allocation, 3D rendering support
- USB Controller: USB 2.0 or 3.0 pass-through to guest
- VMCI Socket: VM communication interface for host-to-VM communication

6.3 VMware Tools — The Critical Component

VMware Tools is a suite of utilities installed inside the guest OS (inside the VM). Without VMware Tools, you're missing critical functionality. Always install it — no exceptions.

What VMware Tools Provides

- Improved video resolution and display performance
- Time synchronization between VM and host (critical for domain-joined machines)
- Quiesced snapshots — application-consistent snapshots for databases
- VM heartbeat — tells ESXi the VM is healthy and responsive
- Graceful shutdown from vCenter — instead of hard power-off
- Improved network performance with VMXNET3 drivers
- Drag and drop, shared clipboard between VM and console
- Guest customization — VMware can automatically configure hostname, IP, domain join

 **Note:** *Never take a snapshot of a running database VM without VMware Tools installed. Without it, you get a crash-consistent snapshot (like yanking the power cord). With Tools, you get a quiesced, application-consistent snapshot where the database flushes its buffers first.*

6.4 Snapshots — Power and Peril

VM snapshots are one of the most misunderstood features in VMware. They're powerful when used correctly and dangerous when misunderstood.

What a Snapshot Is

A snapshot captures the state of a VM at a specific moment in time: the VM's disk, memory (optional), and settings. After a snapshot is taken, all new disk writes go to a 'delta file' — the base disk remains unchanged. If something goes wrong, you can revert to the snapshot state instantly.

What a Snapshot Is NOT

A snapshot is NOT a backup. This is the most common VMware misconception. Snapshots are stored on the same datastore as the VM. If that datastore fails, you lose both the VM and its snapshots. Snapshots also grow over time and can consume massive amounts of storage.

Real World Disaster Story

A sysadmin at a retail company took a 'safety snapshot' before a major patch cycle and never deleted it. Over 6 months, the snapshot grew to 2TB, completely filling the datastore. All VMs on the datastore stopped — they couldn't write to disk. The company had a 4-hour outage on a Monday morning. The lesson: snapshots are temporary tools, not long-term storage. Delete them within 24-72 hours.

☆ **Pro Tip:** Set an alarm in vCenter to alert you when any snapshot is older than 72 hours. This is a production best practice that prevents storage crises.

6.5 VM Templates and Cloning

VM Templates

A template is a master image of a VM that you've configured exactly the way you want — OS installed, patched, VMware Tools installed, company software deployed, hardened per security policy. From one template, you can deploy dozens of identical VMs in minutes.

Content Library

A vCenter Content Library is a centralized repository for templates, ISO files, and OVF packages. It can synchronize between multiple vCenter Servers, making templates available across geographic locations. This is the modern, recommended way to manage templates.

Method	Best For
Clone from VM	Quick one-off clone, maintaining linked snapshots

Deploy from Template	Consistent, repeatable VM deployment for production
OVF/OVA Import	Importing VMs from vendors or between platforms
Instant Clone	Rapid deployment with minimal storage (uses parent VM memory)

Chapter 7: VMware Networking Deep Dive

Networking is often where VMware projects get complicated. Virtual networking in VMware is powerful and flexible, but it requires a solid understanding of both networking fundamentals and VMware-specific concepts. In this chapter, we'll cover everything from basic virtual switches to advanced distributed networking.

7.1 VMware Virtual Switch Architecture

A vSphere Standard Switch (vSS) or Distributed Switch (vDS) is a Layer 2 software switch that operates within the ESXi kernel. It connects VMs to each other and to the physical network, just like a physical switch does for physical servers — but it lives entirely in software.

How Physical NICs Connect to Virtual Switches

Physical NICs on ESXi hosts (called 'uplinks' or 'pNICs') are connected to virtual switches as uplink ports. The virtual switch uses these physical NICs to communicate with the rest of the physical network (other servers, routers, the internet).

Think of it this way: The physical NIC is the 'cable' plugged into your virtual switch. The virtual switch then connects that cable to the rest of your network infrastructure — just like a patch cable connects a physical switch to a router.

Practical Architecture

A typical ESXi host has 4 physical NICs: - 2 NICs → Management vSwitch (for ESXi management + VMkernel ports) - 2 NICs → VM Traffic vSwitch (for virtual machine network traffic) This gives you redundancy on each switch. If one NIC fails, the other takes over instantly.

7.2 Port Groups

A Port Group is a named network configuration on a virtual switch. VMs connect to port groups, not directly to virtual switches. Each port group defines the VLAN, traffic shaping, security policies, and failover policies for VMs connected to it.

Real-world example: Your vSS has one physical NIC uplink. You create three port groups: 'Production VLAN 100', 'DMZ VLAN 200', and 'Management VLAN 300'. Each port group is associated with a different VLAN tag. A web server VM gets connected to 'DMZ VLAN 200' and instantly has access only to that network segment.

7.3 vSphere Distributed Switch (vDS)

The vDS is the enterprise-grade virtual networking solution. Unlike the Standard Switch which is configured per-host, the Distributed Switch is configured once in vCenter and automatically pushed to all ESXi hosts in the cluster.

Key vDS Features

- Centralized management — configure once, apply everywhere
- Port state persistence — VM network settings survive migration to different hosts
- NetFlow — send network traffic statistics to a NetFlow collector for analysis
- Port Mirroring — mirror VM traffic to another port for intrusion detection systems
- LACP — Link Aggregation Control Protocol for NIC teaming with network switches
- Traffic shaping — control ingress and egress bandwidth per port group
- Health checks — validate MTU consistency across uplinks

7.4 Network I/O Control (NIOC)

On a single physical NIC, multiple types of traffic compete: VM traffic, vMotion, management, vSAN, FT logging. During peak times, one type can starve others. Network I/O Control lets you define shares and limits so each traffic type gets appropriate bandwidth.

Traffic Type	Recommended Share Priority
FT Logging	High (50 shares) — Fault Tolerance needs reliable bandwidth
iSCSI Storage	High (50 shares) — Storage latency kills performance
vSAN	High (50 shares) — Hyper-converged storage traffic
vMotion	Normal (50 shares) — Can burst, but shouldn't starve VMs
Virtual Machine Traffic	Normal (100 shares) — This is your production workload
Management	Low (20 shares) — Important but not latency-sensitive

7.5 VMware NSX-T — Software-Defined Networking

NSX-T (now called NSX) is VMware's network virtualization platform. It takes the concept of server virtualization and applies it to networking. Instead of physical firewalls, routers, and load balancers, NSX creates them all in software.

What NSX Provides

- Micro-segmentation: Apply firewall rules at the VM level, not just the network edge
- Logical routing: Create software-defined routers that connect VM networks
- Distributed firewall: Each ESXi host enforces firewall rules at the hypervisor level
- Overlay networking: Create isolated virtual networks on top of physical networks (VXLAN/GENEVE)
- Load balancing: Software-defined load balancers without dedicated hardware

Real-world example: In a traditional network, a firewall sits at the edge of your data center. A compromised VM can freely communicate with other VMs in the same subnet. With NSX micro-segmentation, each VM has its own firewall rules. A compromised web server VM literally cannot even ping the database VM in the same VLAN — the distributed firewall blocks it at the hypervisor.

Chapter 8: VMware Storage

Storage is where virtual machines live — it's the foundation everything else is built on. Understanding VMware storage deeply is what separates good VMware administrators from great ones. Poor storage design causes more VMware performance problems than anything else.

8.1 Storage Concepts Review

Before diving into VMware-specific storage, let's establish the fundamentals. Virtual machines store their files (including large virtual disk files) on storage devices. VMware supports block storage and file storage, accessed via several protocols.

Protocol	Type / Description
Fibre Channel (FC)	Block storage over dedicated fiber network. Fastest, most expensive. Used by large enterprises.
iSCSI	Block storage over standard IP network. Cost-effective alternative to FC.
FCoE	Fibre Channel over Ethernet — combines FC and Ethernet on same NICs
NFS	File-based storage over IP network. Simple, flexible, widely supported.
vSAN	Hyper-converged storage using hosts' local disks

8.2 VMFS — VMware's File System

VMFS (Virtual Machine File System) is VMware's proprietary cluster file system. It's specifically designed for virtualization workloads and allows multiple ESXi hosts to simultaneously access the same datastore — which is what makes VMotion and HA possible.

VMFS Key Characteristics

- Cluster file system — multiple ESXi hosts read/write simultaneously
- Supports large files — virtual disks can be up to 62TB
- Sub-block allocation — efficient storage for small files
- On-disk locking — prevents multiple VMs from writing to same disk simultaneously
- ATS (Atomic Test and Set) — hardware-accelerated locking for modern arrays

 **Note:** VMFS version 6 (VMFS6) is the current version, introduced with vSphere 6.5. It supports automatic space reclamation (UNMAP) which helps thin-provisioned arrays reclaim space when data is deleted from VMs.

8.3 vSAN — Hyper-Converged Storage

vSAN is VMware's software-defined storage product. Instead of buying a separate SAN (Storage Area Network), vSAN uses the local SSDs and HDDs inside your ESXi hosts and pools them together into a shared storage cluster. This is called hyper-converged infrastructure (HCI).

How vSAN Works

vSAN requires a minimum of 3 ESXi hosts, each with at least one SSD (for caching) and one HDD or additional SSD (for capacity). Each host contributes its local storage to the vSAN cluster, and vSAN stripes data across all hosts with configurable redundancy.

vSAN Data Protection — FTT

vSAN protects data using Failures to Tolerate (FTT) policies: FTT=1: Data survives 1 host failure (requires 3 hosts minimum) FTT=2: Data survives 2 host failures (requires 5 hosts minimum) FTT=3: Data survives 3 host failures (requires 7 hosts minimum) VMs on vSAN are automatically protected according to VM Storage Policies you assign.

vSAN Advantages

- Eliminates dedicated SAN hardware — significant cost savings
- Scales linearly — add a host, add compute AND storage simultaneously
- Policy-based storage — assign storage SLAs per VM, not per LUN
- All-NVMe configurations support extremely high IOPS
- Stretched cluster support — vSAN across two data centers for disaster recovery

vSAN Requirements

- Minimum 3 ESXi hosts per vSAN cluster
- 10GbE networking strongly recommended (25GbE for all-flash)
- Dedicated vmkernel port for vSAN traffic
- Compatible hardware from VMware's vSAN HCL
- Each host needs a cache tier (SSD) and capacity tier (SSD/HDD)

8.4 Storage Performance Best Practices

The Storage Stack — Where Latency Hides

When a VM writes data, it travels through multiple layers before reaching physical storage. Each layer adds latency. Understanding the stack helps you diagnose and eliminate bottlenecks.

- 55. Guest OS issues write to file system buffer
- 56. VMDK — virtual disk driver in ESXi kernel
- 57. Storage driver (SCSI, NVMe)
- 58. VMkernel storage stack
- 59. Host Bus Adapter (HBA) or NIC
- 60. Physical storage array / vSAN

☆ **Pro Tip:** The single biggest storage performance killer in VMware environments is snapshots. Every snapshot adds a delta file that the VM must write through. 3 stacked snapshots means 3x write amplification. Delete snapshots you don't need!

Symptom	Likely Storage Cause
VM disk latency > 20ms	Datastore overloaded or snapshot stack
Disk queue depth > 32	Too many VMs on a single LUN/datastore
Datastore >80% full	Risk of thin disk expansion failure
vMotion takes >30 minutes	Storage I/O contention during migration
Database slow after snapshot	Delta file I/O overhead from active snapshot

Chapter 9: VMotion — The Magic of Live Migration

VMotion is arguably the most impressive technology VMware ever created. The ability to move a running virtual machine from one physical host to another — while it continues to run, without the users noticing, without dropping a single network packet — is nothing short of remarkable. In this chapter, we go deep on everything VMotion.

9.1 What Is VMotion?

VMotion (Virtual Motion) is VMware's technology for live migration of virtual machines. 'Live migration' means the VM keeps running throughout the entire process. Users connected to the VM don't notice anything. Applications running in the VM don't pause. Network connections are maintained. Transactions complete normally.

Real-world analogy: Imagine performing open-heart surgery on a patient while they're fully conscious and having a conversation. VMotion does the equivalent for virtual machines — migrating them without interrupting their operation.

Real-world use case: It's 10am on a Monday. Your most important database server is running on ESXi Host A. Host A is due for an emergency security patch that requires a reboot. Without VMotion, you'd have to schedule maintenance, notify users, shut down the database, patch, reboot, verify, and restart the database — causing an outage. With VMotion, you move the database VM to Host B (takes 30-60 seconds), patch and reboot Host A (takes 10-15 minutes), move the VM back if desired — zero downtime, zero user impact.

VMotion in Numbers

A typical VMotion migration takes:

- 10-60 seconds for most VMs
- Memory-intensive VMs (128GB+ RAM) may take 2-5 minutes
- The VM experiences less than 1 second of 'stun' (pause) during the final cutover
- Modern environments with 25GbE+ networks can migrate VMs almost instantaneously
- VMware guarantees the VM remains operational throughout

9.2 How VMotion Works — The Technical Process

Understanding the VMotion process step by step helps you design networks and storage correctly, troubleshoot failures, and explain the technology to stakeholders. Here's exactly what happens during a VMotion migration:

Phase 1: Pre-Copy Memory Transfer

VMotion begins by copying the VM's memory (RAM) from the source host to the destination host while the VM keeps running. This happens over the VMotion network.

61. VMotion initiates — source ESXi opens a network connection to destination ESXi over the VMotion vmkernel network

62. Source host begins transmitting the VM's memory pages to the destination host
63. Since the VM is still running, memory is changing as it's being copied — dirty pages (recently changed) are tracked
64. VMotion iteratively copies dirty pages — pages that changed during the copy are re-sent
65. With each iteration, fewer pages are dirty (the VM's memory 'converges')

Phase 2: The Stun and Cutover

Once memory convergence reaches a threshold (usually less than 100ms worth of changes remain), VMotion executes the cutover:

66. Source VM is briefly 'stunned' (paused) — this pause is typically under 1 second
67. Remaining dirty memory pages are transmitted to destination host
68. Destination host receives the final memory state
69. Destination host begins executing the VM from the exact point where source stopped
70. Network state is transferred — the VM's MAC address and IP address move to the destination host
71. The physical network switch learns the new port for the VM's MAC address (via RARP/gratuitous ARP)
72. Source VM is deleted — its memory is freed
73. Storage access: the VM's disk files remain in place on shared storage; only the pointer changes

 **Note:** The 'stun time' is the only moment users might notice VMotion. In practice, it's so brief (usually well under 1 second) that TCP connections stay alive and users experience at most a momentary slowdown. Most users never notice.

9.3 VMotion Requirements

VMotion has specific requirements that must be met for it to function. Meeting these requirements is not optional — VMotion simply won't work without them. This is one of the most important topics for VMware certification exams AND real-world deployments.

1. Shared Storage — The Most Critical Requirement

The VM's disk files must be accessible from BOTH the source and destination host. This is why shared storage (SAN, NAS, vSAN) is fundamental to any VMware infrastructure that uses VMotion. Both hosts must be able to read and write to the same datastore simultaneously.

Why: During VMotion, only memory is transferred — the VM's disk files don't move. They stay on the shared storage. The destination host simply starts reading the disk from the same location. If the disk isn't accessible from the destination host, VMotion cannot complete.

Exception: Storage VMotion

Storage VMotion (svMotion) allows you to move a VM's disk files to a different datastore while the VM keeps running. Regular VMotion moves the VM to a different host. You can even do both simultaneously — moving the VM to a different host AND different storage at the same time. This is called 'Enhanced VMotion' or just a combined migration.

2. vMotion Network — Dedicated and Fast

VMotion traffic needs its own dedicated vmkernel port and sufficient bandwidth. The amount of bandwidth needed depends on the VM's memory size and how quickly you need migrations to complete.

- Minimum: 1 Gbps dedicated NIC for VMotion (functional but slow for large VMs)
- Recommended: 10 Gbps or higher for production environments
- Enterprise: 25 Gbps or 40 Gbps for mission-critical environments with large VMs
- VMotion vmkernel must have 'VMotion' checkbox enabled in its settings
- VMotion vmkernel must be able to reach the VMotion vmkernel on the destination host

★ **Pro Tip:** Create a dedicated VLAN for VMotion traffic. This prevents VMotion from competing with VM traffic on the same NIC. In large environments with frequent migrations (DRS moves), VMotion traffic can consume significant bandwidth.

3. Compatible CPU Features

The source and destination ESXi hosts must have CPUs with compatible instruction sets. A VM can't run on a CPU that's missing features the VM's OS is using.

The good news: VMware has a feature called Enhanced VMotion Compatibility (EVC) that solves this problem in clusters. EVC masks advanced CPU features so all hosts in the cluster appear to have the same CPU baseline. This allows VMotion between hosts with different CPU generations.

EVC Mode	Description
Intel Merom EVC	Lowest common denominator — works with Core 2 Duo era and newer
Intel Westmere EVC	Xeon 5600 era features available
Intel Sandy Bridge EVC	AVX instructions added
Intel Broadwell EVC	AVX2, ADCX, RDSEED added
Intel Skylake EVC	AVX-512 (partial), MPX instructions
AMD Opteron EVC	First-gen AMD compatible baseline

AMD Naples EVC	EPYC 1st generation baseline
AMD Rome EVC	EPYC 2nd generation baseline, AVX2

 **Note:** Set EVC mode when you create the cluster — before adding hosts. Changing EVC mode later requires powering off all VMs in the cluster. EVC mode can only be raised (to a newer generation) if all running VMs support the new feature set.

4. vCenter Server

VMotion between two ESXi hosts requires vCenter Server. VMotion is NOT available without vCenter — this is a critical point. You can't initiate VMotion from the ESXi host's DCUI or standalone web interface. vCenter orchestrates the entire process.

5. Matching Networks on Destination Host

The destination host must have the same network port groups as the source host, or VMotion won't know where to put the VM's network adapters. If the VM has a NIC on 'Production VLAN 100,' the destination host must also have a port group named 'Production VLAN 100' (on a Standard Switch) or the same distributed switch port group.

This is one of the biggest advantages of vSphere Distributed Switches — because networks are configured centrally and applied to all hosts, you automatically have matching networks on every host in the cluster. With Standard Switches, you must manually ensure consistency across hosts.

6. Sufficient Resources on Destination Host

The destination host must have enough free CPU, memory, and any reserved resources to accommodate the VM being migrated. If the destination is already at capacity, VMotion will fail or may not even be attempted (DRS won't suggest the migration).

9.4 Types of VMotion Migration

Migration Type	What Moves
VMotion (Compute)	The running VM moves to a different ESXi host. Disks stay on shared storage.
Storage VMotion (Storage)	VM disks move to a different datastore. VM can stay on same host or move.
Enhanced VMotion (Both)	VM moves to different host AND disks move to different datastore simultaneously.

Cross-vCenter VMotion	VM moves between two different vCenter Server instances (even different data centers).
Cross-Cloud VMotion	VM moves to/from VMware Cloud on AWS or other cloud providers.

9.5 Performing VMotion — Step by Step

Here's how to perform a VMotion migration in the vSphere Client:

Method 1: Migrate Wizard

74. In vSphere Client, navigate to the VM you want to migrate
75. Right-click the VM > Migrate
76. Select migration type:
 - 'Change compute resource only' — VMotion only (host migration)
 - 'Change storage only' — Storage VMotion only
 - 'Change both compute resource and storage' — Enhanced VMotion
77. Select destination host or cluster (for compute migrations)
78. Select network — confirm the VM's network adapters will connect to the right port group on the destination
79. Select destination datastore (for storage migrations)
80. Select VM storage policy (if using vSAN)
81. Review the migration summary
82. Click Finish — VMotion begins immediately
83. Monitor progress in the vSphere Client 'Recent Tasks' panel

Method 2: Drag and Drop

In the vSphere Client host and cluster view, you can literally drag a VM from one host and drop it onto another host. This triggers a VMotion migration using the default settings. Quickest method for ad-hoc migrations.

Method 3: PowerCLI (Automated VMotion)

For migrating many VMs or automating migrations, VMware PowerCLI (PowerShell for VMware) is the tool of choice:

PowerCLI Example

```
# Connect to vCenter Connect-VIServer -Server vcenter.company.com # Migrate a single VM Move-VM -VM 'WebServer01' -Destination (Get-VMHost 'esxi02.company.com') # Migrate all VMs from a host (prep for maintenance) $vms = Get-VMHost 'esxi01.company.com' | Get-VM foreach ($vm in $vms) { Move-VM -VM $vm -Destination (Get-VMHost 'esxi02.company.com') }
```

9.6 What VMotion Is Required For

VMotion is not just a 'nice to have' — it's the foundation of several critical enterprise VMware capabilities. Without VMotion, most of what makes VMware valuable in the enterprise doesn't work.

High Availability (HA)

Wait — HA doesn't actually use VMotion to restart VMs after a host failure. HA simply starts VMs on other hosts (a cold restart). However, VMotion is used to evacuate VMs from hosts that might be having issues — for example, when you detect hardware degradation before a complete failure and want to proactively move VMs off a troubled host.

DRS — Distributed Resource Scheduler

DRS uses VMotion constantly. When DRS determines that workloads are imbalanced (one host is at 90% CPU while another is at 20%), it uses VMotion to move VMs from the overloaded host to the underloaded host — automatically, without any administrator intervention. VMotion is the mechanism that makes DRS work.

Maintenance Mode

Putting an ESXi host into Maintenance Mode evacuates all VMs from it using VMotion before any maintenance activity begins. This is the safe way to patch, upgrade, or perform hardware maintenance on hosts with zero downtime.

Host Profiles

When applying host profiles to standardize configurations across hosts, VMotion is used to evacuate VMs before the profile is applied.

vSphere Lifecycle Manager (vLCM)

When updating ESXi hosts with patches or new versions, vLCM puts hosts into maintenance mode (triggering VMotion) and then applies updates. This is how entire clusters are patched with zero VM downtime.

VMotion Requirement Summary

VMotion is required for:

- Zero-downtime host maintenance (patching, hardware repair)
- DRS automatic load balancing
- Proactive HA before host failure
- Scheduled maintenance

windows • VM migration during disaster recovery testing • Cross-data center workload balancing • Cloud bursting (moving VMs to public cloud)

9.7 VMotion Troubleshooting

VMotion can fail for various reasons. Here's how to diagnose common problems:

VMotion Error	Likely Cause and Fix
No compatible hosts found	Destination hosts may lack required CPU features — enable EVC on the cluster
The VM is using host devices	VMs with physical DVD drives, USB pass-through, or RDMs may have VMotion restrictions
Insufficient resources on destination	Destination host doesn't have enough free RAM — free up resources or add memory
Migration network not available	VMotion vmkernel port not configured on source or destination — add VMotion vmkernel
VMotion fails partway through	Network interruption during migration — check network stability and speed
Migration takes too long (>30 min)	VM has high memory write rate (memory-intensive app) or insufficient VMotion bandwidth
Cannot migrate: device not supported	CD/DVD connected to client device prevents VMotion — disconnect virtual media first

☆ **Pro Tip:** The most common VMotion issue is the VM having a CD/DVD drive connected to a client device (ISO mounted from your laptop via the web console). Always disconnect client devices before migrating a VM.

Chapter 10: High Availability (HA)

Server failures happen. Disks die, power supplies fail, network cards burn out, entire servers crash. VMware High Availability (HA) ensures that when a physical ESXi host fails, the virtual machines that were running on it automatically restart on other healthy hosts in the cluster — with minimal disruption.

10.1 What Is VMware HA?

VMware HA is a cluster feature that monitors the health of ESXi hosts. When HA detects that a host has failed (it's no longer responding), it automatically restarts the VMs that were running on that host on the surviving hosts. This happens automatically, without administrator intervention, typically within 1-3 minutes of detecting the failure.

Real-world scenario: It's 3:07am on a Sunday. Your ESXi host 'PROD-ESXi-03' has a PSU failure and crashes. Your email server, HR application, and file server were all running on it. Without HA, you'd need an on-call administrator to notice the failure, log in, and manually restart those VMs on another host — potentially an hour or more of downtime. With HA, vCenter detects the failure within seconds, and within 1-2 minutes, all three VMs are back online on 'PROD-ESXi-01' and 'PROD-ESXi-02' automatically.

10.2 How HA Works

HA Master and Slave Hosts

When HA is enabled on a cluster, vCenter elects one ESXi host as the HA Master. All other hosts become HA Slaves. The master monitors the slaves. If the master loses contact with a slave, it checks whether the VMs on that slave are reachable via storage heartbeating before declaring the host failed.

HA Detection Mechanisms

- Network heartbeats: ESXi hosts send heartbeats to each other every second. Missing heartbeats trigger failure investigation.
- Datastore heartbeating: Hosts write heartbeat data to shared datastores. This differentiates host failures from network partitions.
- Management network isolation: If a host can still write to datastores but not communicate on the network, it may be an isolated host, not a failed host.

10.3 HA Admission Control

Admission Control is HA's way of ensuring there are always enough resources reserved to restart VMs if a host fails. HA won't let VMs power on if doing so would prevent successful failover.

Admission Control Policy	How It Works
Percentage based	Reserve X% of cluster CPU and RAM for failover capacity
Slot policy	Calculate 'slots' based on largest VM reservation; reserve enough slots for host failure
Dedicated failover hosts	Designate specific hosts as failover-only — they run no production VMs normally
Cluster resource percentage	Modern default — simple percentage-based resource reservation

 **Note:** If Admission Control is preventing VMs from powering on, it means the cluster is over-committed relative to your failover protection level. Either add hosts, reduce VM reservations, or (carefully) lower the failover capacity percentage.

10.4 HA vs. Fault Tolerance (FT)

HA provides restart protection (VMs restart after failure). Fault Tolerance provides zero-downtime protection (VMs never go down, even during a host failure). FT maintains a live shadow copy of a VM on a different host, mirroring every instruction in real time.

Feature	HA vs Fault Tolerance
Recovery Type	HA: Restart (1-3 min downtime) FT: Zero downtime (continuous)
VM Size Support	HA: Unlimited VM sizes FT: Up to 8 vCPUs, specific limits
Performance Impact	HA: None during normal operations FT: ~10-20% CPU overhead
Storage Impact	HA: Shared storage only FT: Shared storage + FT logging network
Use Case	HA: All production VMs FT: Only the most critical workloads

 **Pro Tip:** Reserve FT for applications where even 2 minutes of downtime is catastrophic — trading systems, payment processing, emergency services. For everything else, HA + application clustering (like SQL Always On) is more cost-effective.

Chapter 11: Distributed Resource Scheduler (DRS)

DRS is VMware's automated workload balancing system. It monitors CPU and memory utilization across all hosts in a cluster and automatically moves VMs (using VMotion) to balance the load. Think of DRS as a highly intelligent air traffic controller for your virtual machines — constantly adjusting where workloads run to prevent congestion.

11.1 How DRS Works

DRS runs as part of vCenter Server and continuously monitors resource utilization across all hosts in a cluster. Every 5 minutes (by default), DRS calculates whether the current VM placement is optimal. If it detects imbalance, it recommends migrations (or automatically performs them, depending on DRS mode).

DRS Imbalance Score

DRS uses a normalized imbalance metric to determine if migration is warranted. It calculates a 'target load' for each host (what the cluster average would be if load were perfectly balanced) and compares each host's actual load to that target. If the deviation exceeds DRS's threshold, it recommends moves.

11.2 DRS Automation Levels

Mode	Behavior
Manual	DRS suggests VM placements and migrations but does nothing automatically. Admin approves each one.
Partially Automated	DRS automatically places VMs when they power on (initial placement). Migration recommendations require approval.
Fully Automated	DRS automatically places VMs at power-on AND automatically migrates VMs for balance. Recommended for production.

★ **Pro Tip:** Start with Partially Automated to see what DRS recommends before trusting it fully. After a few weeks of observing good recommendations, switch to Fully Automated. Most mature VMware environments use Fully Automated.

11.3 DRS Affinity and Anti-Affinity Rules

Sometimes you want to control where VMs run relative to each other. DRS rules let you define these preferences.

VM-VM Affinity Rules — Keep VMs Together

Use case: Your three-tier application has a web server, app server, and cache server that communicate constantly. Keeping them on the same host eliminates network latency between them (VM-to-VM traffic on the same host goes through virtual switch memory, not physical network).

VM-VM Anti-Affinity Rules — Keep VMs Apart

Use case: You're running two SQL Server nodes in an Always On Availability Group for high availability. If both nodes are on the same host and that host fails, you lose both nodes — defeating the purpose of clustering. Anti-affinity rules ensure they're always on different hosts.

VM-Host Affinity Rules

Force VMs to run on specific hosts. Use case: Software with per-socket licensing that requires running on specific hosts, or VMs needing access to specific hardware (GPU, specialized NIC).

11.4 Storage DRS (sDRS)

Storage DRS does for datastores what compute DRS does for hosts — it automatically balances VM disk files across multiple datastores based on space utilization and I/O load. It can automatically move VM disks between datastores using Storage VMotion.

Storage DRS Use Cases

- Prevent any single datastore from filling up
- Balance I/O load across multiple datastores
- Automatic initial placement of VM disks on best datastore
- Proactive space management — move VMs before storage is full

Chapter 12: VMware Security

Security in a VMware environment has multiple layers. Physical security of hosts, ESXi hardening, VM isolation, network security, access control, and encryption all contribute to a comprehensive security posture. This chapter covers the essential security configurations every VMware administrator should implement.

12.1 ESXi Security Hardening

Lockdown Mode

When Lockdown Mode is enabled on an ESXi host, local access is disabled. Only vCenter Server can manage the host — direct root SSH or API access is blocked. This prevents someone from bypassing vCenter's access controls by connecting directly to the host.

- Normal Lockdown Mode: vCenter manages the host; DCUI still accessible for emergencies
- Strict Lockdown Mode: vCenter manages the host; DCUI is also disabled (maximum security)

 **Note:** Always enable at least Normal Lockdown Mode in production. The exception is hosts you're actively maintaining or troubleshooting. Re-enable it when done.

SSH and ESXi Shell

SSH provides command-line access to ESXi. It's a powerful troubleshooting tool, but it should be disabled when not actively needed. Every security framework (CIS, NIST, PCI DSS) recommends keeping SSH disabled on ESXi hosts and enabling it only for specific maintenance windows.

 **Pro Tip:** Create an alarm in vCenter that alerts you if SSH is enabled on any ESXi host. This catches both intentional changes and unauthorized access.

12.2 VM Encryption

vSphere VM Encryption encrypts virtual machine files at rest — including the VM's memory files, disk files, and configuration files. Even if someone steals the physical disk from a server, they can't read the VM data without the encryption keys.

How vSphere Encryption Works

84. A Key Management Server (KMS) stores the encryption keys — VMware partners with vendors like Thales, HashiCorp, and others
85. vCenter requests a Data Encryption Key (DEK) from the KMS when encrypting a VM
86. The DEK is used to encrypt the VM's data

87. The DEK itself is encrypted with a Key Encryption Key (KEK) from the KMS

88. Without access to the KMS, the encrypted VM data is unreadable

12.3 vSphere Trust Authority

vSphere Trust Authority (vTA) is an advanced security feature that provides hardware root of trust attestation for ESXi hosts. It ensures that ESXi hosts are running trusted, unmodified software before they're allowed to decrypt sensitive VM workloads.

12.4 Role-Based Access Control Best Practices

- Principle of least privilege: Give users only the permissions they need
- Never use the default 'Administrator' role for day-to-day operations — create custom roles
- Create service accounts for automation tools (Backup software, monitoring, etc.)
- Integrate vCenter with Active Directory for centralized user management
- Audit access control changes regularly — vCenter logs all permission changes
- Use two-factor authentication (2FA) for vCenter access, especially for administrator accounts

Access Control Best Practice	Why It Matters
Separate admin accounts	Administrators should have separate privileged accounts from daily-use accounts
AD integration	Centralized password policies, group management, and audit trails
Object-level permissions	Assign permissions at the lowest necessary object, not always at the global level
No sharing credentials	Every admin needs their own named account — no shared 'admin' passwords
Regular permission audits	Check quarterly that permissions still match job roles

Chapter 13: Monitoring and Performance

A VMware administrator who doesn't monitor their environment is flying blind. Performance problems, capacity issues, and emerging failures all manifest in metrics before they cause outages. This chapter covers VMware's monitoring tools and how to interpret what they're telling you.

13.1 vCenter Performance Charts

vCenter Server collects performance metrics for every object in the environment — hosts, clusters, VMs, datastores, and networks. These metrics are stored in the vCenter database and viewable through performance charts.

Key Performance Metrics to Monitor

Metric	What It Tells You
CPU Ready (ms)	Time a VM's vCPU spent waiting for physical CPU. >5ms is a concern; >10ms is a problem.
Memory Balloon (KB)	ESXi reclaiming memory from VMs. High balloon = host is memory-constrained.
Memory Swap (KB/s)	VMs using disk-backed swap. This is severe — disk is 100x slower than RAM.
Disk Latency (ms)	Average disk read/write response time. >20ms for VMs = investigate storage.
Disk Queue Depth	Number of pending I/O operations. High queue depth = storage overloaded.
Network Dropped Packets	Packets dropped due to NIC saturation. Any dropped packets need investigation.
CPU Co-stop (ms)	Multi-vCPU VMs waiting for all vCPUs to be scheduled simultaneously. High value = over-provisioned vCPUs.

13.2 vRealize Operations (vROps)

vRealize Operations (now called VMware Aria Operations) is VMware's advanced monitoring and analytics platform. It goes far beyond the built-in performance charts, providing predictive analysis, capacity planning, and actionable recommendations.

What vROps Does

- Analyzes performance trends and predicts when resources will run out
- Identifies 'oversized' VMs consuming more resources than they use
- Identifies 'undersized' VMs that need more resources
- Provides workload optimization recommendations
- Tracks compliance against hardening standards
- Generates capacity planning reports for budget planning

Rightsizing — Saving Real Money

vROps can analyze VM utilization over time and recommend rightsizing. A VM with 16 vCPUs and 128GB RAM that actually uses 2 vCPUs and 12GB RAM can be safely resized — freeing enormous resources for other VMs or reducing hardware costs. Many organizations find 20-40% of VMs are oversized when they first do a vROps analysis.

13.3 vCenter Alarms

vCenter Server has a comprehensive alarm system. Alarms trigger when monitored conditions exceed defined thresholds. You can configure alarms to send email, send SNMP traps, or execute scripts.

Essential Alarms to Configure

- Host CPU usage > 85% for 15 minutes
- Host memory usage > 85%
- Datastore space usage > 75%
- VM CPU Ready > 5% average
- VM memory balloon > 0 (any balloon is worth knowing about)
- ESXi host connection lost
- HA failover event (know when HA restarts a VM)
- SSH enabled on any host
- Snapshot age > 72 hours
- Certificate expiration within 30 days

Chapter 14: Backup and Disaster Recovery

Backups are not optional. Disaster recovery planning is not optional. Every production VMware environment must have both — and they must be tested regularly. An untested backup is not a backup; it's a hope.

14.1 VMware Backup Strategies

vSphere APIs for Data Protection (VADP)

Modern VMware backup solutions use VADP — VMware's official backup API. VADP allows backup software to take consistent VM snapshots and back them up without installing agents inside each VM. This is the recommended backup method for VMware environments.

How VADP Backup Works

89. Backup software (Veeam, Zerto, Commvault, etc.) communicates with vCenter via VADP
90. Backup triggers a VMware snapshot of the VM — this quiesces the disk
91. The backup software reads the snapshot's data, identifying only changed blocks since last backup (incremental)
92. Data is written to the backup target (NAS, tape, cloud)
93. Snapshot is deleted after backup completes
94. The VM continues running throughout — no downtime, no agents needed

14.2 Backup Best Practices

Best Practice	Why It Matters
3-2-1 Rule	3 copies of data, 2 different media types, 1 offsite copy. Survives almost any disaster.
Test restores monthly	An untested backup gives false confidence. Test restore random VMs regularly.
Immutable backups	Backups attackers can't delete or encrypt. Critical protection against ransomware.
RPO and RTO defined	Know your Recovery Point Objective and Recovery Time Objective before disaster.
Application-consistent backups	Use VMware Tools quiescing for databases — ensure clean backup state.

Encrypt backup data

Backup data contains all your sensitive VM data — encrypt it in transit and at rest.

14.3 Site Recovery Manager (SRM)

VMware Site Recovery Manager (SRM) is VMware's disaster recovery orchestration platform. It automates the process of failing over to a secondary site when disaster strikes. Instead of an administrator manually powering on VMs at a DR site and re-configuring networking, SRM does it all automatically.

SRM Key Concepts

- Protection Groups: VMs grouped together for coordinated failover
- Recovery Plans: Step-by-step orchestrated failover instructions
- Test Failover: Test your DR plan without affecting production
- Failover: Actual failover to DR site
- Failback: Return operations to primary site after disaster

SRM Test Failover

SRM's most powerful feature is the Test Failover capability. You can test your entire disaster recovery plan — actually starting up every VM at the DR site — without impacting production. SRM spins up VMs in an isolated network bubble so they don't interfere with production, runs your validation scripts, then tears it all down. Run this quarterly to prove your DR plan actually works.

Chapter 15: VMware Cloud and Modern Trends

VMware has evolved far beyond on-premises data center virtualization. Today, VMware technologies span private clouds, public clouds, and hybrid environments. Understanding VMware's cloud strategy is essential for any VMware professional working in modern IT.

15.1 VMware Cloud on AWS (VMC on AWS)

VMware Cloud on AWS is a jointly engineered service from VMware and AWS that delivers VMware's SDDC software stack on AWS's bare-metal infrastructure. Organizations can run VMware workloads in AWS using the same vSphere/vSAN/NSX tools they use on-premises.

Key Benefits

- Extend your data center to AWS in hours, not months
- Use familiar VMware tools to manage cloud workloads
- VMotion workloads between on-premises and AWS (Cross-Cloud VMotion)
- Native access to AWS services (S3, RDS, Lambda) from VMware VMs
- Elastic capacity — scale up and down with AWS's massive resource pool

15.2 VMware Tanzu — Kubernetes on vSphere

VMware Tanzu brings Kubernetes (container orchestration) directly into vSphere. Instead of building separate Kubernetes infrastructure, vSphere administrators can provision Kubernetes clusters directly from vCenter.

The vision: A single platform manages both traditional VMs and modern containerized applications. The same networking (NSX), storage (vSAN), and management (vCenter) tools work for both workload types.

15.3 VMware Cloud Foundation (VCF)

VCF is VMware's complete SDDC (Software-Defined Data Center) platform. It bundles vSphere (compute), vSAN (storage), NSX-T (networking), and SDDC Manager (lifecycle management) into one integrated, pre-validated solution.

VCF Component	What It Provides
vSphere (ESXi + vCenter)	Compute virtualization — run VMs

vSAN	Hyper-converged storage — no separate SAN needed
NSX-T	Software-defined networking and security
SDDC Manager	Automates deployment, patching, and lifecycle of the entire stack
VMware Aria Suite	Operations, automation, and logging across the entire environment

15.4 The Broadcom Era — What's Changing

Following Broadcom's 2023 acquisition of VMware, the product lineup has been simplified and restructured. Key changes include the shift to subscription-based licensing for enterprise customers, the consolidation of many products into VMware Cloud Foundation, and some product discontinuations.

The core technology remains the same — ESXi, vCenter, VMotion, HA, DRS, vSAN, and NSX are all still here and continue to be developed. The change is primarily commercial, not technical.

★ **Pro Tip:** Stay current with Broadcom's VMware product announcements at blogs.vmware.com and the official Broadcom documentation portal. The product naming and licensing landscape has been changing rapidly post-acquisition.

Chapter 16: Hands-On Labs and Exercises

Theory without practice is incomplete. This chapter provides step-by-step lab exercises to solidify your VMware knowledge. Each lab builds on previous chapters and gives you real hands-on experience.

Lab 1: Set Up a Home VMware Lab

You don't need enterprise hardware to learn VMware. This lab sets up a nested VMware environment on a single workstation — perfect for learning and certification preparation.

Requirements

- PC or Mac with at least 32GB RAM (64GB preferred)
- Modern CPU with Intel VT-x or AMD-V and hardware virtualization extensions
- 500GB+ SSD (NVMe preferred)
- VMware Workstation Pro 17 or Fusion 13 (licensed or evaluation)
- ESXi 8.0 ISO (60-day evaluation from VMware Customer Connect)
- VCSA 8.0 ISO (60-day evaluation)

Lab Steps

95. Install VMware Workstation Pro on your PC
96. Create a new VM in Workstation: 16 vCPUs, 32GB RAM, 200GB disk — this is your nested ESXi host
97. Enable nested virtualization in VM settings (Processors > Enable Intel VT-x/EPT)
98. Install ESXi 8.0 on this nested VM
99. Create a second VM with the same specs — second nested ESXi host
100. Create a third VM: 4 vCPUs, 16GB RAM, 100GB disk — this is for vCenter
101. Deploy VCSA on the first ESXi host
102. Add both ESXi hosts to vCenter
103. Create a shared NFS datastore (can use a Windows share or FreeNAS VM)
104. Create a cluster with both hosts
105. Enable HA and DRS on the cluster
106. Test VMotion between the two ESXi hosts

 **Note:** Nested ESXi performance is limited — don't expect production-grade results. But it's perfect for learning all the concepts, procedures, and troubleshooting. The VCP-DCV certification exam accepts nested lab experience.

Lab 2: VMotion Step-by-Step

Prerequisites: Two ESXi hosts in the same vCenter cluster with shared storage.

Exercise 2A: Manual VMotion

107. Create a test VM on Host A (4 vCPUs, 8GB RAM, 40GB thin disk)
108. Install Windows Server and VMware Tools
109. Start a continuous ping from the VM to its gateway: `ping -t 8.8.8.8`
110. Open a second browser window to vCenter
111. Right-click the VM > Migrate > Change compute resource only
112. Select Host B as the destination
113. Select VMotion for migration priority
114. Click Finish and watch the progress bar
115. Return to the ping window — did any pings drop? (They shouldn't!)
116. Verify the VM is now on Host B in vCenter

Exercise 2B: DRS-Initiated VMotion

117. Create 10-15 VMs on your cluster
118. Enable Fully Automated DRS on the cluster
119. Use a stress testing tool inside several VMs to generate CPU load
120. Watch vCenter as DRS automatically migrates VMs to balance the load
121. Review the DRS History in the cluster's Monitor tab

Lab 3: Configure HA and Test Failover

Exercise

122. Enable HA on the cluster with FTT=1 (N+1 failover capacity)
123. Power on several VMs and note which host each runs on
124. In VMware Workstation, suspend (simulate crash) the ESXi VM running some of those VMs
125. Watch vCenter — within 1-3 minutes, VMs should restart on the remaining host
126. Resume the 'failed' host and add it back to the cluster

★ **Pro Tip:** Test HA in your lab before you rely on it in production. Understanding exactly how it behaves — detection time, restart order, behavior during isolation — prepares you to respond correctly during a real incident.

Lab 4: Storage VMotion and Snapshot Management

Exercise 4A: Storage VMotion

127. Create two datastores in your lab environment
128. Deploy a VM on Datastore 1
129. Right-click the VM > Migrate > Change storage only
130. Select Datastore 2 as the destination
131. Choose 'Same format as source' for disk format
132. Click Finish and watch the VM's disk migrate live

Exercise 4B: Snapshot Operations

133. Take a snapshot of a running VM (include memory, quiesce guest FS)
134. Make significant changes inside the VM (install software, create files)
135. Take a second snapshot
136. Make more changes
137. Revert to the first snapshot — see all changes disappear
138. Use 'Delete All Snapshots' to consolidate the snapshot tree
139. Verify disk consolidation completed successfully in Recent Tasks

Appendix: Glossary and Quick Reference

A.1 VMware Glossary

Term	Definition
ESXi	VMware's bare-metal Type 1 hypervisor. The foundation of vSphere.
vCenter Server	Centralized management platform for ESXi hosts and VMs.
vSphere	VMware's complete virtualization platform (ESXi + vCenter + suite).
VM	Virtual Machine — a software emulation of a physical computer.
vCPU	Virtual CPU assigned to a VM. Scheduled on physical CPU by VMware.
VMDK	Virtual Machine Disk — the virtual hard disk file format used by VMware.
VMotion	Live migration of a running VM between ESXi hosts with zero downtime.
Storage VMotion	Live migration of a VM's disk files between datastores.
HA	High Availability — automatic VM restart after host failure.
DRS	Distributed Resource Scheduler — automatic VM load balancing via VMotion.
FT	Fault Tolerance — zero-downtime protection via live shadow VM.
vSAN	Hyper-converged storage using ESXi hosts' local disks.
NSX-T	VMware's software-defined networking and security platform.
vDS	vSphere Distributed Switch — centrally managed virtual network switch.

vSS	vSphere Standard Switch — per-host managed virtual network switch.
vmkernel	The core of ESXi that manages hardware resources and VMs.
vmkernel port	A network adapter used for ESXi host traffic (management, VMotion, etc.)
VMFS	Virtual Machine File System — VMware's cluster file system for datastores.
Datastore	Storage repository where VM files are stored.
OVA/OVF	Open Virtualization Archive/Format — standard VM packaging format.
EVC	Enhanced VMotion Compatibility — CPU feature masking for cross-generation VMotion.
VADP	vSphere APIs for Data Protection — backup integration API.
RDM	Raw Device Mapping — direct access to physical storage from a VM.
Snapshot	Point-in-time capture of VM state. NOT a backup.
Template	Master VM image used for rapid, consistent VM deployment.
Content Library	Centralized repository for templates and ISOs across vCenters.
SRM	Site Recovery Manager — DR orchestration platform.
VCF	VMware Cloud Foundation — complete SDDC stack bundle.
DCUI	Direct Console User Interface — text-based console on ESXi hosts.
SSO	Single Sign-On — vCenter's authentication service.
HCL	Hardware Compatibility List — VMware's list of tested/certified hardware.

FTT	Failures to Tolerate — vSAN's redundancy configuration.
NIOC	Network I/O Control — bandwidth allocation between traffic types.
vLCM	vSphere Lifecycle Manager — centralized patching and upgrade management.
RBAC	Role-Based Access Control — permission management system.
Thin Provisioning	Virtual disk that only uses space as data is written.
Thick Provisioning	Virtual disk with pre-allocated space on the datastore.
HBA	Host Bus Adapter — physical adapter connecting a server to storage network.

A.2 VMotion Quick Reference

VMotion Requirements Checklist

✓ vCenter Server installed and hosts added to cluster ✓ Shared storage accessible from both source and destination hosts ✓ VMotion vmkernel port configured on both hosts ✓ VMotion vmkernel can reach destination host's VMotion vmkernel ✓ Compatible CPU features (or EVC enabled on cluster) ✓ Matching network port groups on destination host ✓ Sufficient RAM available on destination host ✓ VM has no local device connections (physical CD/DVD, USB pass-through)

A.3 Common VMware Ports

Port	Purpose
443	vCenter Server HTTPS (web interface, API)
902	ESXi host communication, VMware Tools heartbeat
8443	vSphere Client HTTPS
5989	CIM (hardware monitoring) over HTTPS

8000	VMotion traffic (legacy; now uses 902)
2049	NFS datastore traffic
3260	iSCSI storage traffic
12301-12321	FT (Fault Tolerance) logging traffic
161/162	SNMP monitoring
514	Syslog traffic from ESXi hosts

A.4 VMware Certification Path

Certification	Description
VCP-DCV	VMware Certified Professional — Data Center Virtualization. The most common VMware cert.
VCAP-DCV Deploy	Advanced Deployment skills — hands-on lab exam.
VCAP-DCV Design	Advanced Design skills — scenario-based design exam.
VCDX-DCV	VMware Certified Design Expert — highest VMware certification. Defense required.
VCP-NV	VCP for Network Virtualization (NSX-T focus)
VCP-Cloud	VCP for Cloud Management and Automation
VCP-SEC	VCP for Security (Carbon Black, NSX Security)

A.5 Essential PowerCLI Commands

Daily PowerCLI Reference

```
# Connect to vCenter Connect-VIServer vcenter.company.com # Get all VMs Get-VM # Get all ESXi hosts Get-VMHost # Get VMs on a specific host Get-VMHost 'esxi01' | Get-VM # VMotion a VM Move-VM -VM 'WebServer01' -Destination (Get-VMHost 'esxi02') # Get VM snapshot info Get-VM | Get-Snapshot | Select VM, Name, Created # Get datastore usage Get-Datastore | Select Name, CapacityGB, FreeSpaceGB # Put host in maintenance mode
```

```
Set-VMHost 'esxi01' -State Maintenance # Get cluster info Get-Cluster | Select Name,
HAEnabled, DrsEnabled
```

End of VMware Complete Course Guide

Author: Gautam Shirsath

<https://www.linkedin.com/in/gautam-s2024>